

Convolutional Neural Networks

Elena Congeduti, 20-11-2024



Lecture's Agenda

1. Structured data
2. Convolutions
3. Convolutional Layers
4. Relation to Fully Connected Layers

What makes images different?



What makes images different?

1. High dimensionality

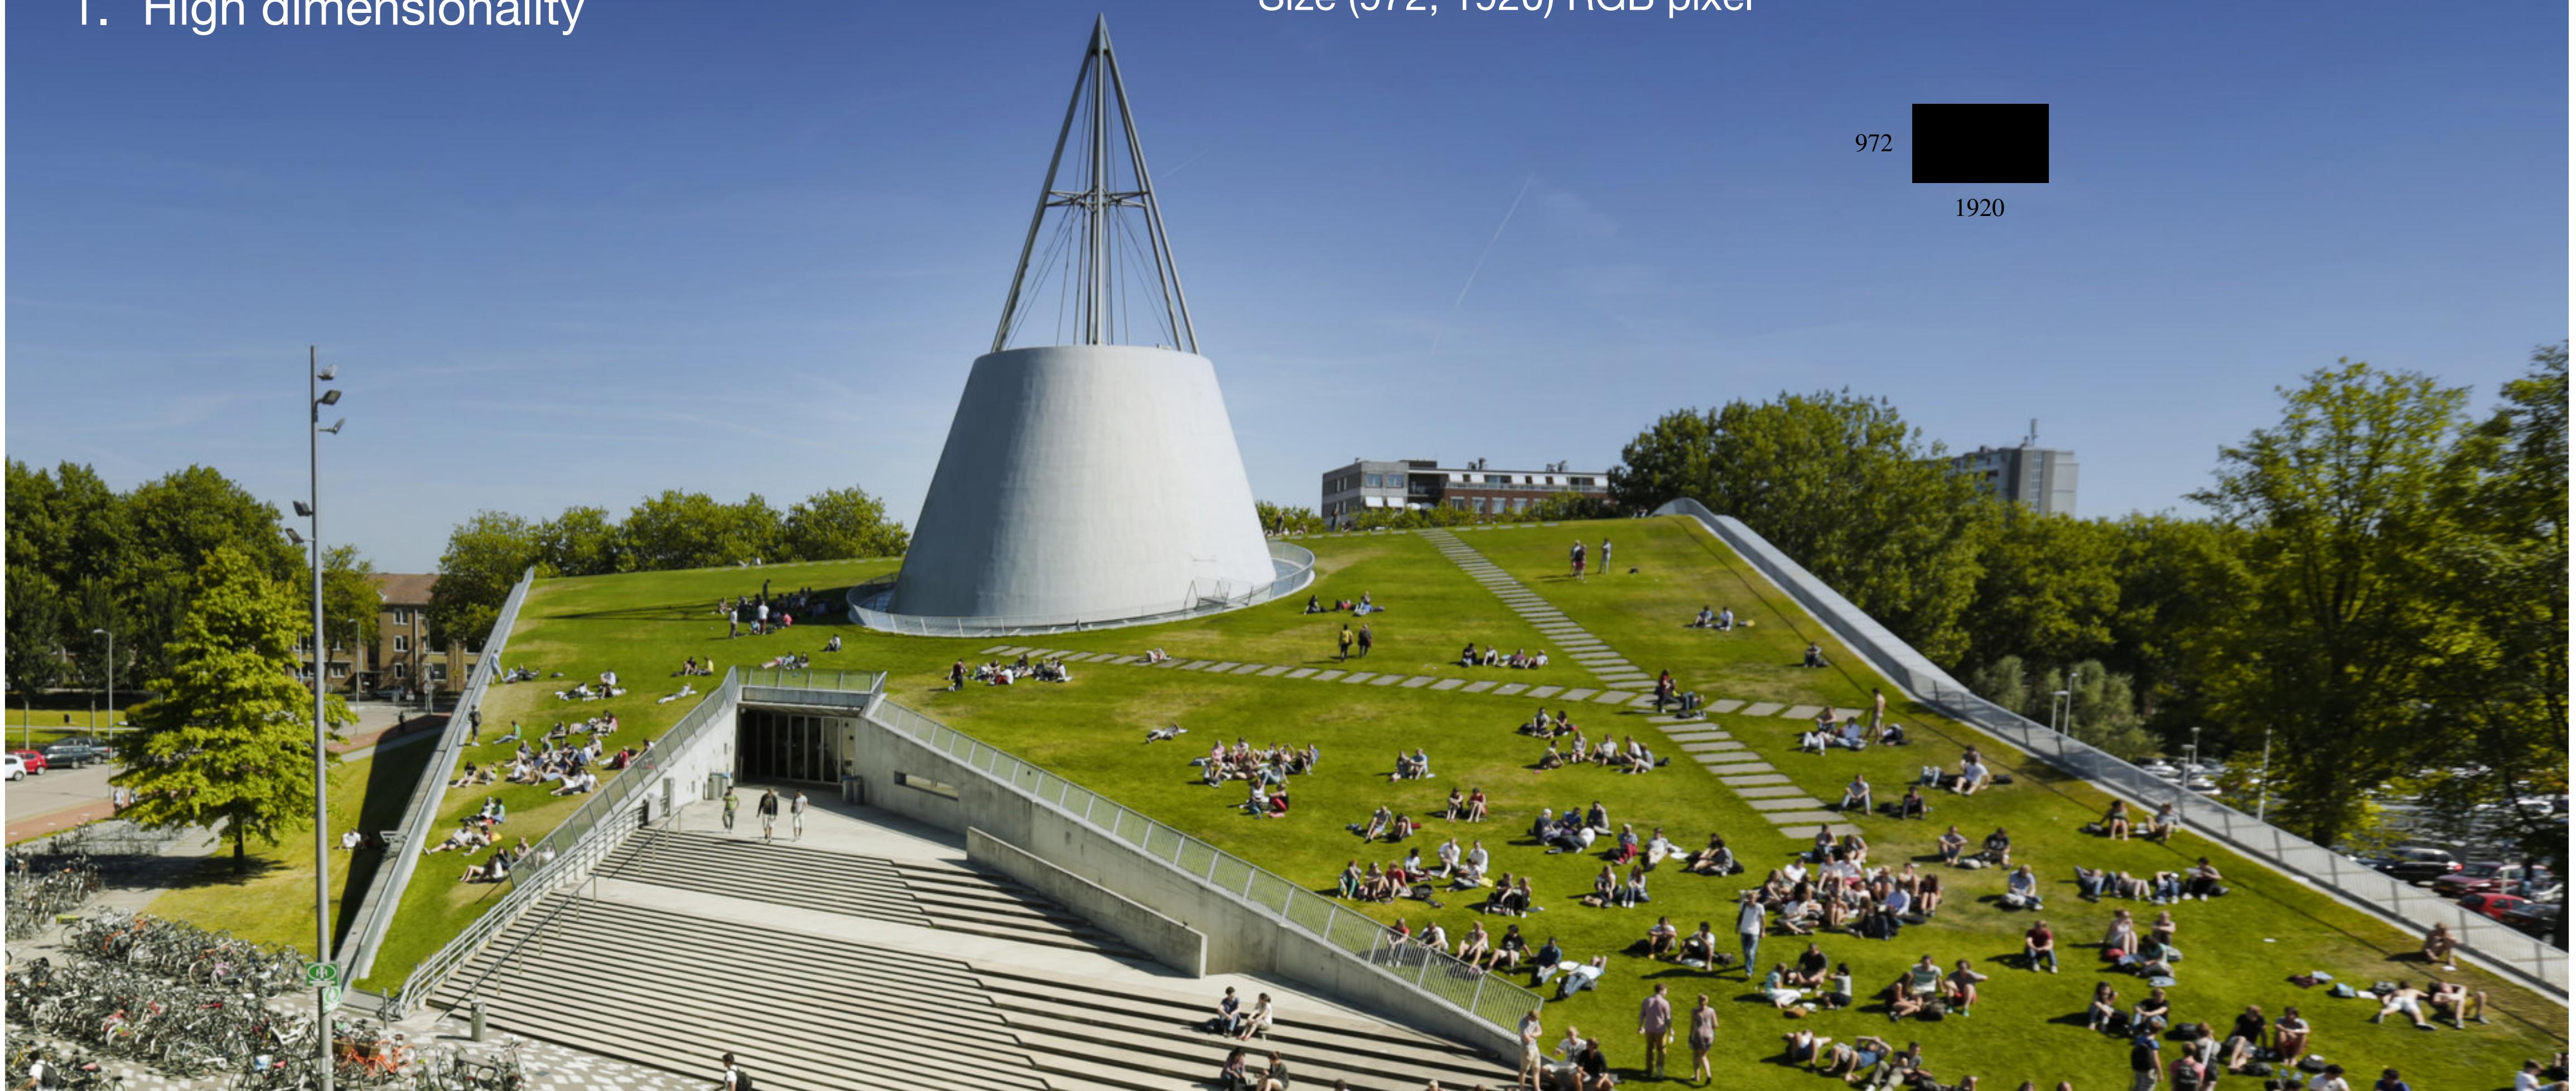


What makes images different?

1. High dimensionality

Size (972, 1920) RGB pixel

972
1920



What makes images different?

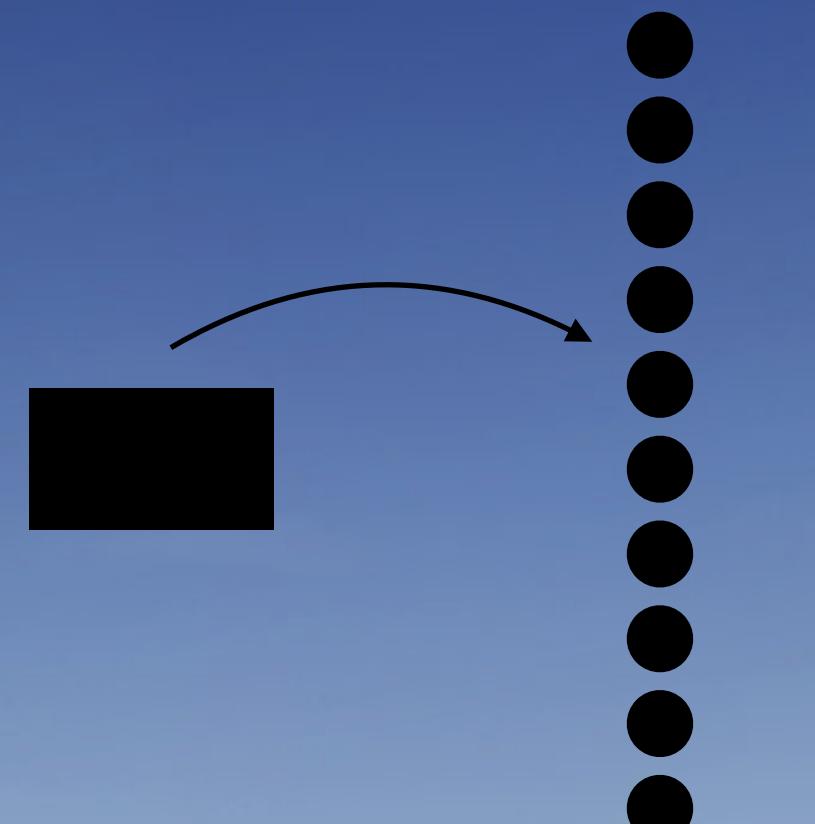
1. High dimensionality



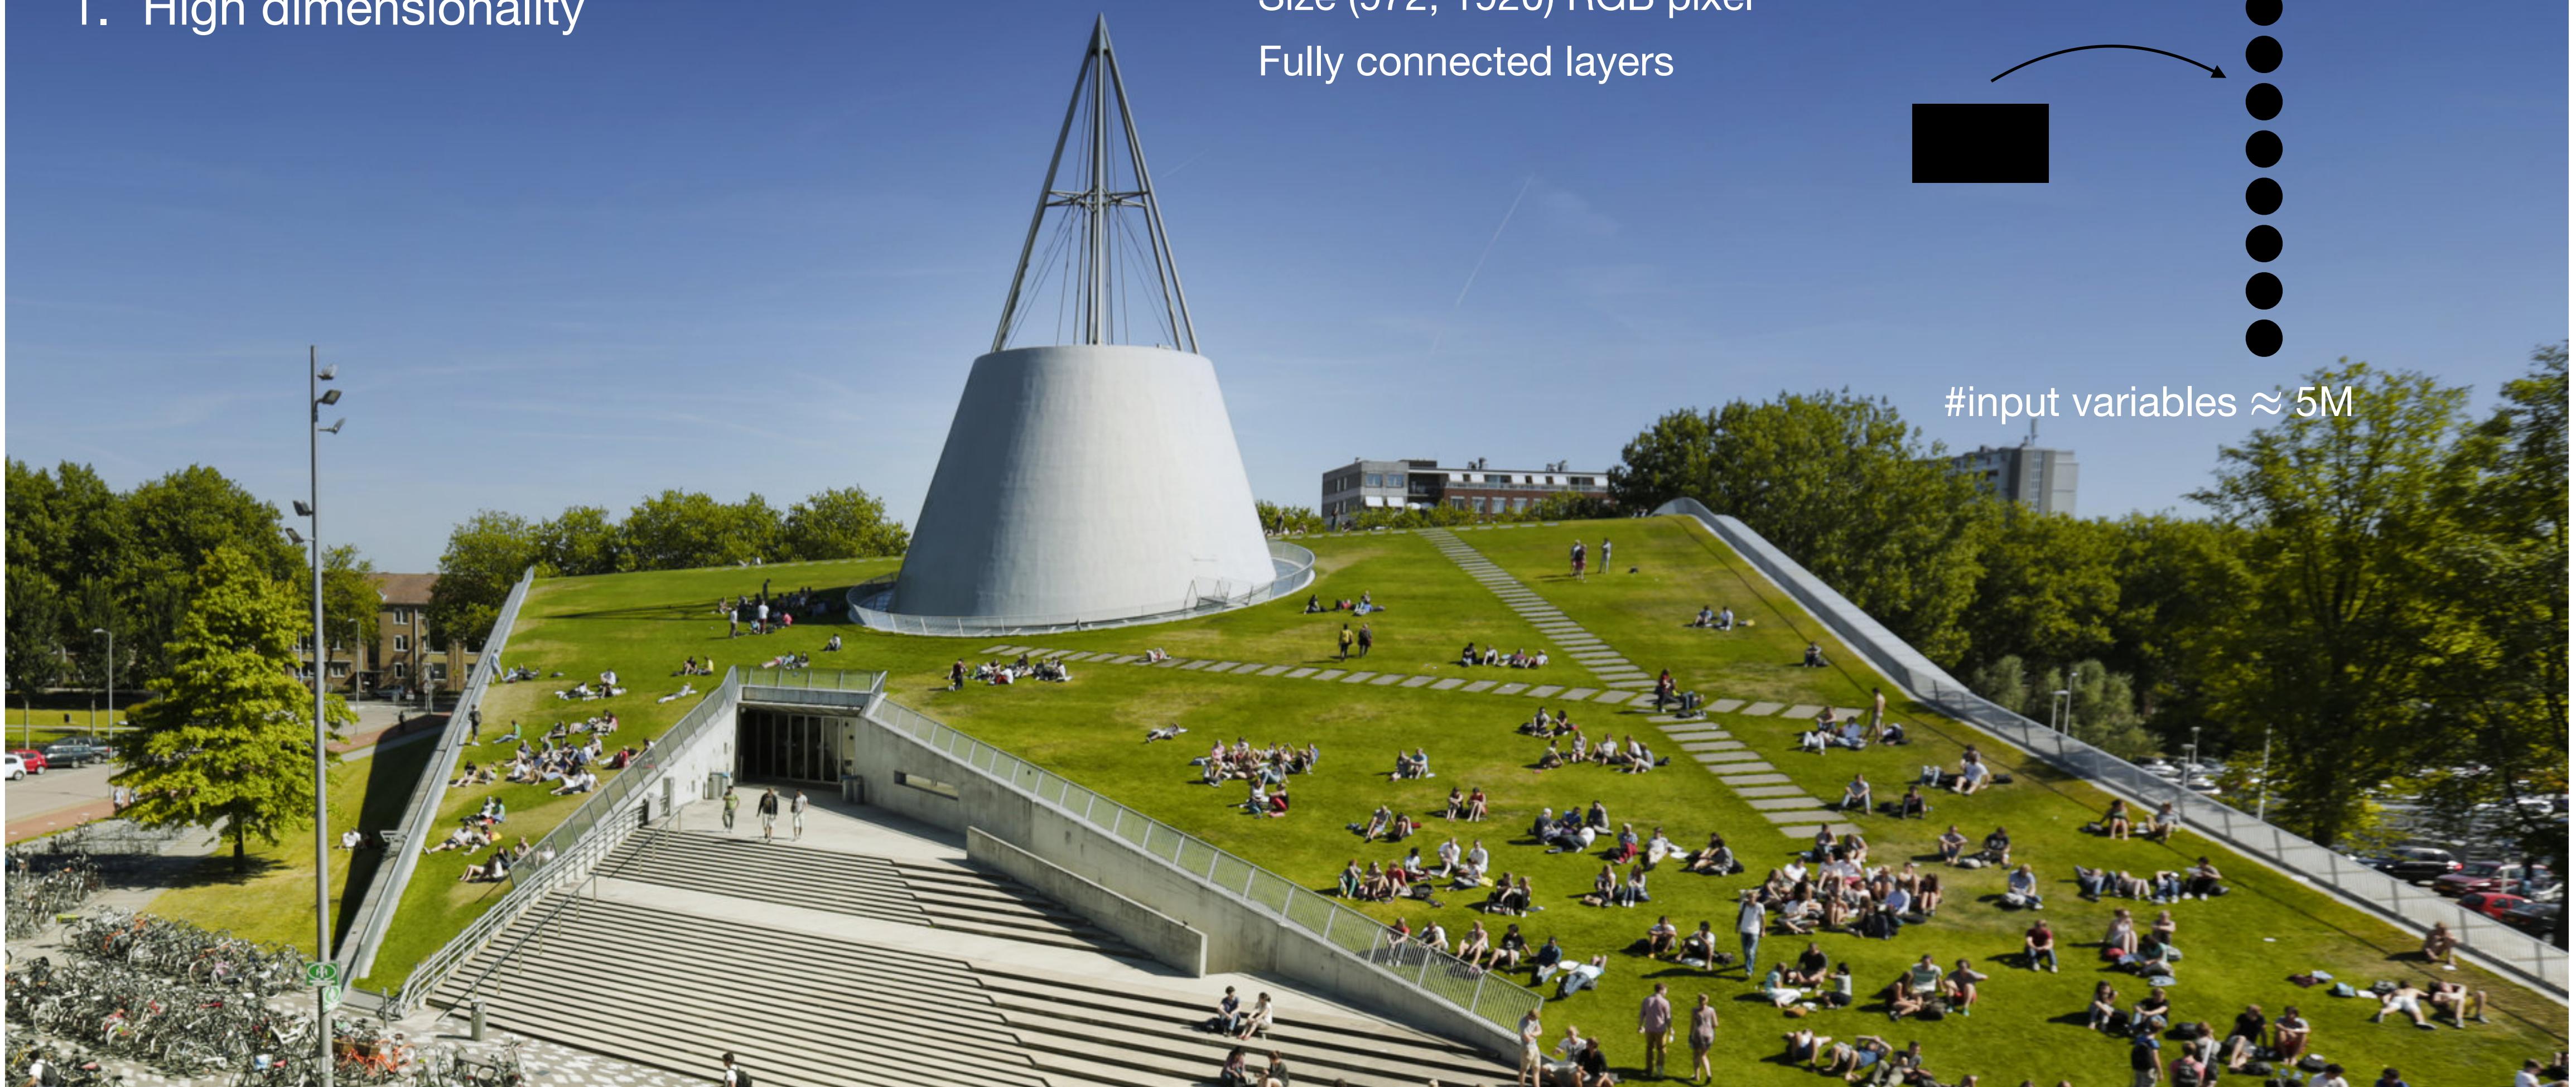
What makes images different?

1. High dimensionality

Size (972, 1920) RGB pixel
Fully connected layers



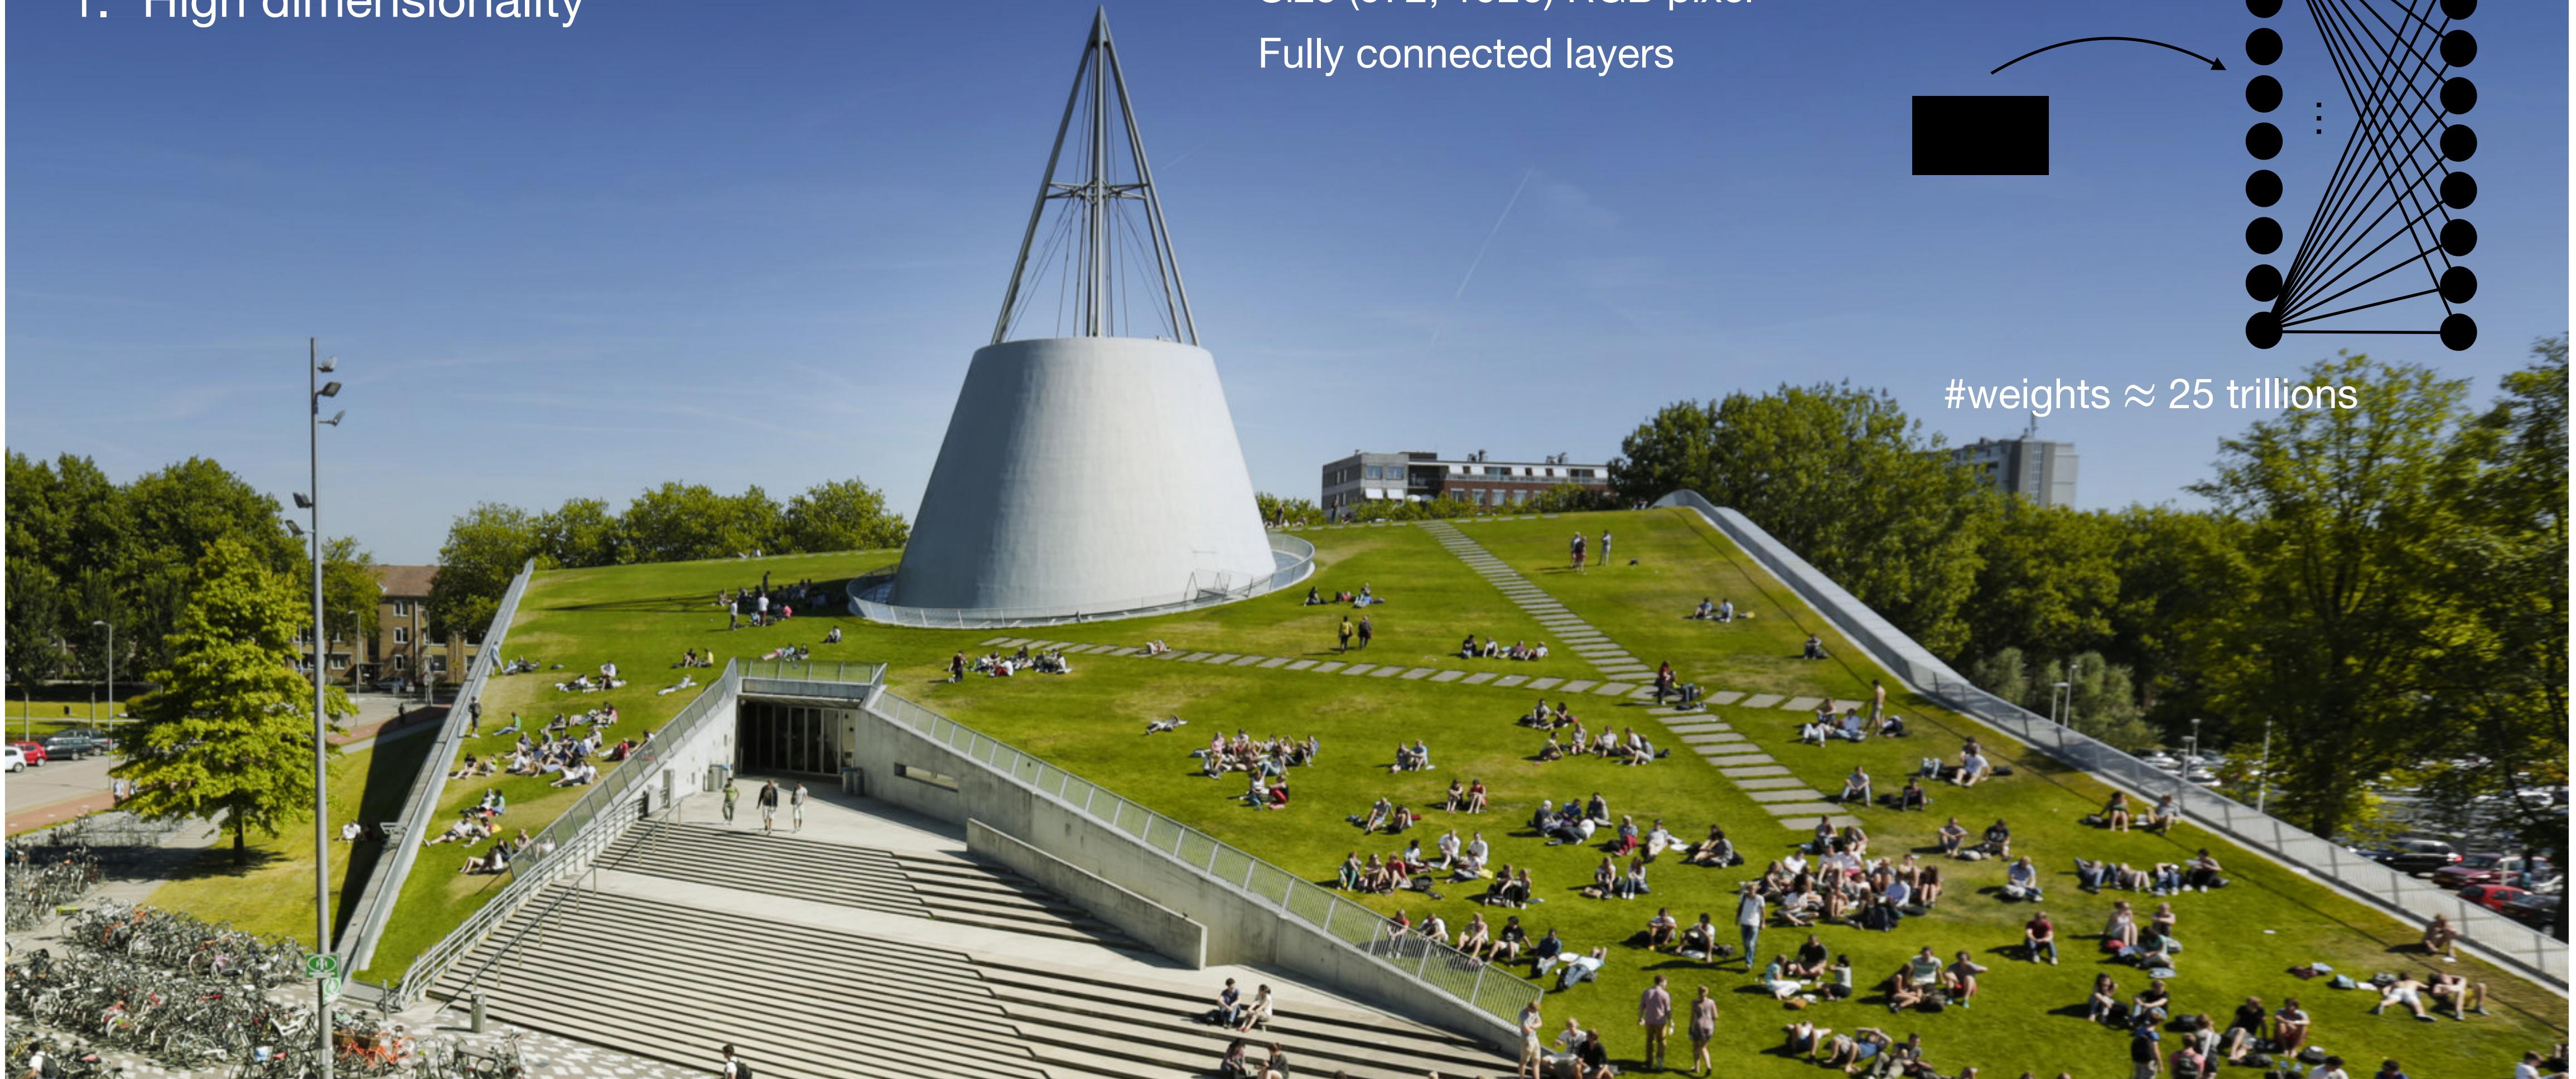
#input variables \approx 5M



What makes images different?

1. High dimensionality

Size (972, 1920) RGB pixel
Fully connected layers



#weights \approx 25 trillions

What makes images different?

1. High dimensionality
2. Spacial structure



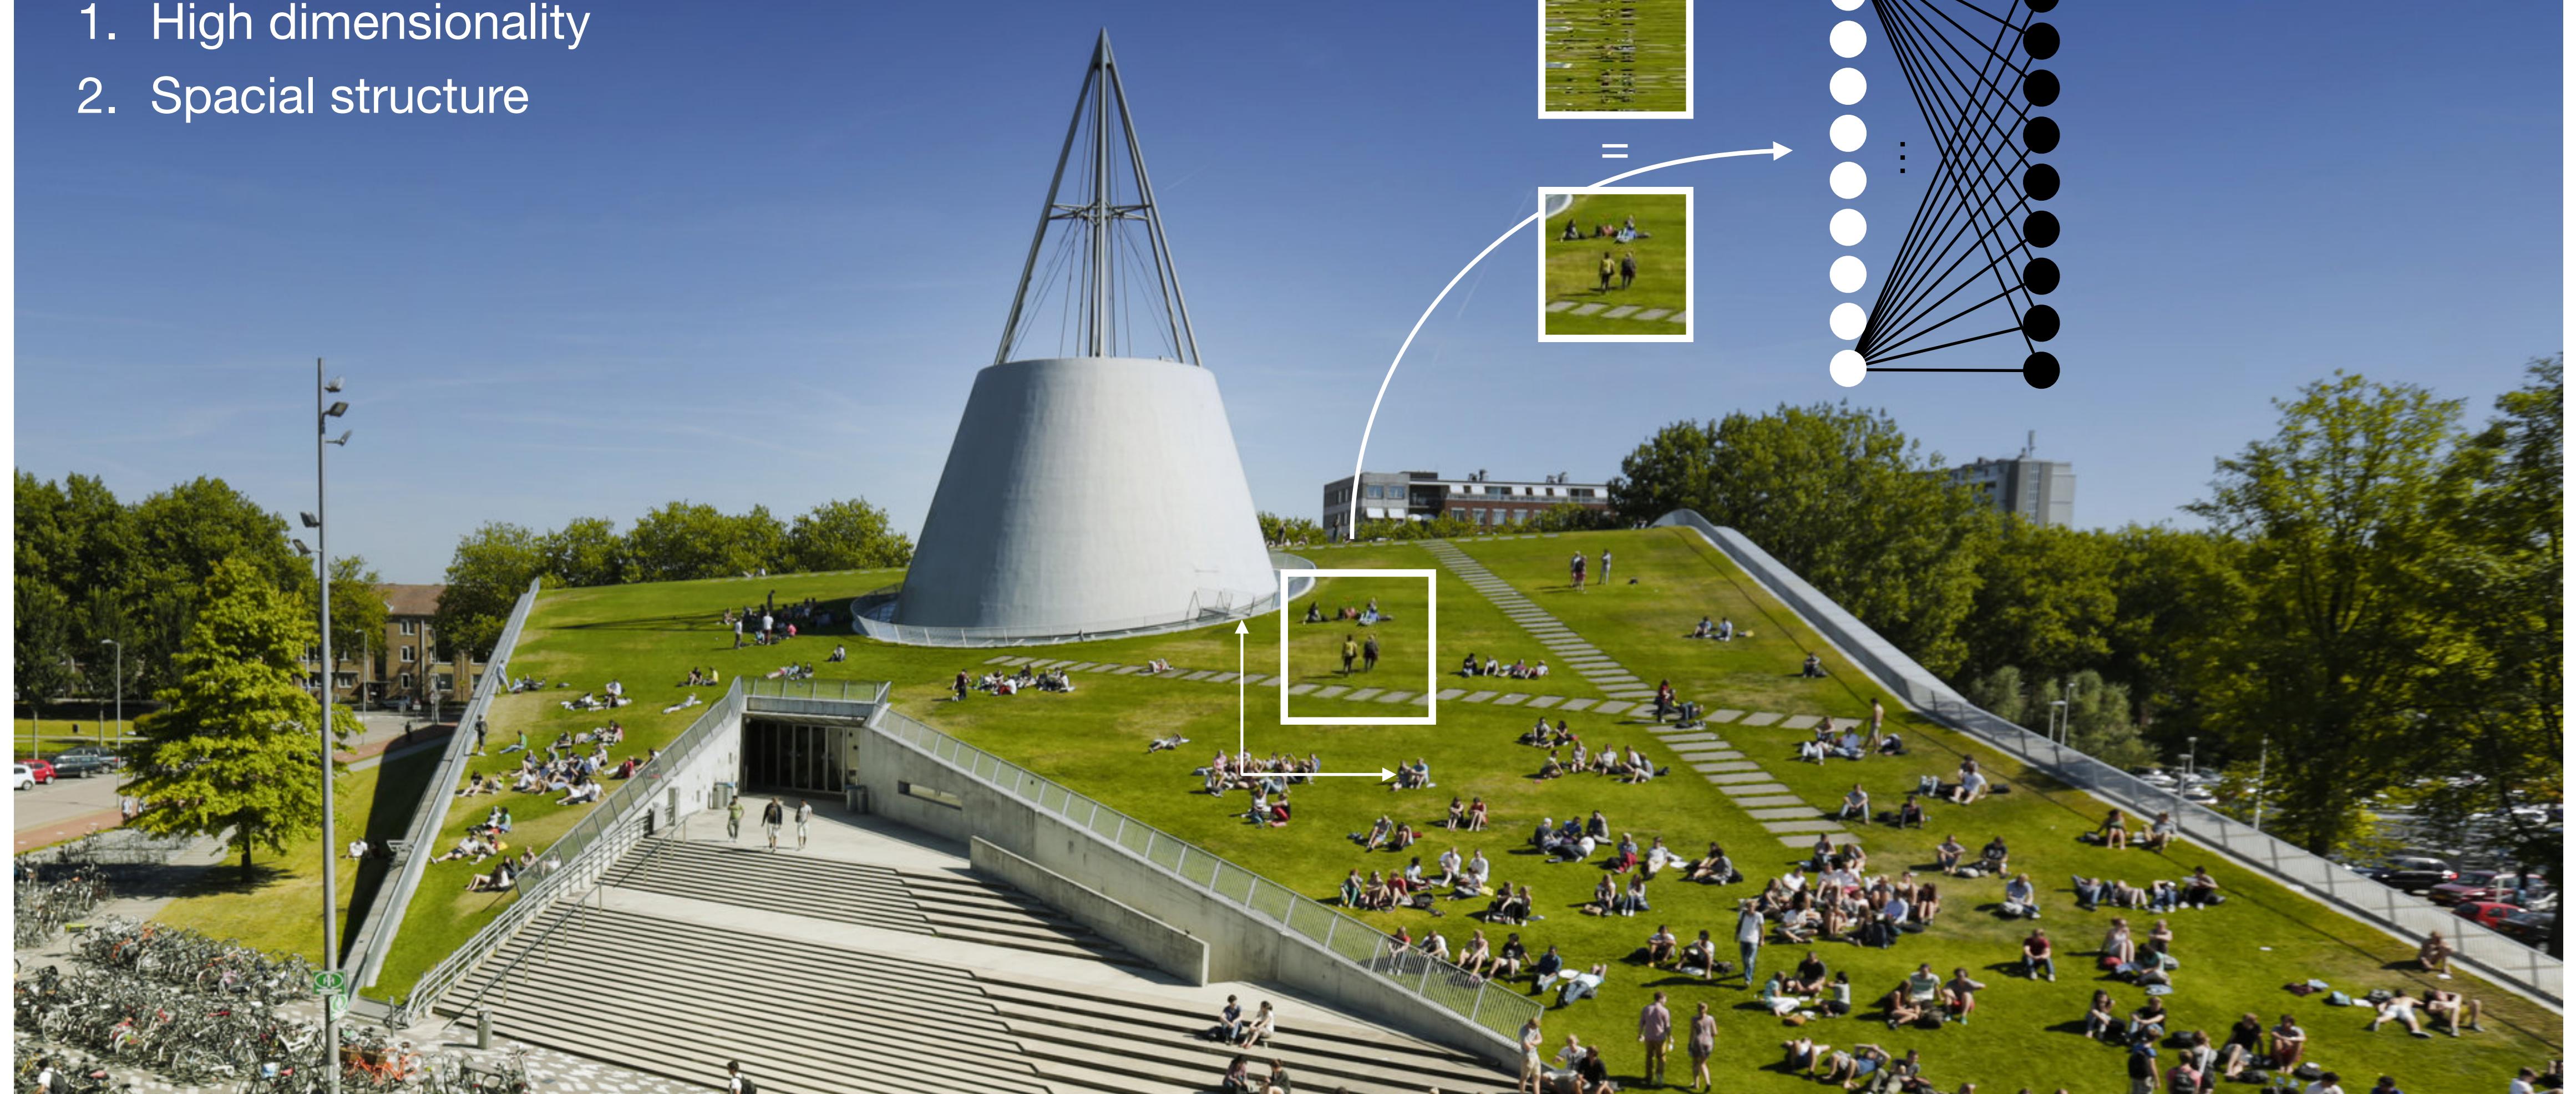
What makes images different?

1. High dimensionality
2. Spacial structure



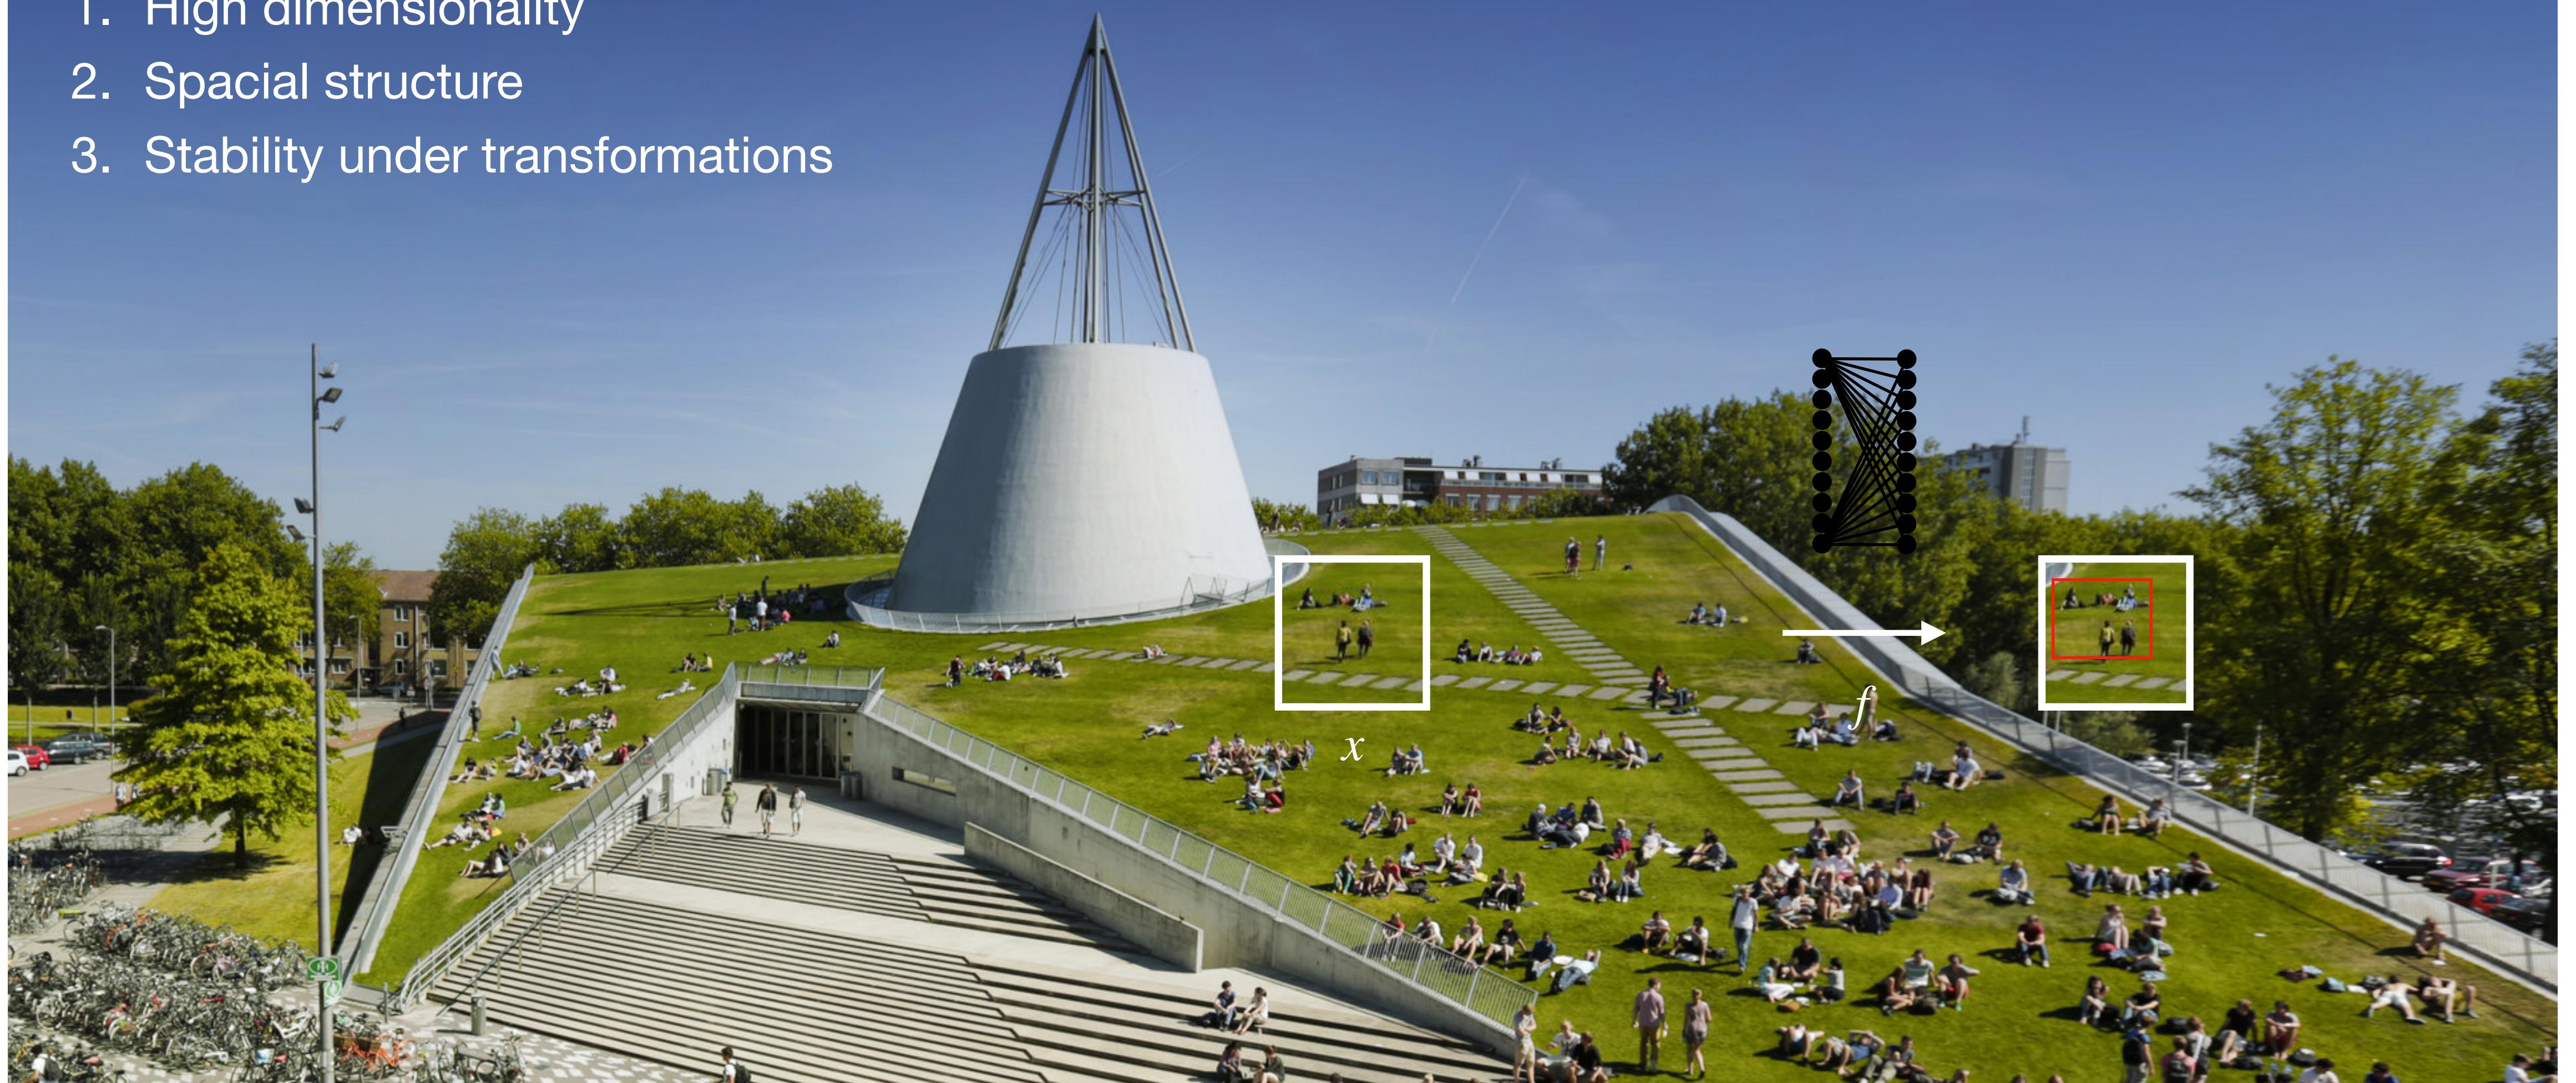
What makes images different?

1. High dimensionality
2. Spacial structure



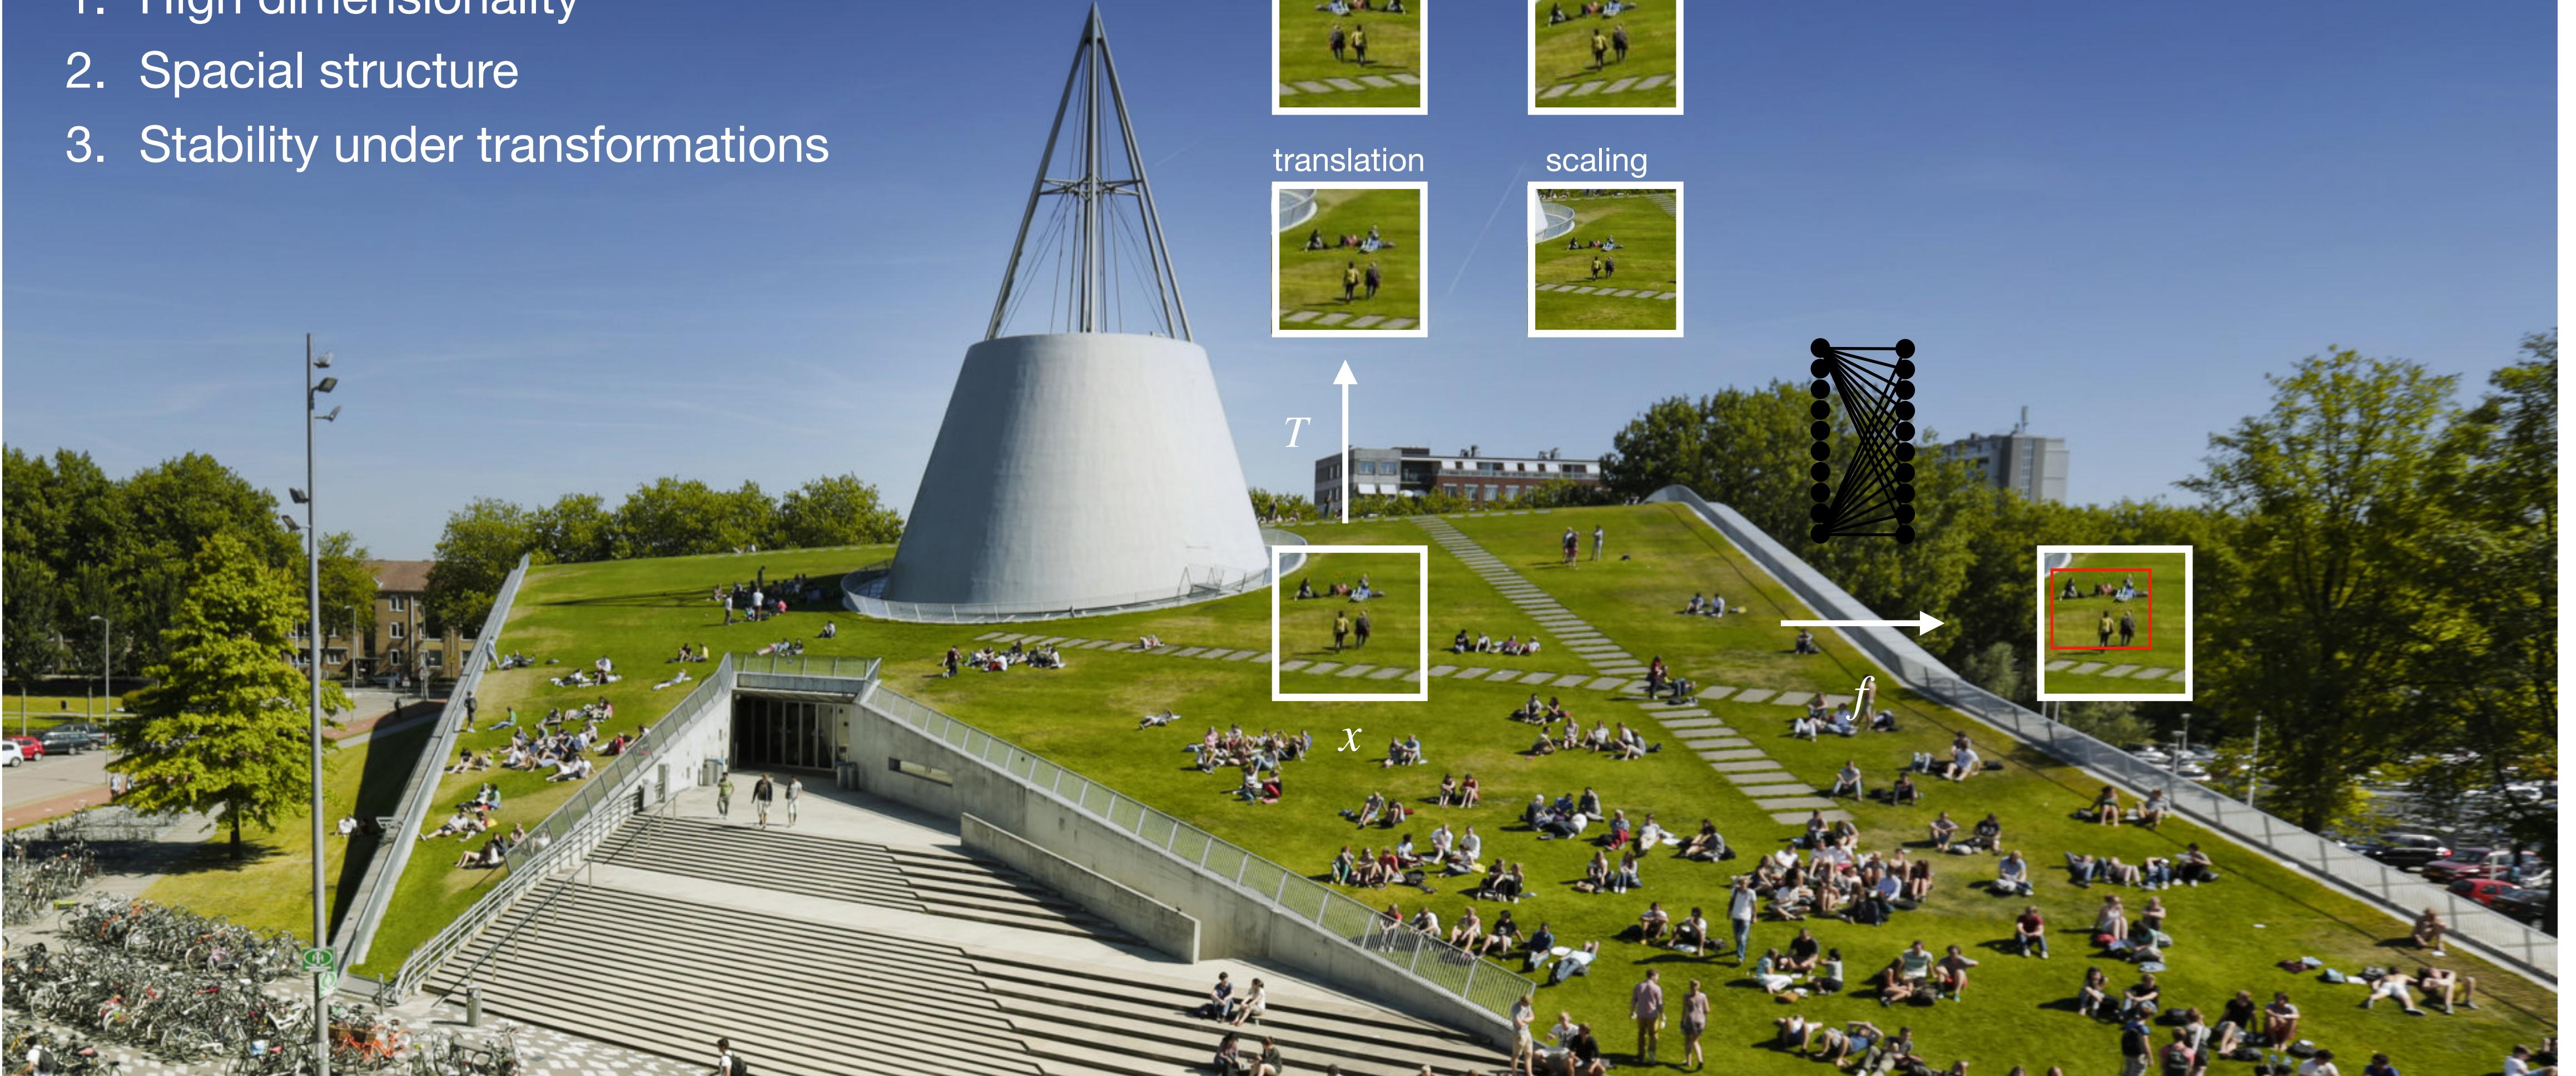
What makes images different?

1. High dimensionality
2. Spacial structure
3. Stability under transformations



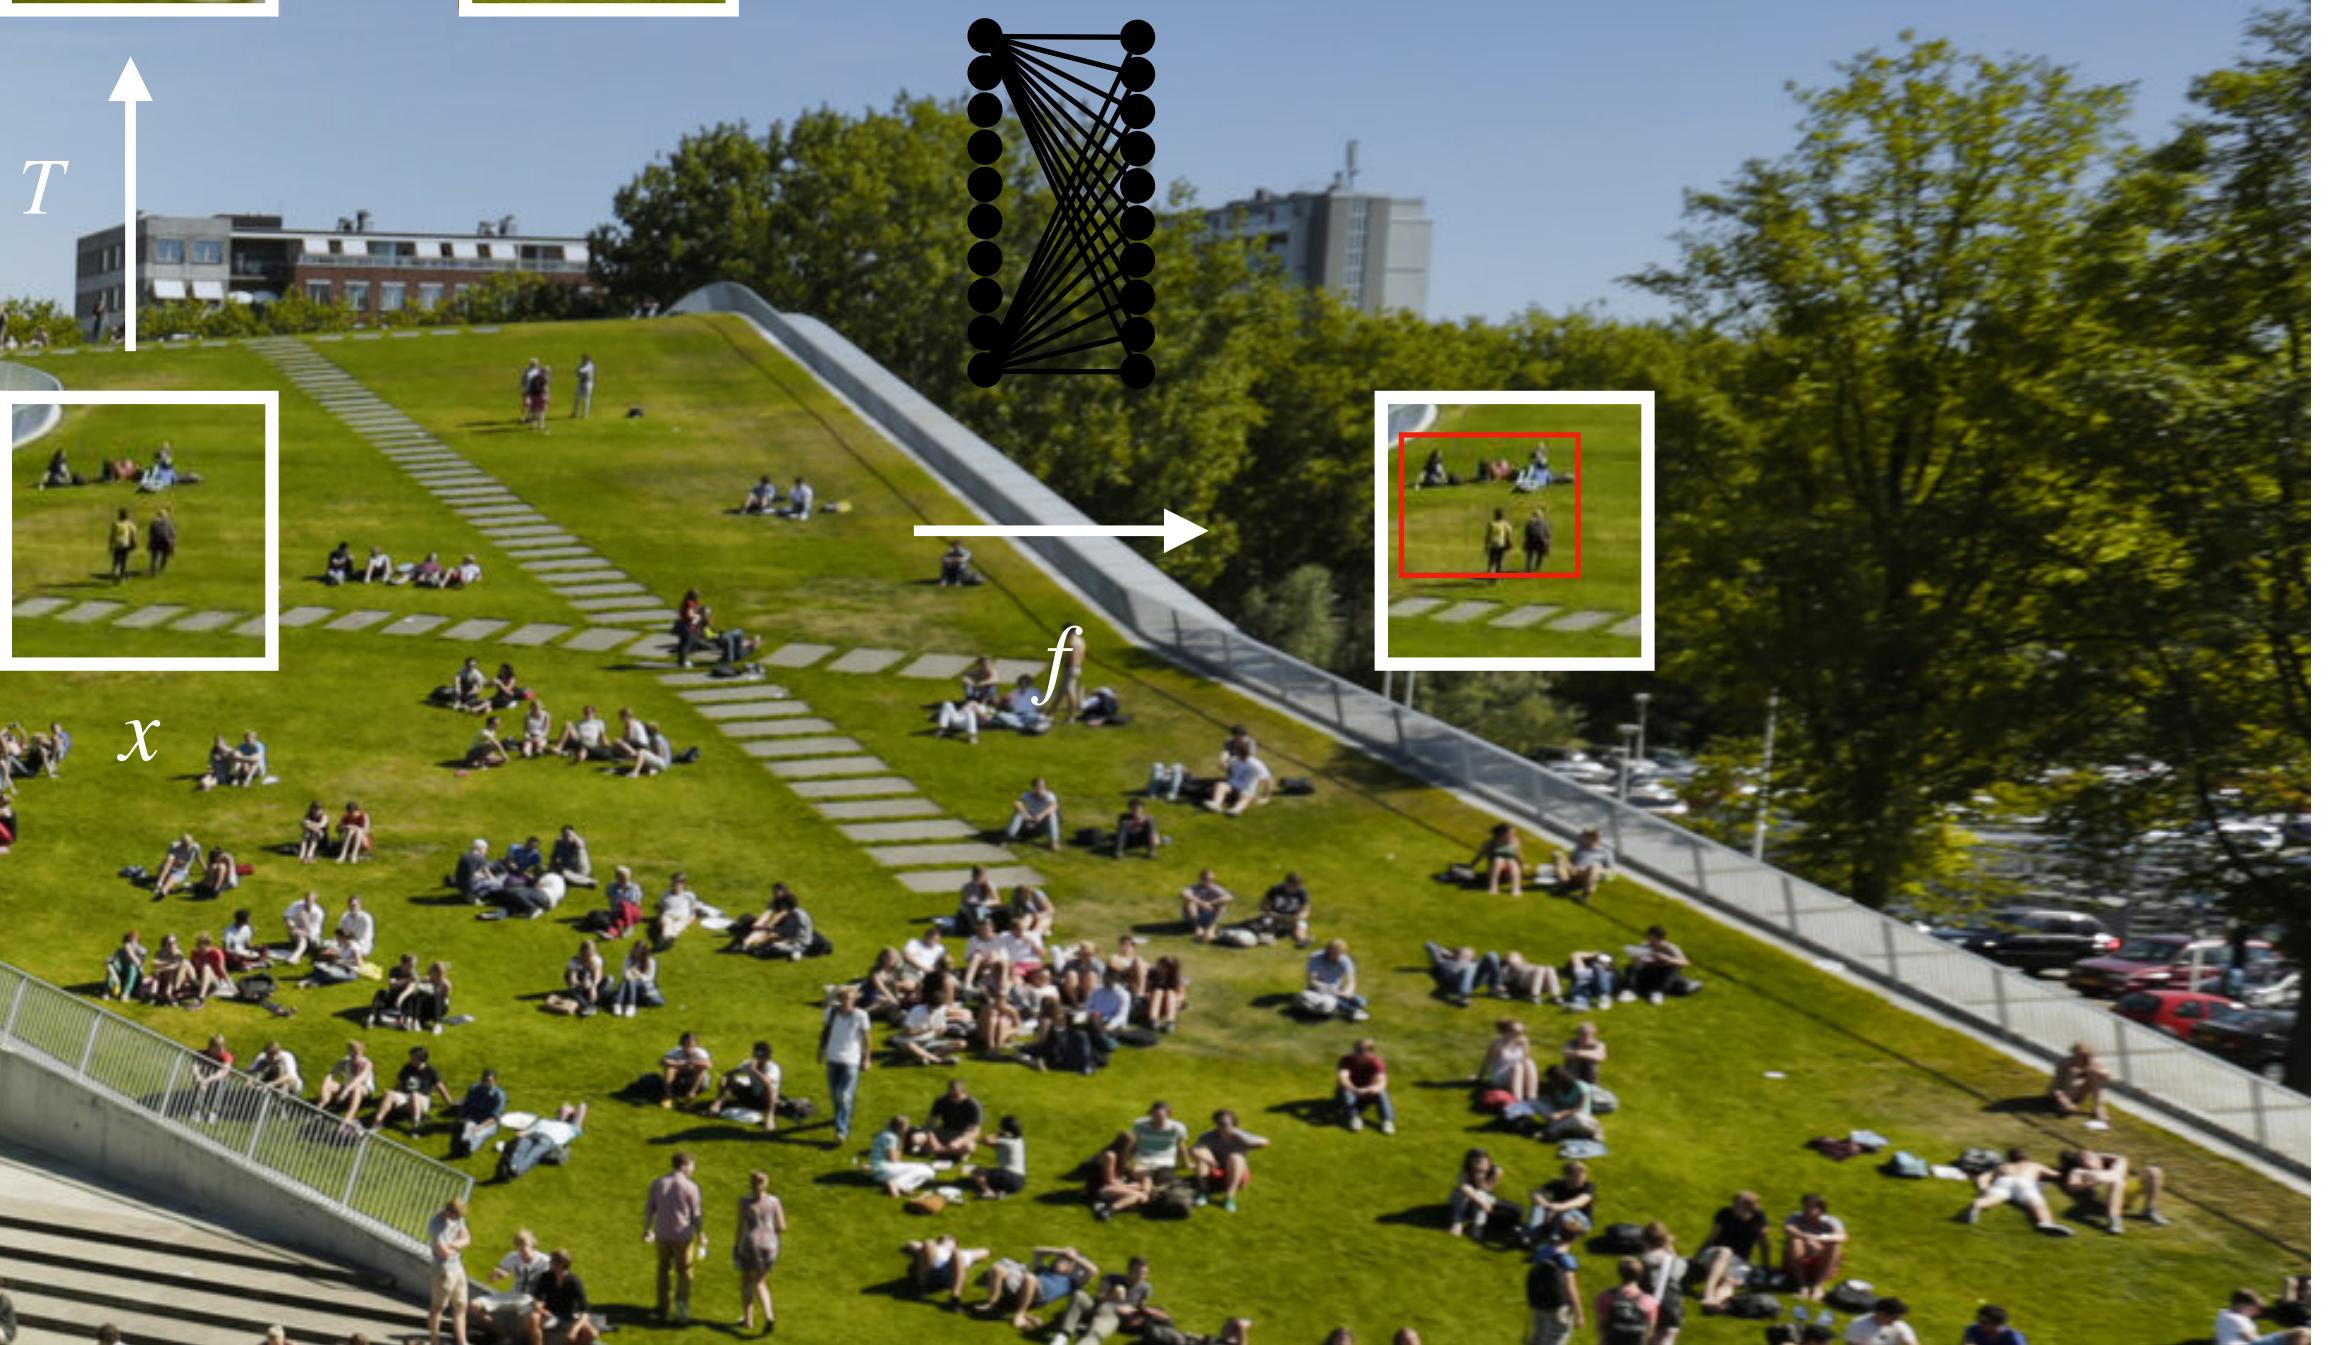
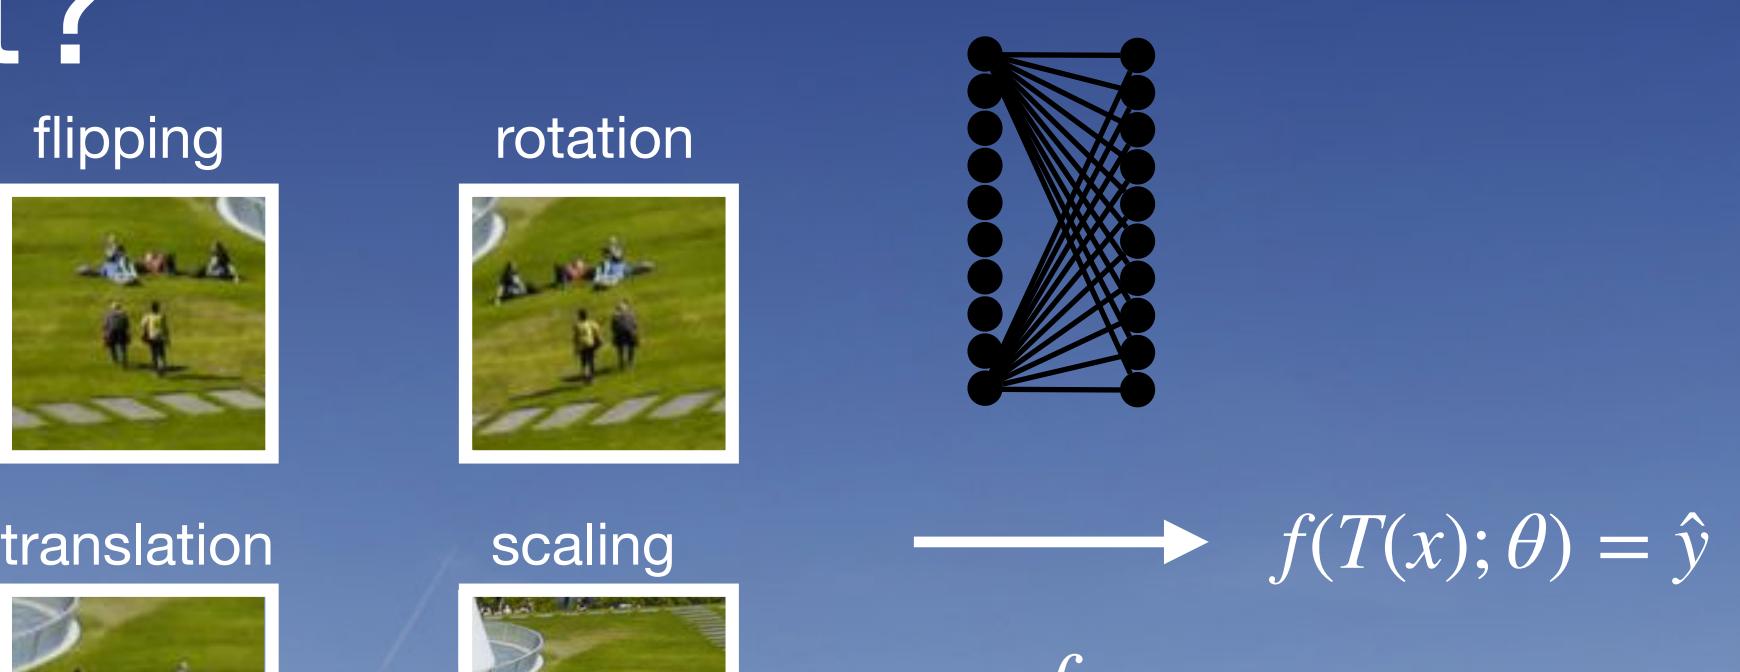
What makes images different?

1. High dimensionality
2. Spacial structure
3. Stability under transformations



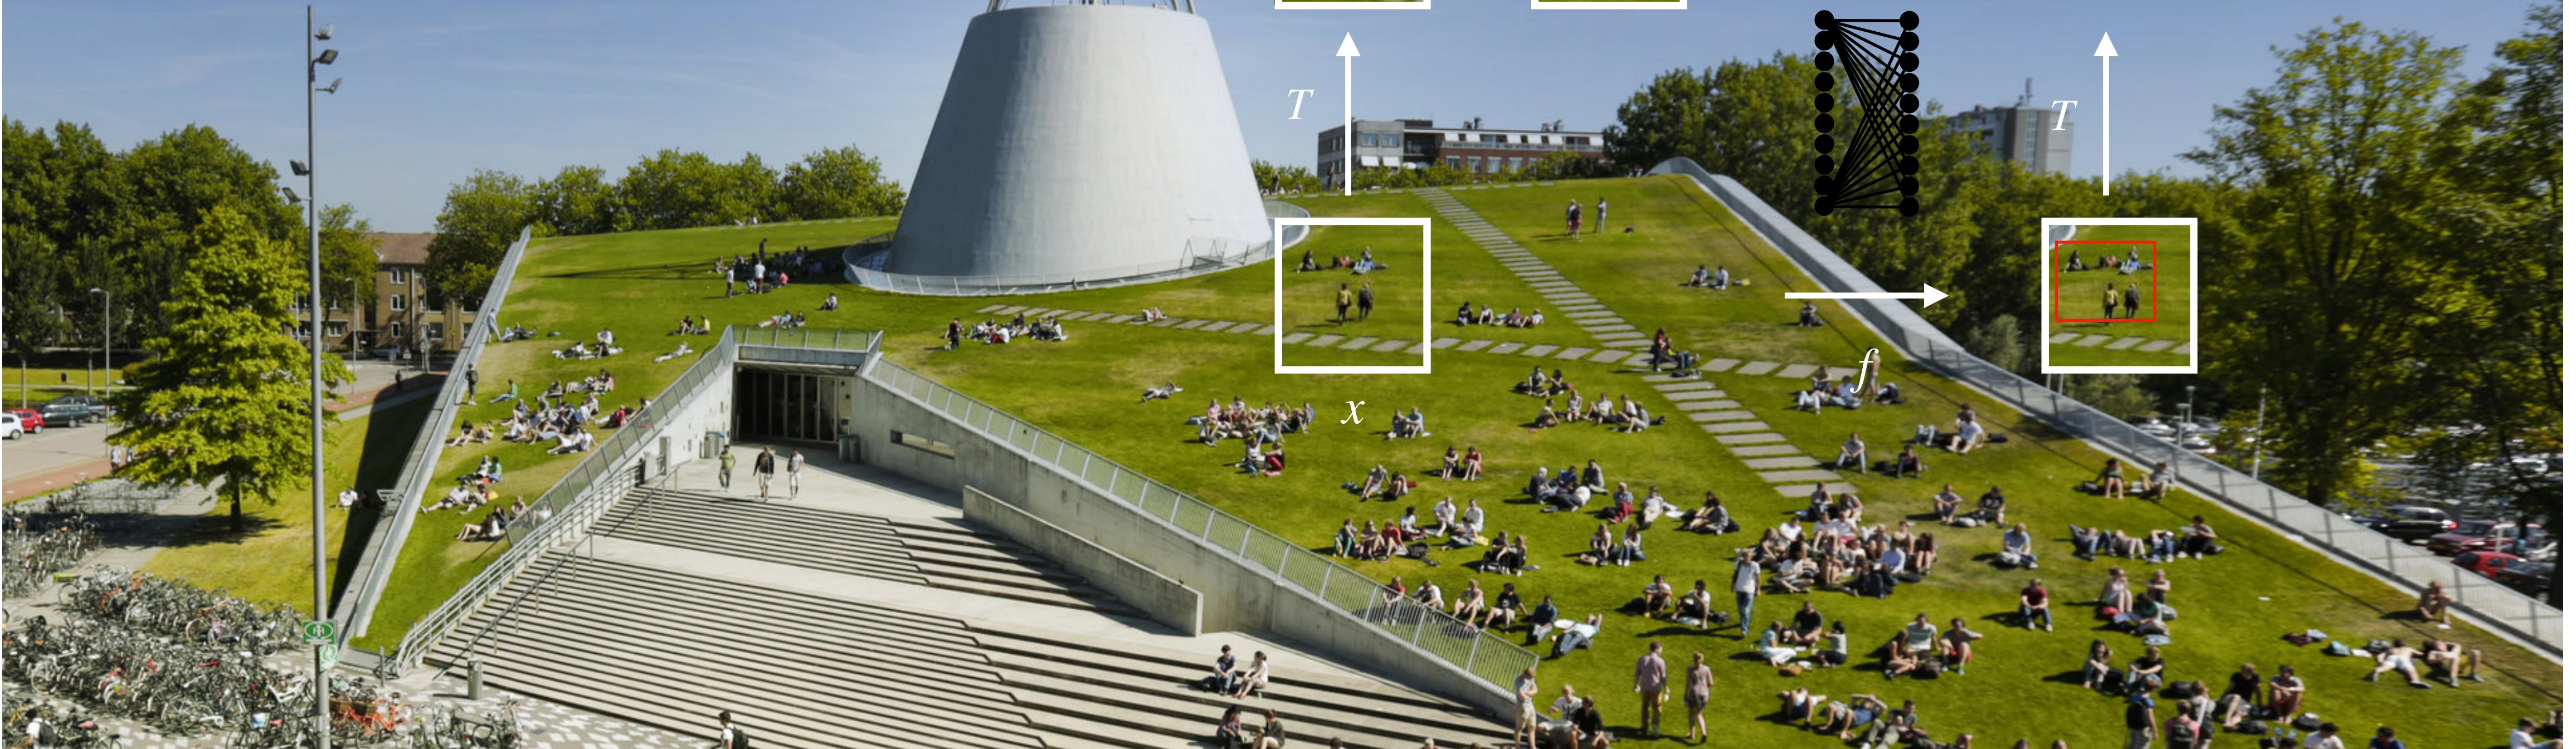
What makes images different?

1. High dimensionality
2. Spacial structure
3. Stability under transformations



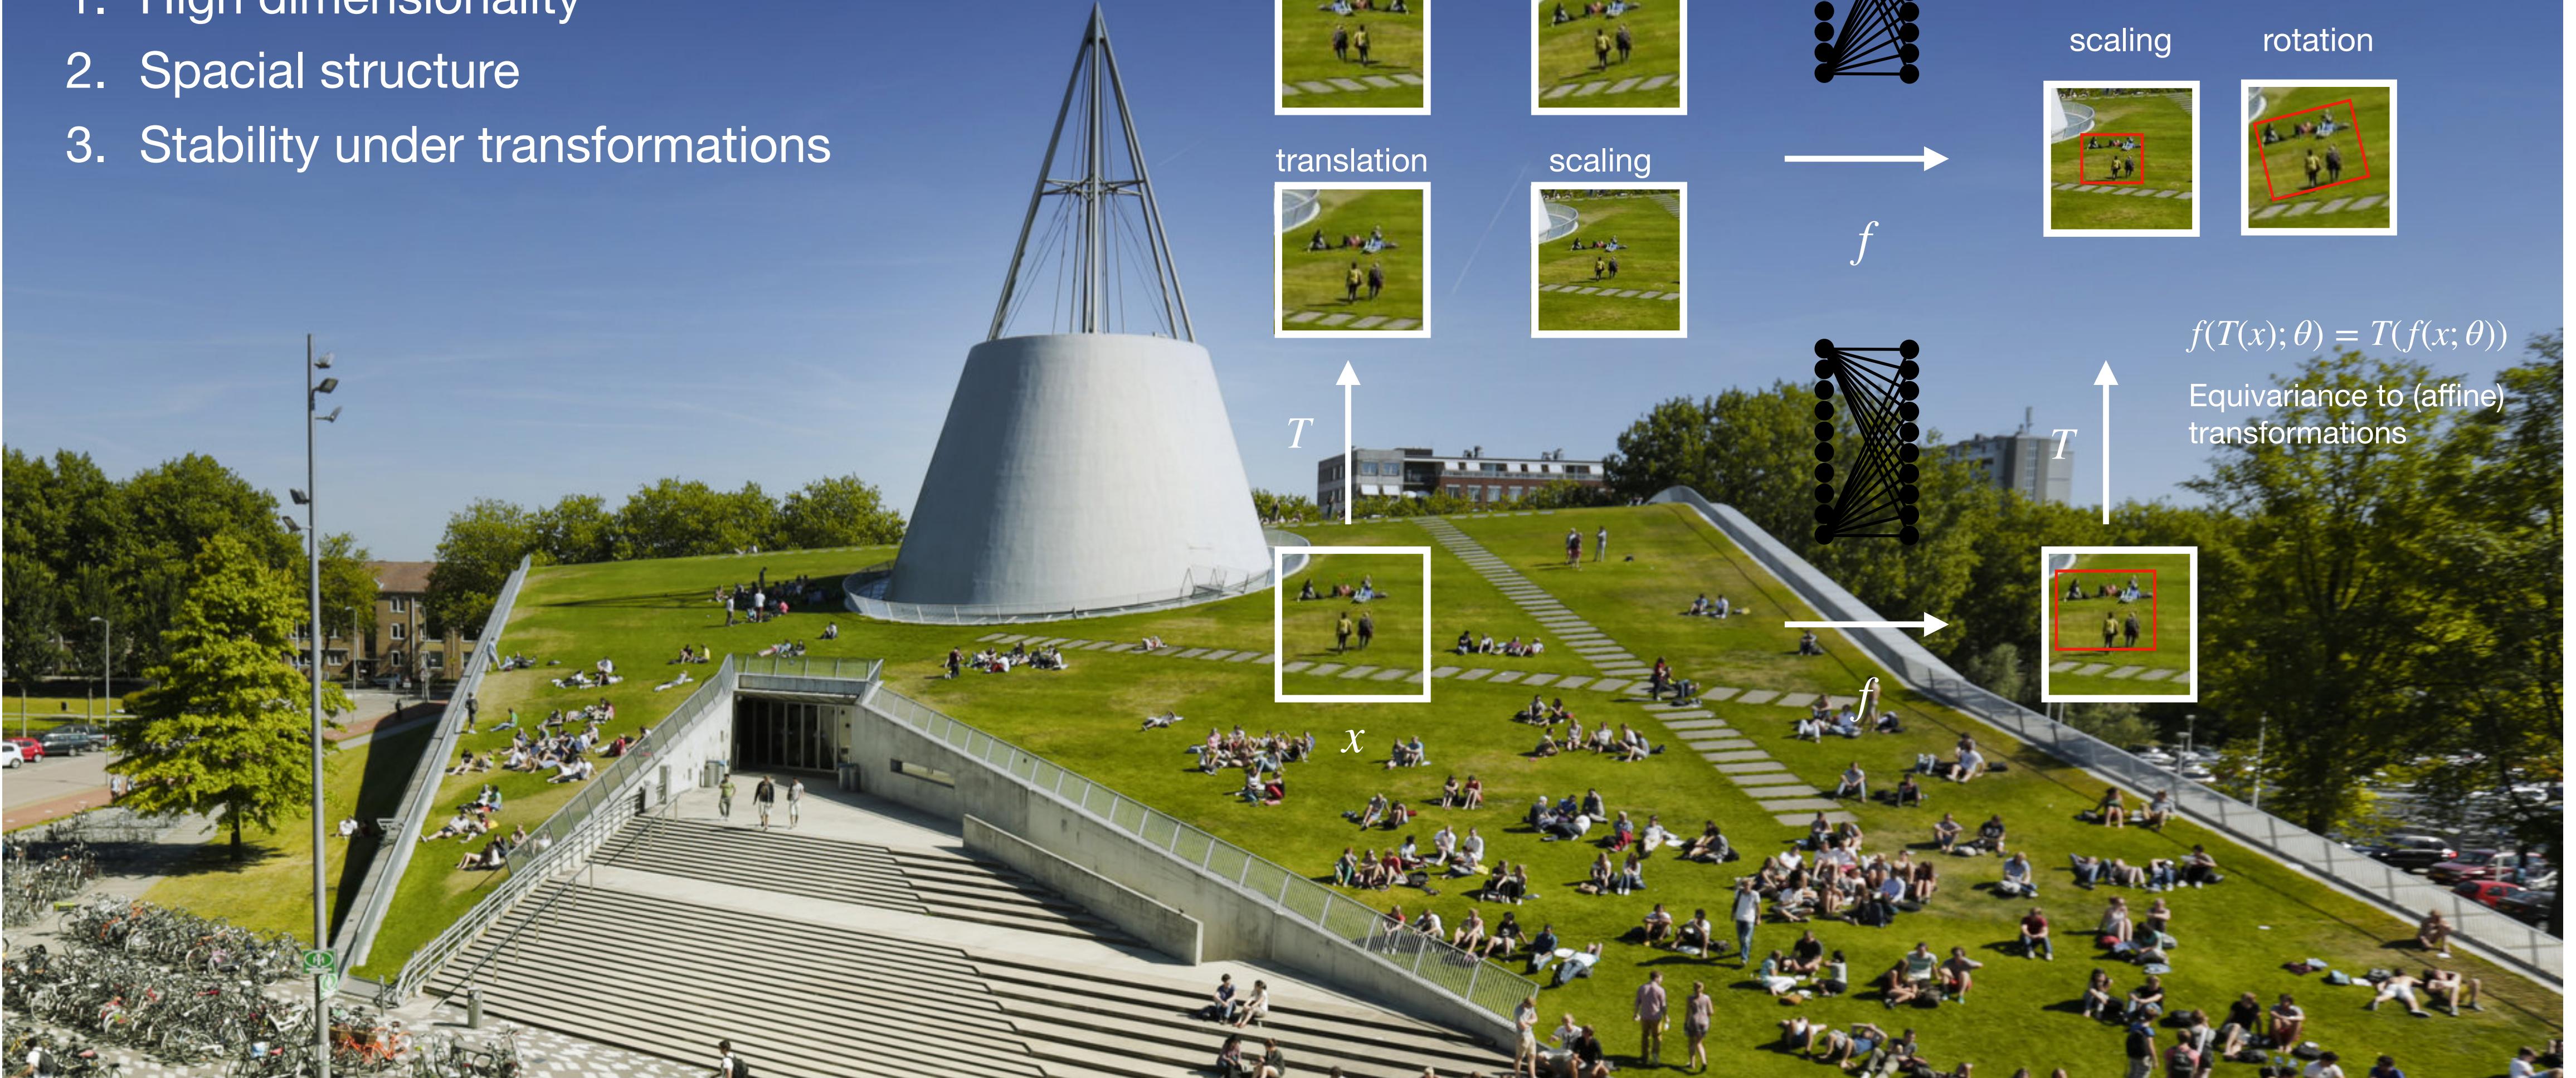
What makes images different?

1. High dimensionality
2. Spacial structure
3. Stability under transformations



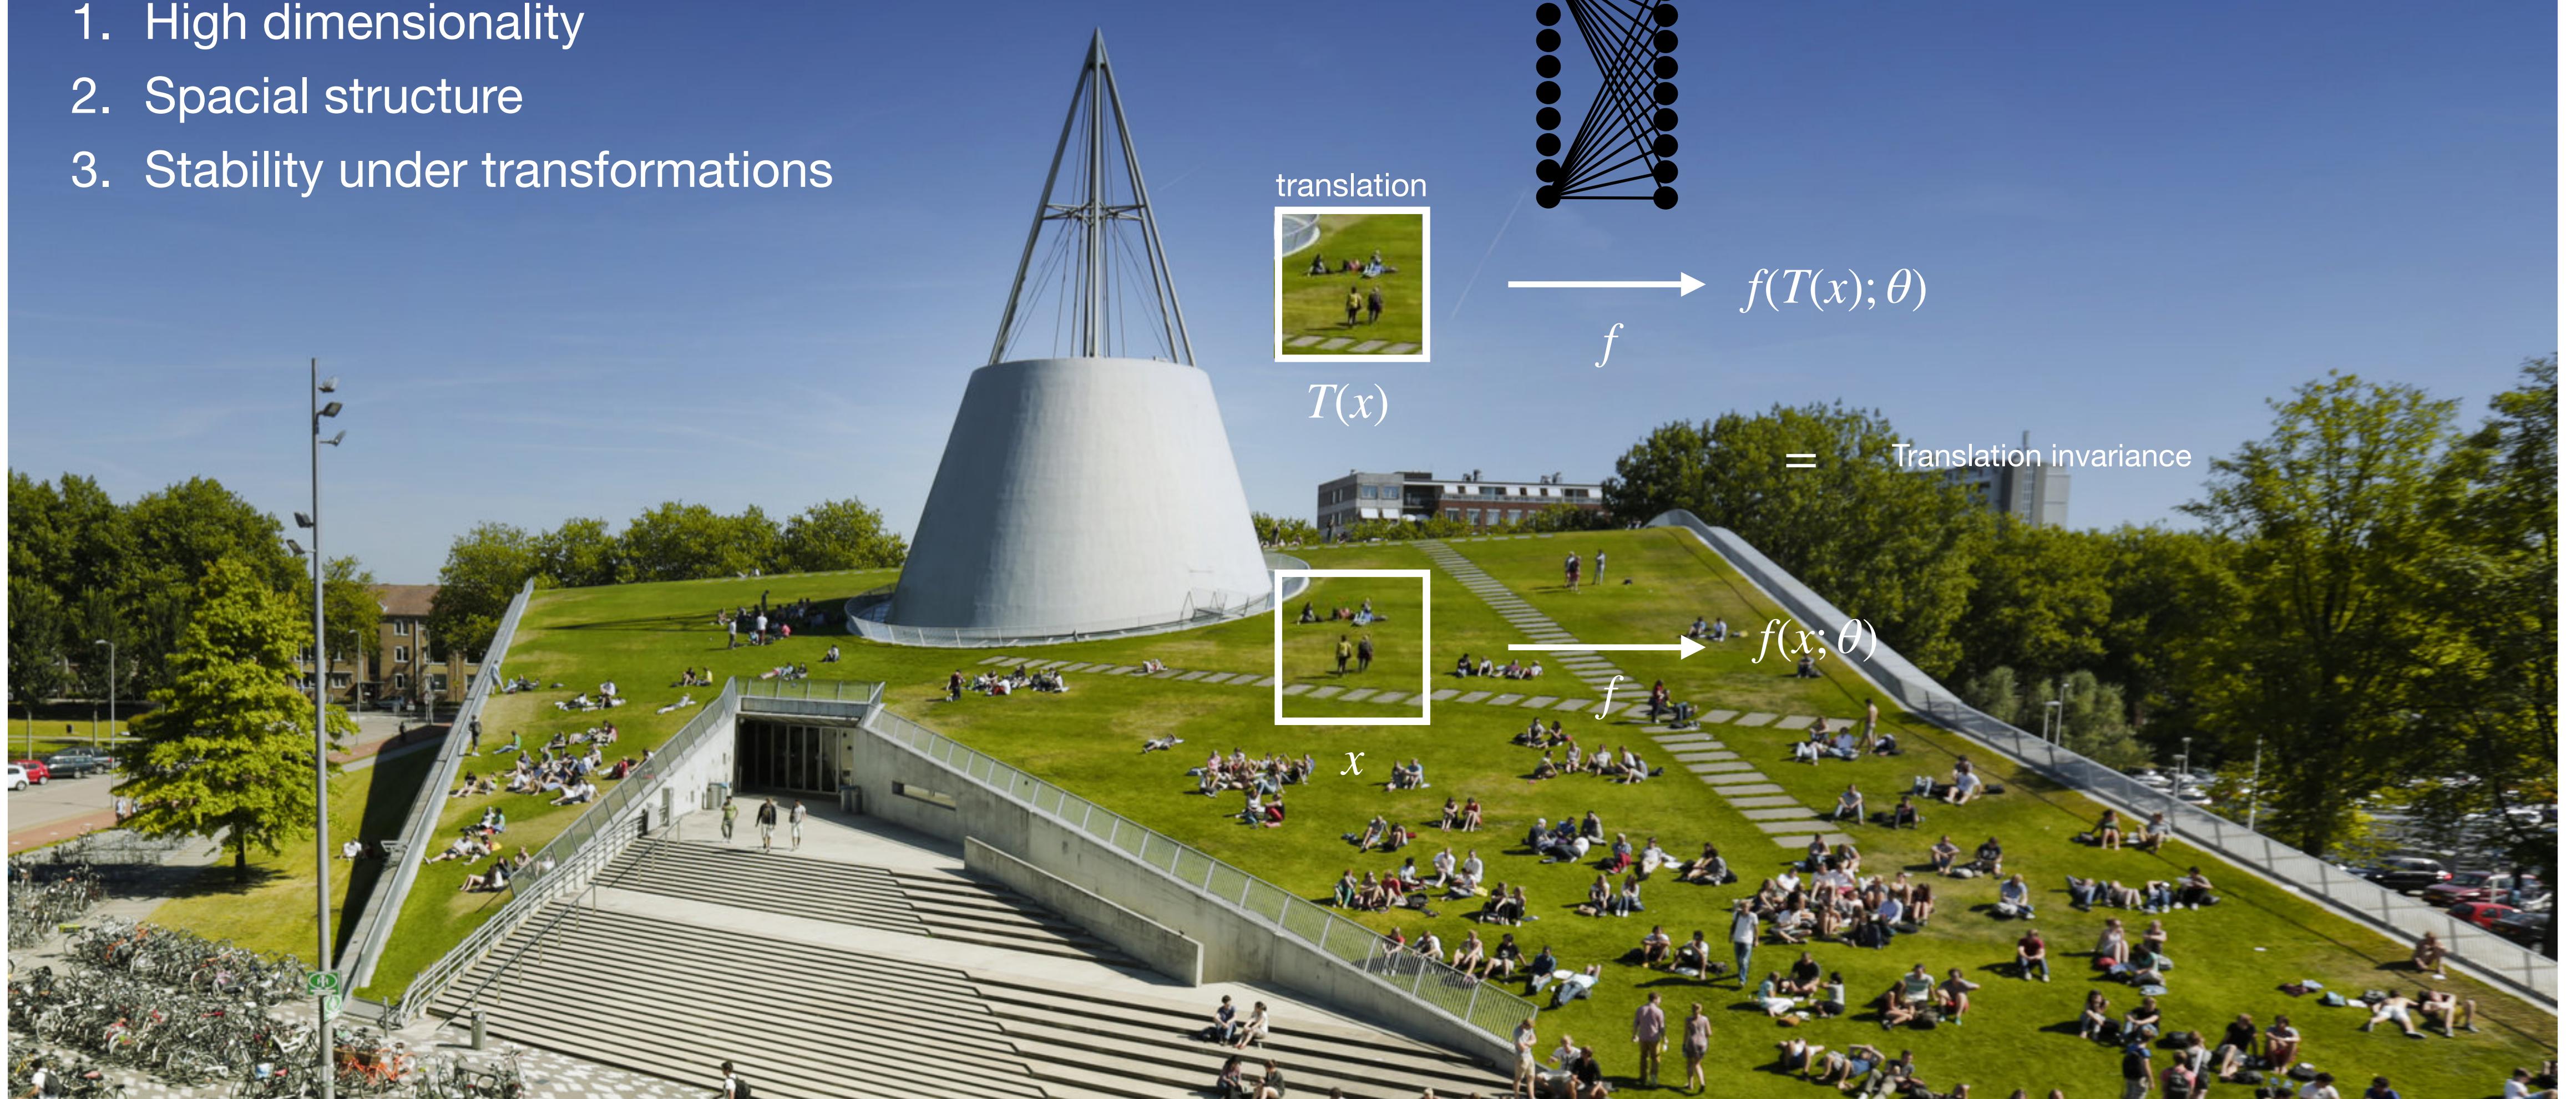
What makes images different?

1. High dimensionality
2. Spacial structure
3. Stability under transformations



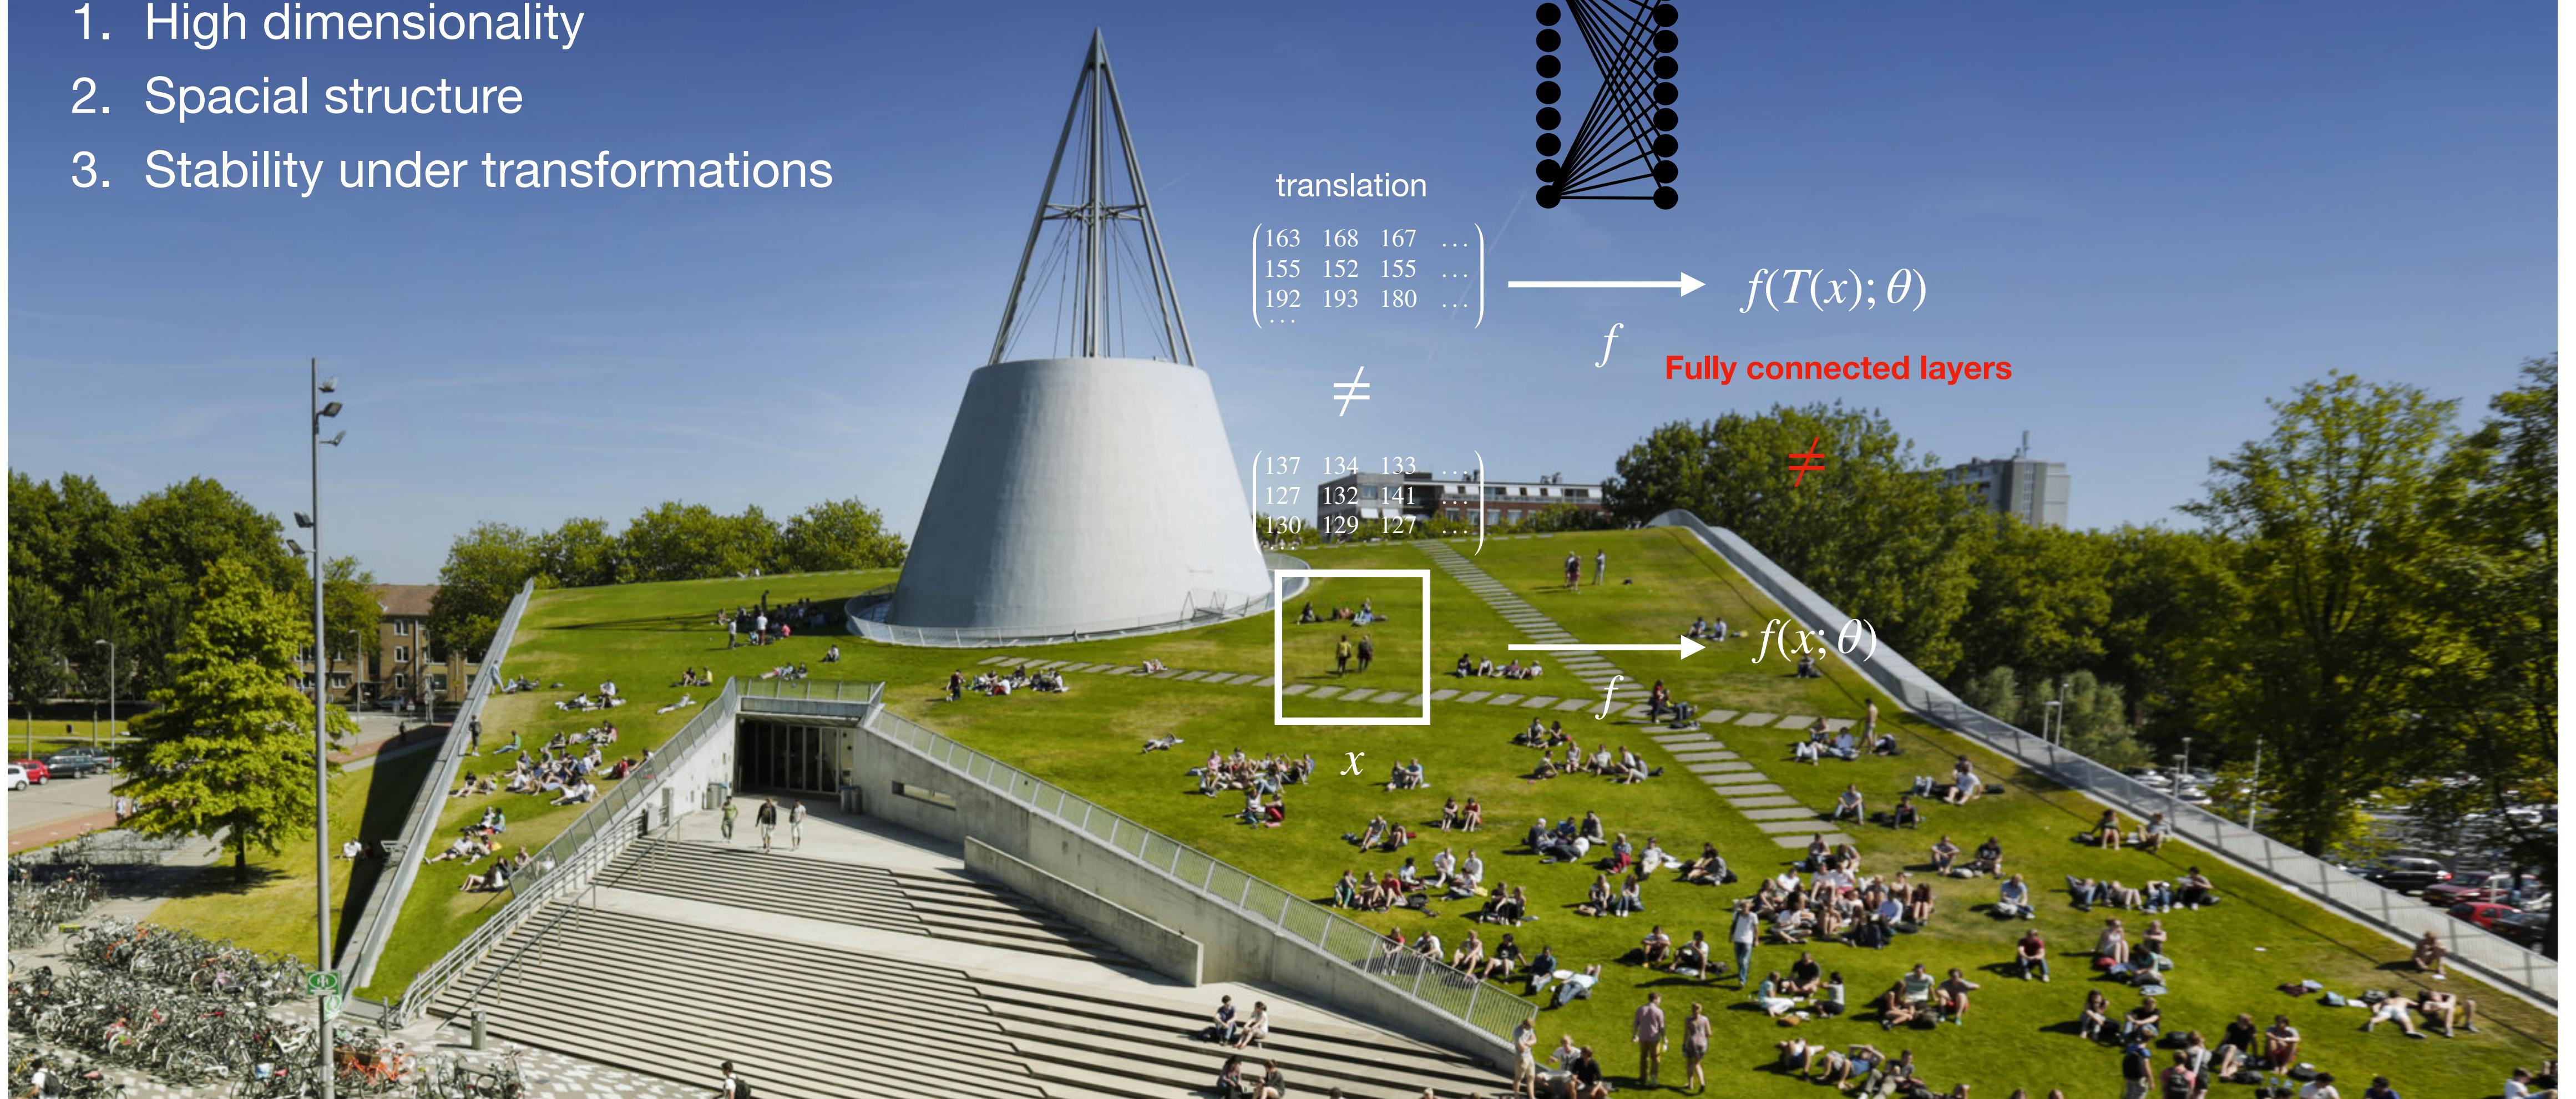
What makes images different?

1. High dimensionality
2. Spacial structure
3. Stability under transformations



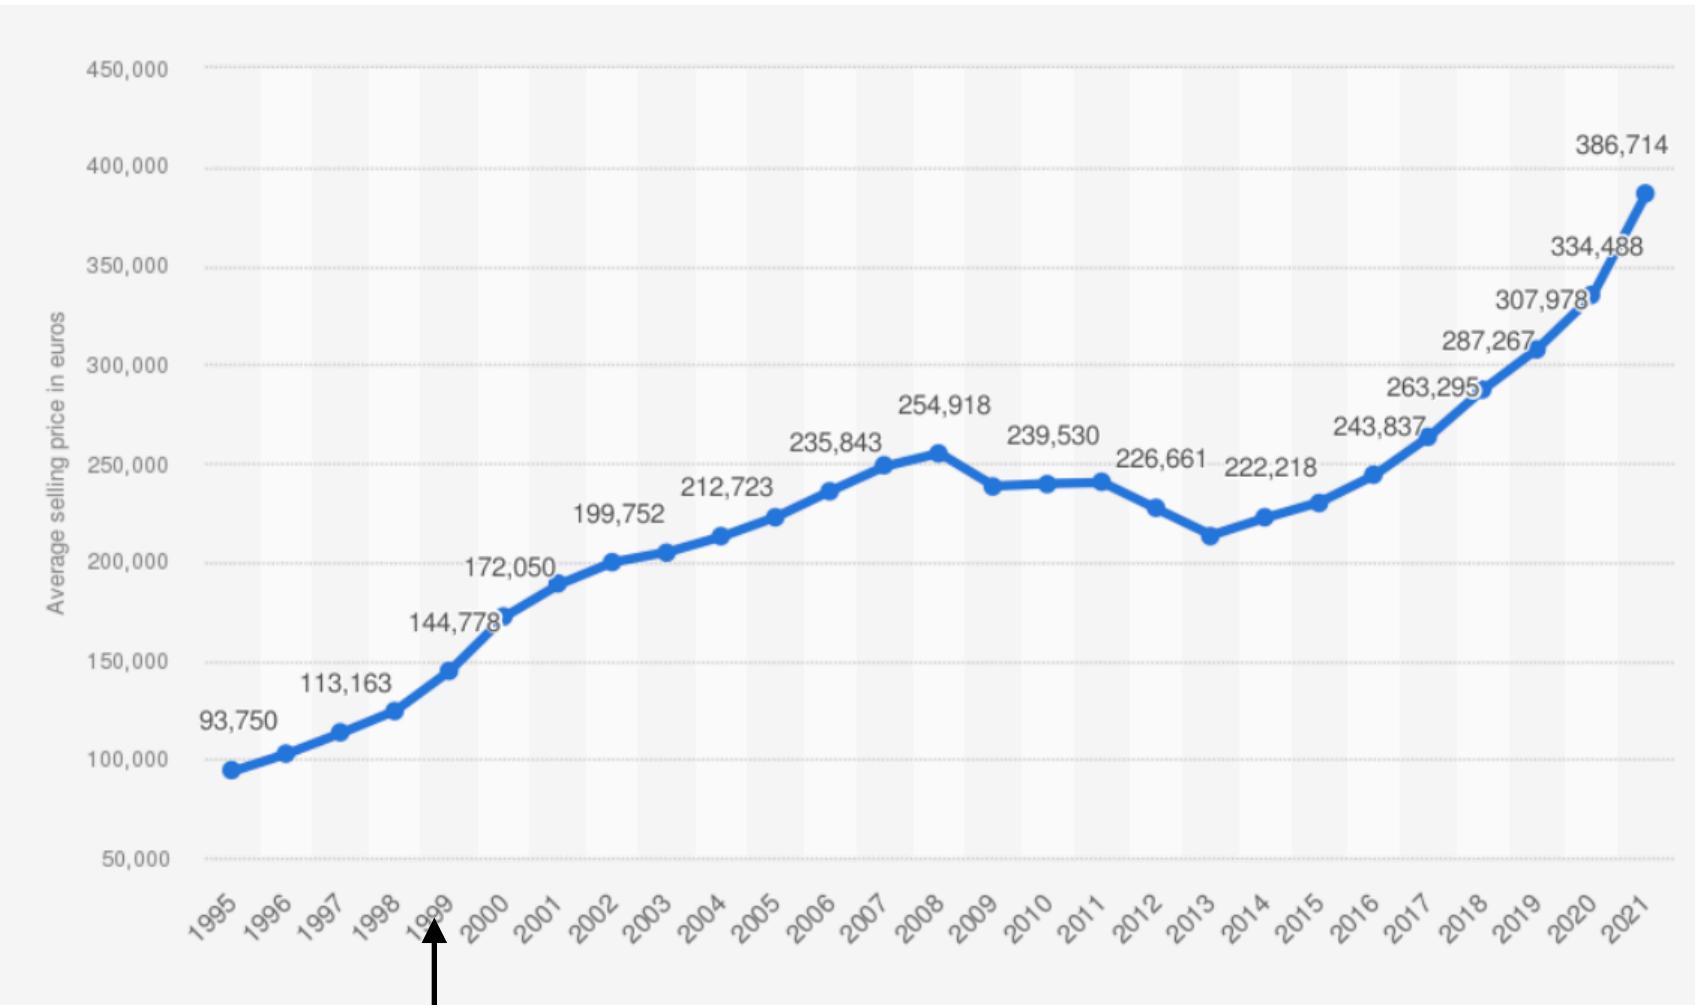
What makes images different?

1. High dimensionality
2. Spacial structure
3. Stability under transformations



Grid topology

1. Time series



(93750, 113163, 144778, 172050, ...)

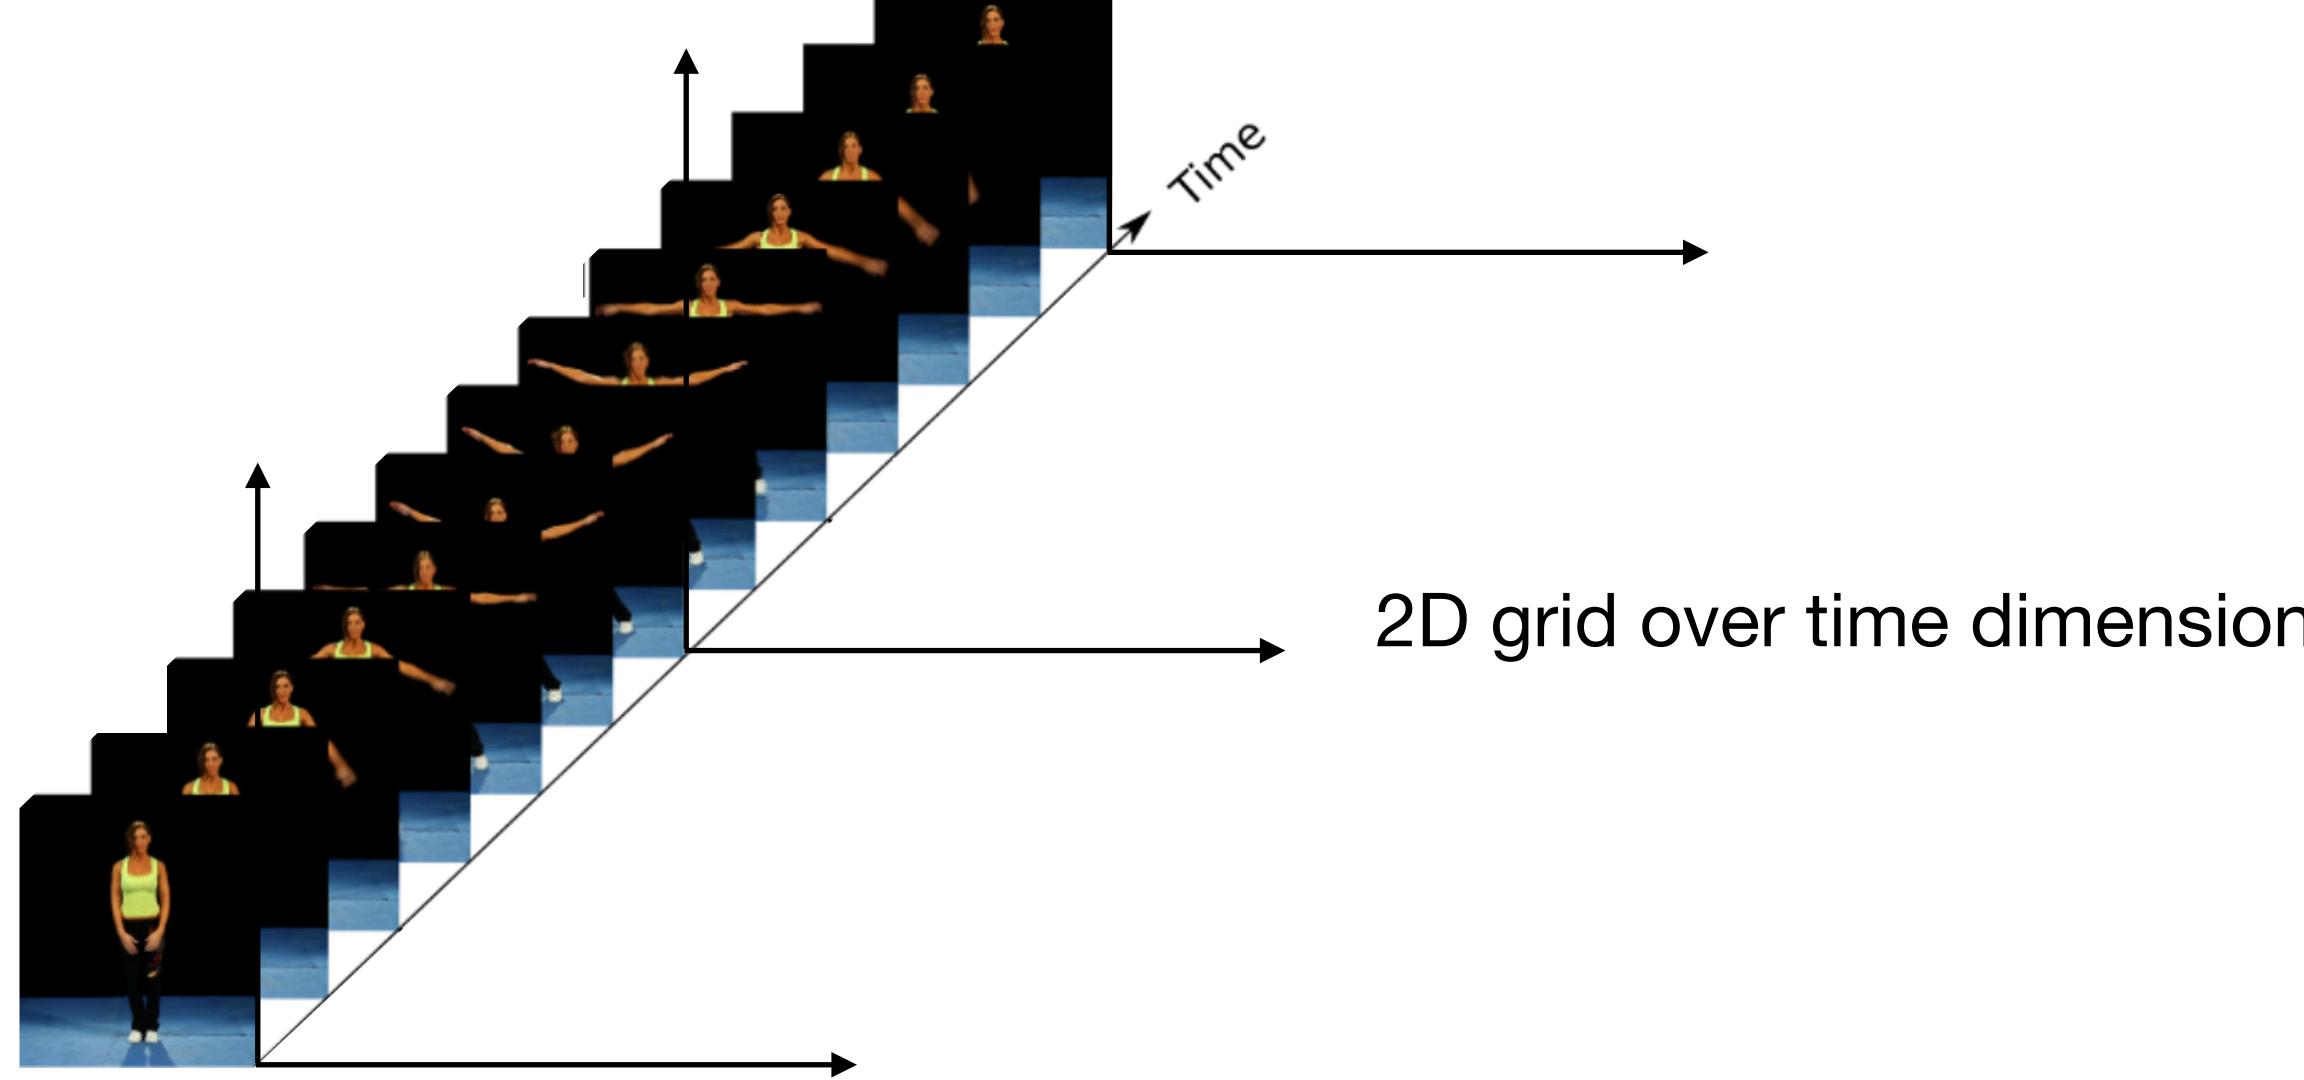
\xrightarrow{t}

1D grid - time dimension

2. Video

3. Audio

4. Text



Convolutions

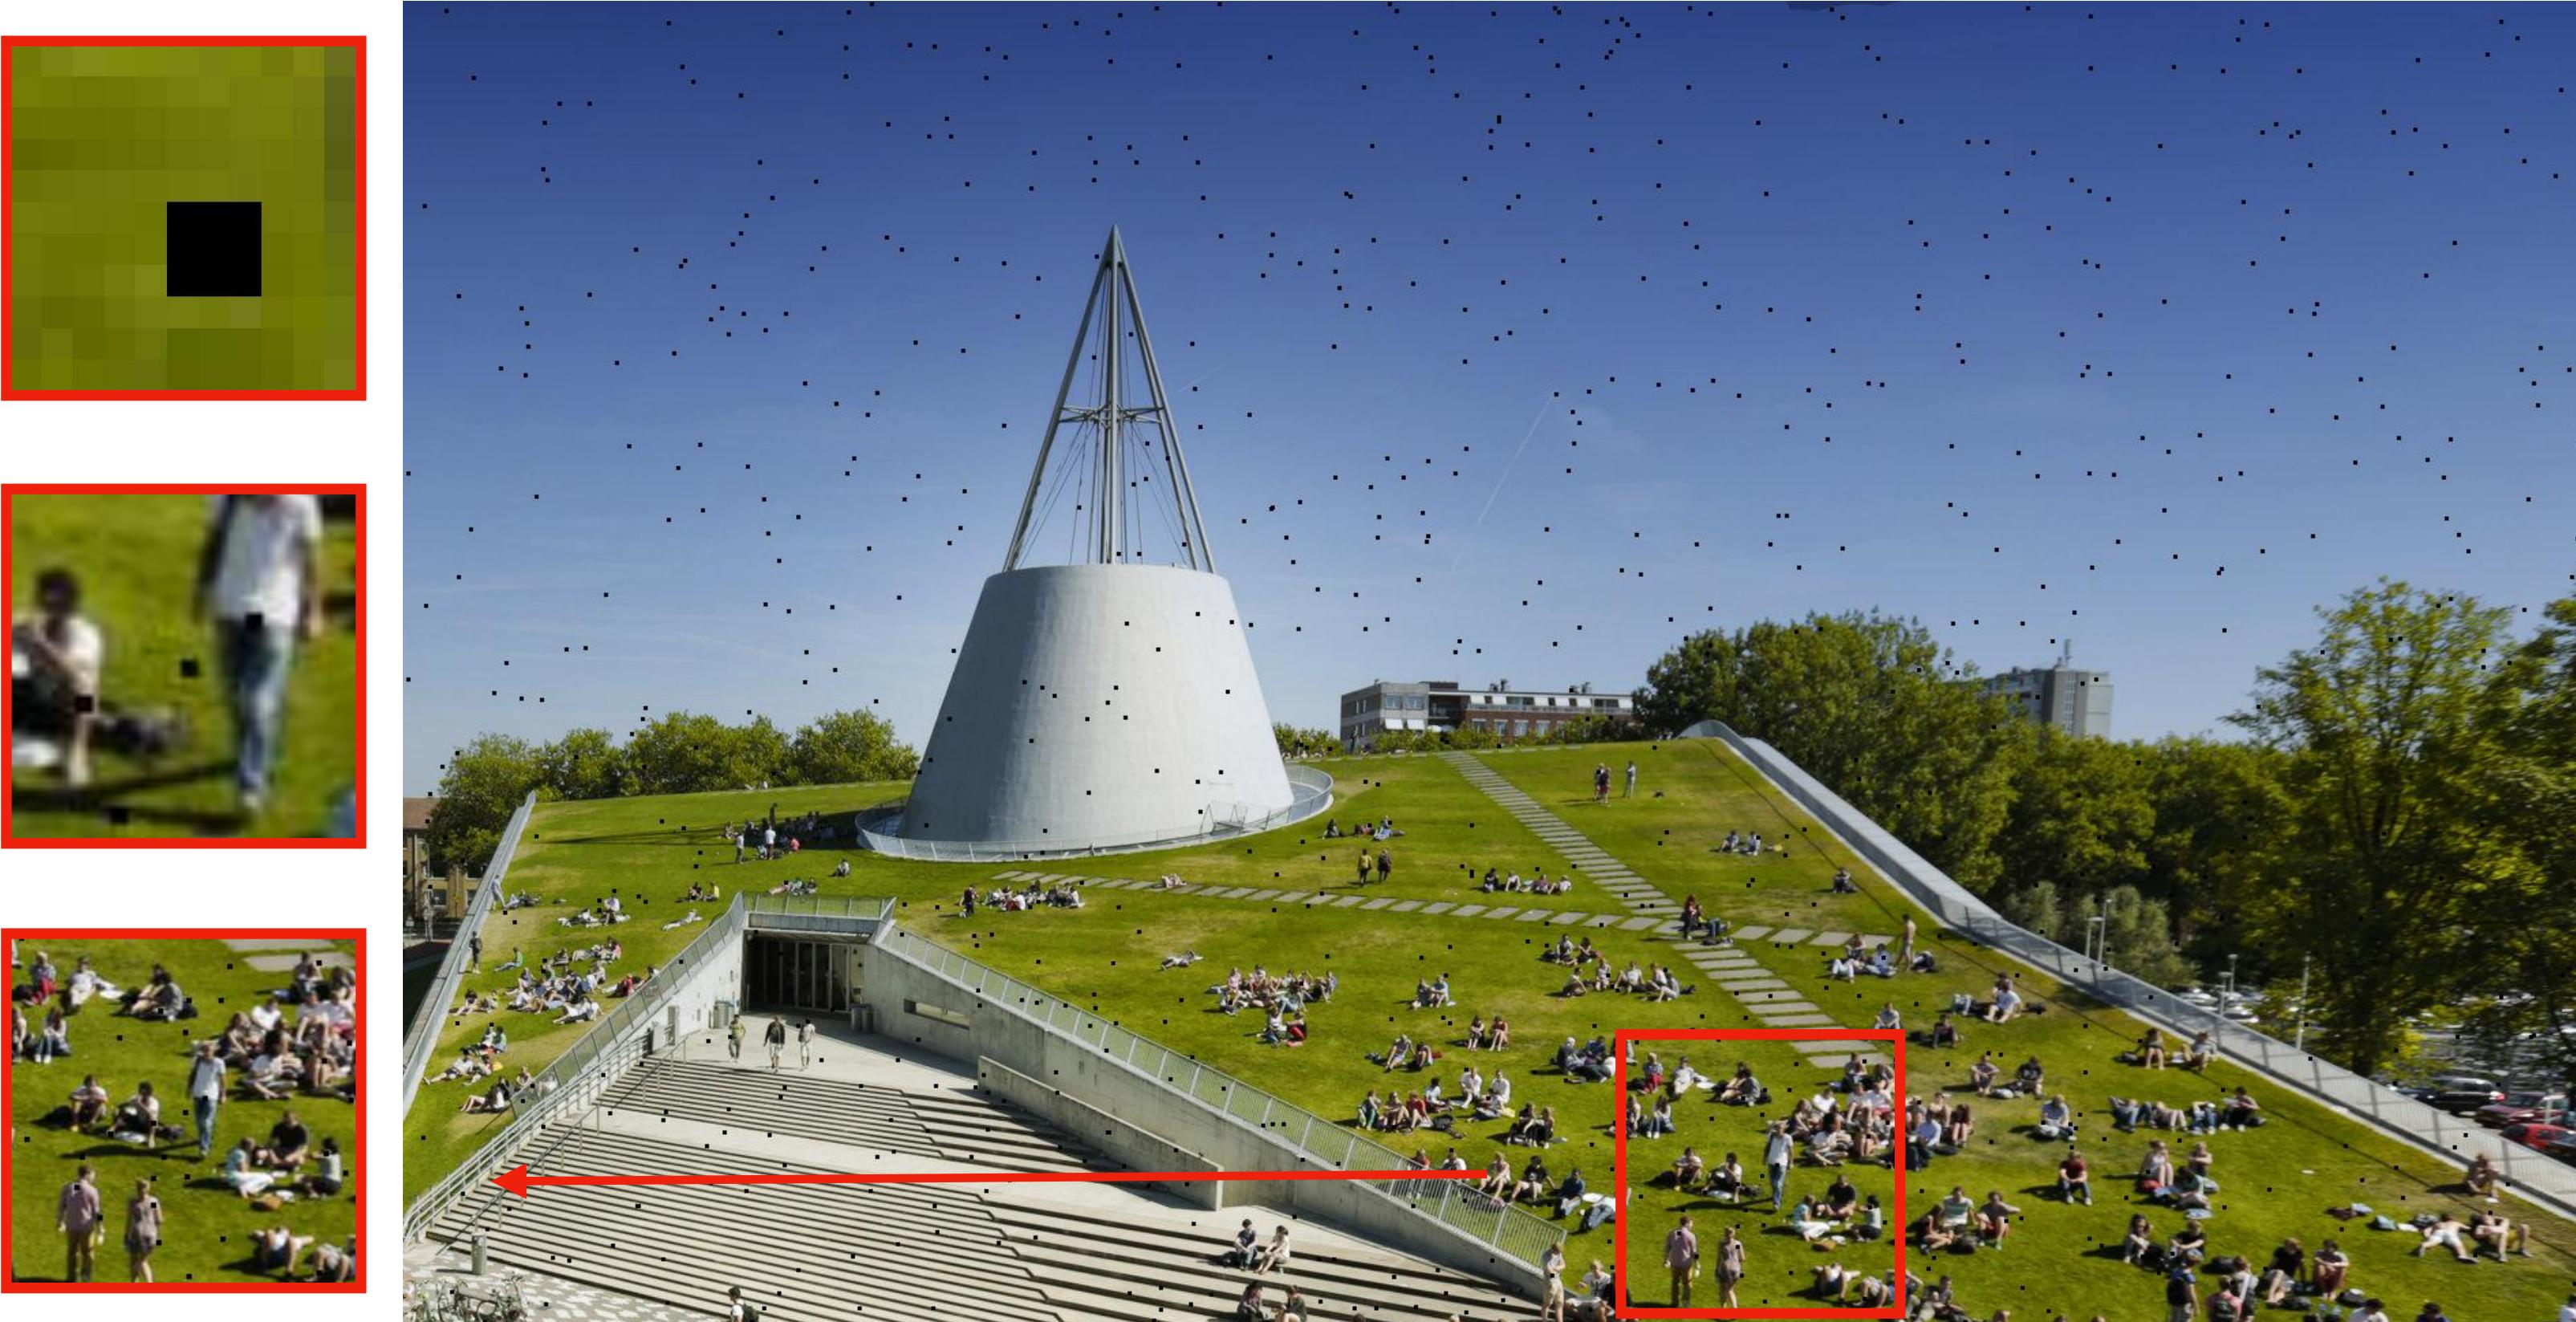
Image processing example



Simple solution for noise removal → neighbourhood average to replace the lost pixels

Convolutions

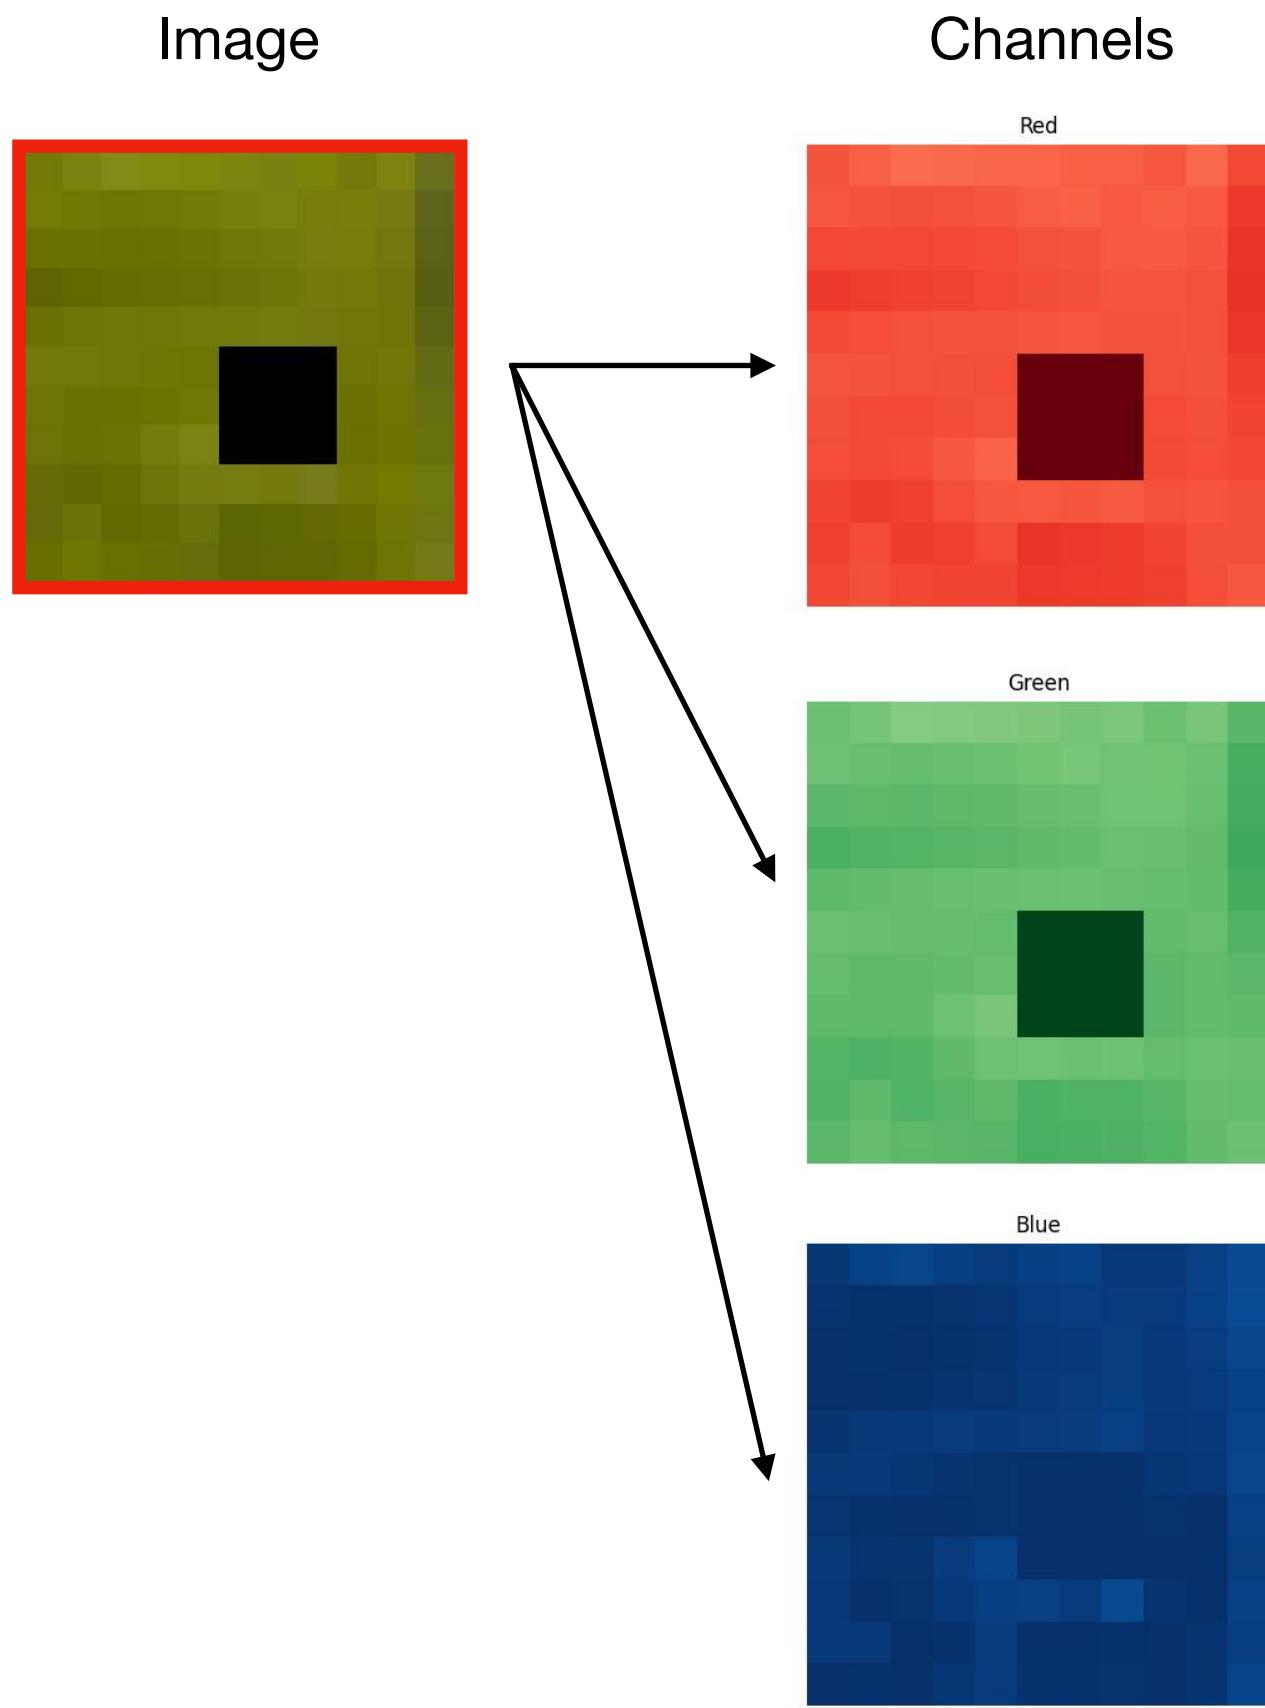
Image processing example



Simple solution for noise removal → neighbourhood average to replace the lost pixels

Convolutions

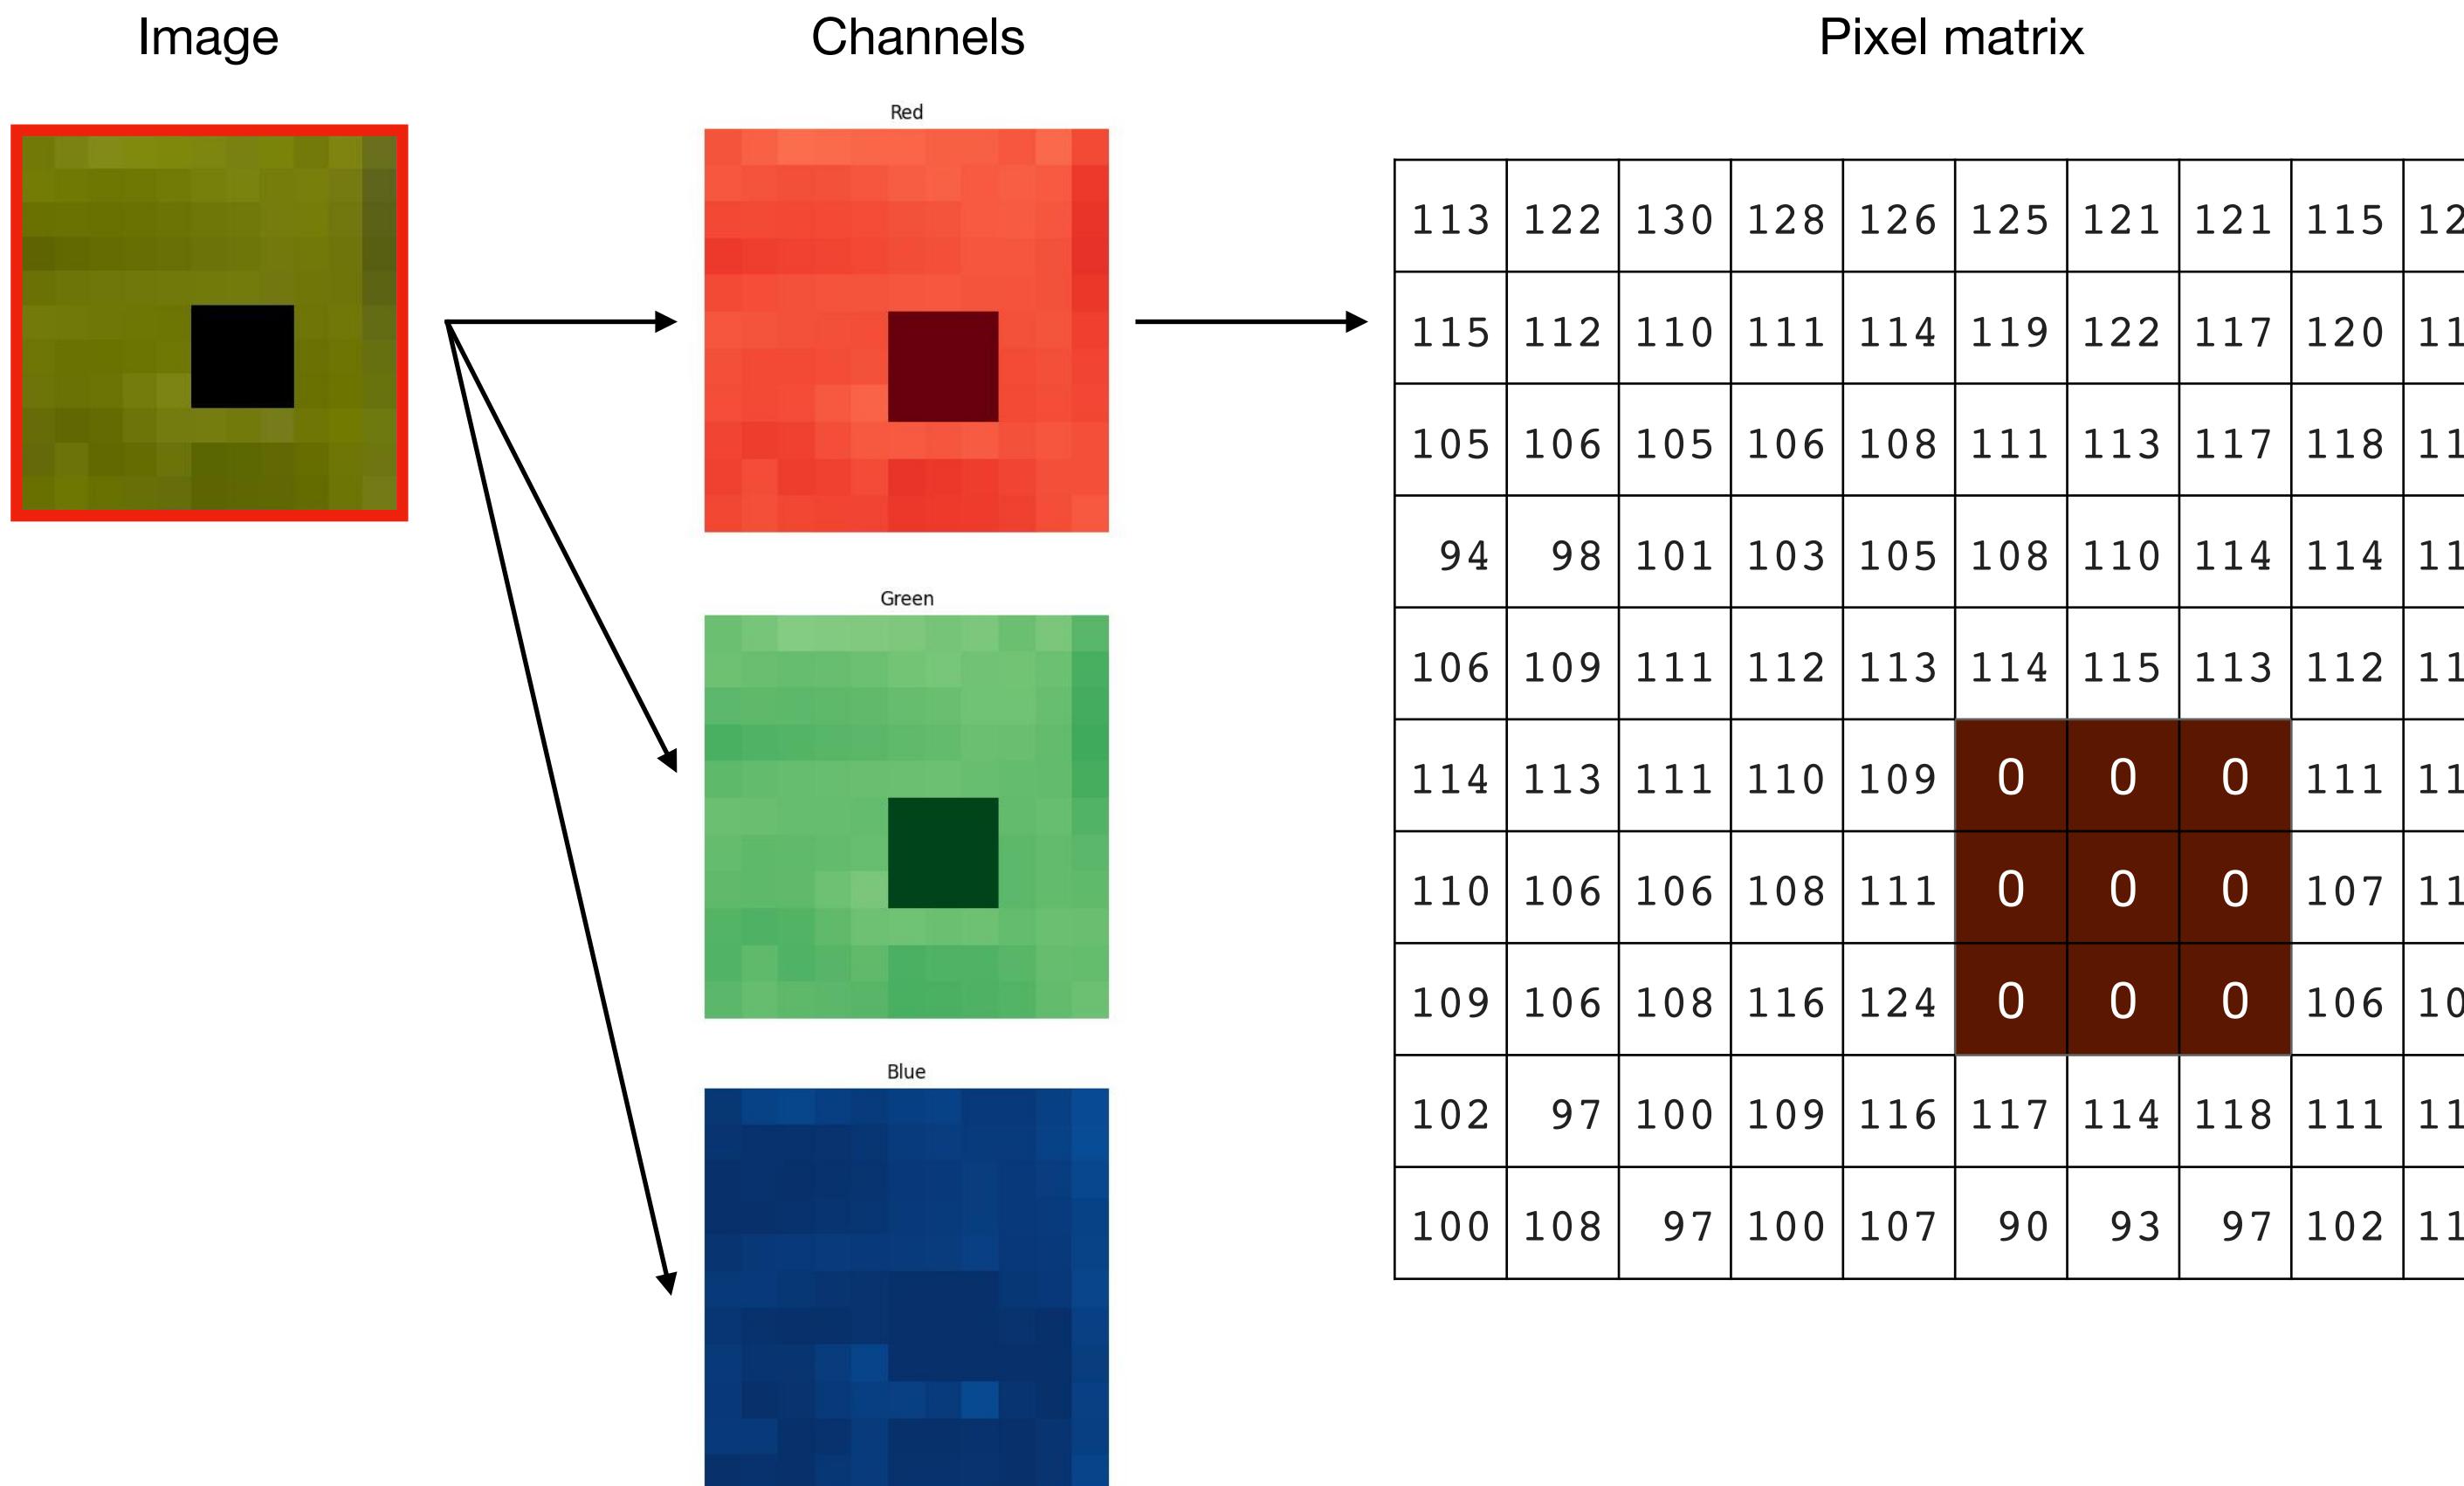
Image processing example



Simple solution for noise removal → neighbourhood average to replace the lost pixels

Convolutions

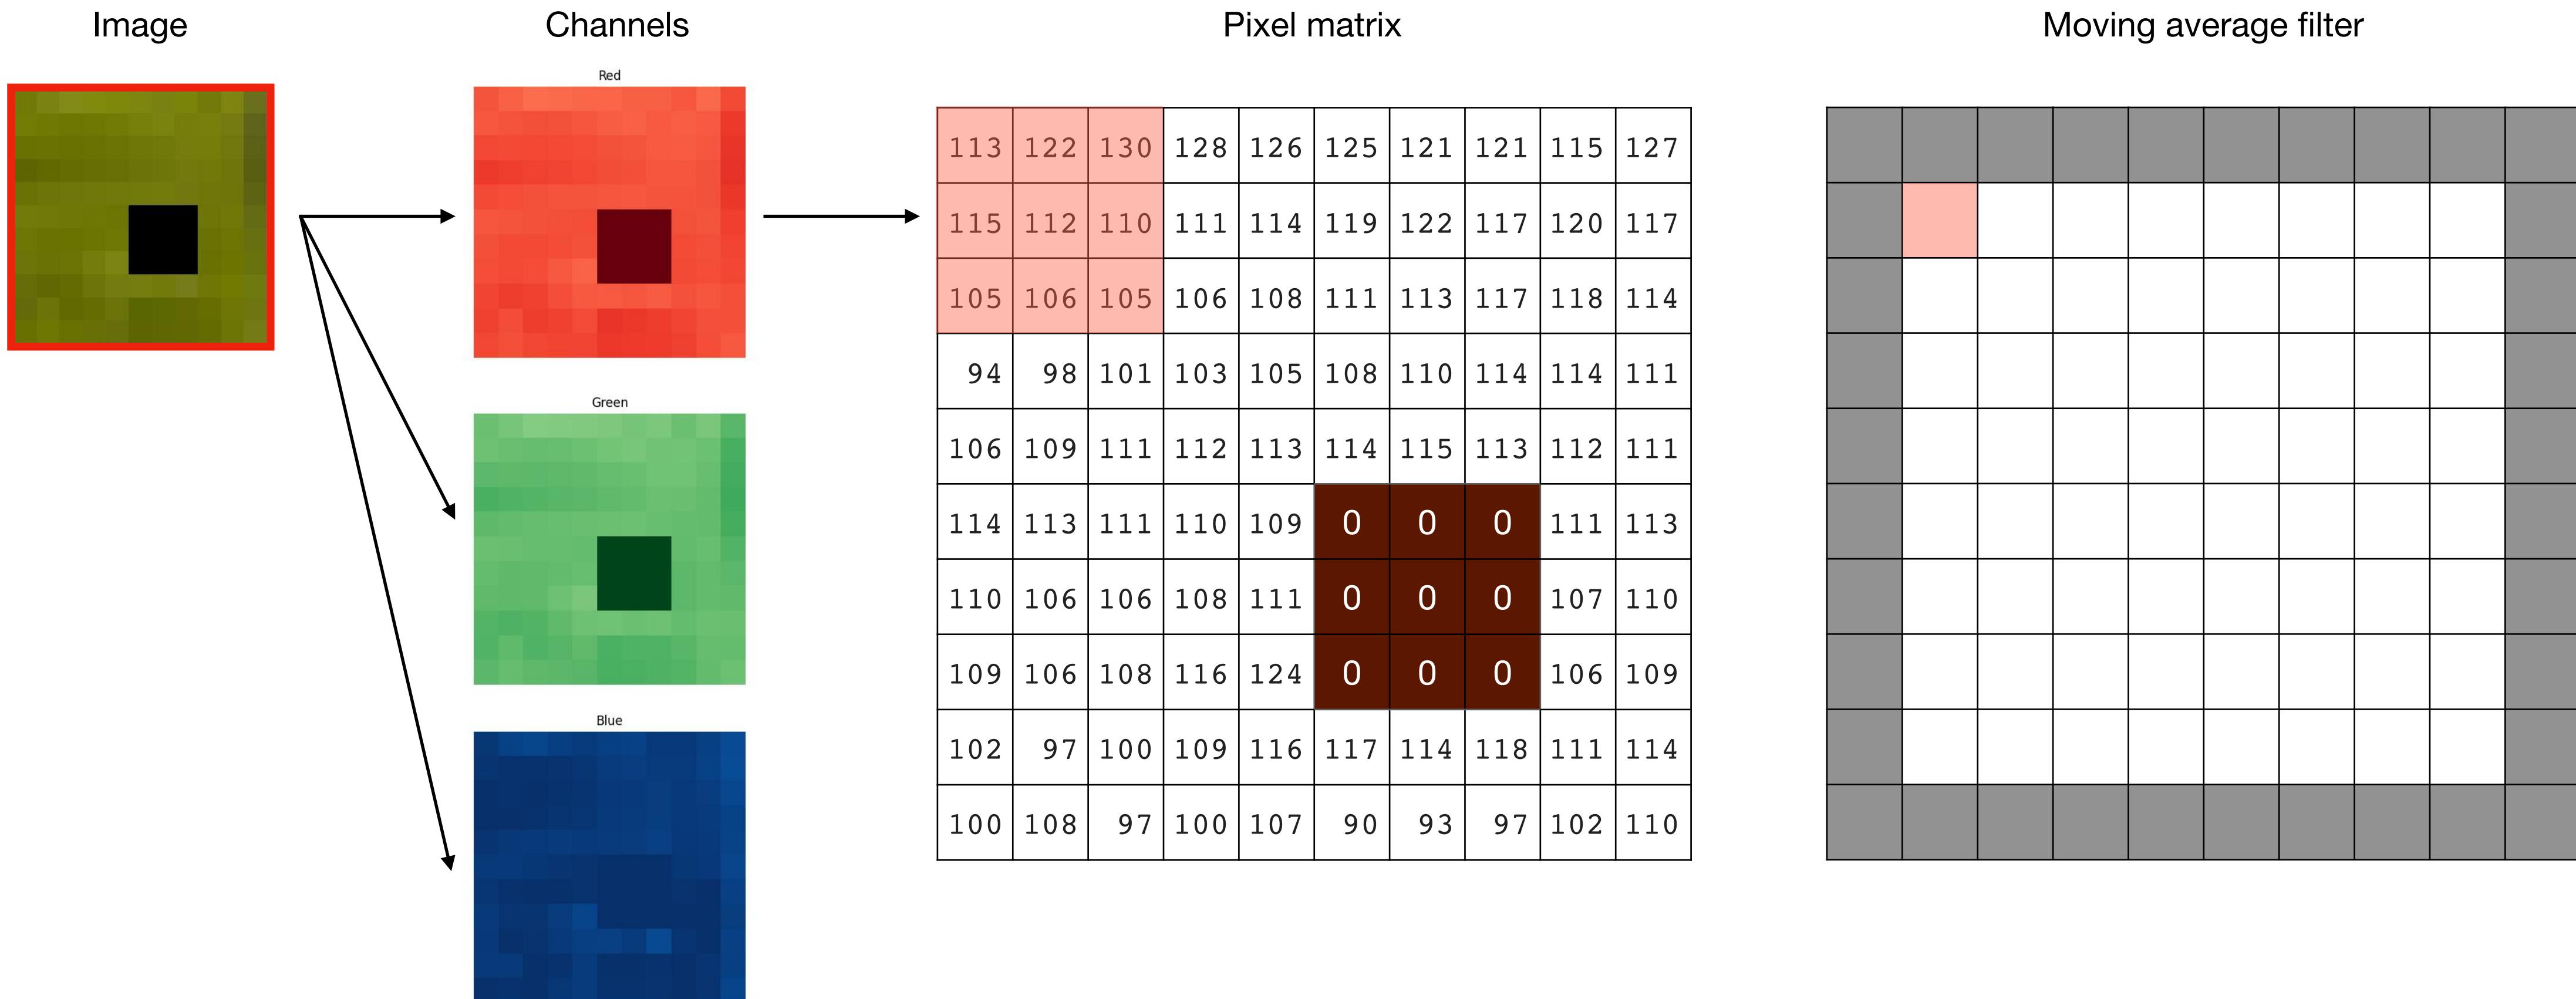
Image processing example



Simple solution for noise removal → neighbourhood average to replace the lost pixels

Convolutions

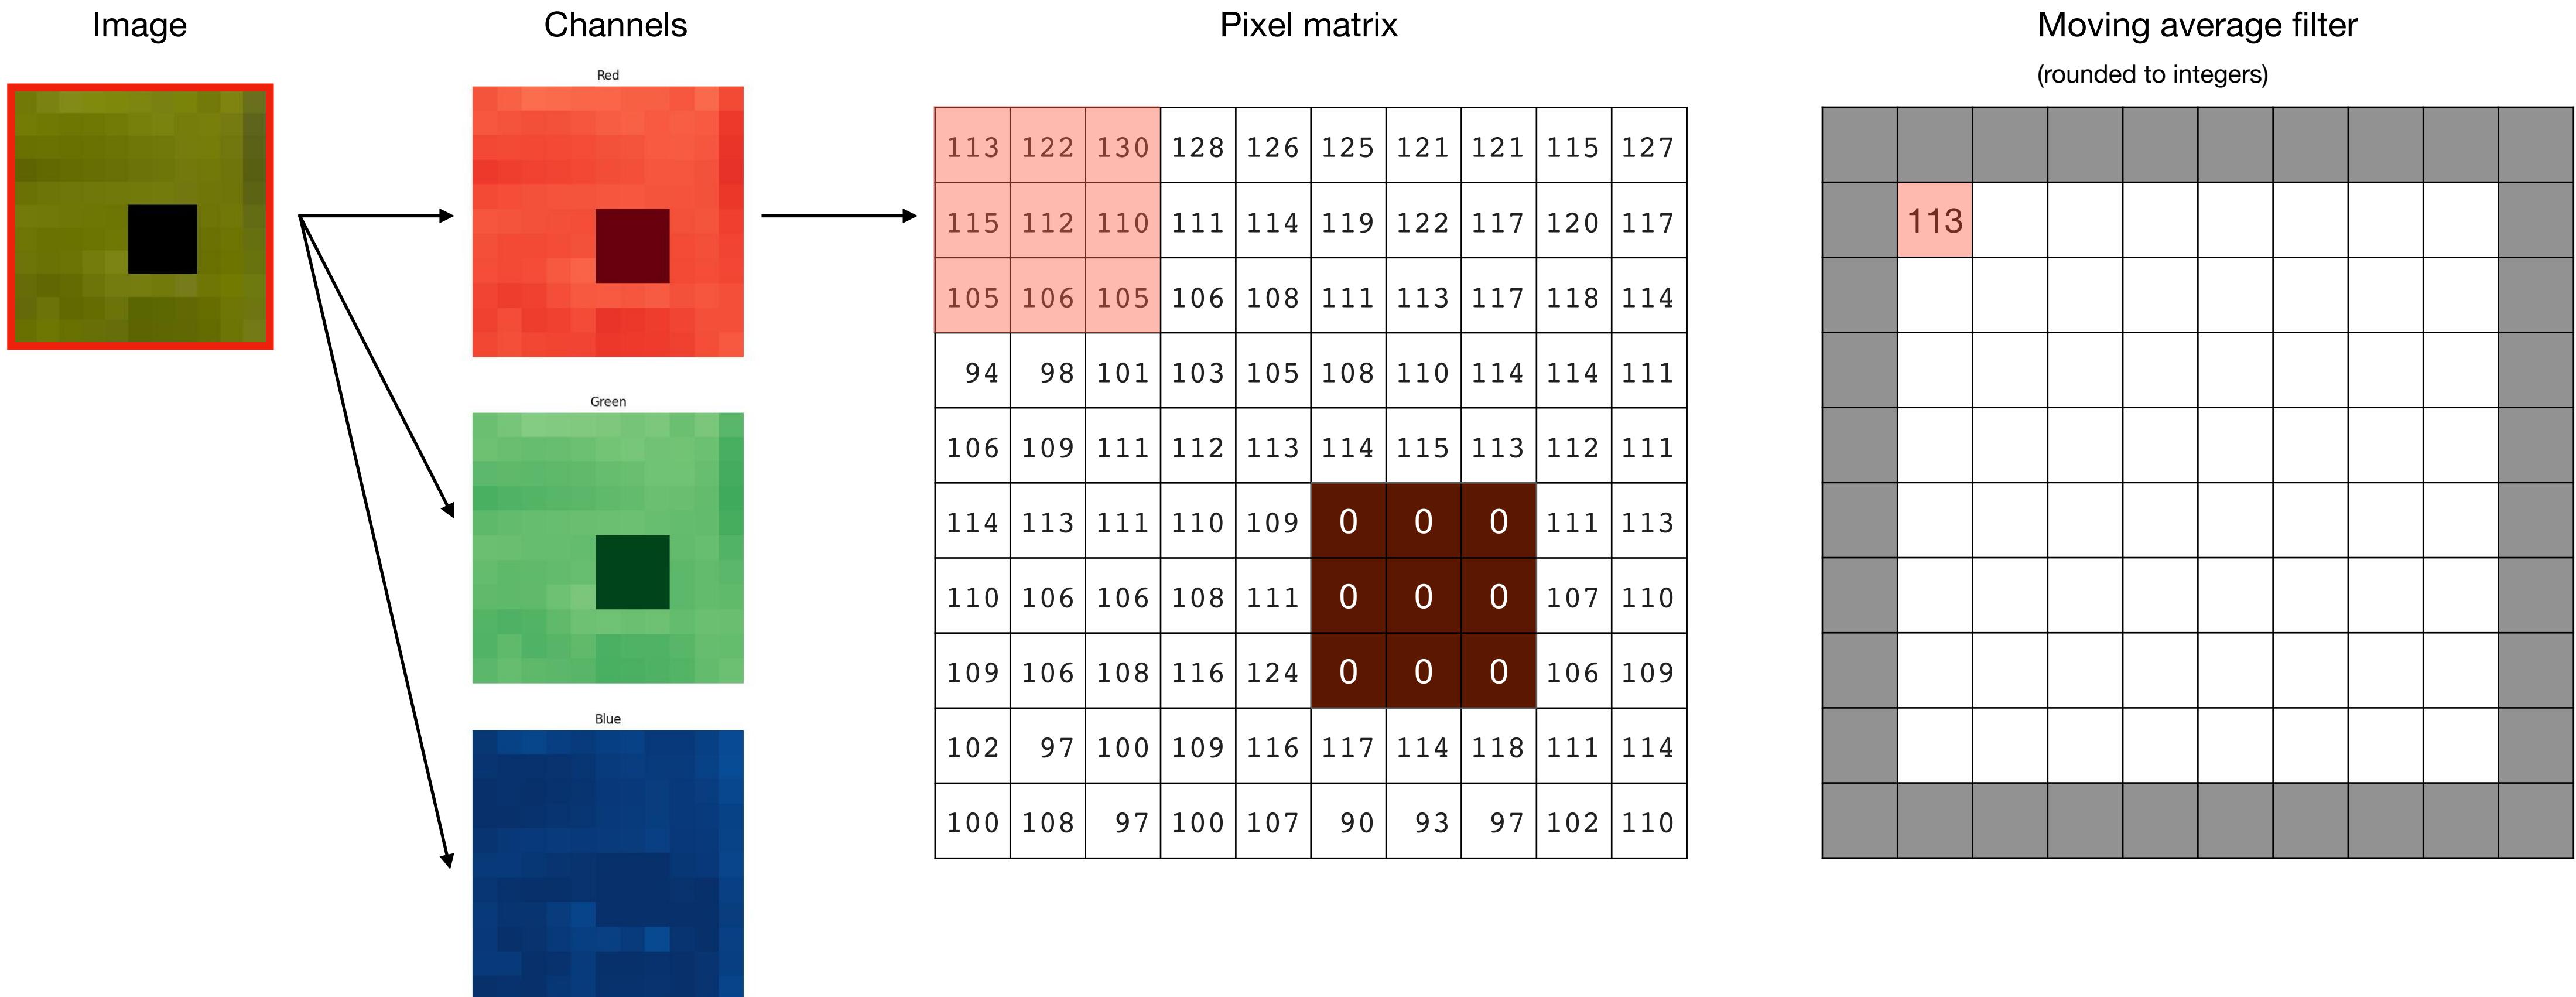
Image processing example



Simple solution for noise removal → neighbourhood average to replace the lost pixels

Convolutions

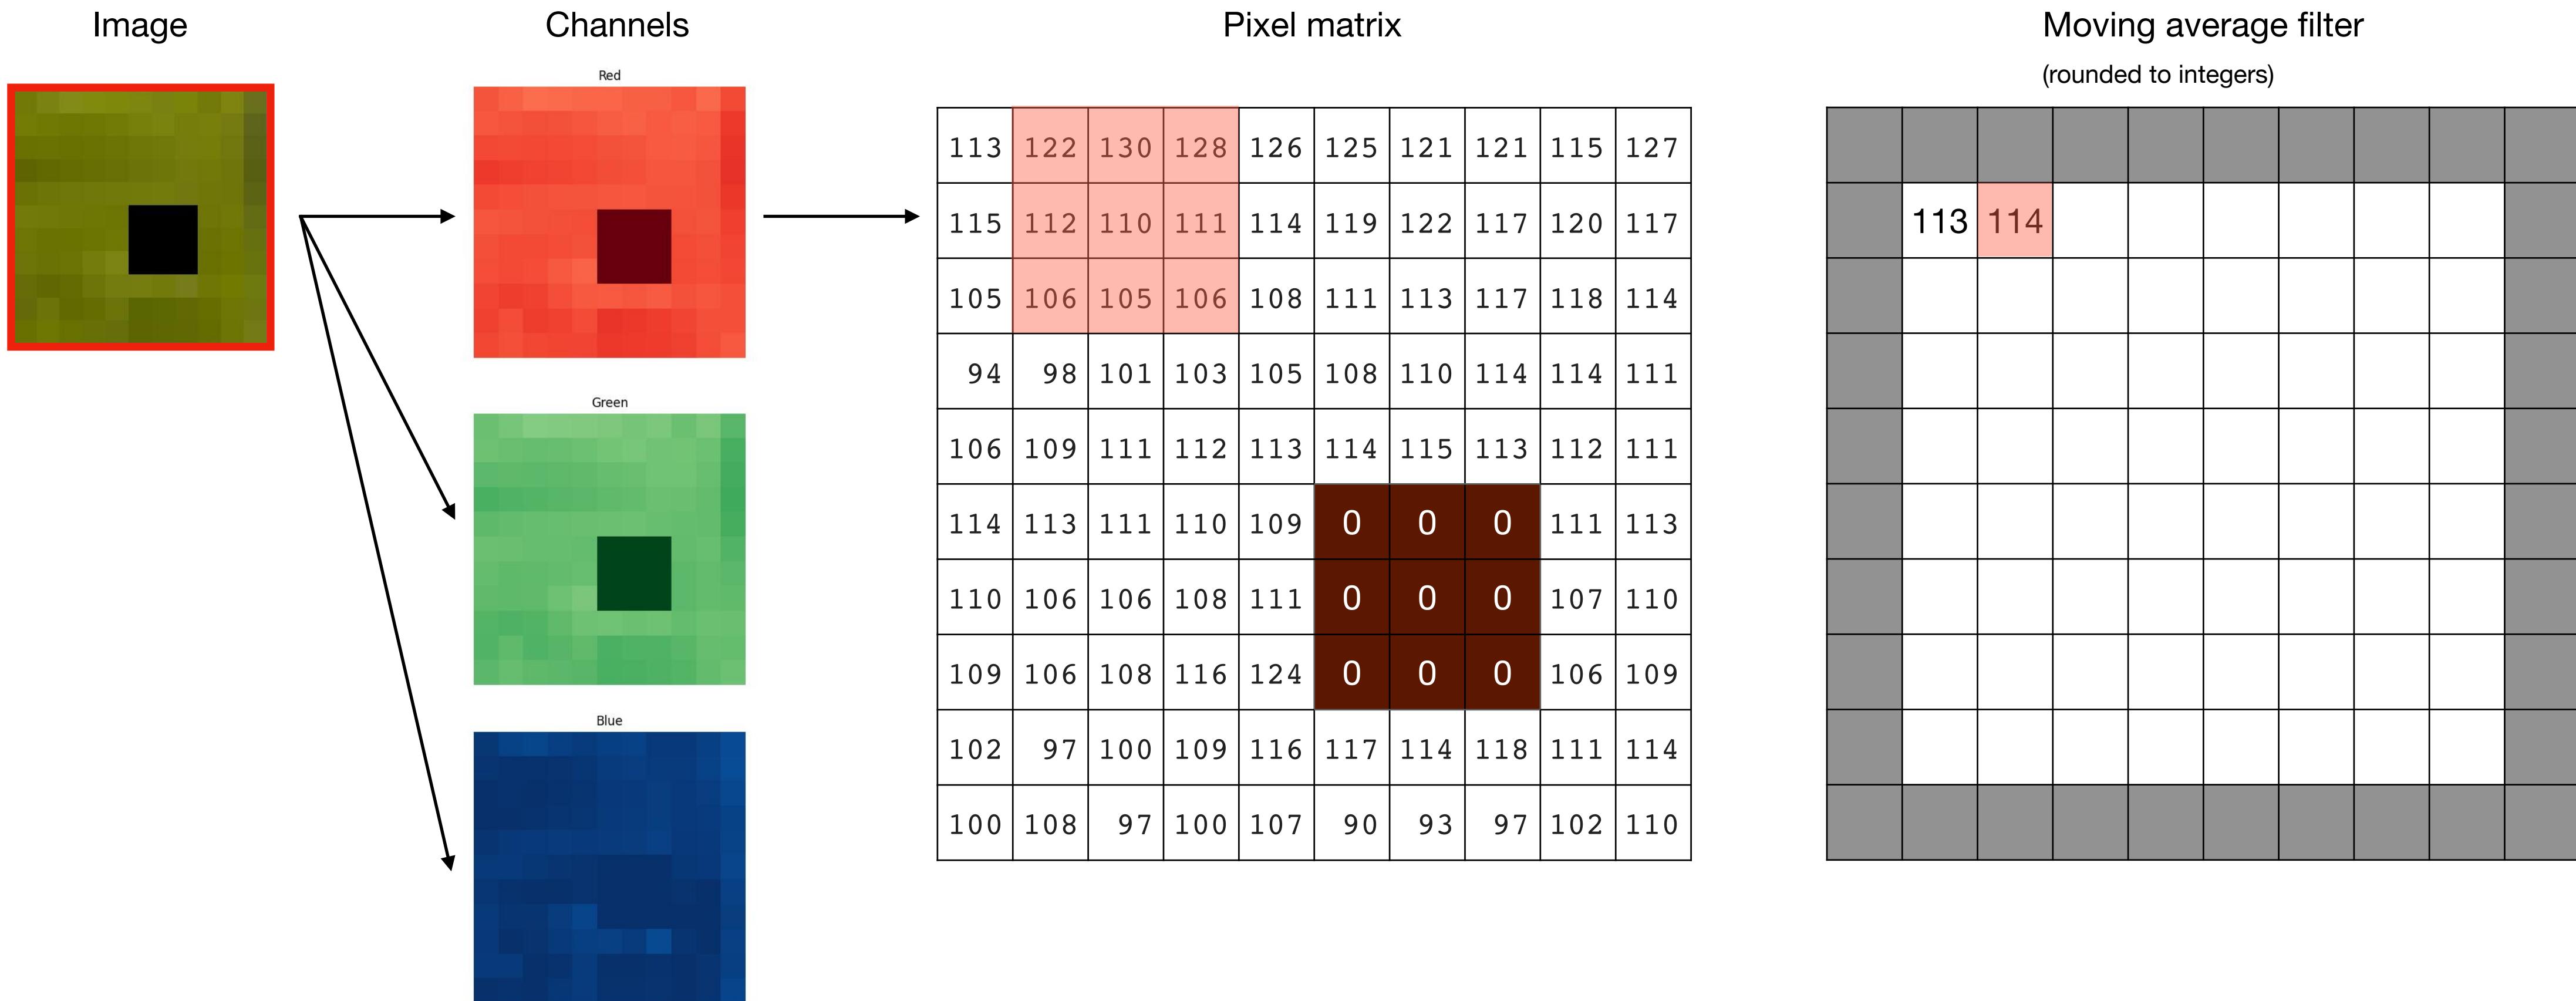
Image processing example



Simple solution for noise removal → neighbourhood average to replace the lost pixels

Convolutions

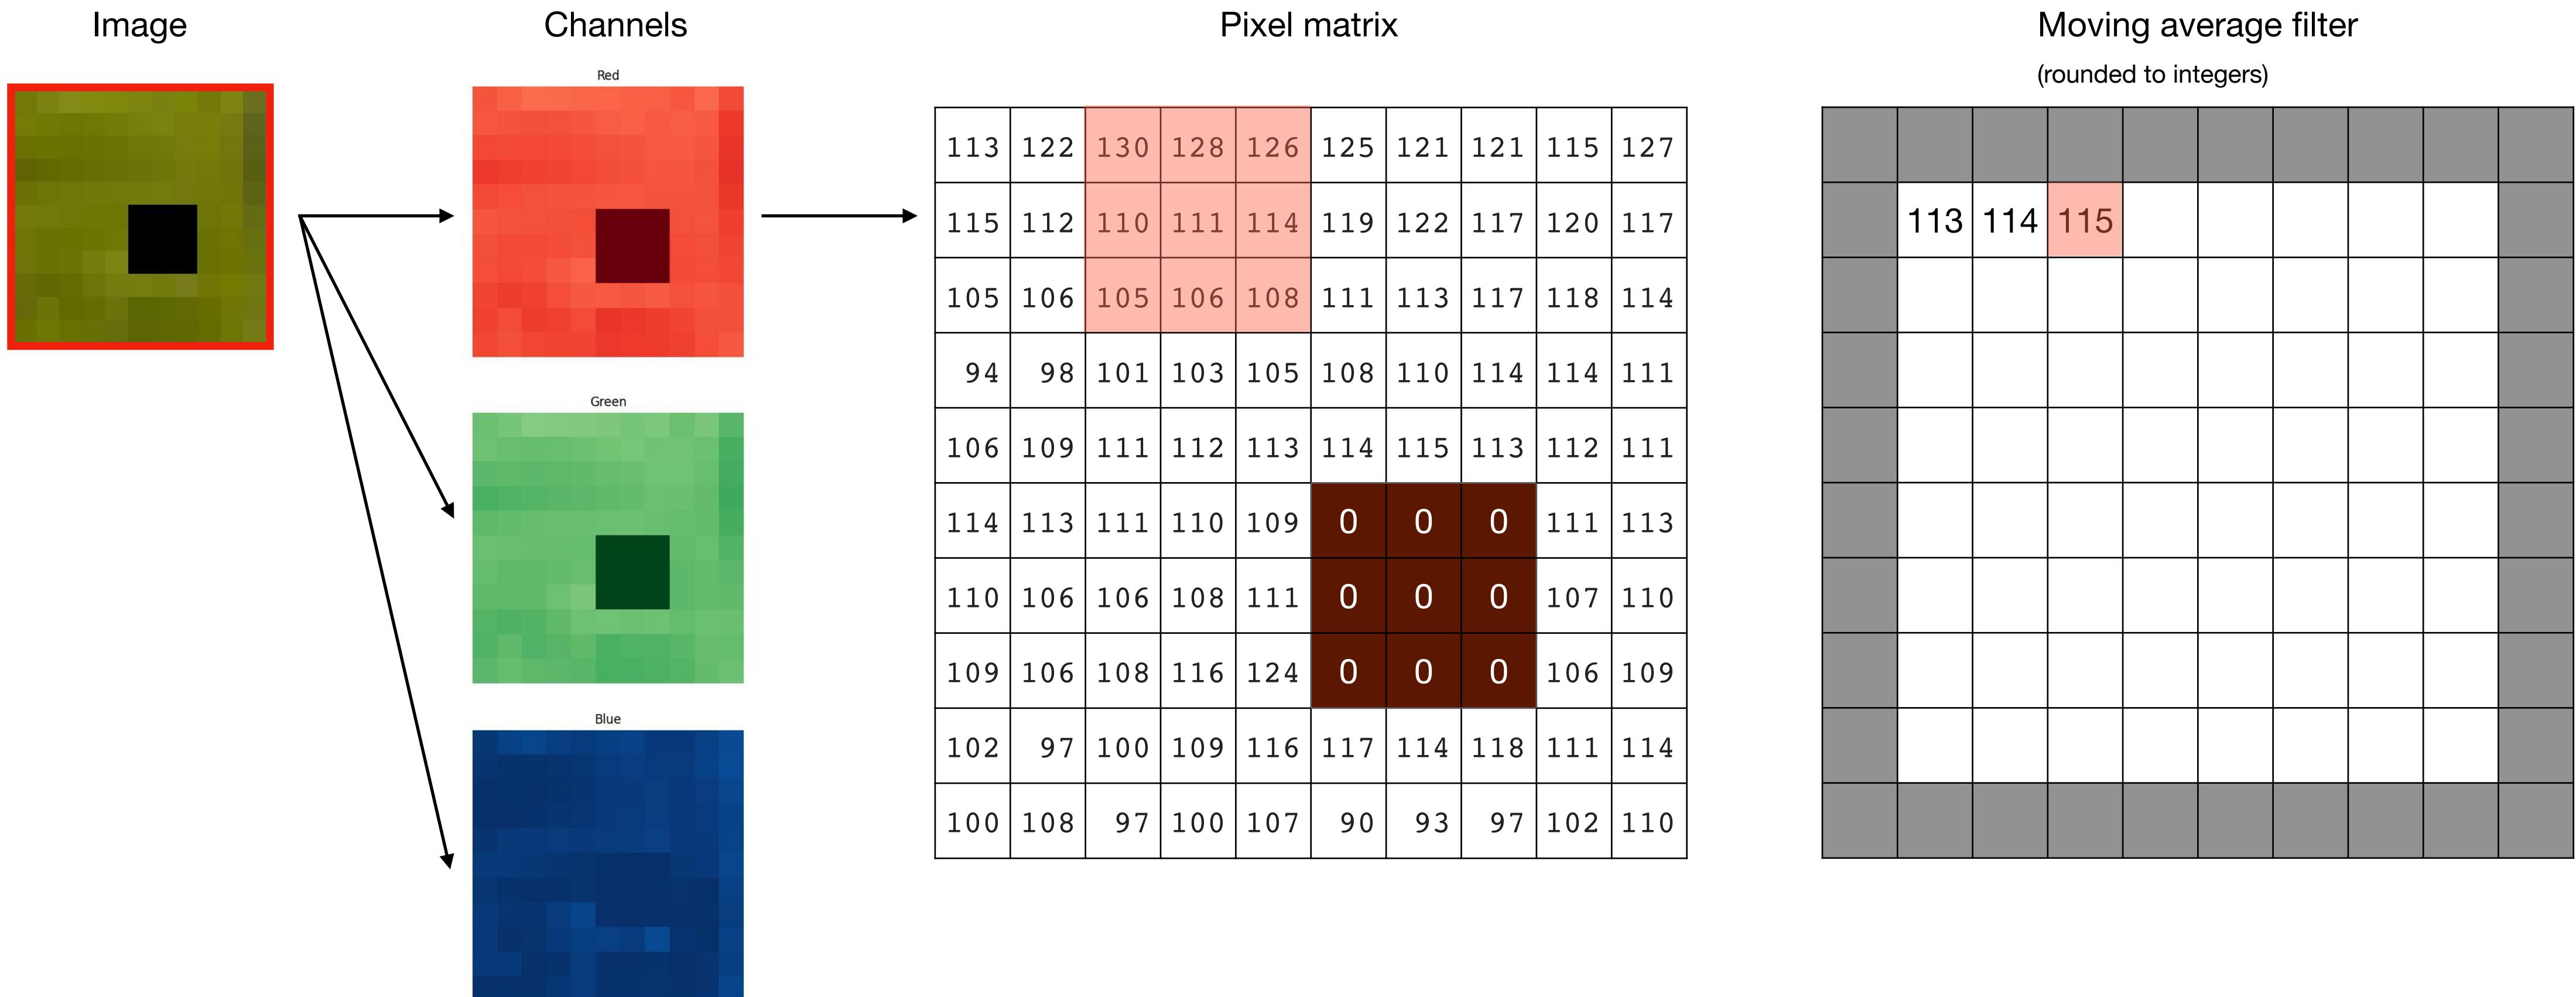
Image processing example



Simple solution for noise removal → neighbourhood average to replace the lost pixels

Convolutions

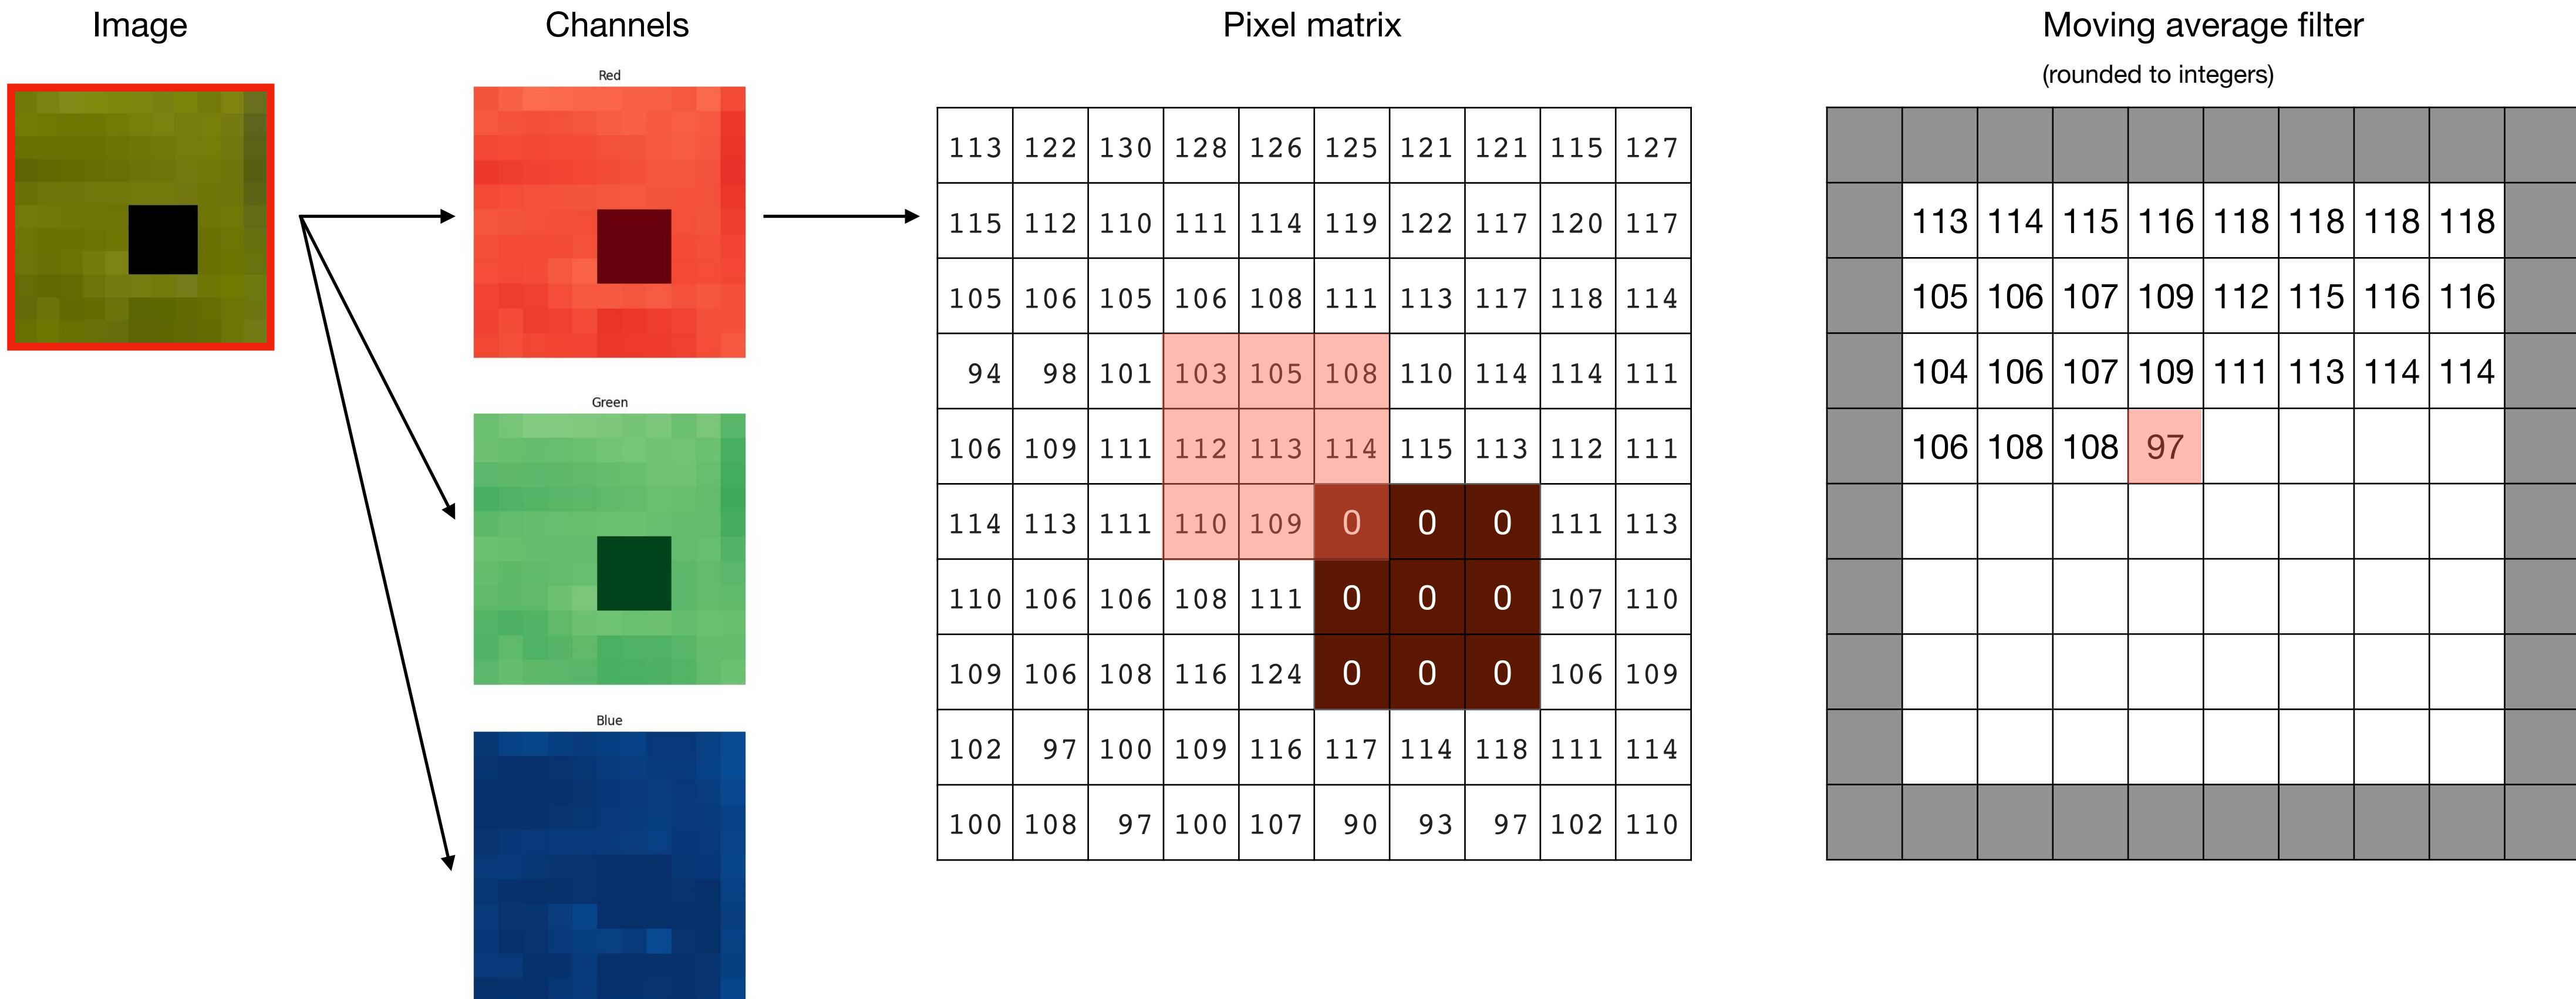
Image processing example



Simple solution for noise removal → neighbourhood average to replace the lost pixels

Convolutions

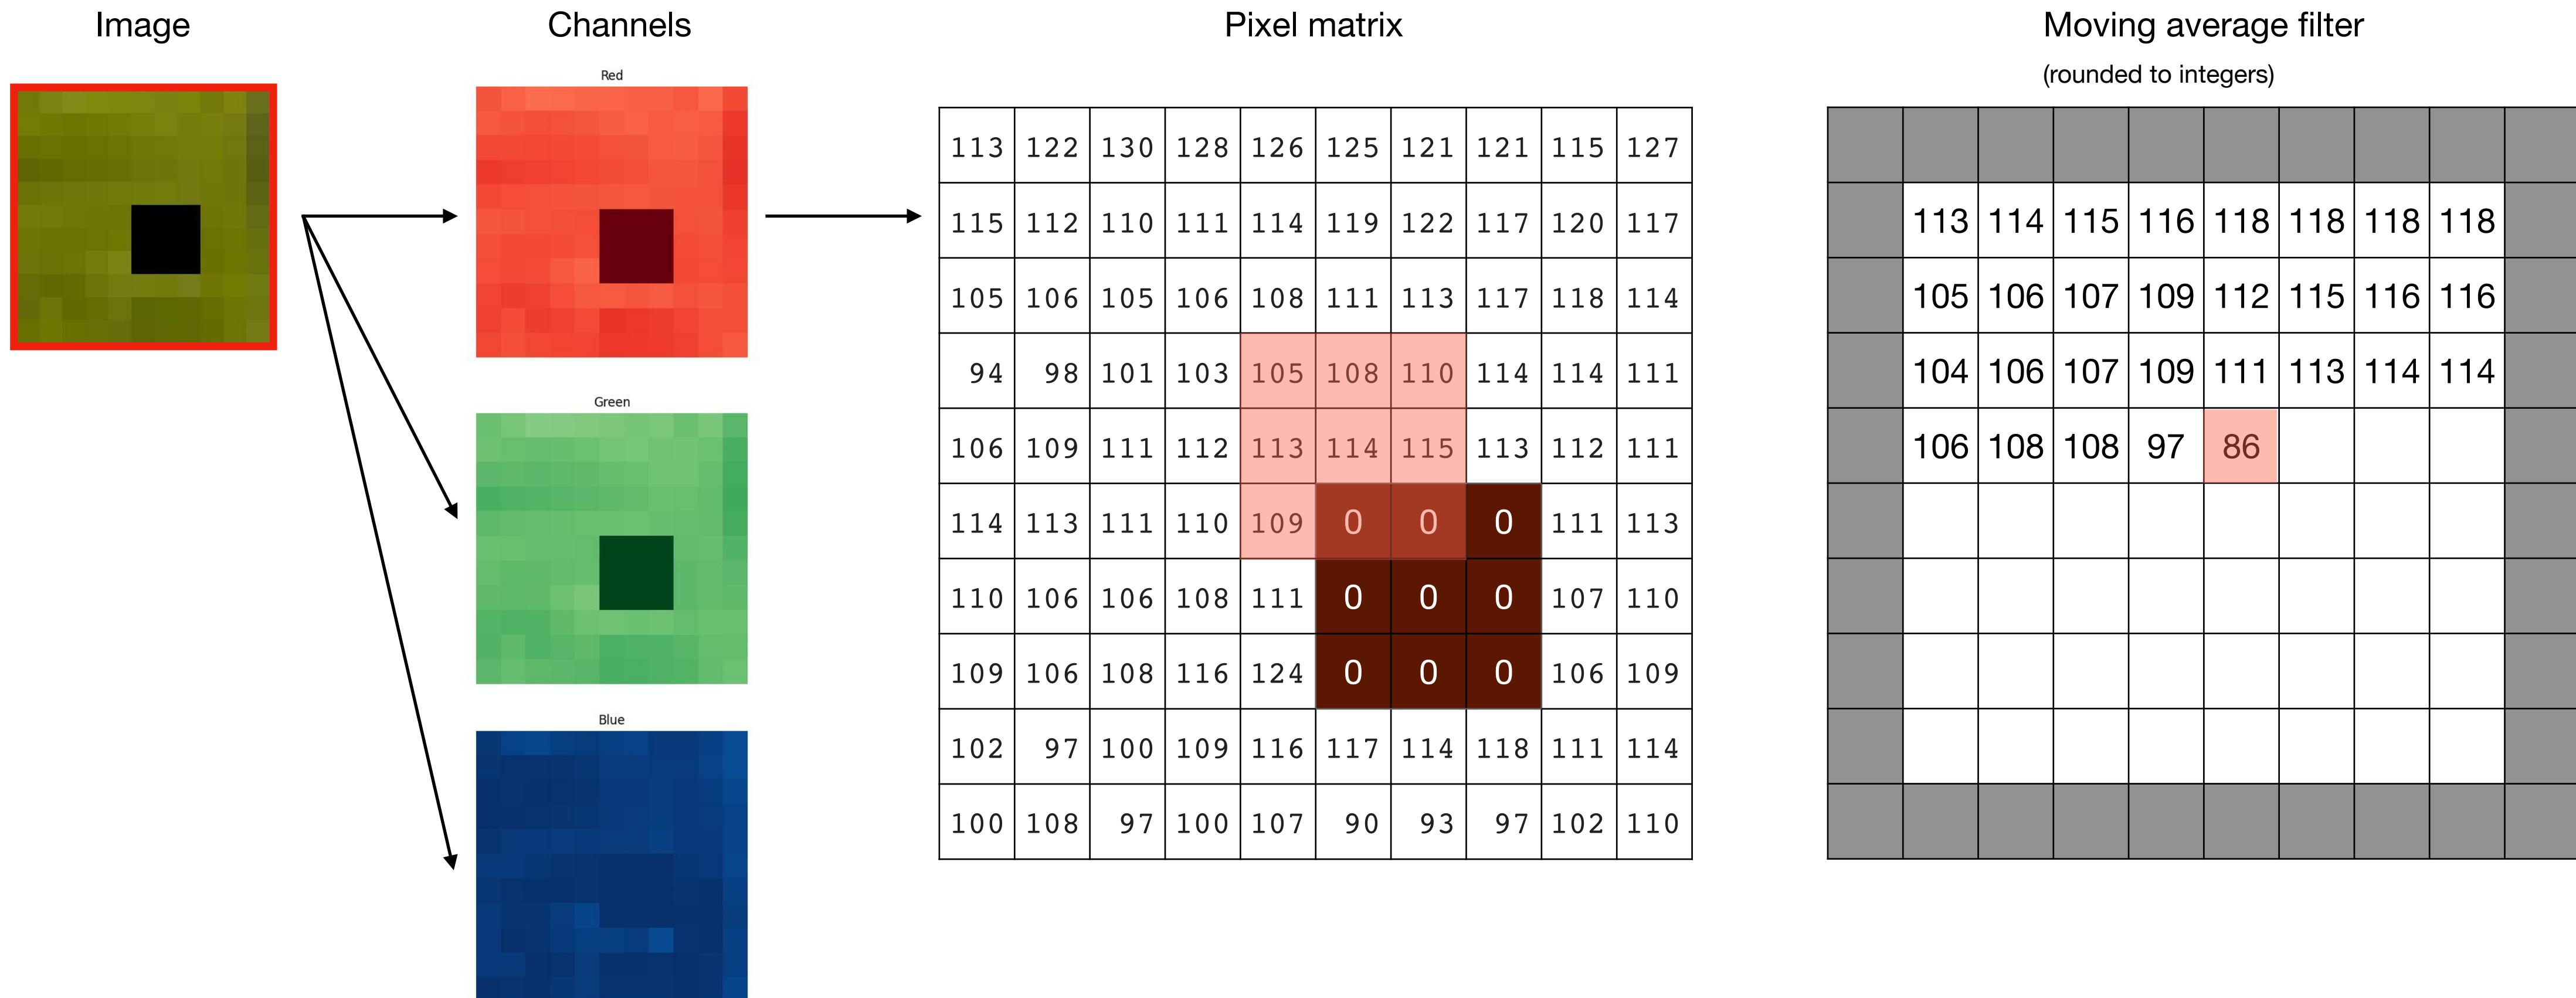
Image processing example



Simple solution for noise removal → neighbourhood average to replace the lost pixels

Convolutions

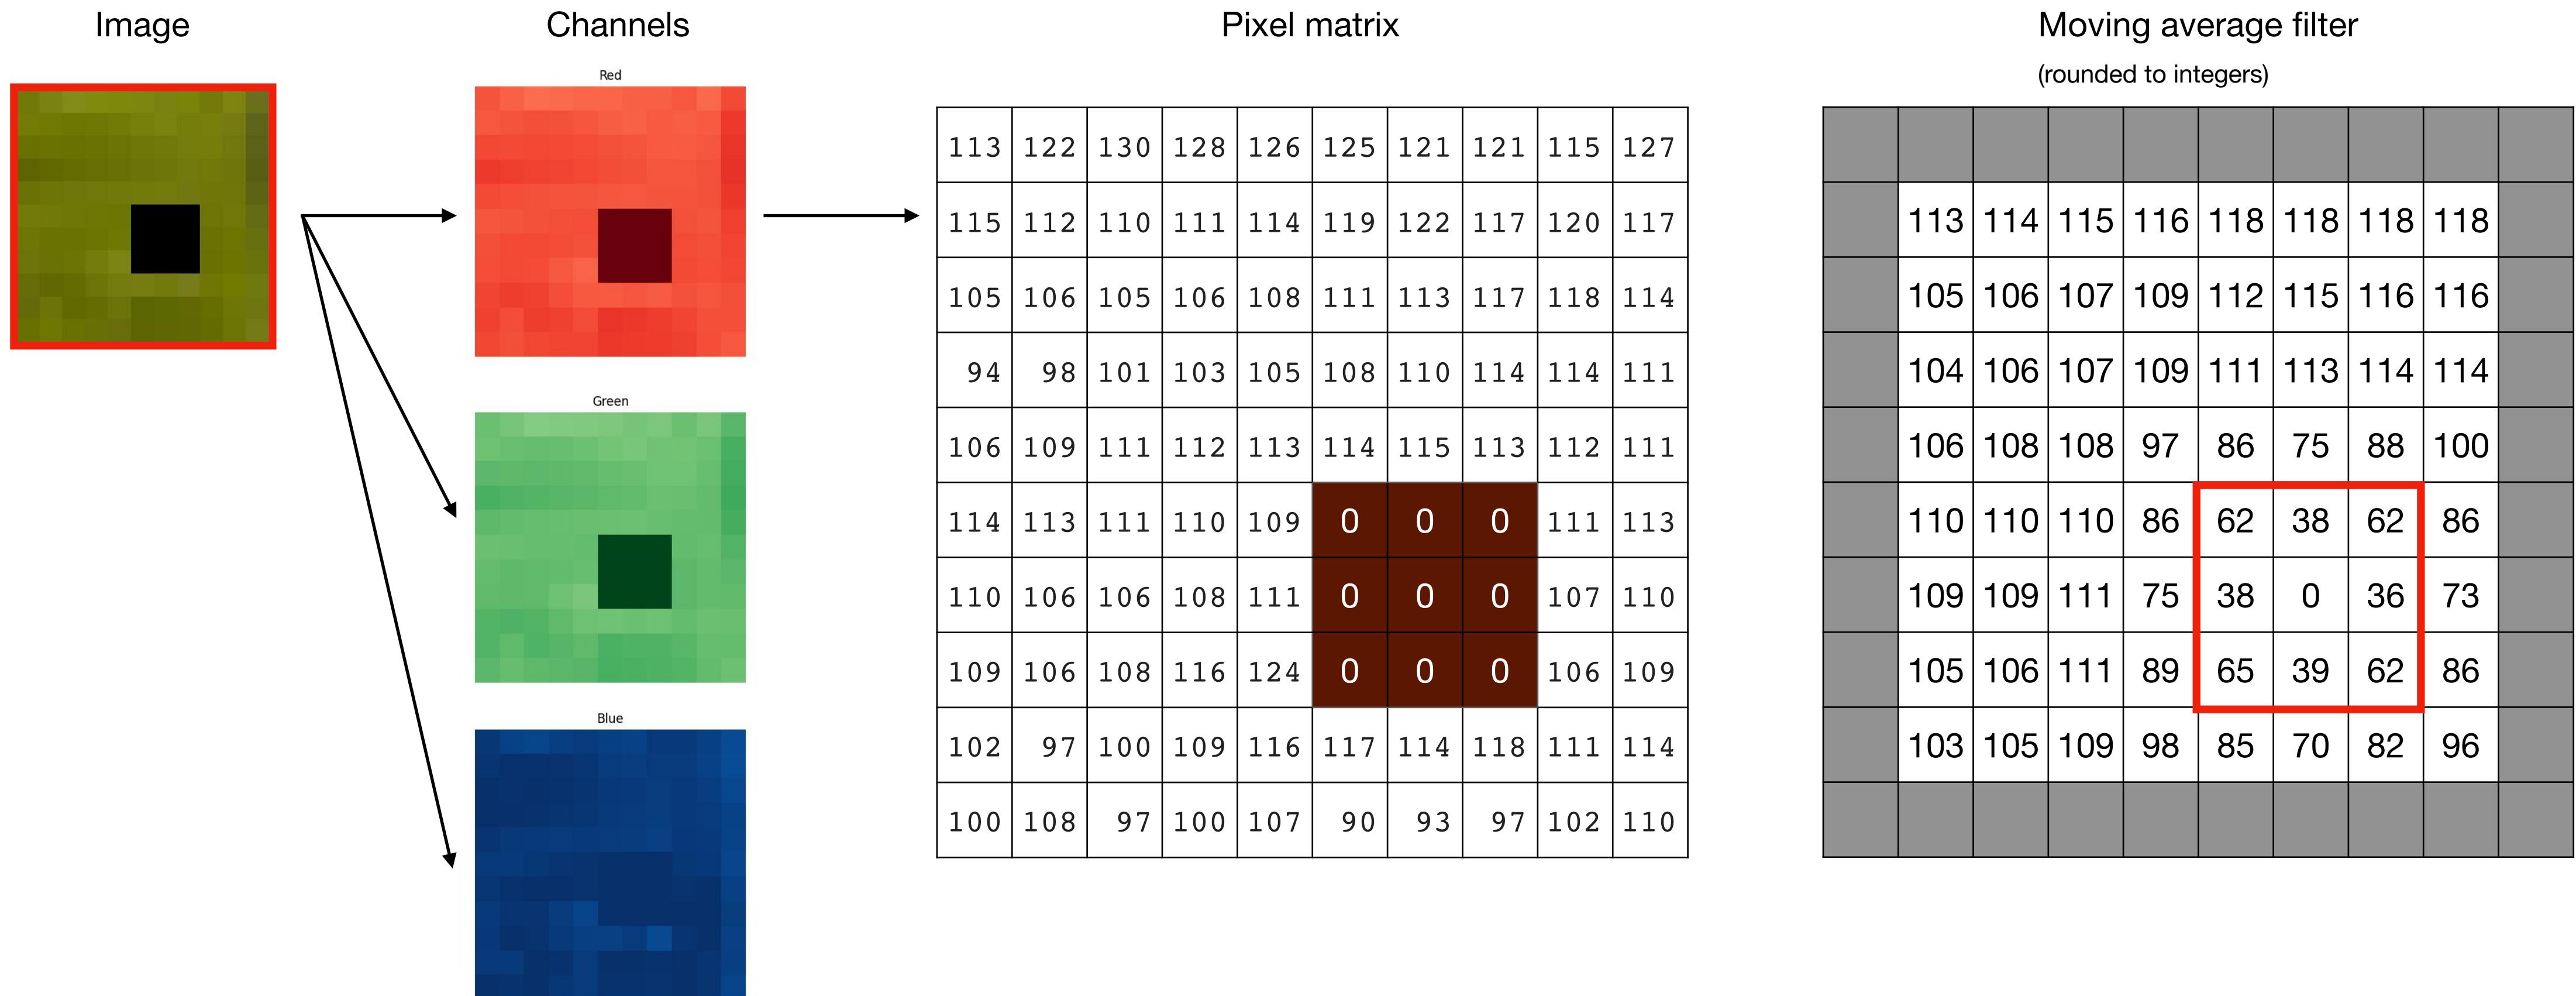
Image processing example



Simple solution for noise removal → neighbourhood average to replace the lost pixels

Convolutions

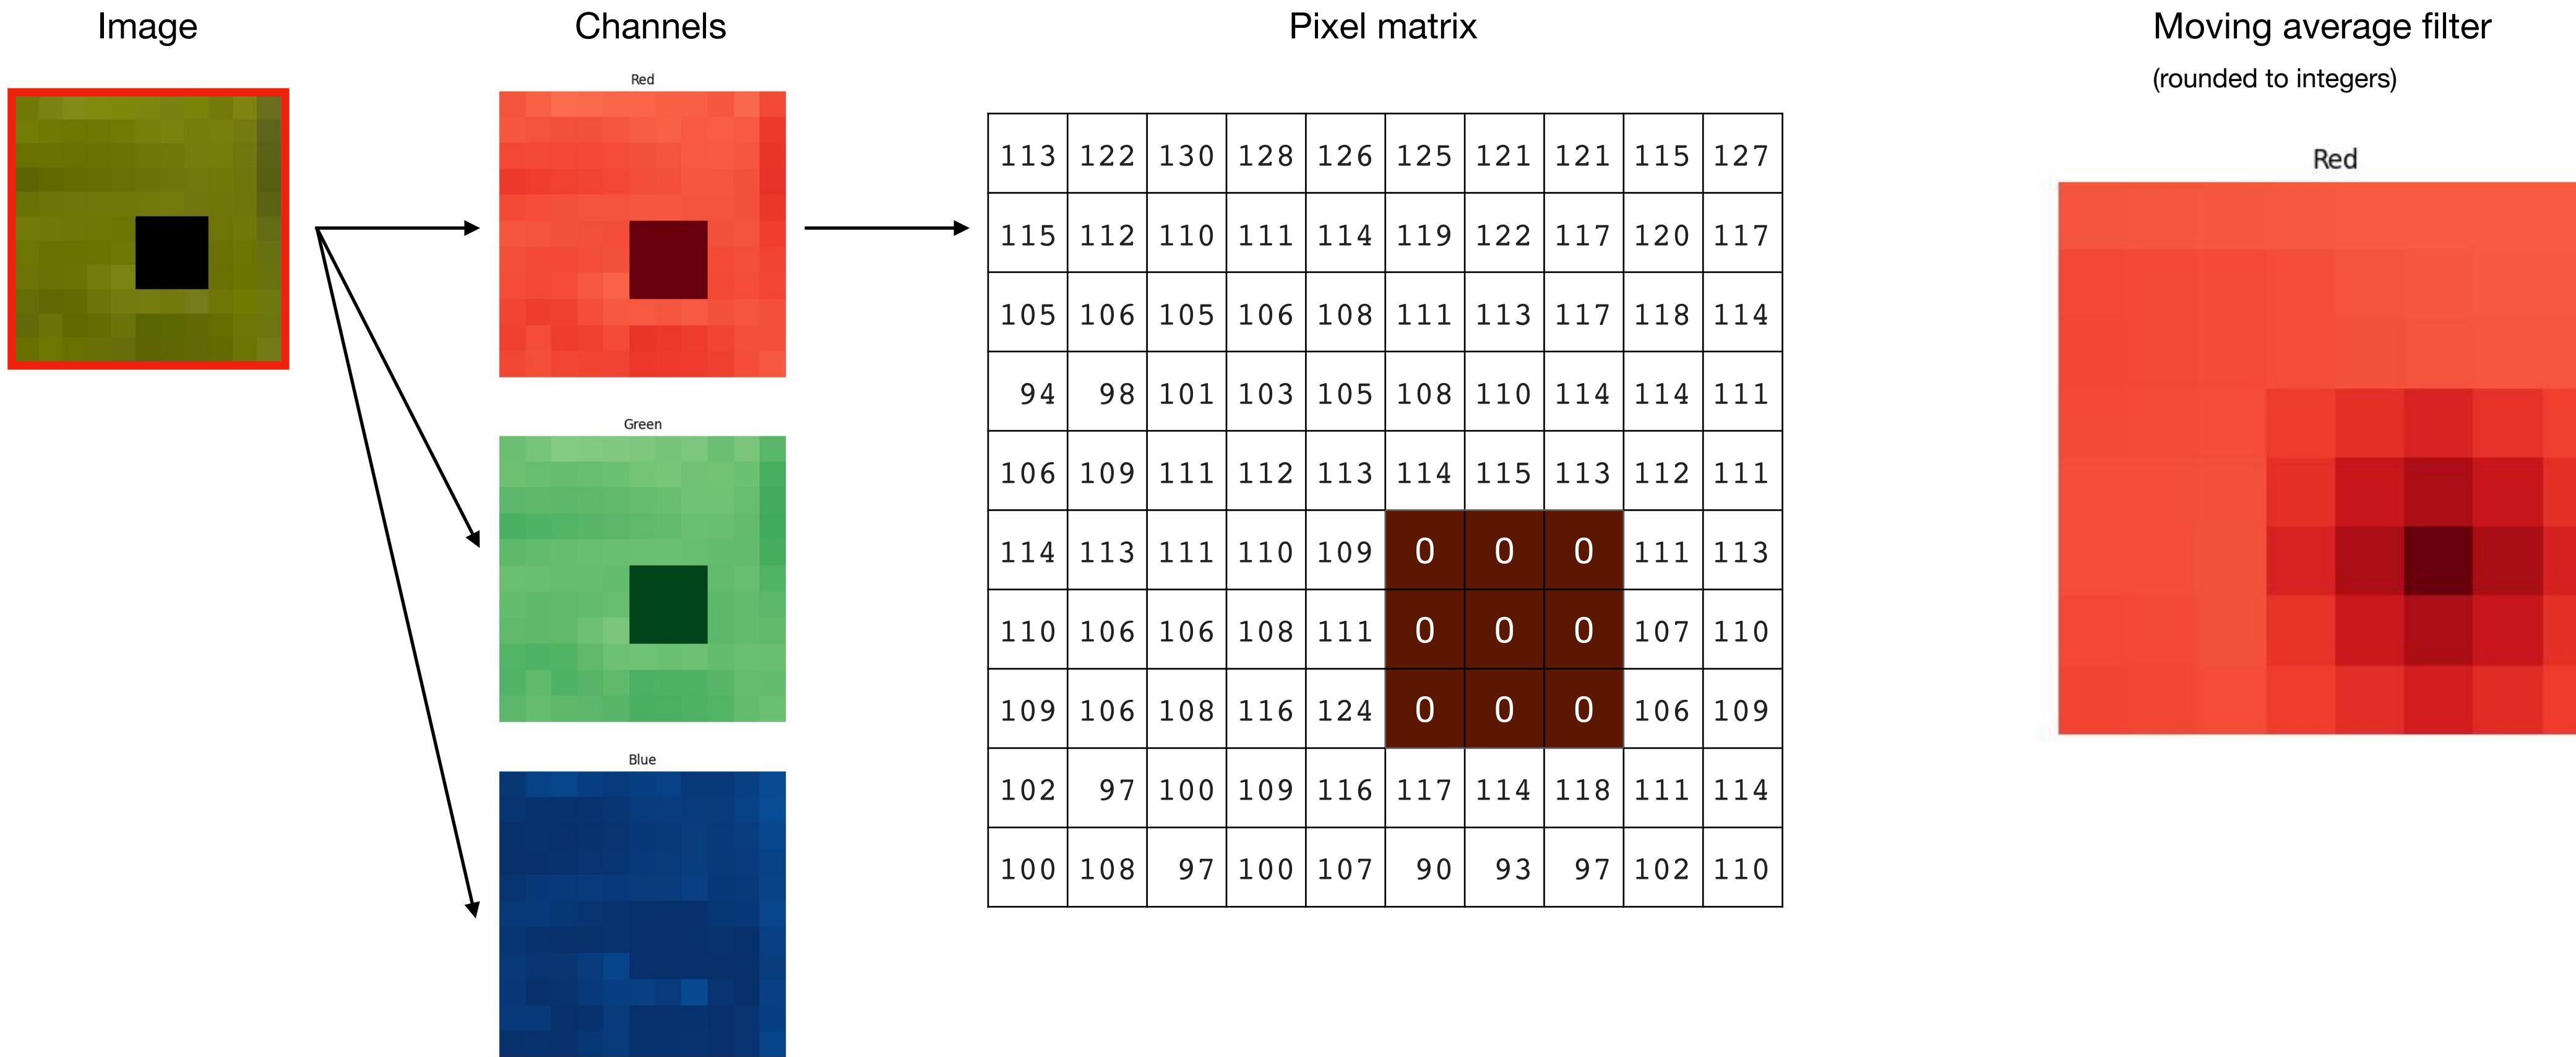
Image processing example



Simple solution for noise removal → neighbourhood average to replace the lost pixels

Convolutions

Image processing example

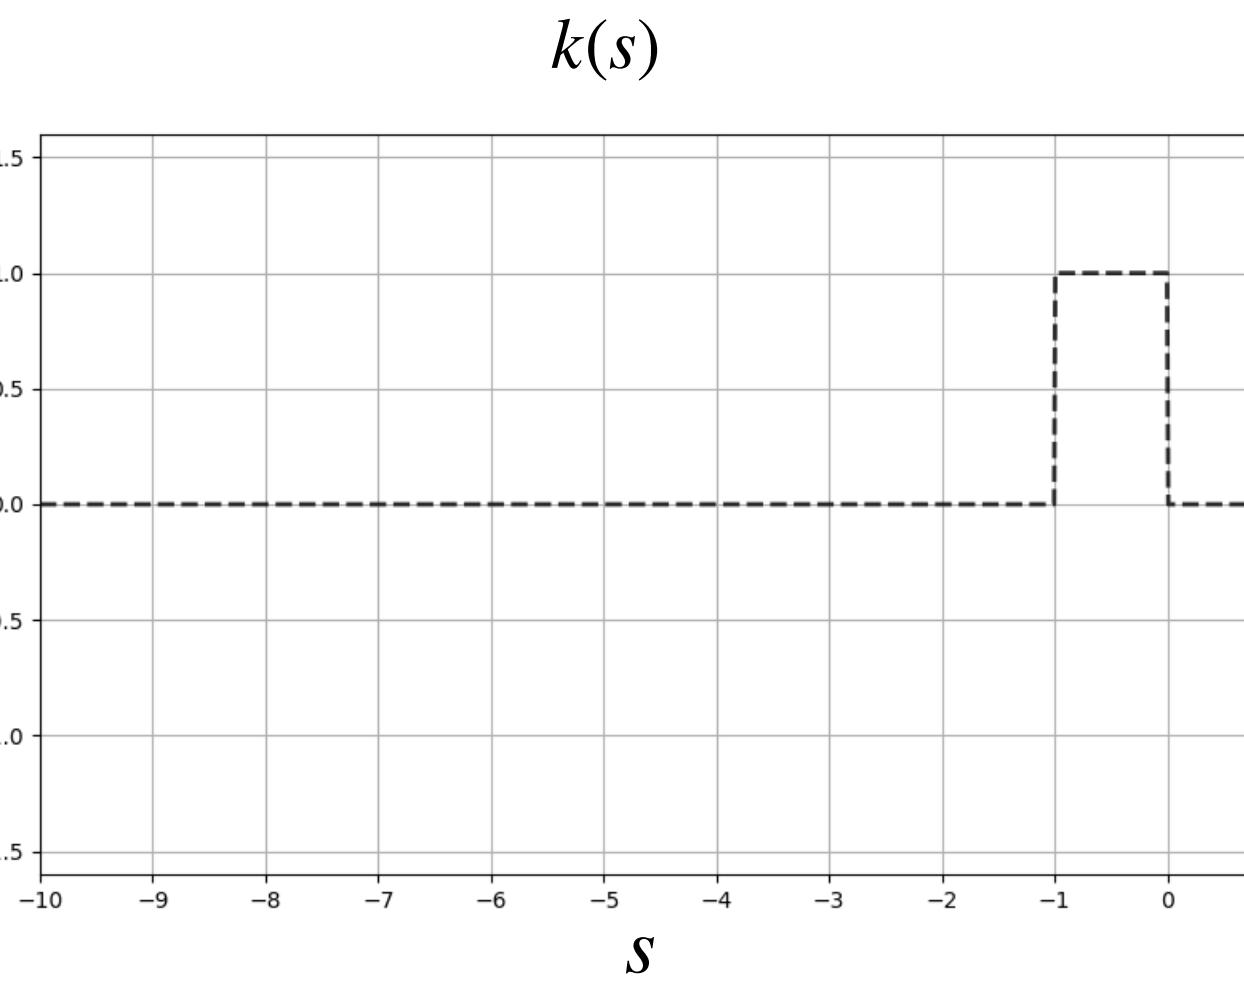
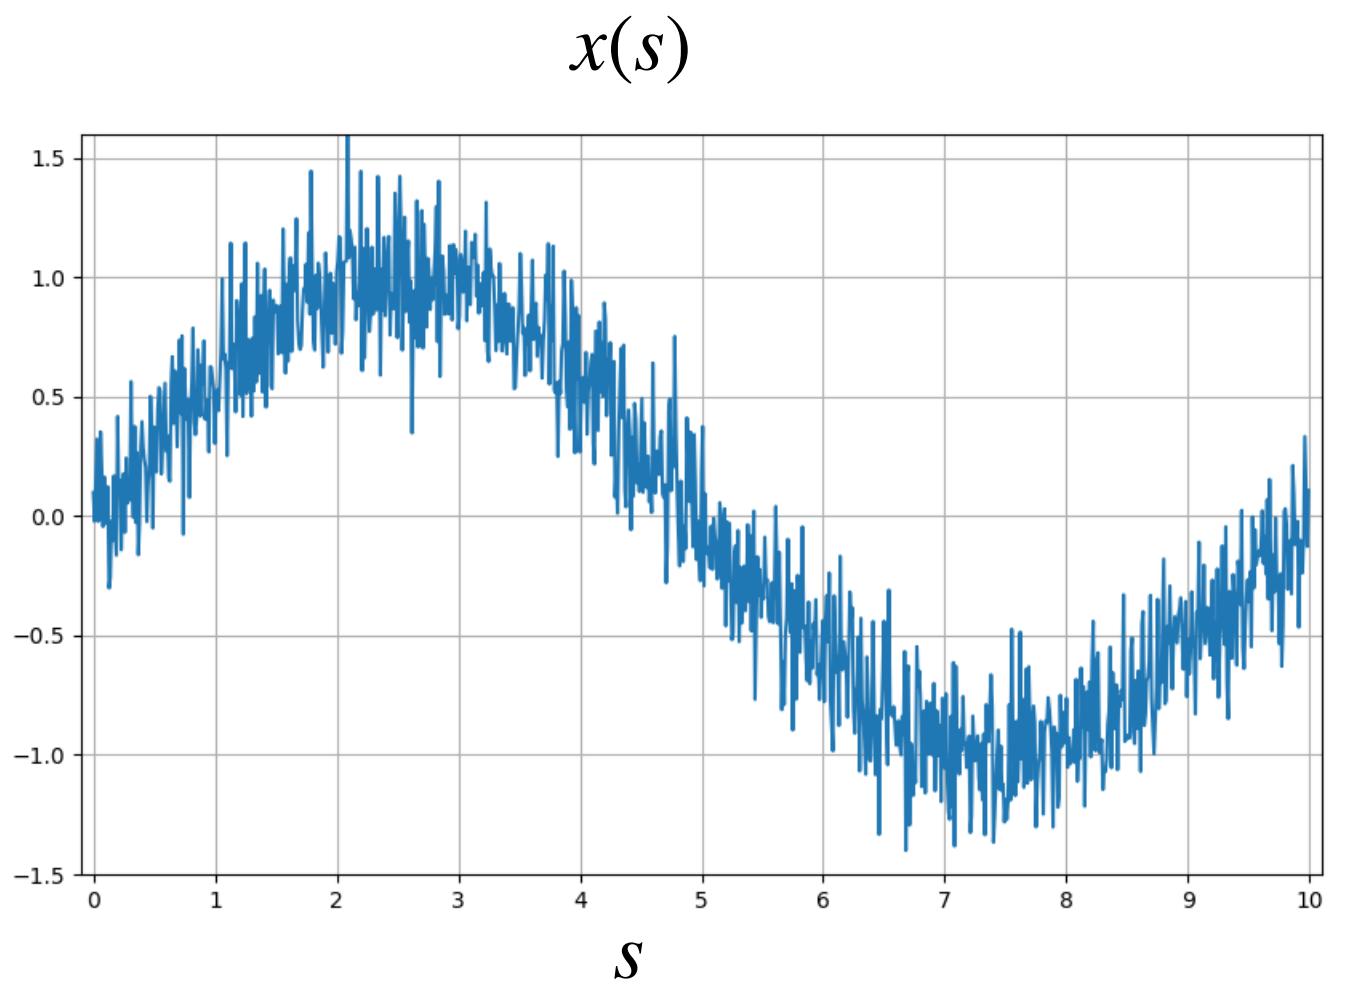


Simple solution for noise removal → neighbourhood average to replace the lost pixels

Convolutions

General definition of convolution between an input signal x and the kernel k

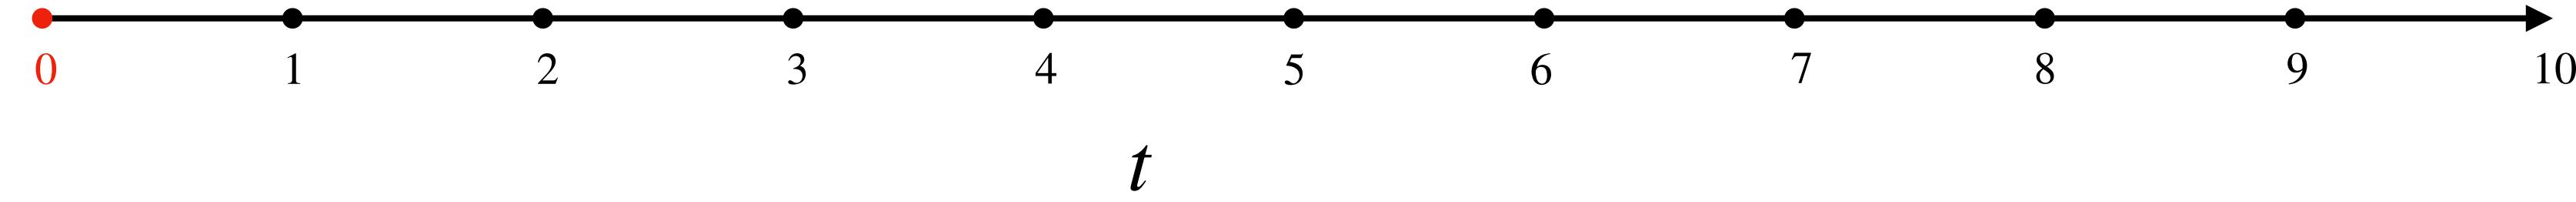
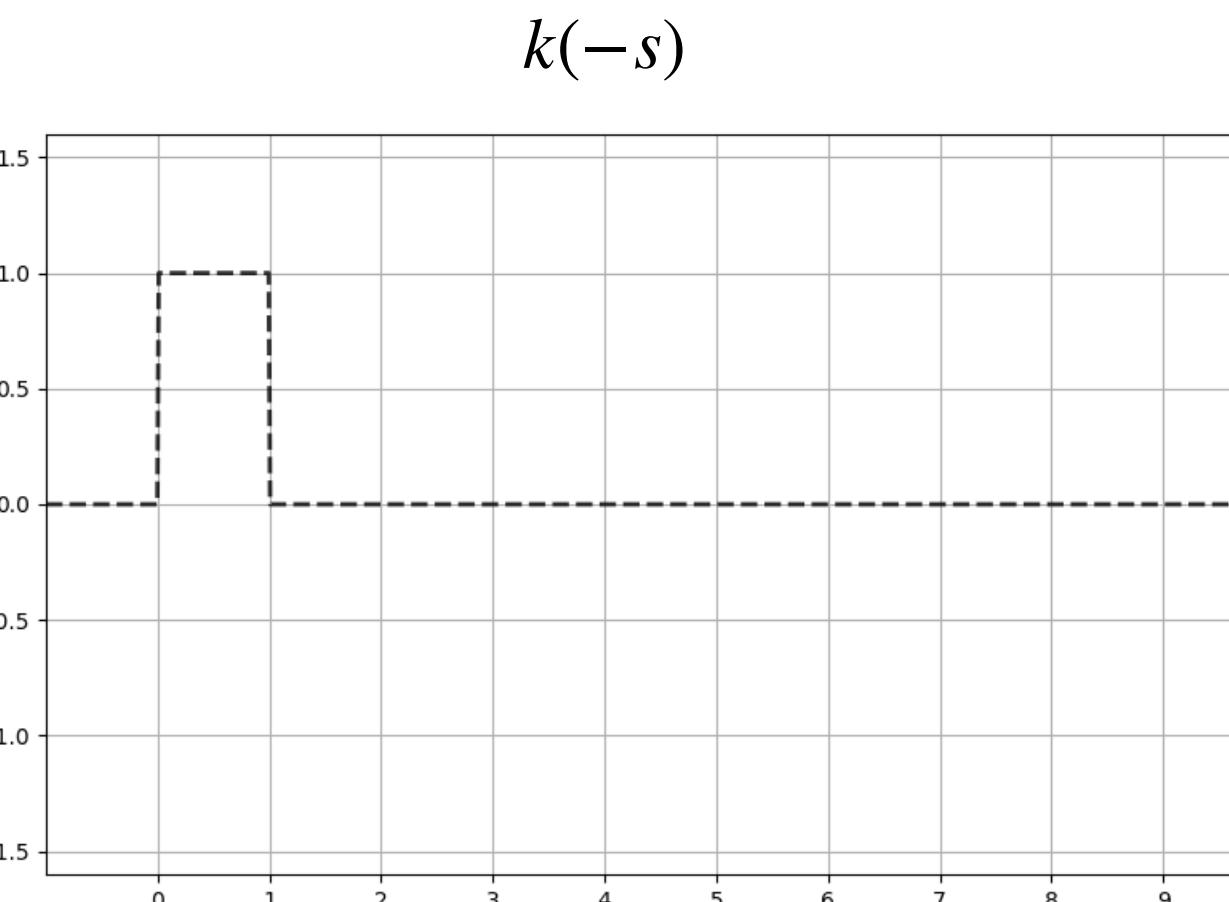
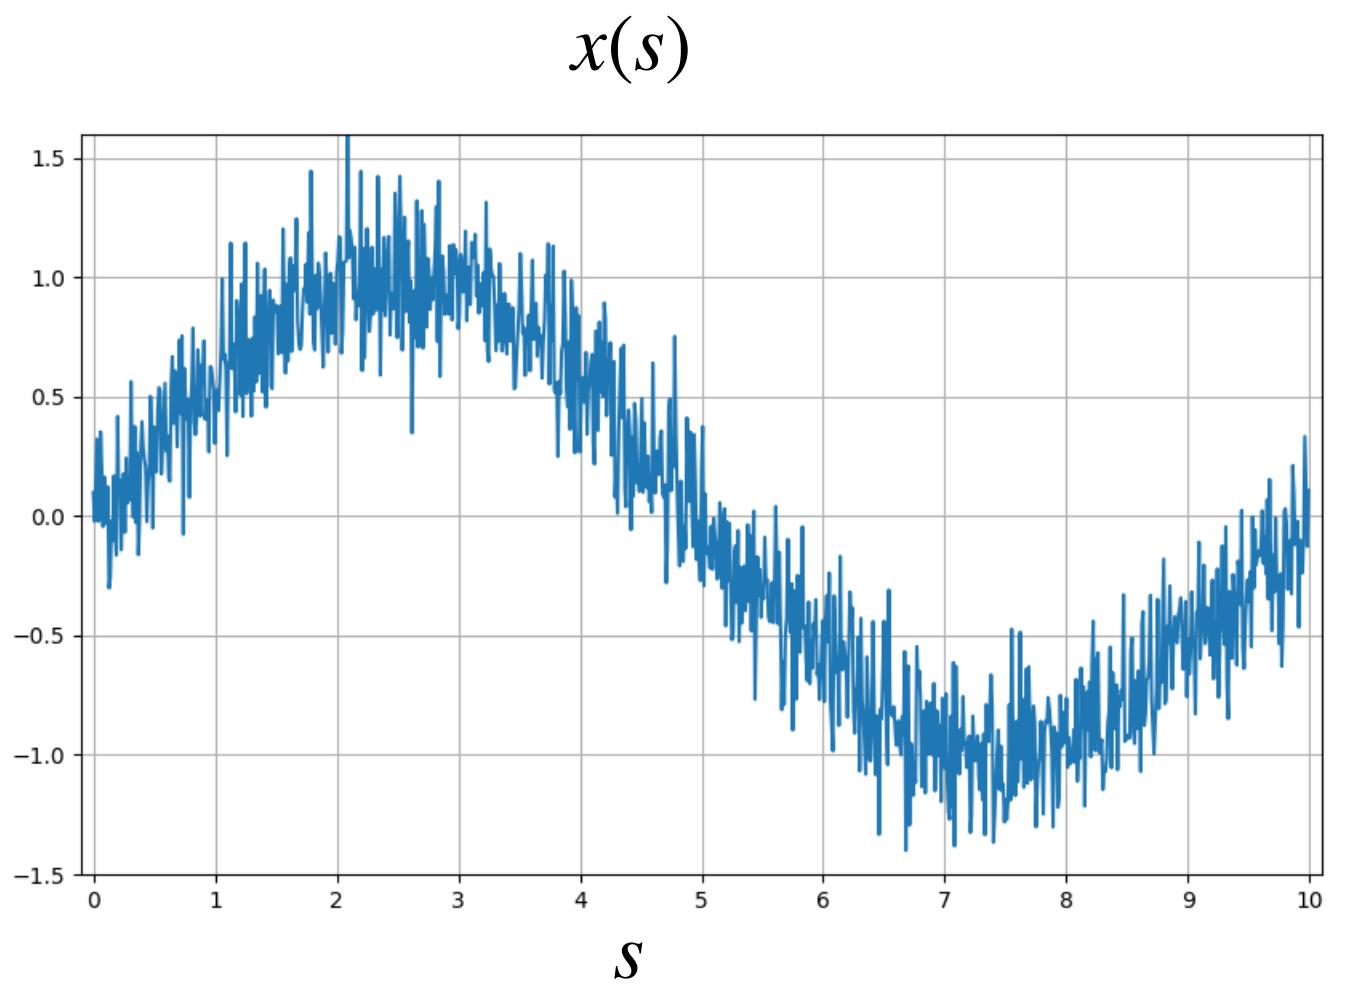
$$(x * k)(t) = \int x(s) k(t - s) ds$$



Convolutions

General definition of convolution between an input signal x and the kernel k

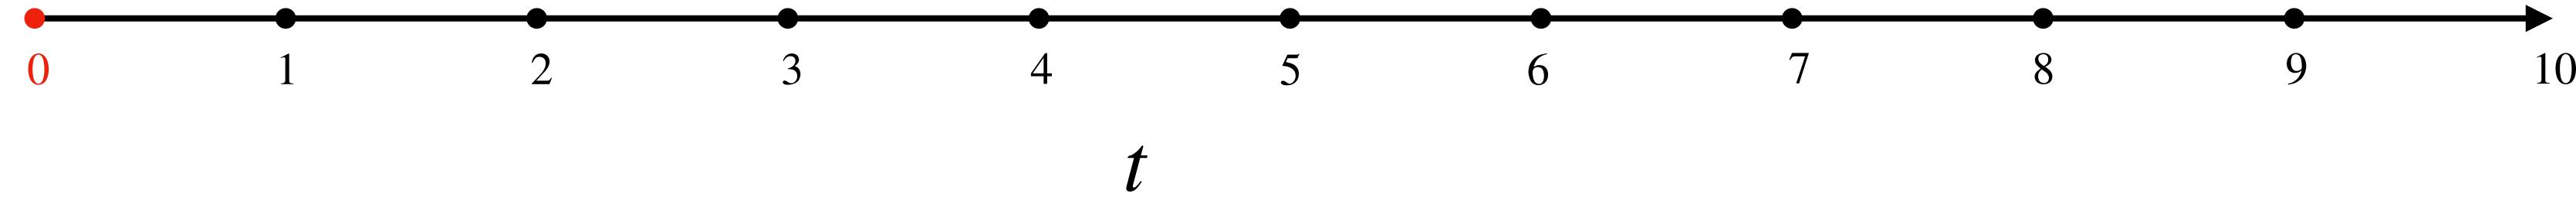
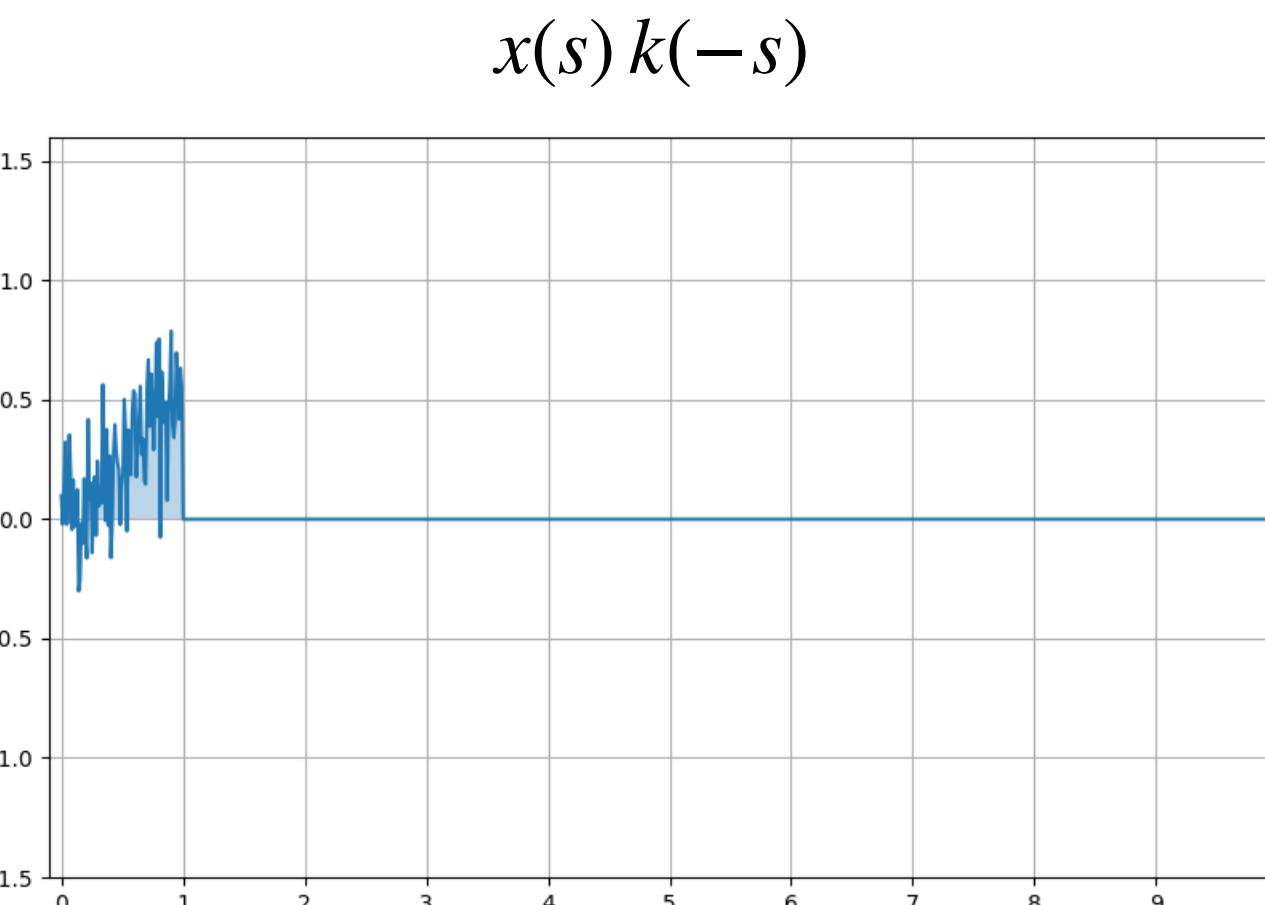
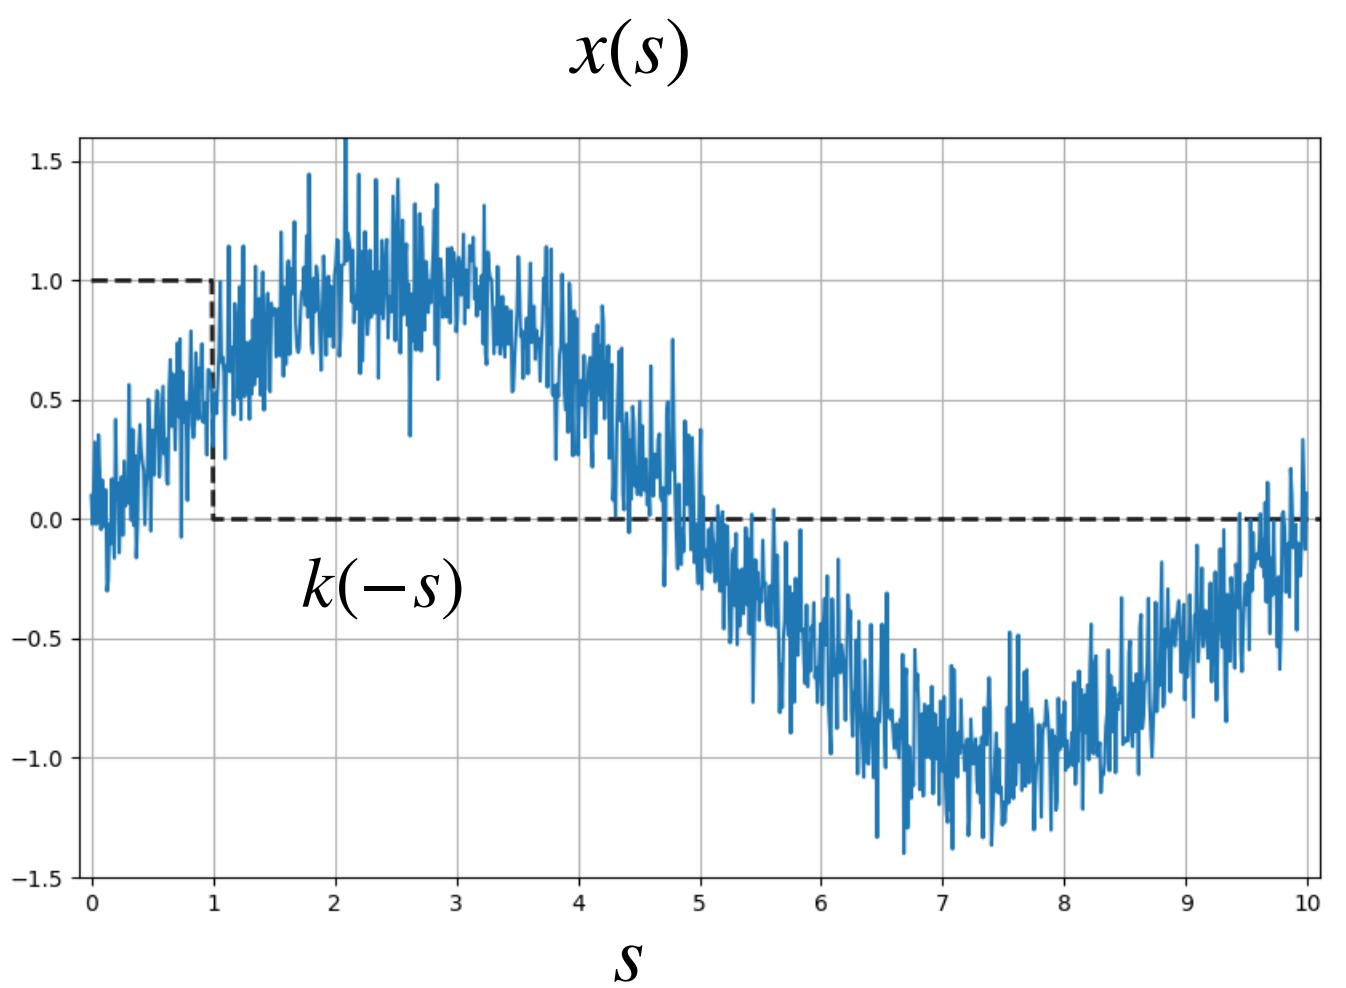
$$(x * k)(t) = \int x(s) k(t - s) ds$$



Convolutions

General definition of convolution between an input signal x and the kernel k

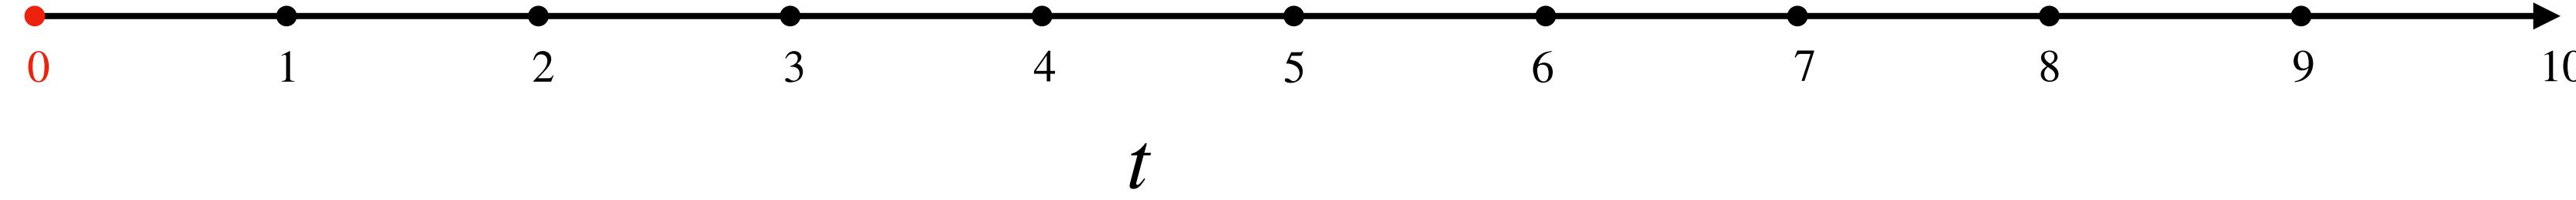
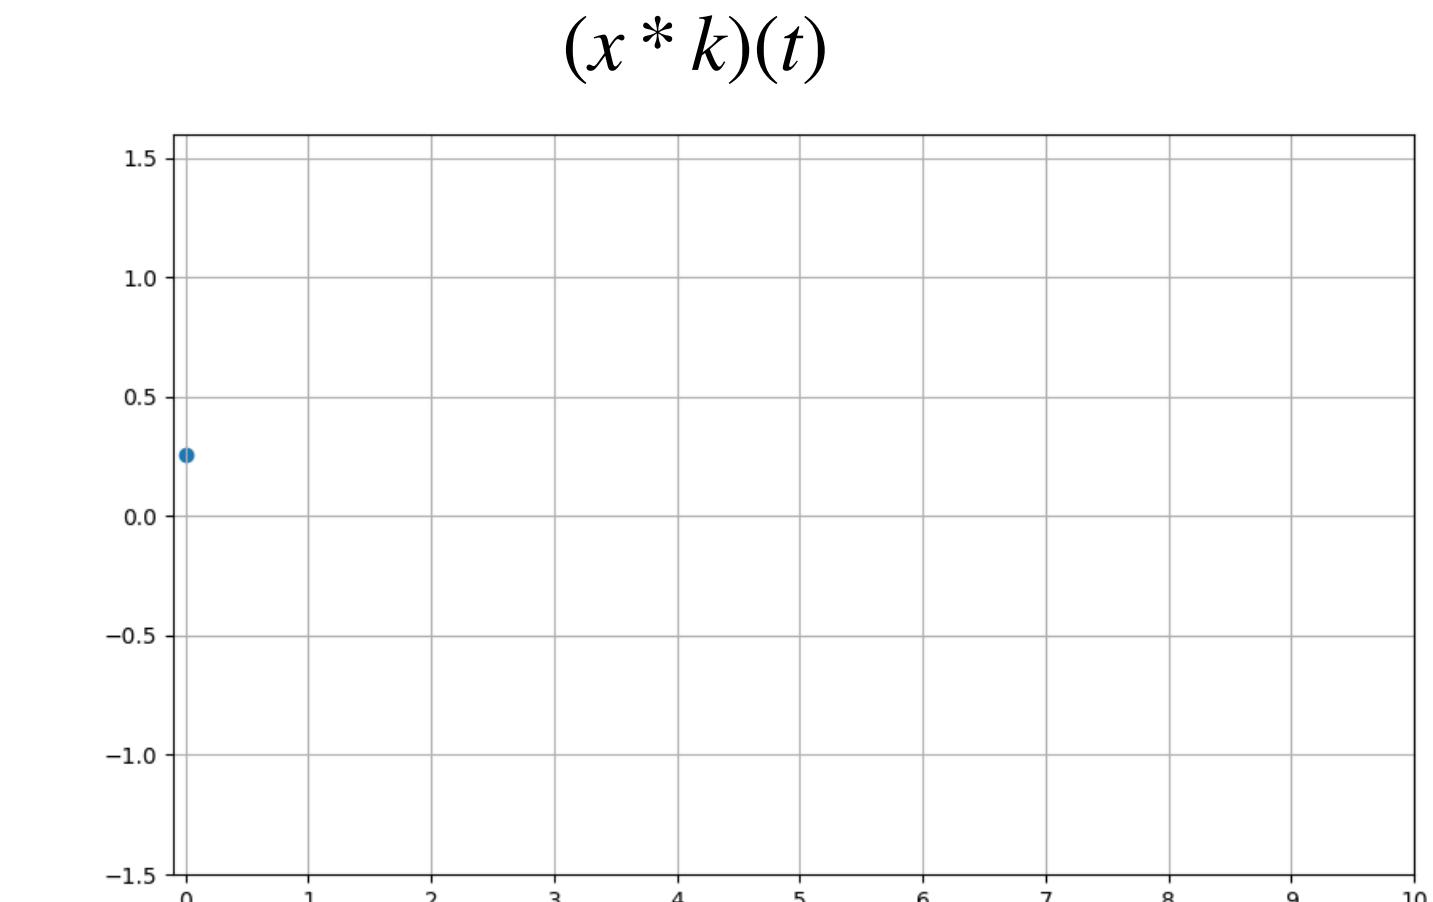
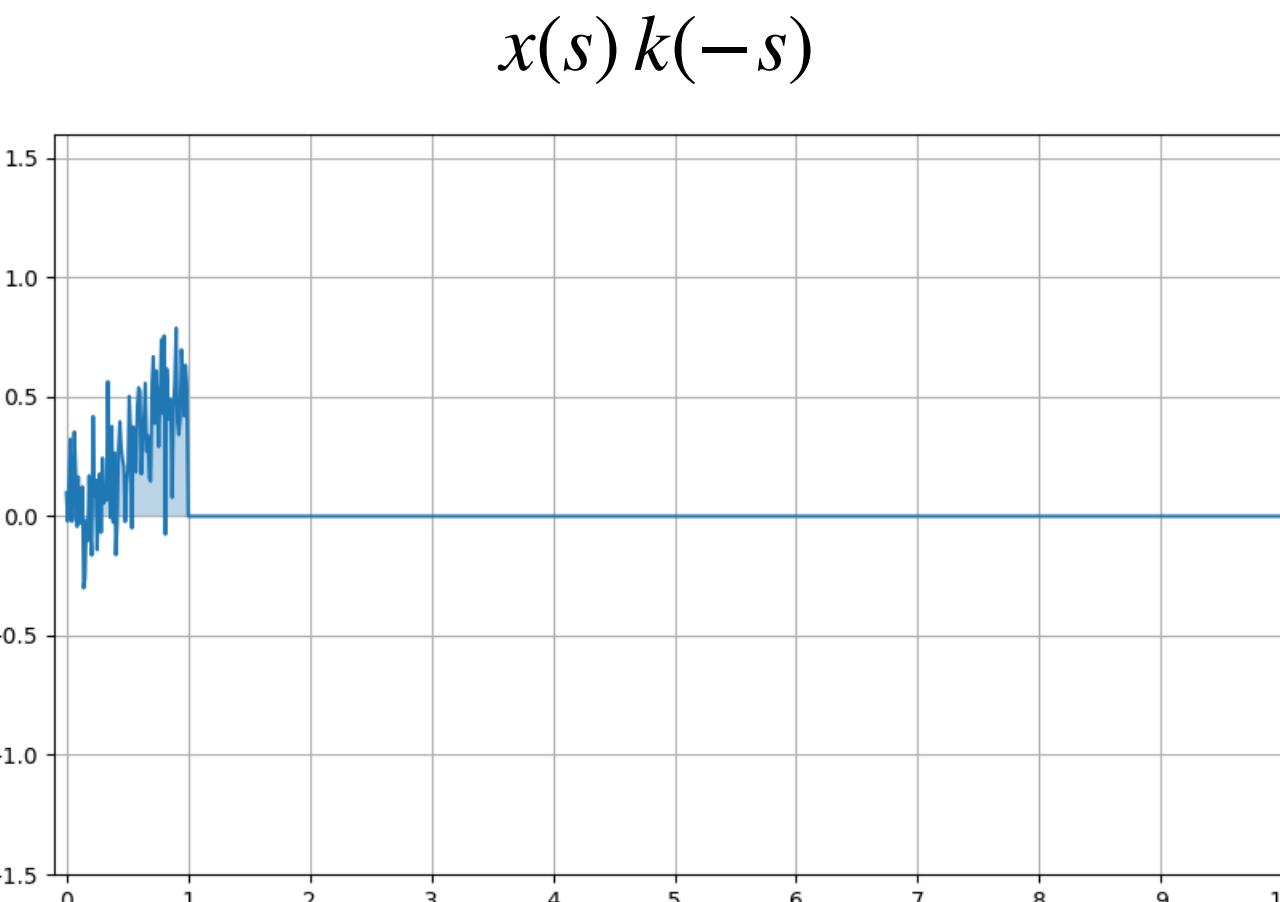
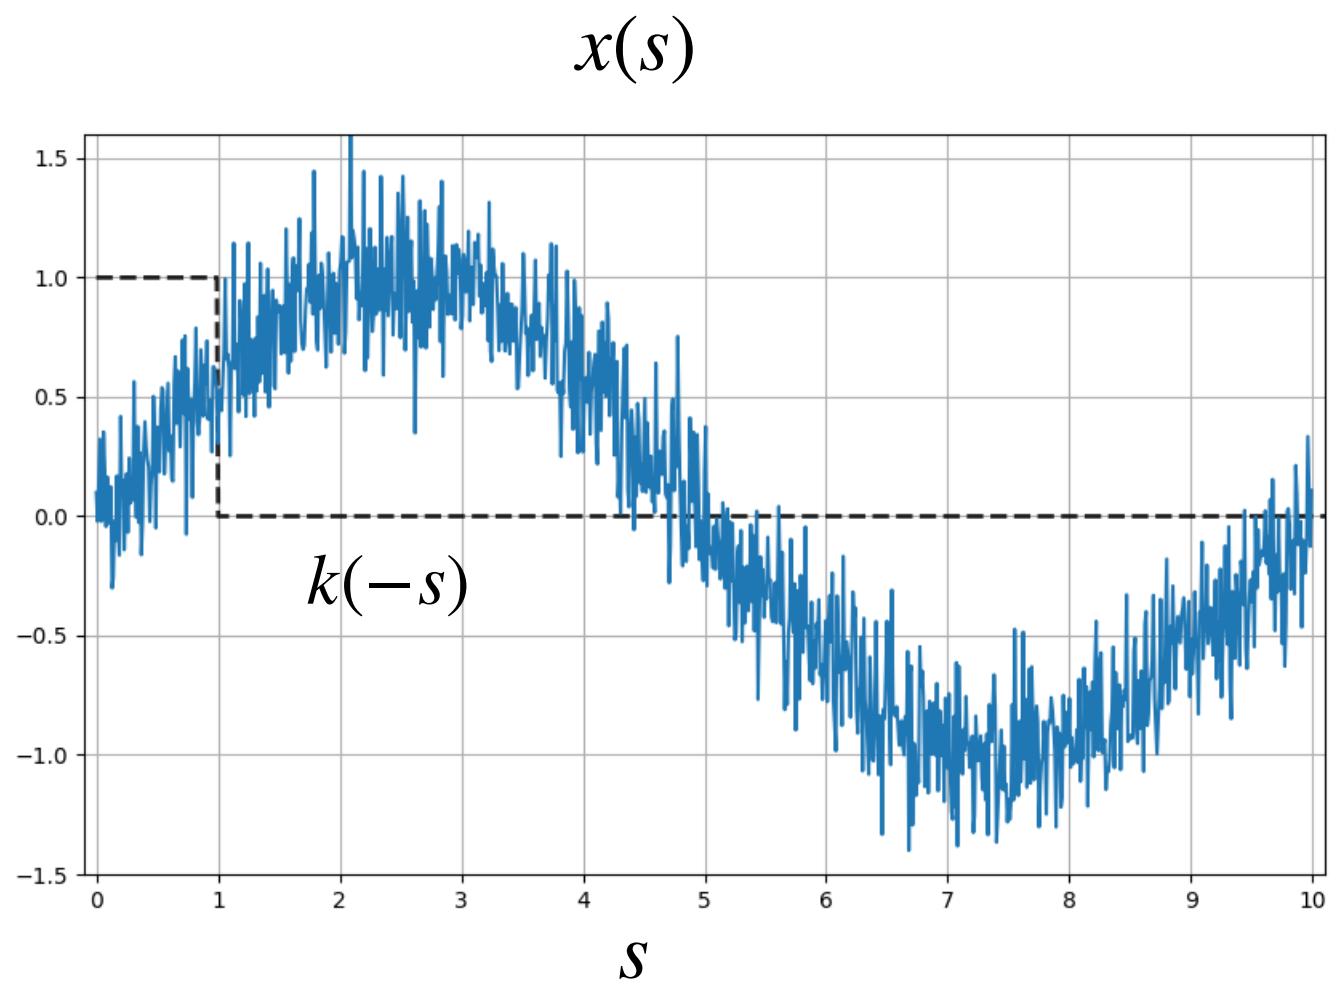
$$(x * k)(t) = \int x(s) k(t - s) ds$$



Convolutions

General definition of convolution between an input signal x and the kernel k

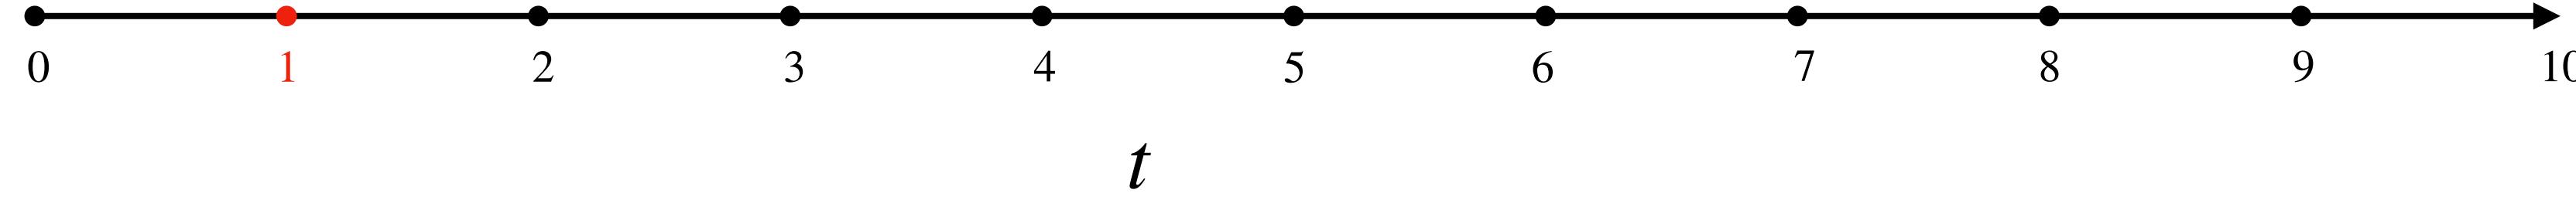
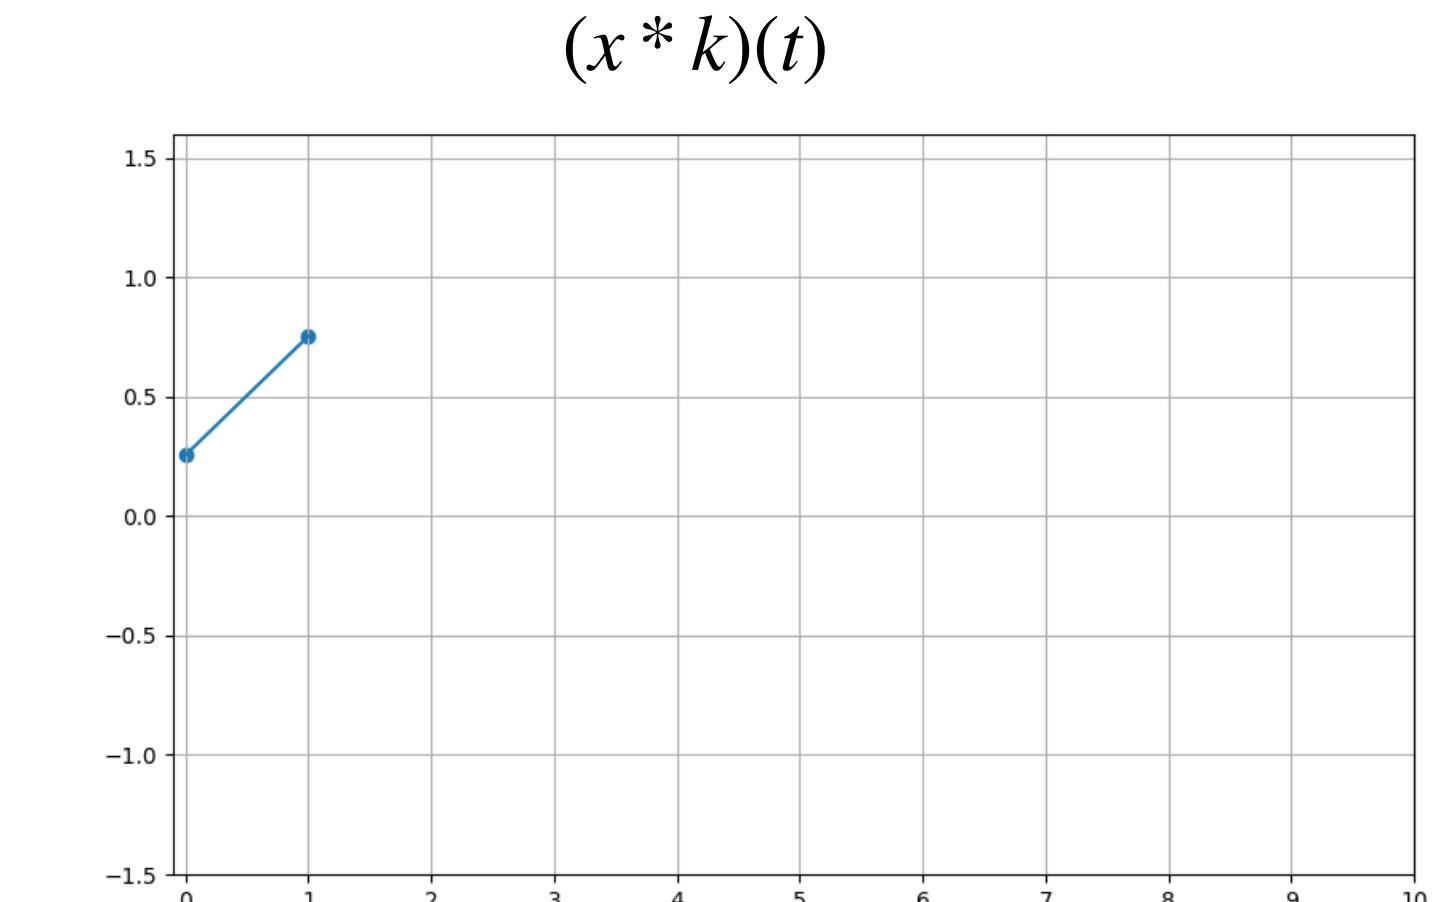
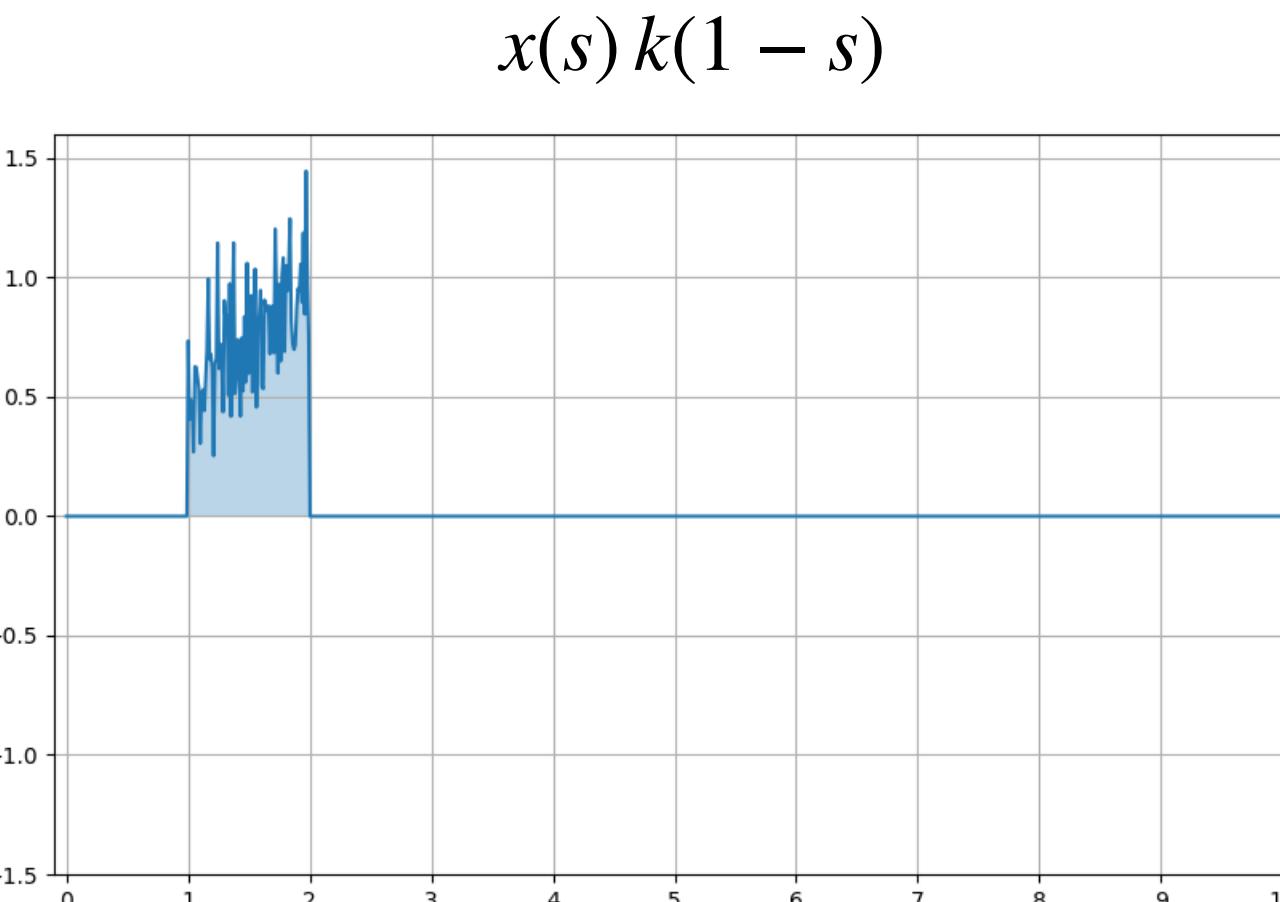
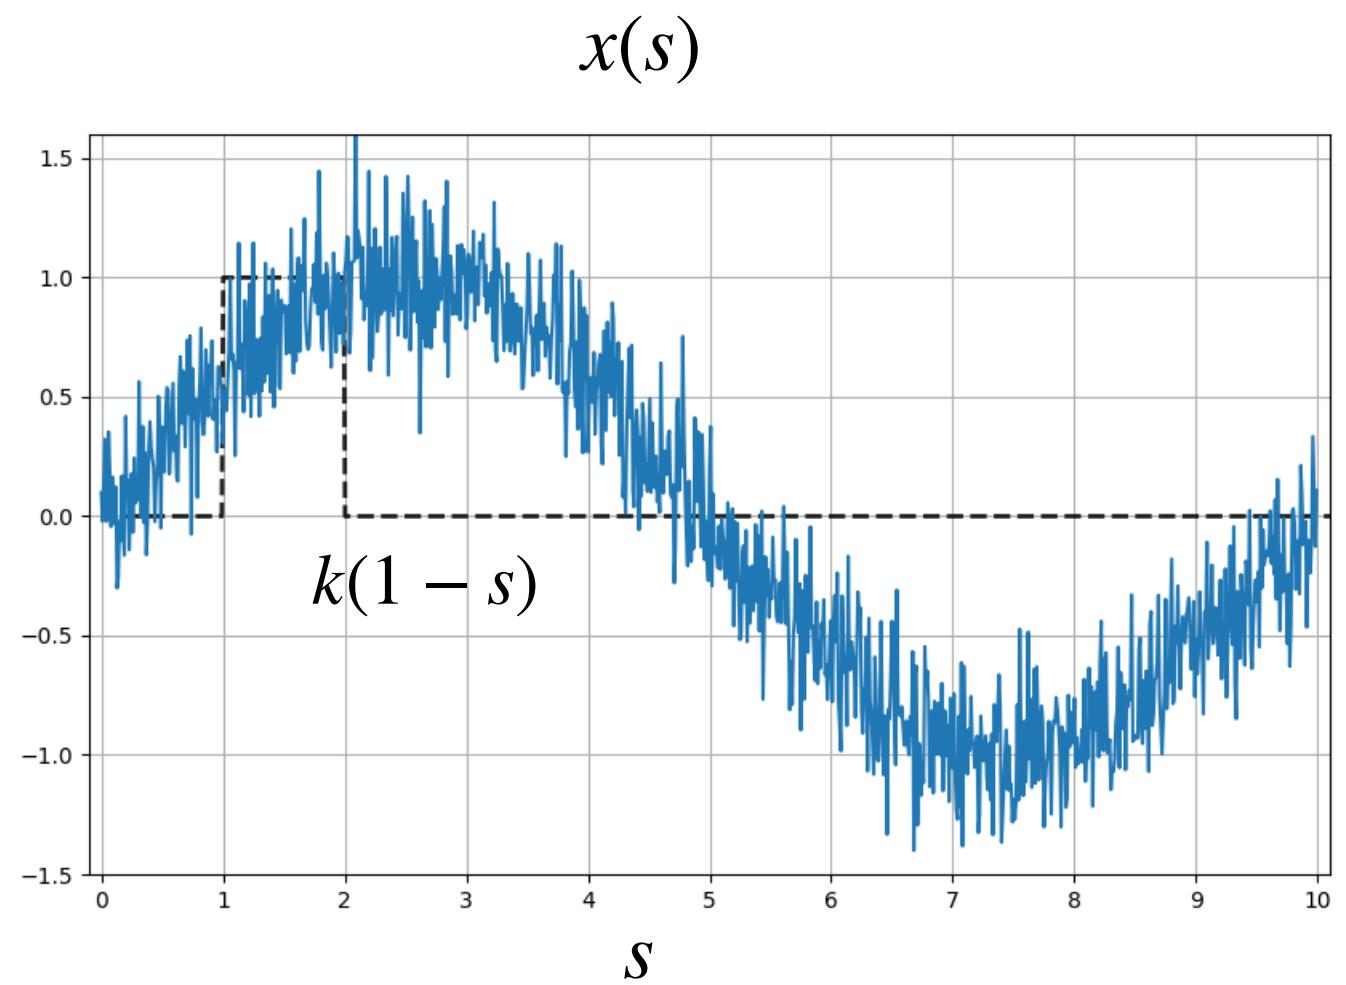
$$(x * k)(t) = \int x(s) k(t - s) ds$$



Convolutions

General definition of convolution between an input signal x and the kernel k

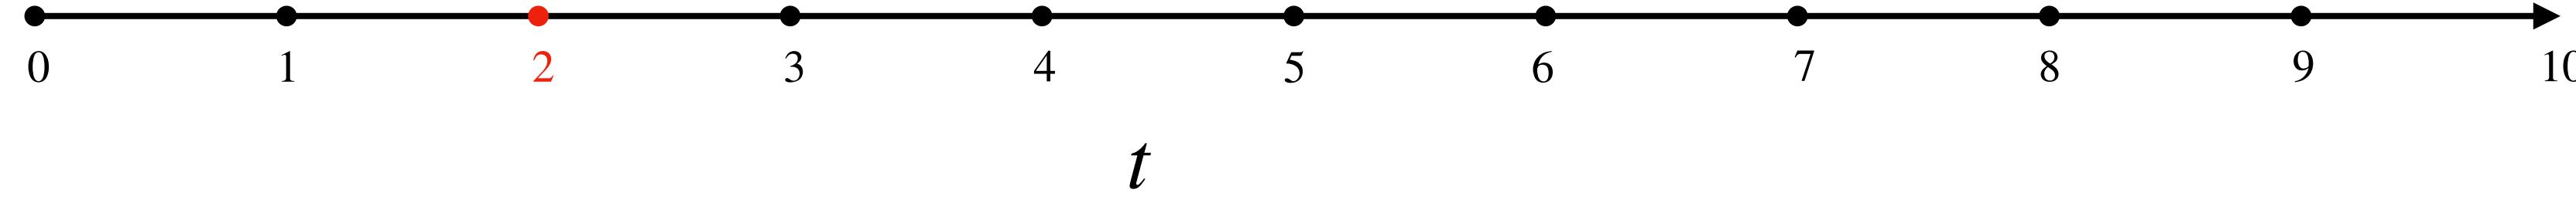
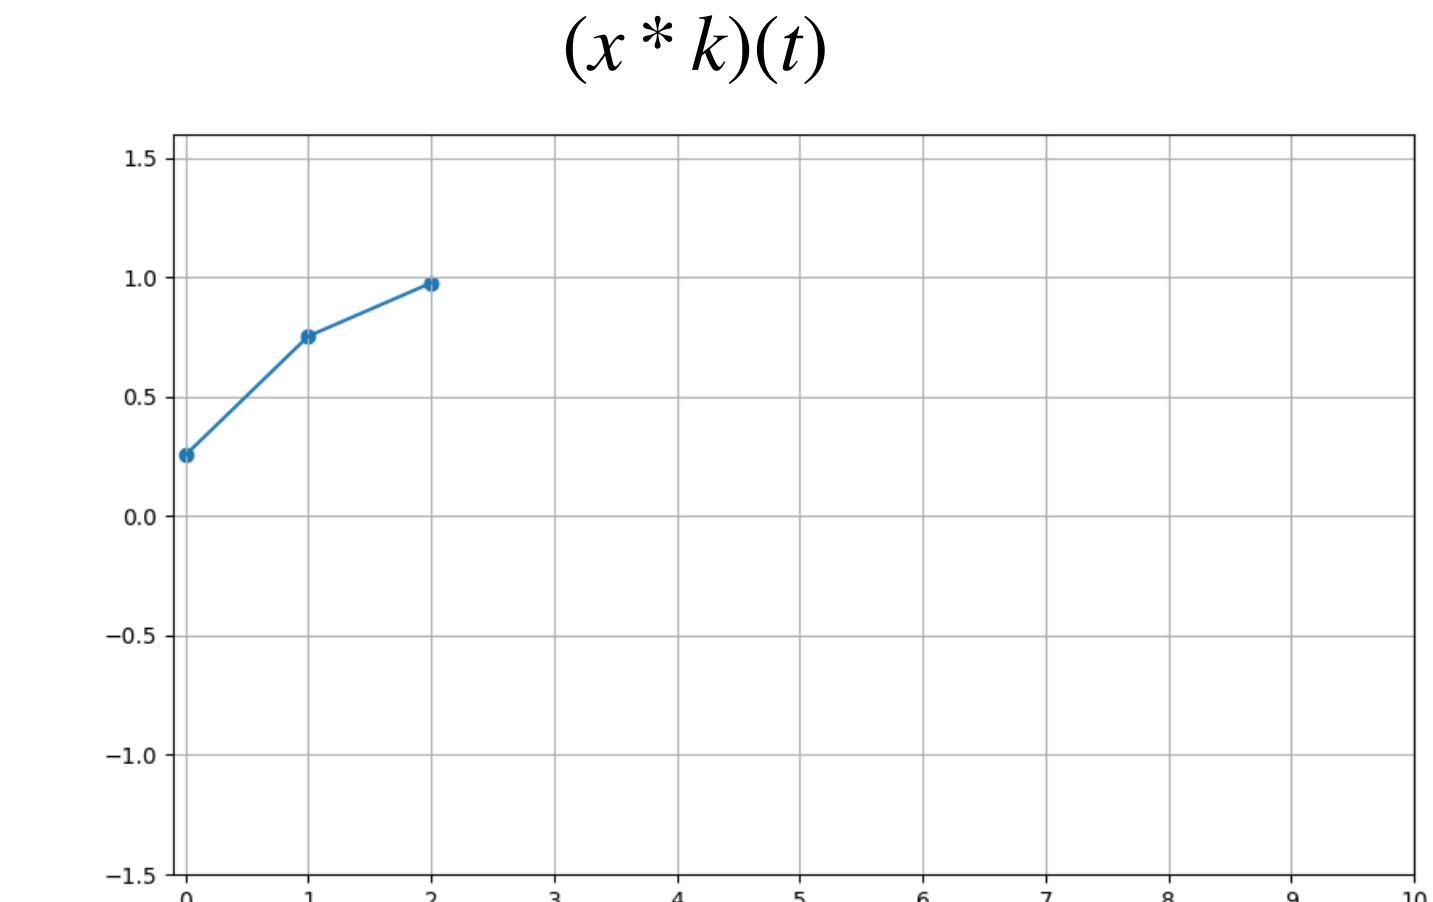
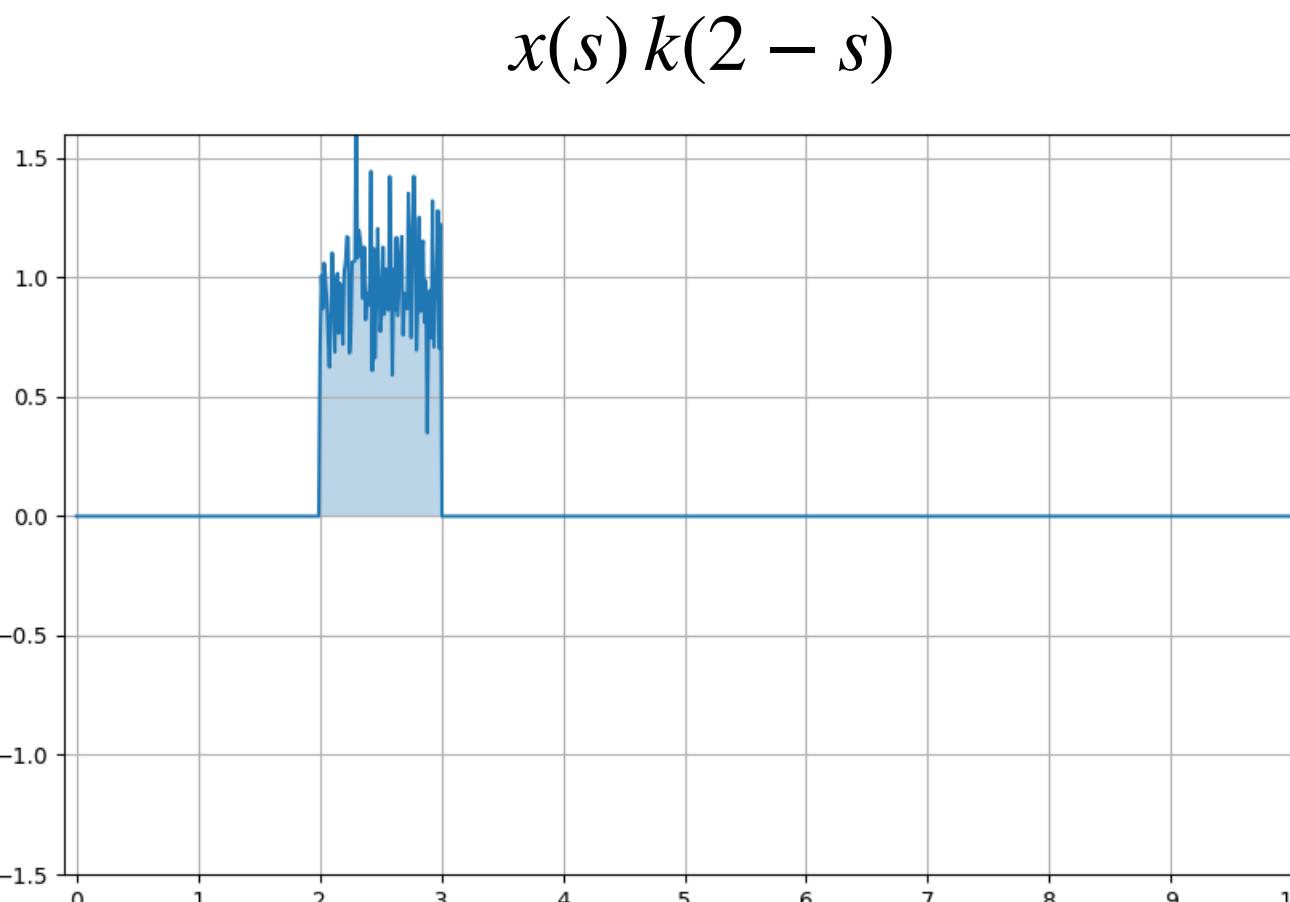
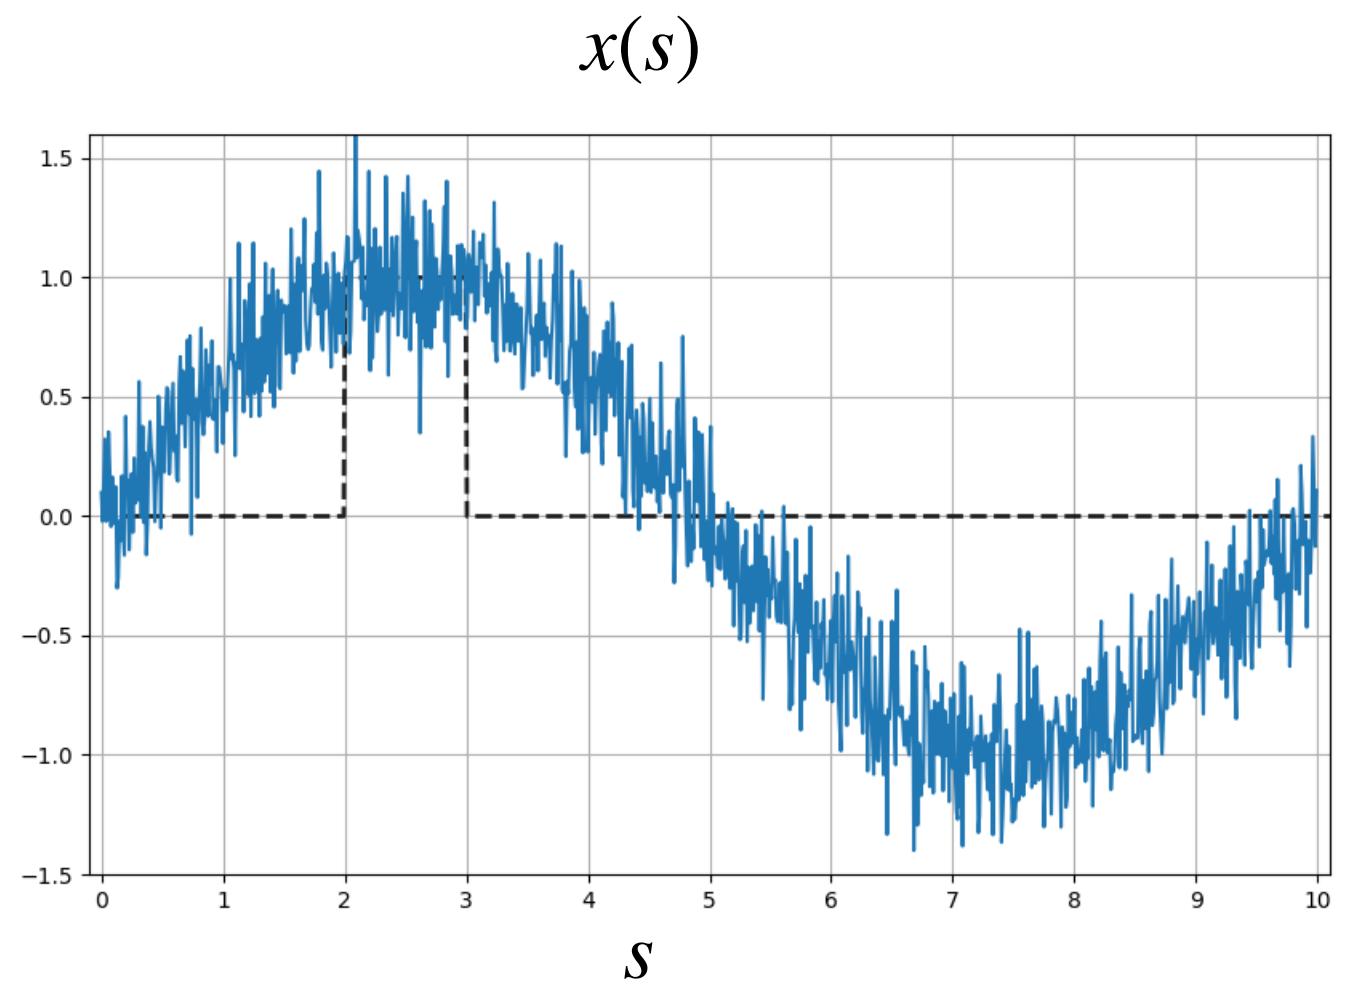
$$(x * k)(t) = \int x(s) k(t - s) ds$$



Convolutions

General definition of convolution between an input signal x and the kernel k

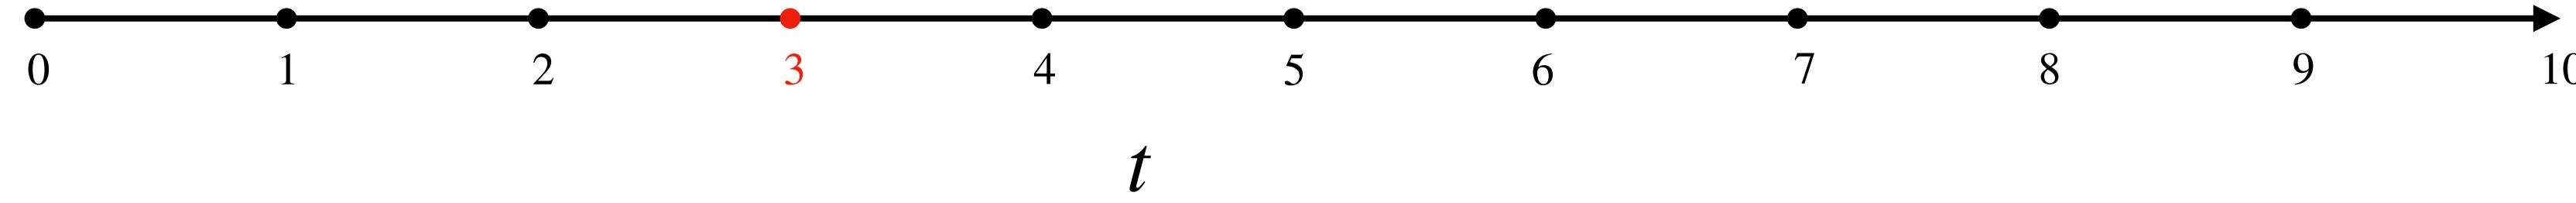
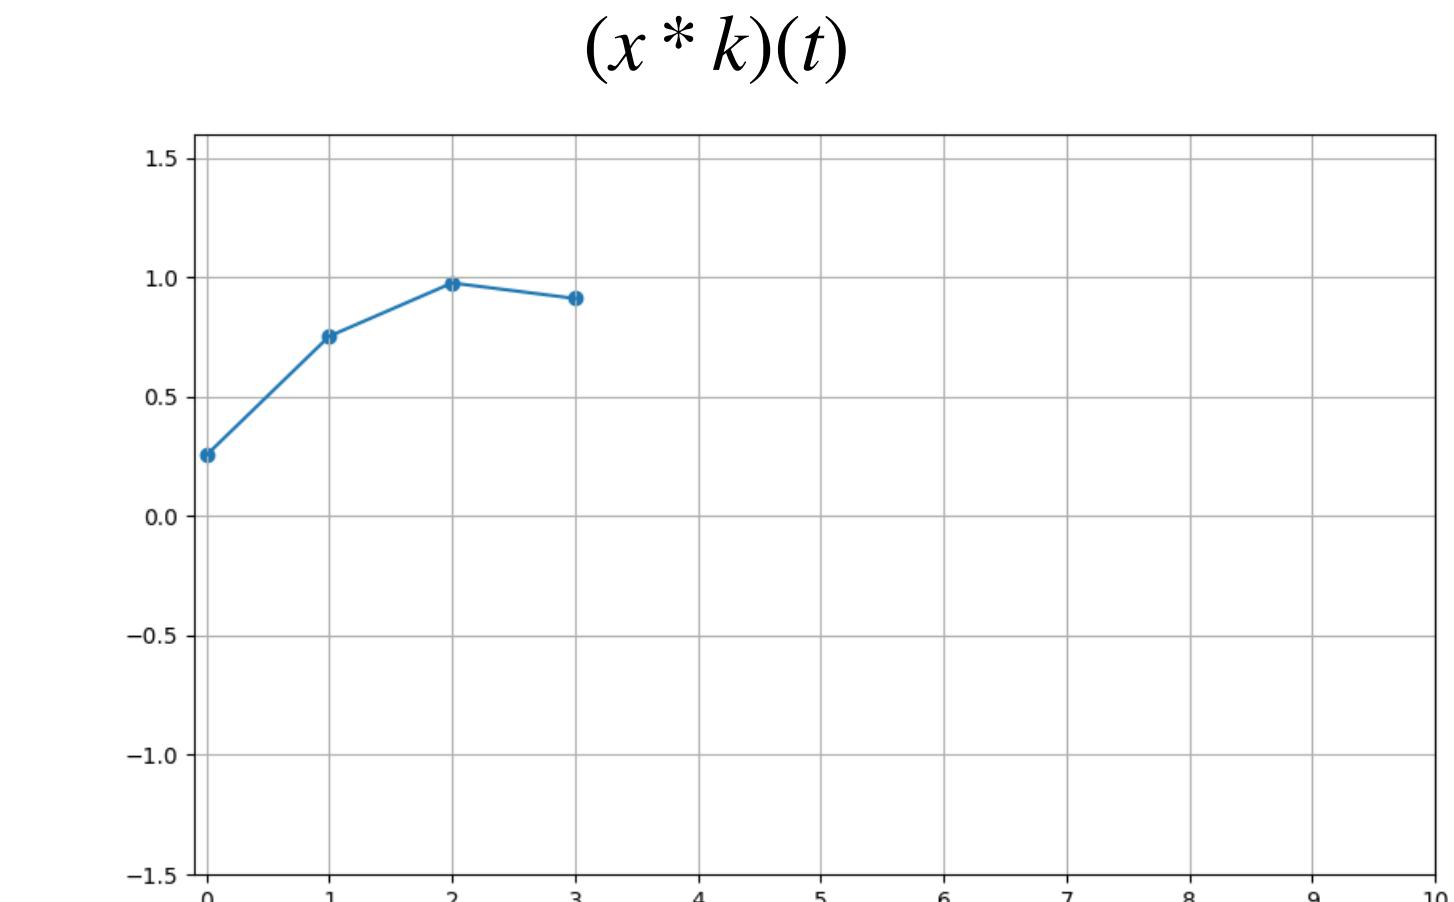
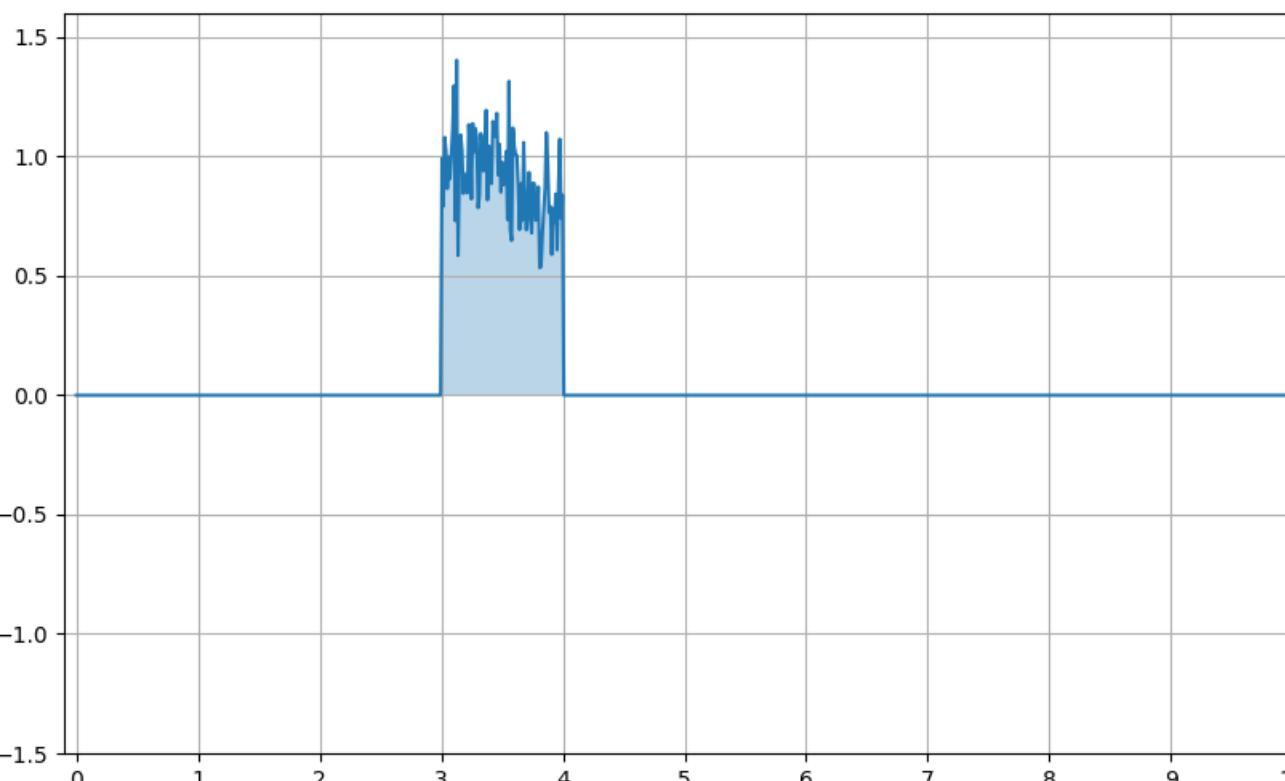
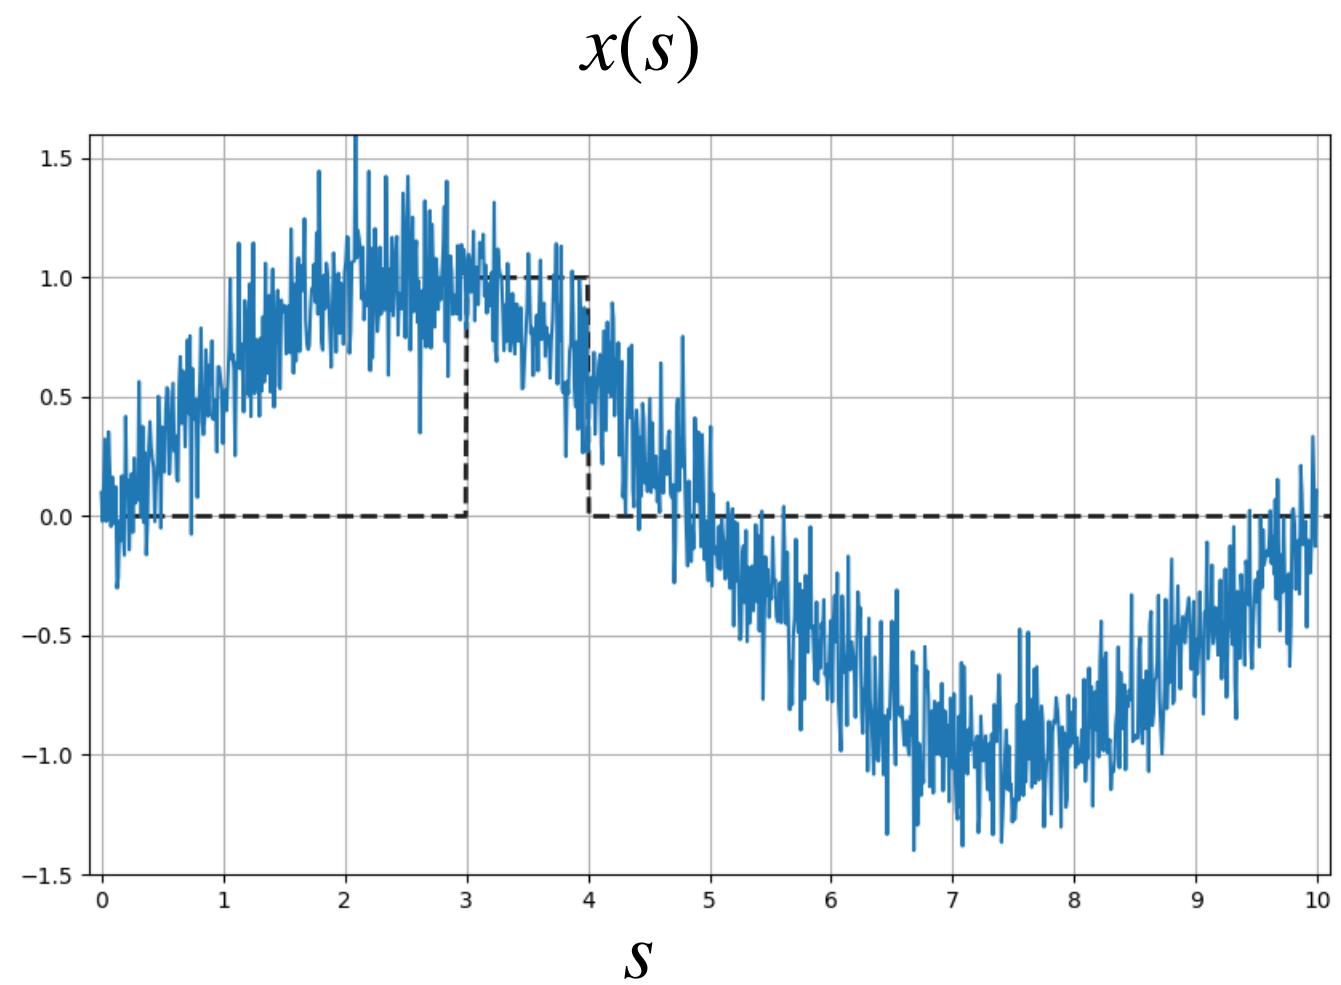
$$(x * k)(t) = \int x(s) k(t - s) ds$$



Convolutions

General definition of convolution between an input signal x and the kernel k

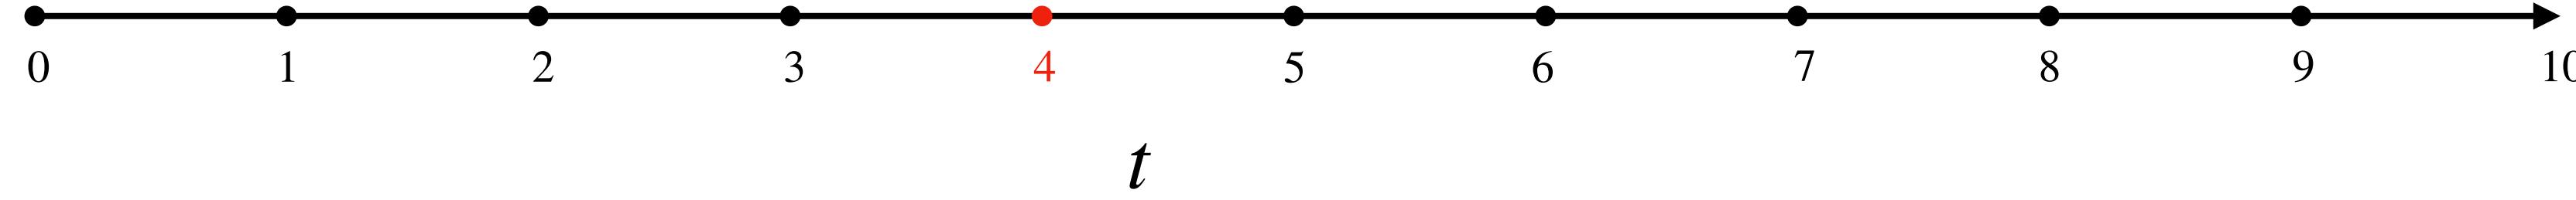
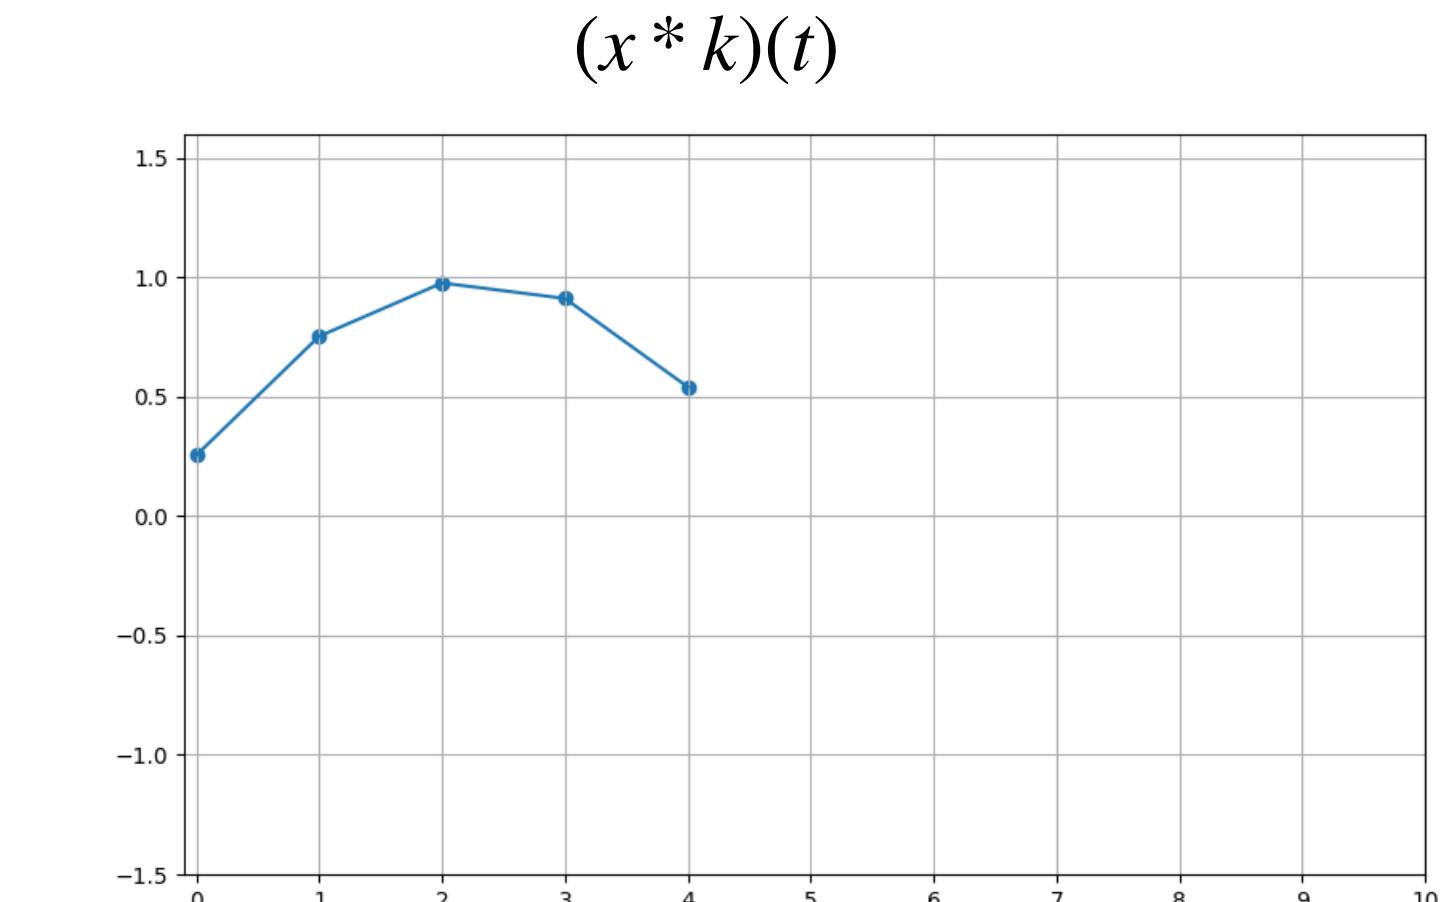
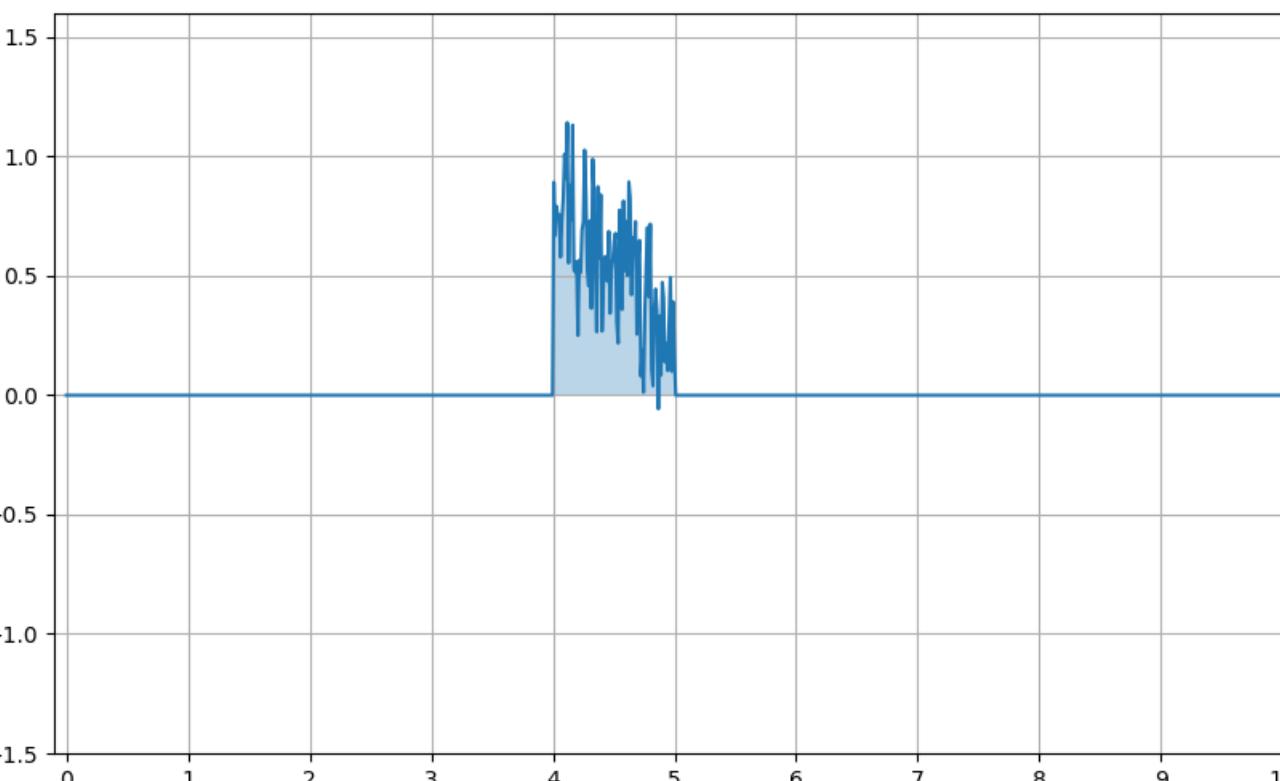
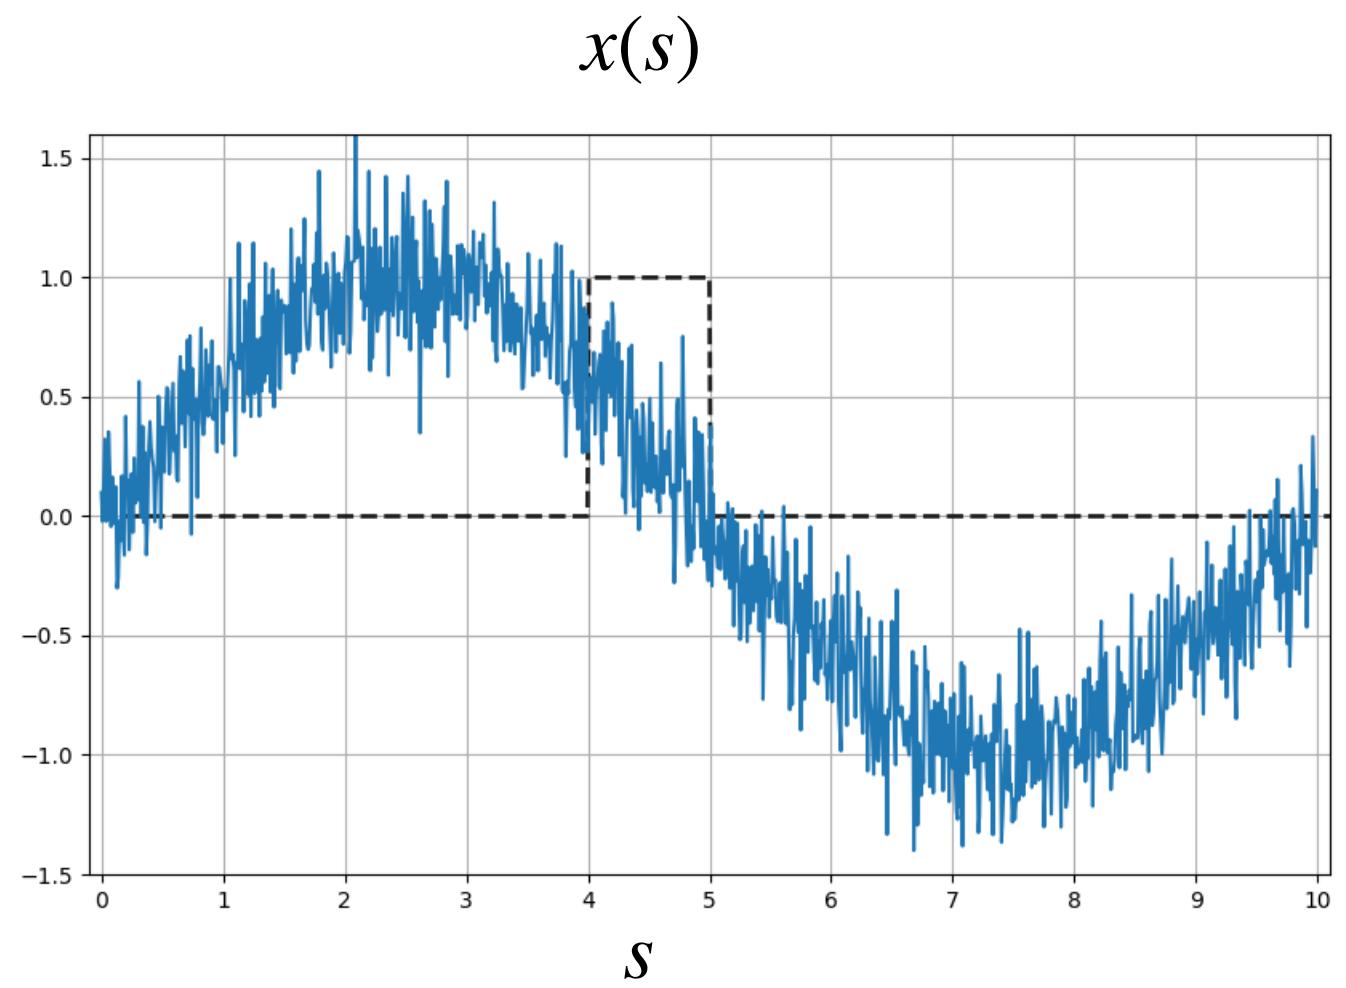
$$(x * k)(t) = \int x(s) k(t - s) ds$$



Convolutions

General definition of convolution between an input signal x and the kernel k

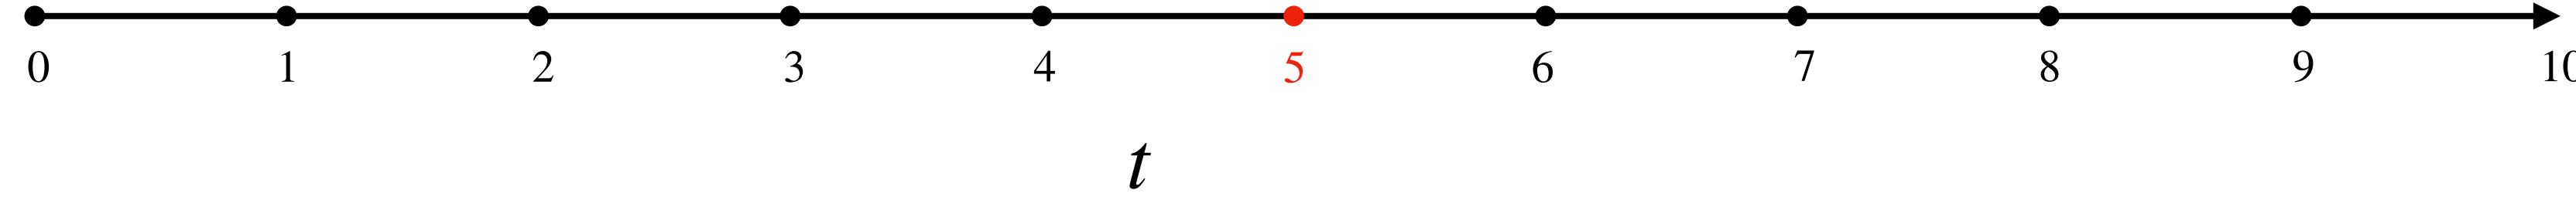
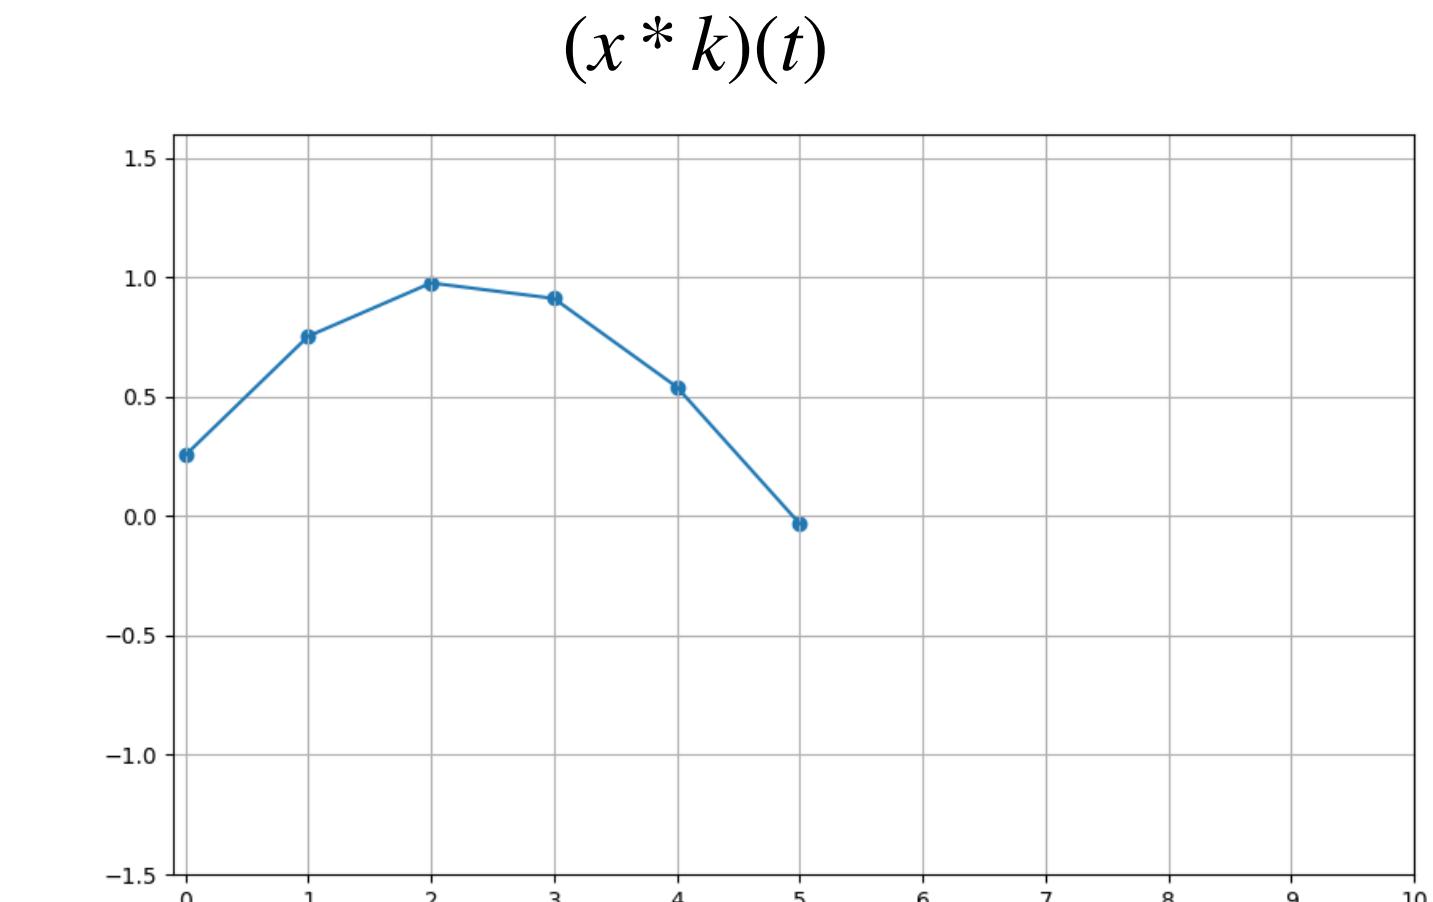
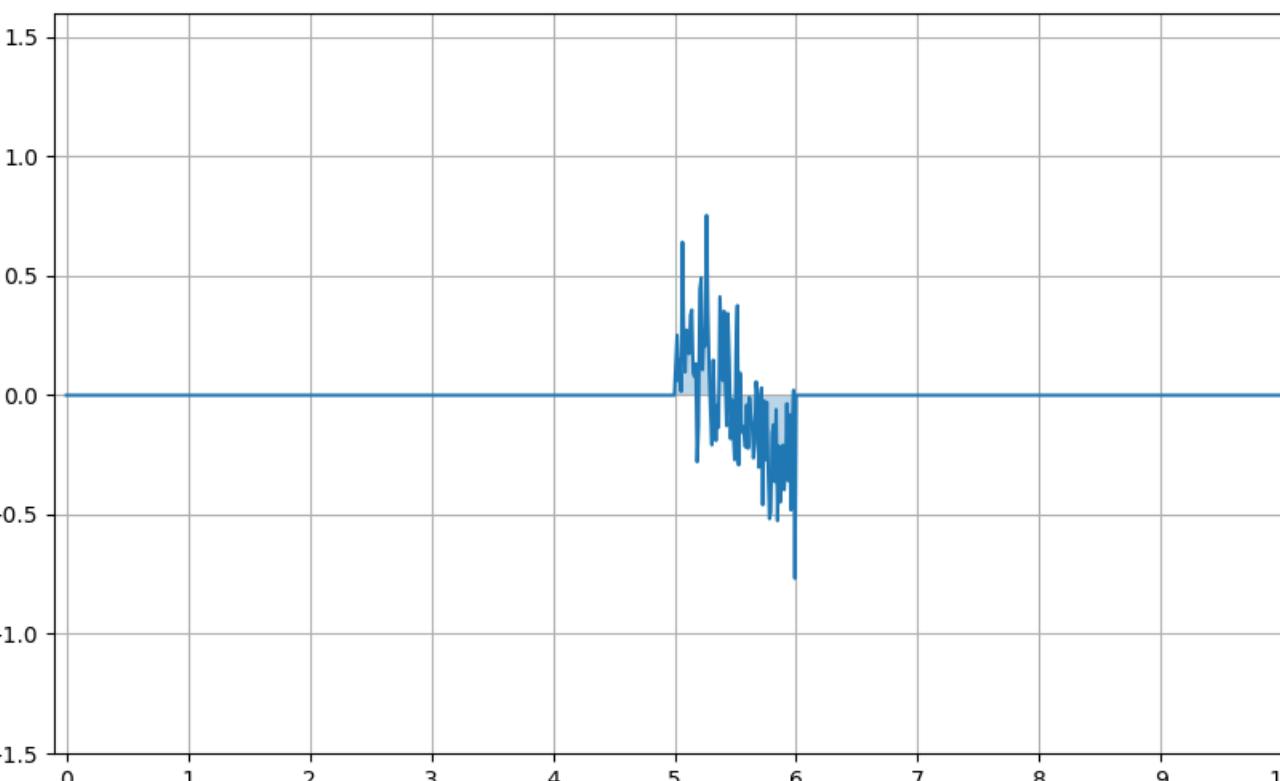
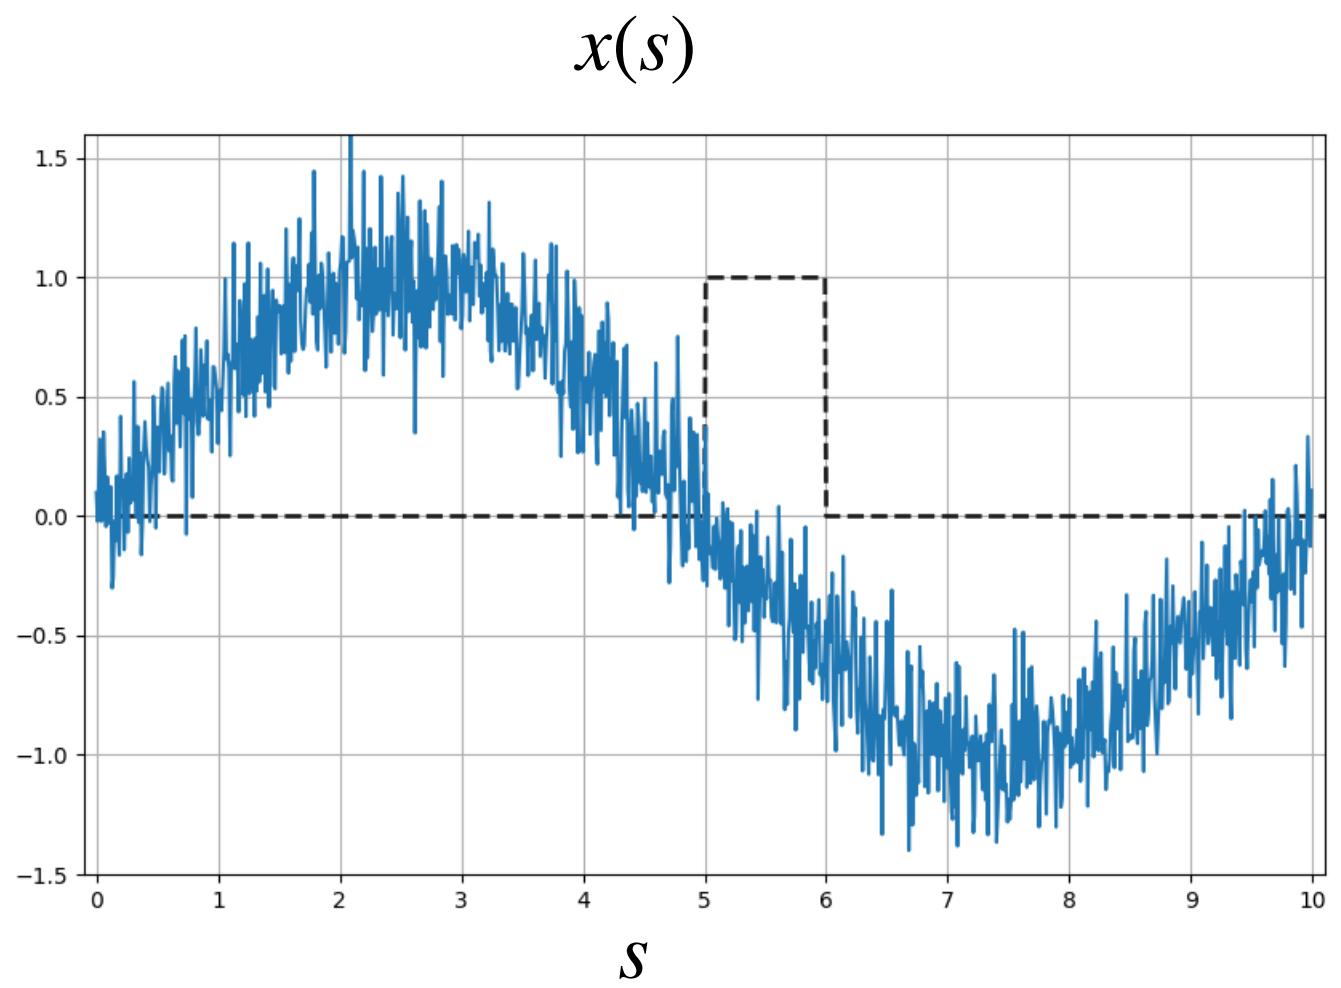
$$(x * k)(t) = \int x(s) k(t - s) ds$$



Convolutions

General definition of convolution between an input signal x and the kernel k

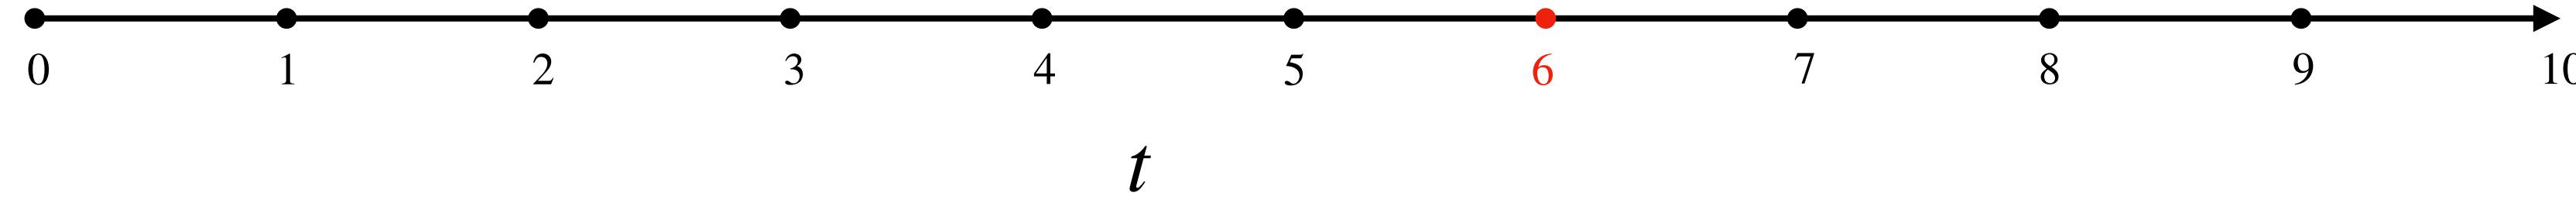
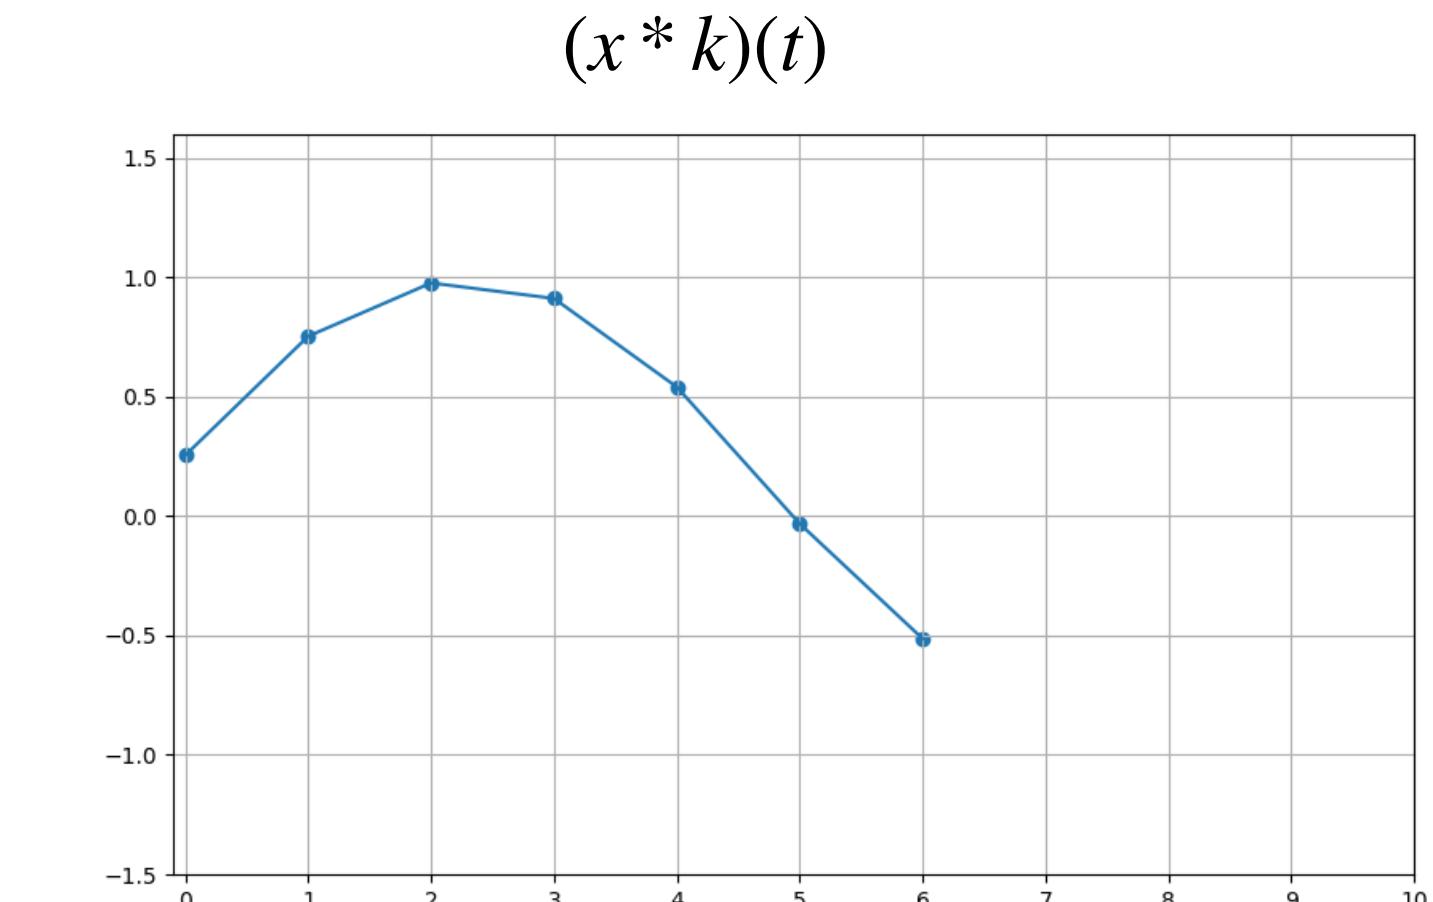
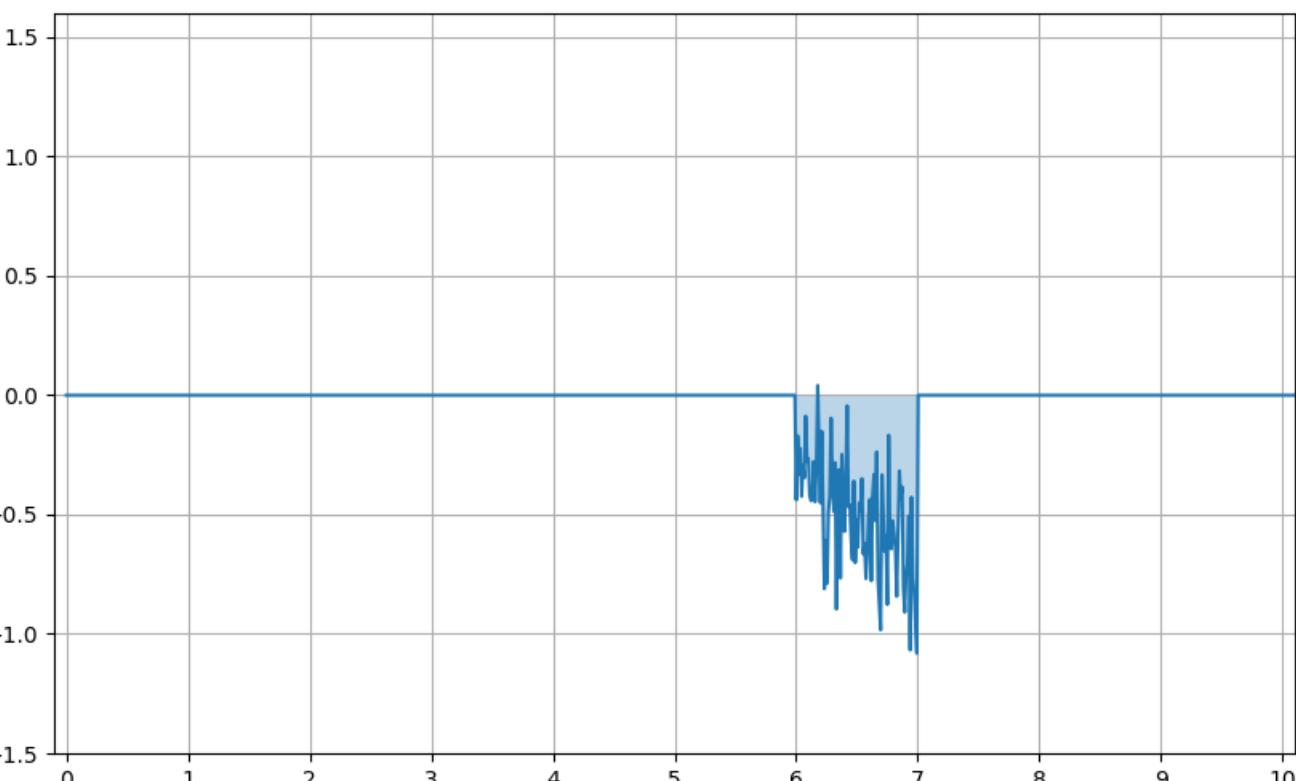
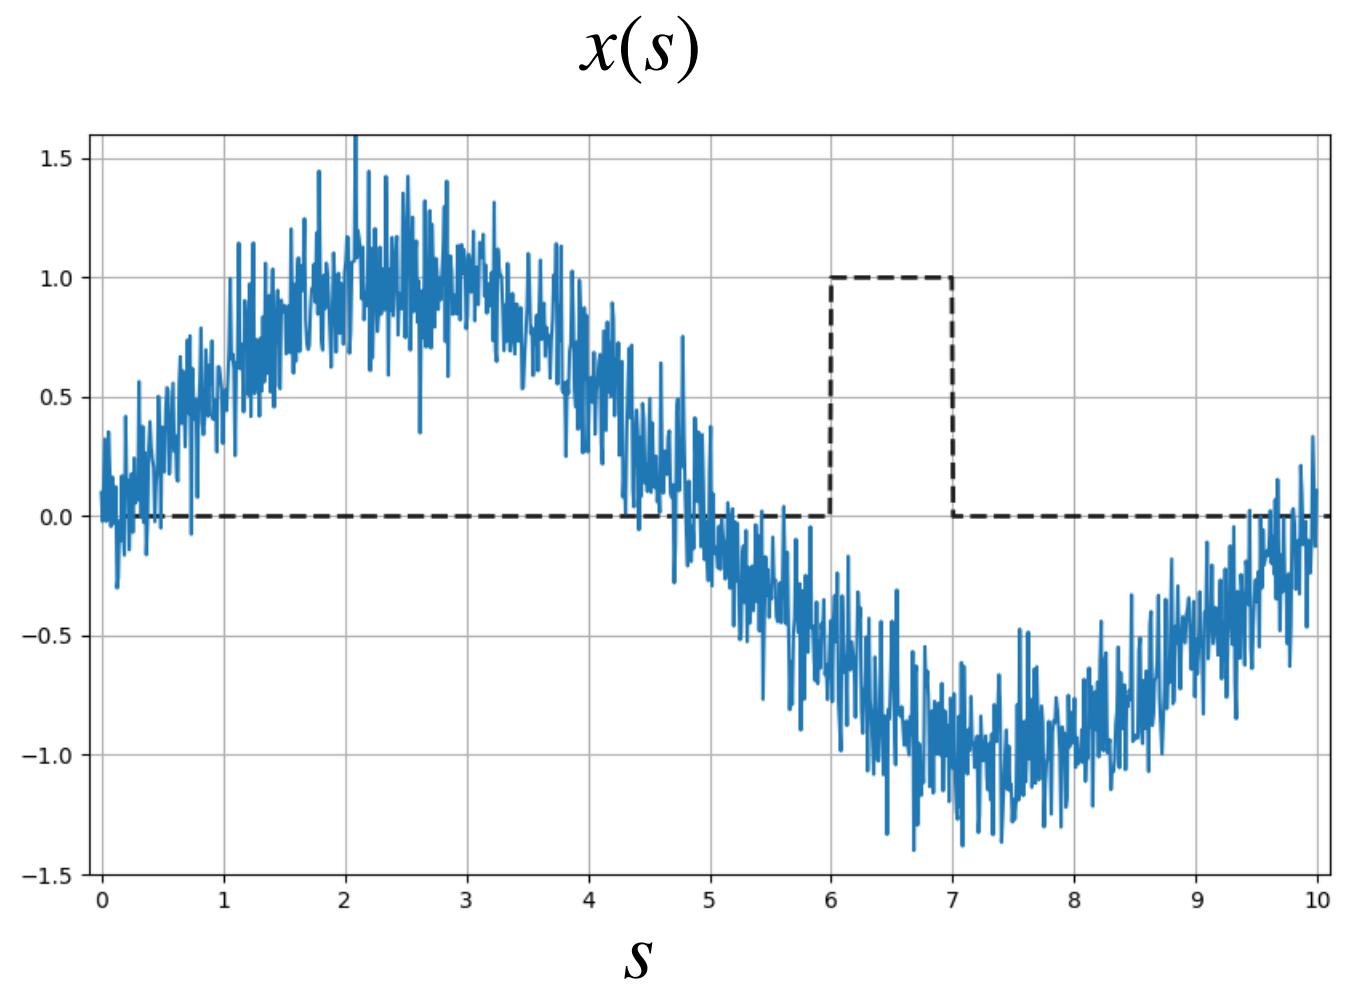
$$(x * k)(t) = \int x(s) k(t - s) ds$$



Convolutions

General definition of convolution between an input signal x and the kernel k

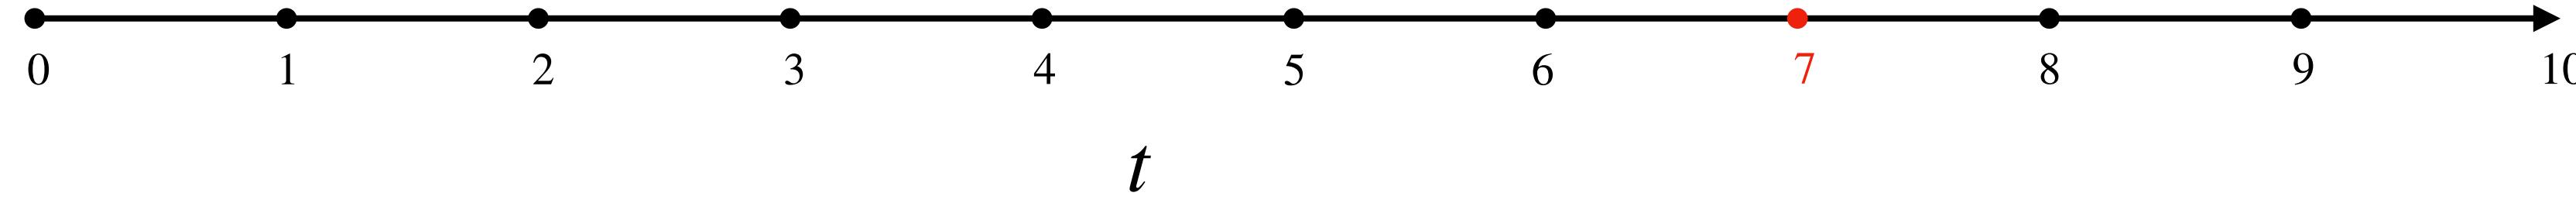
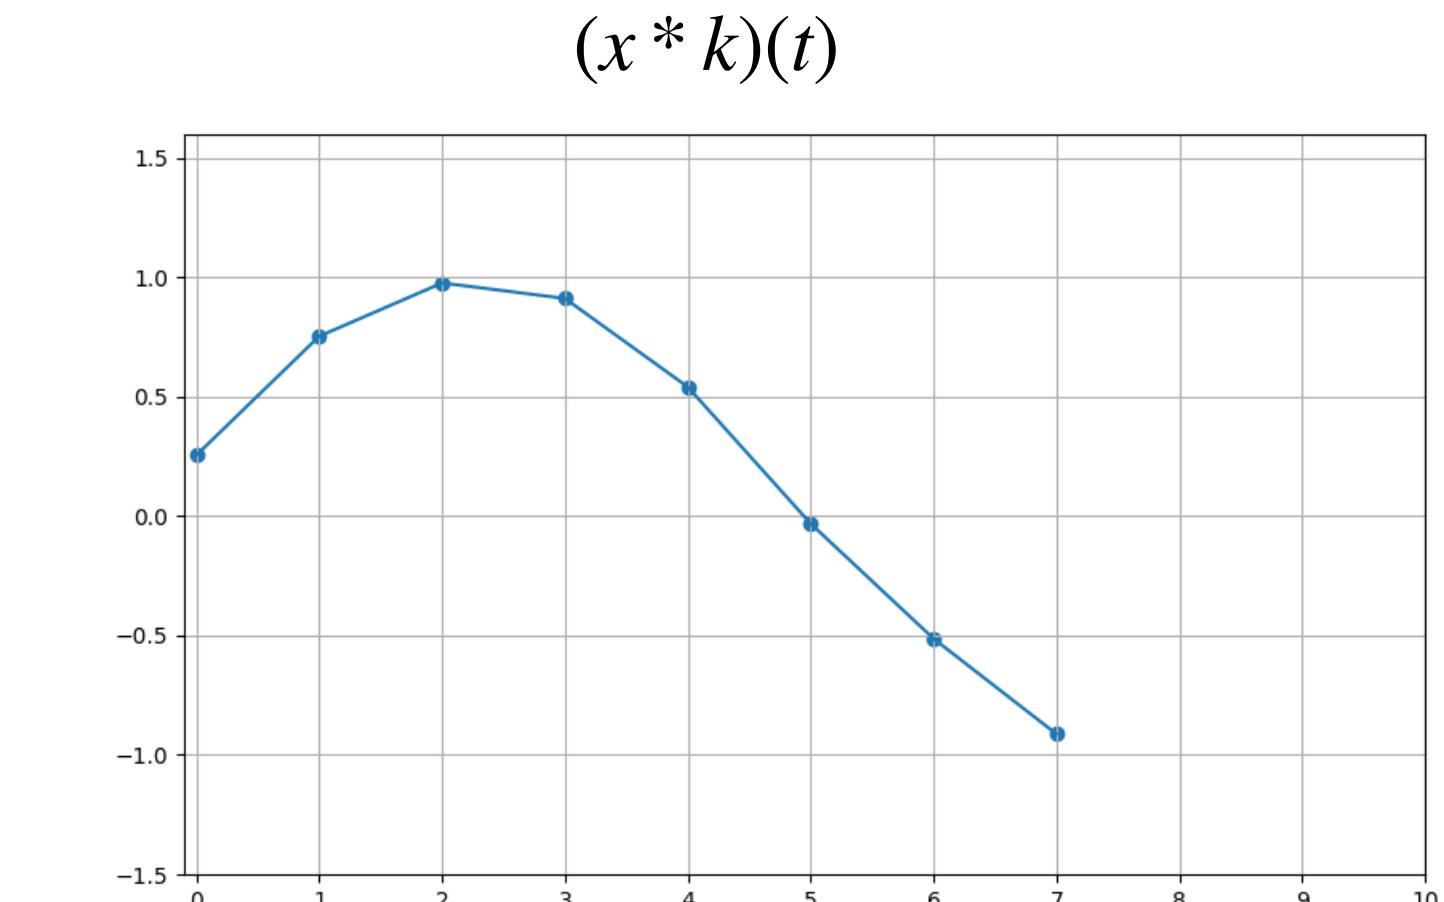
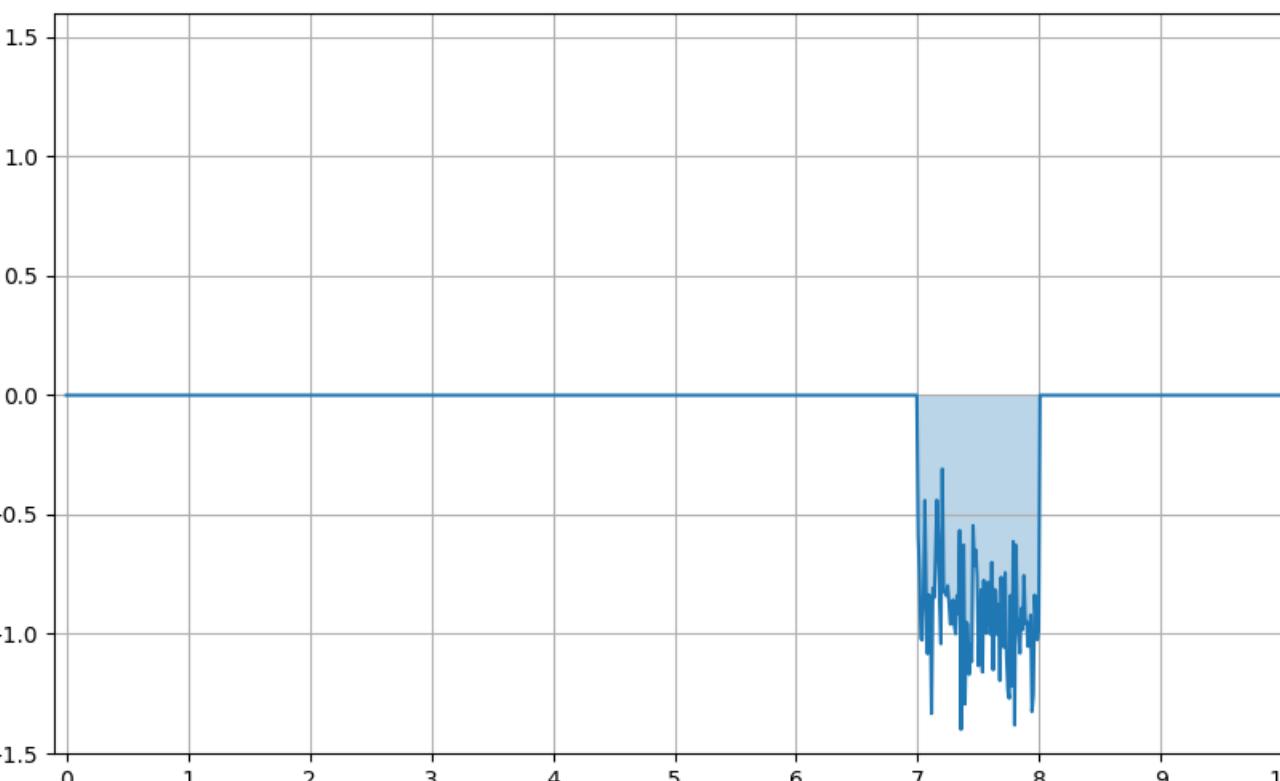
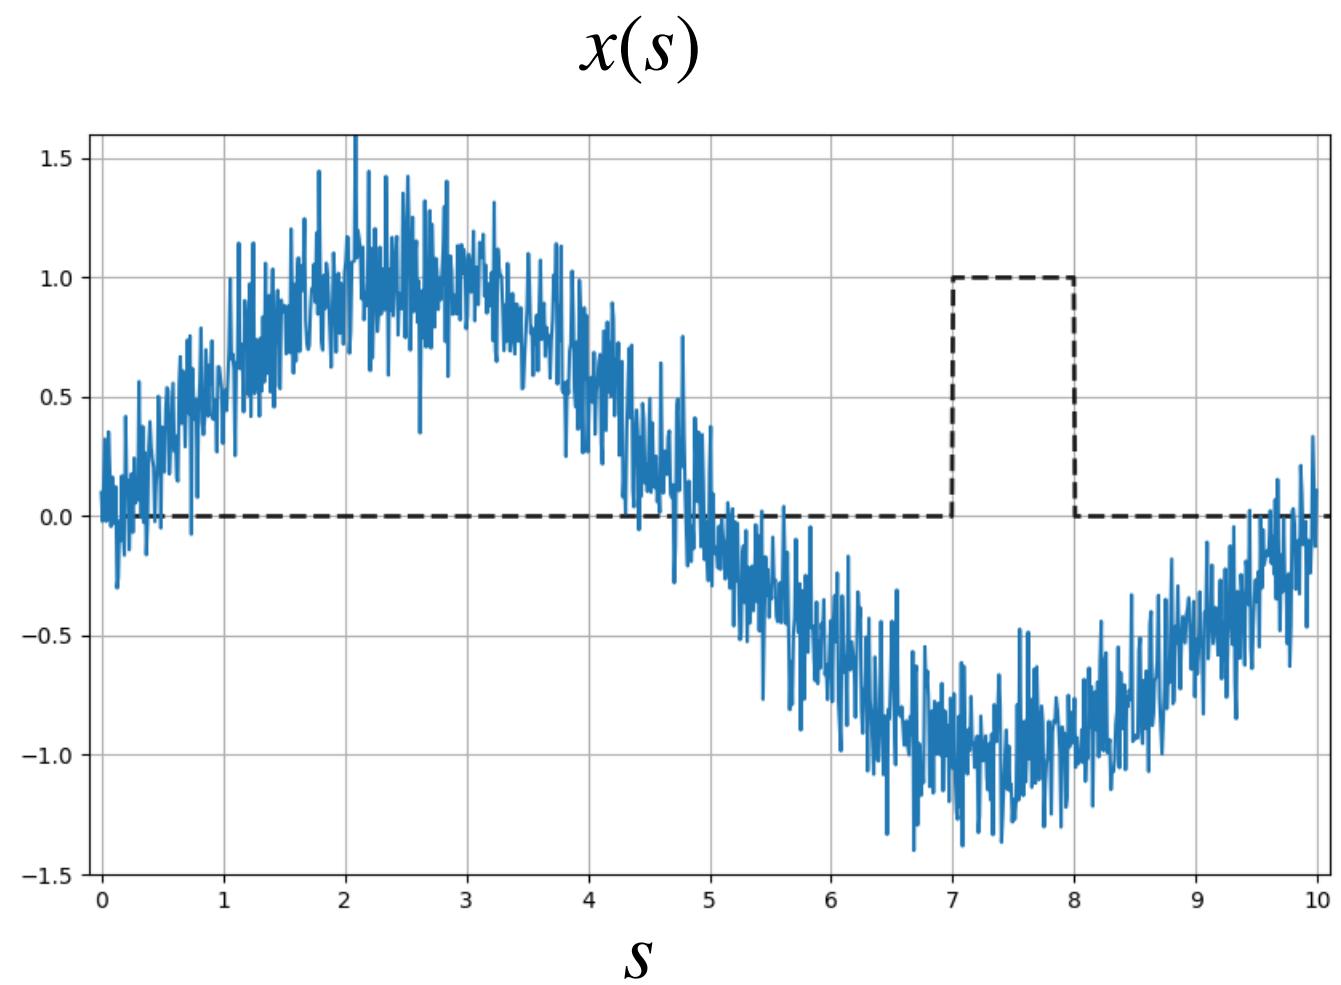
$$(x * k)(t) = \int x(s) k(t - s) ds$$



Convolutions

General definition of convolution between an input signal x and the kernel k

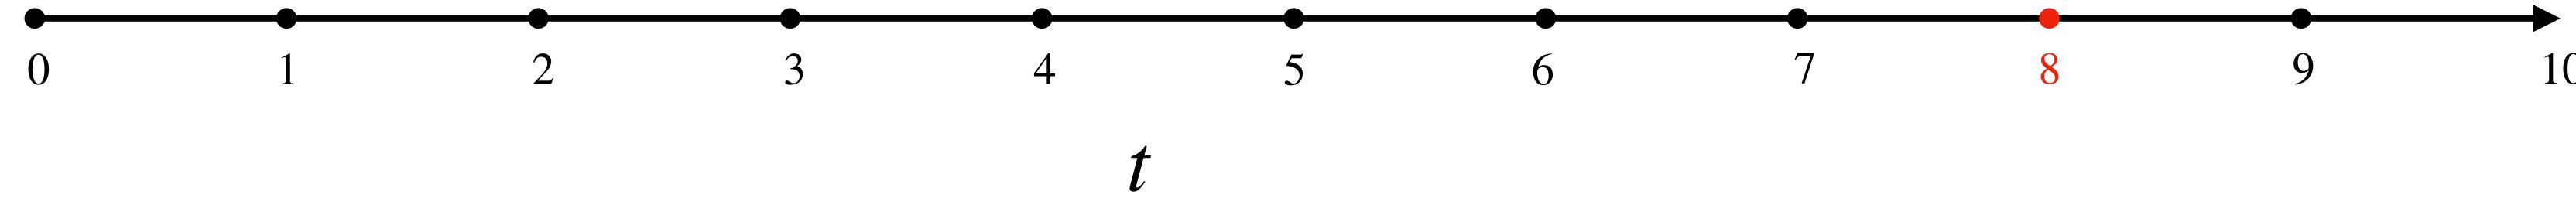
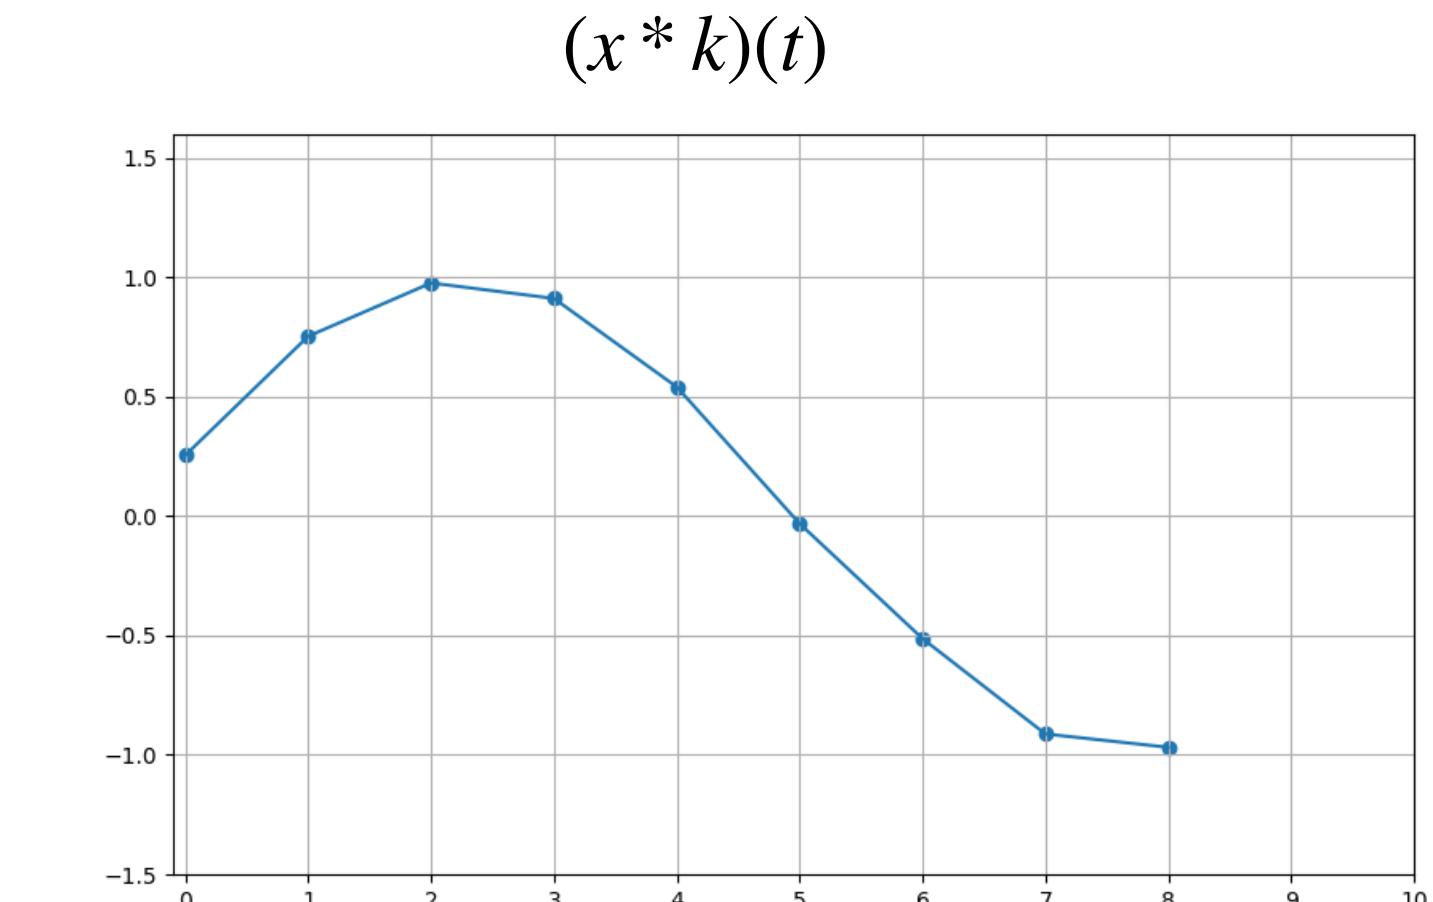
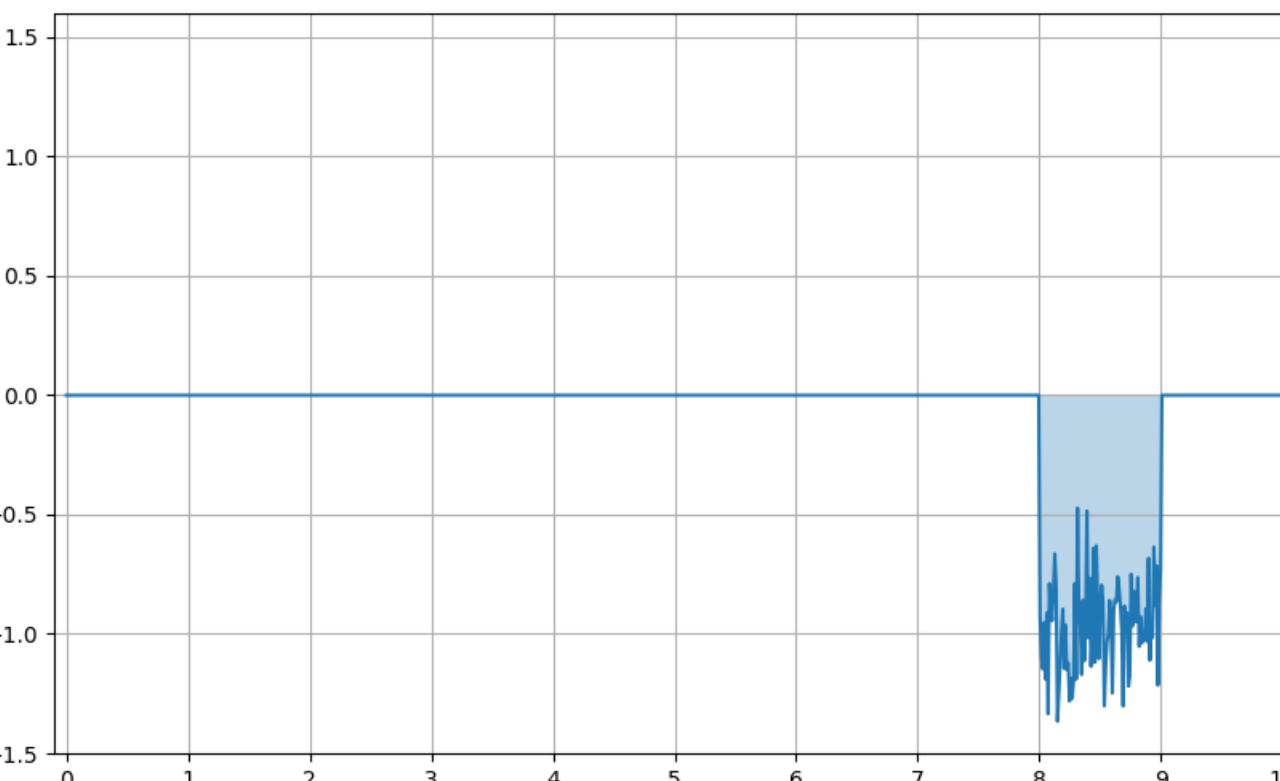
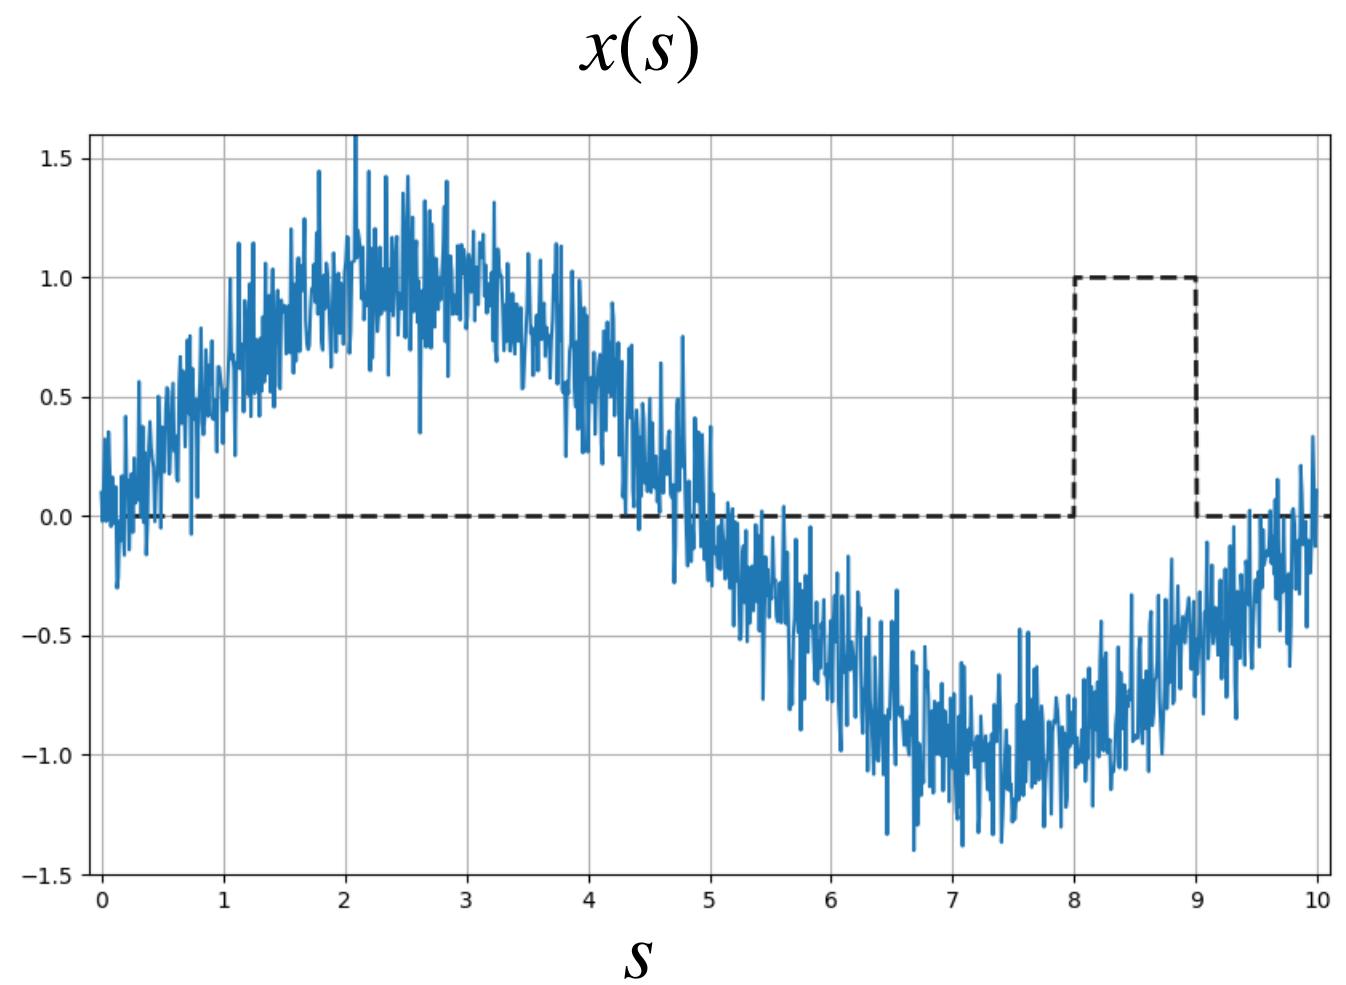
$$(x * k)(t) = \int x(s) k(t - s) ds$$



Convolutions

General definition of convolution between an input signal x and the kernel k

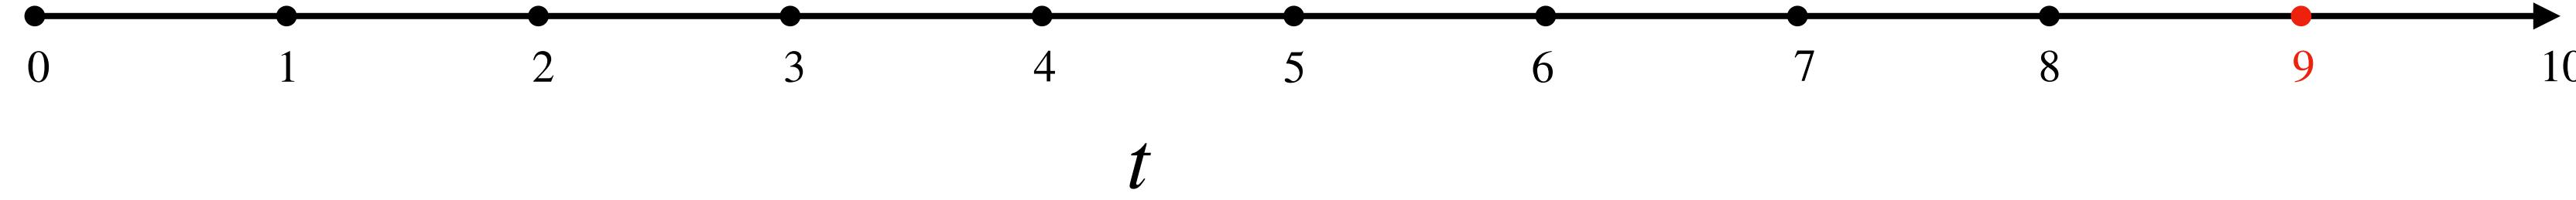
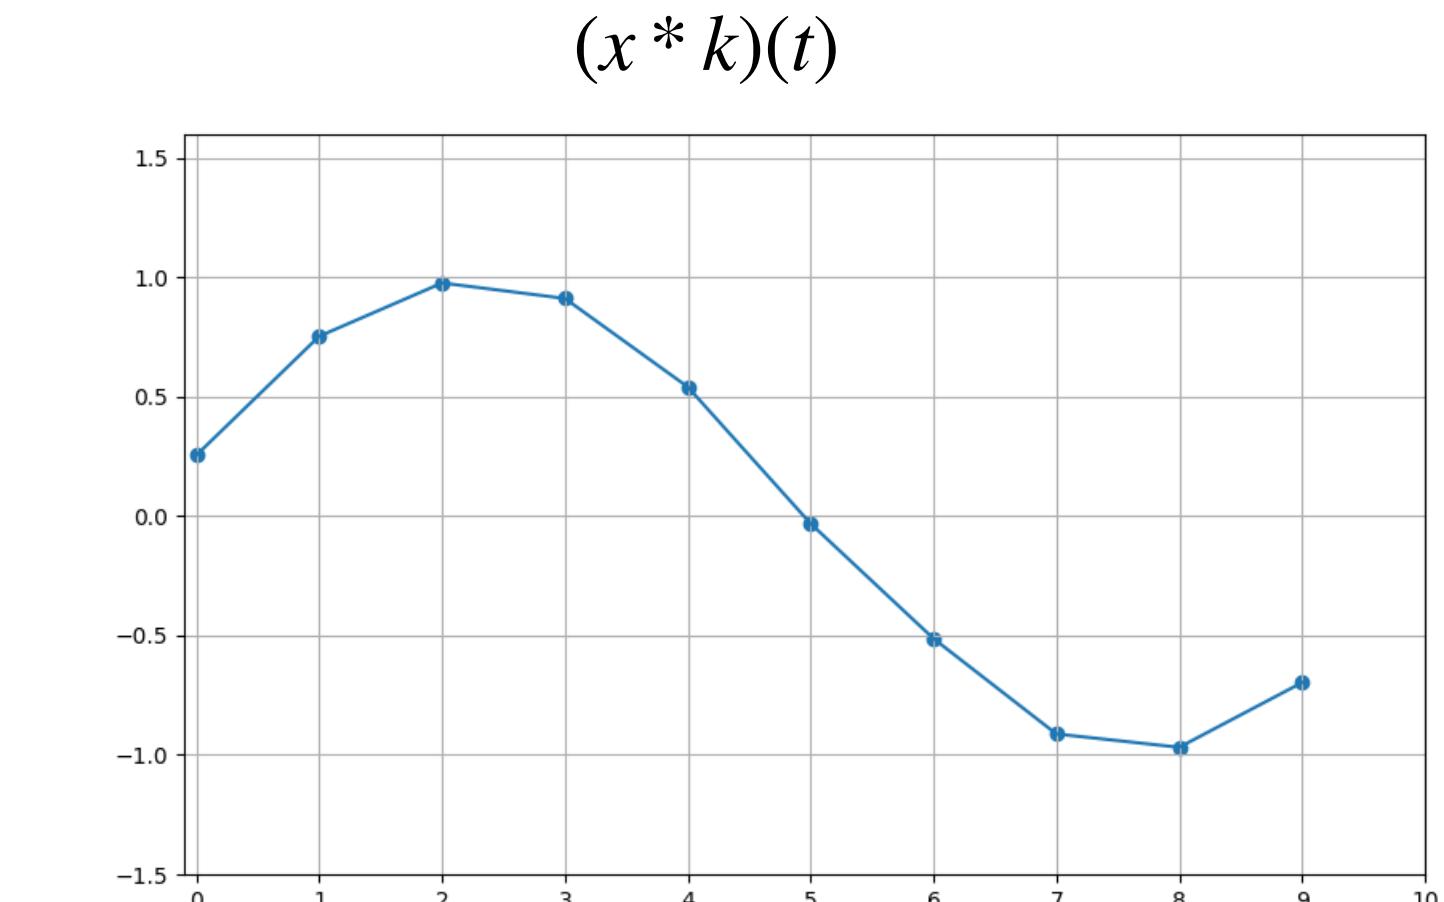
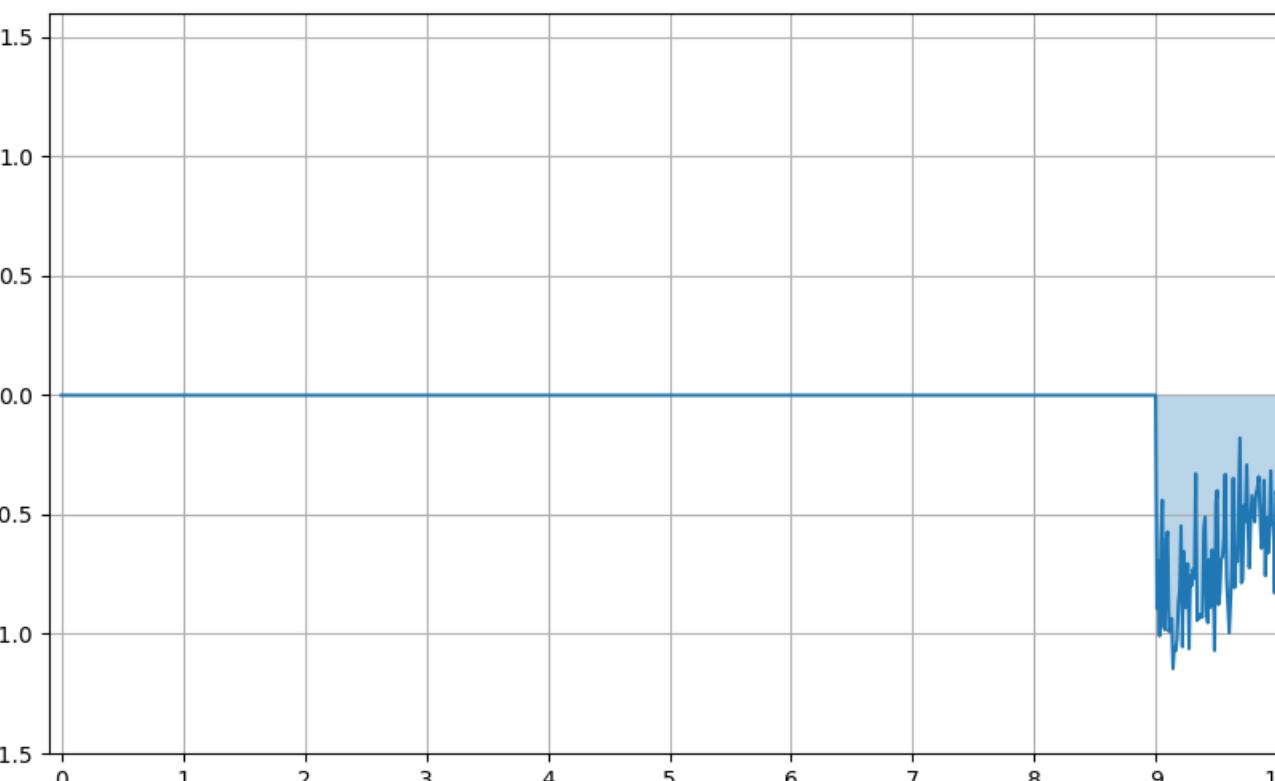
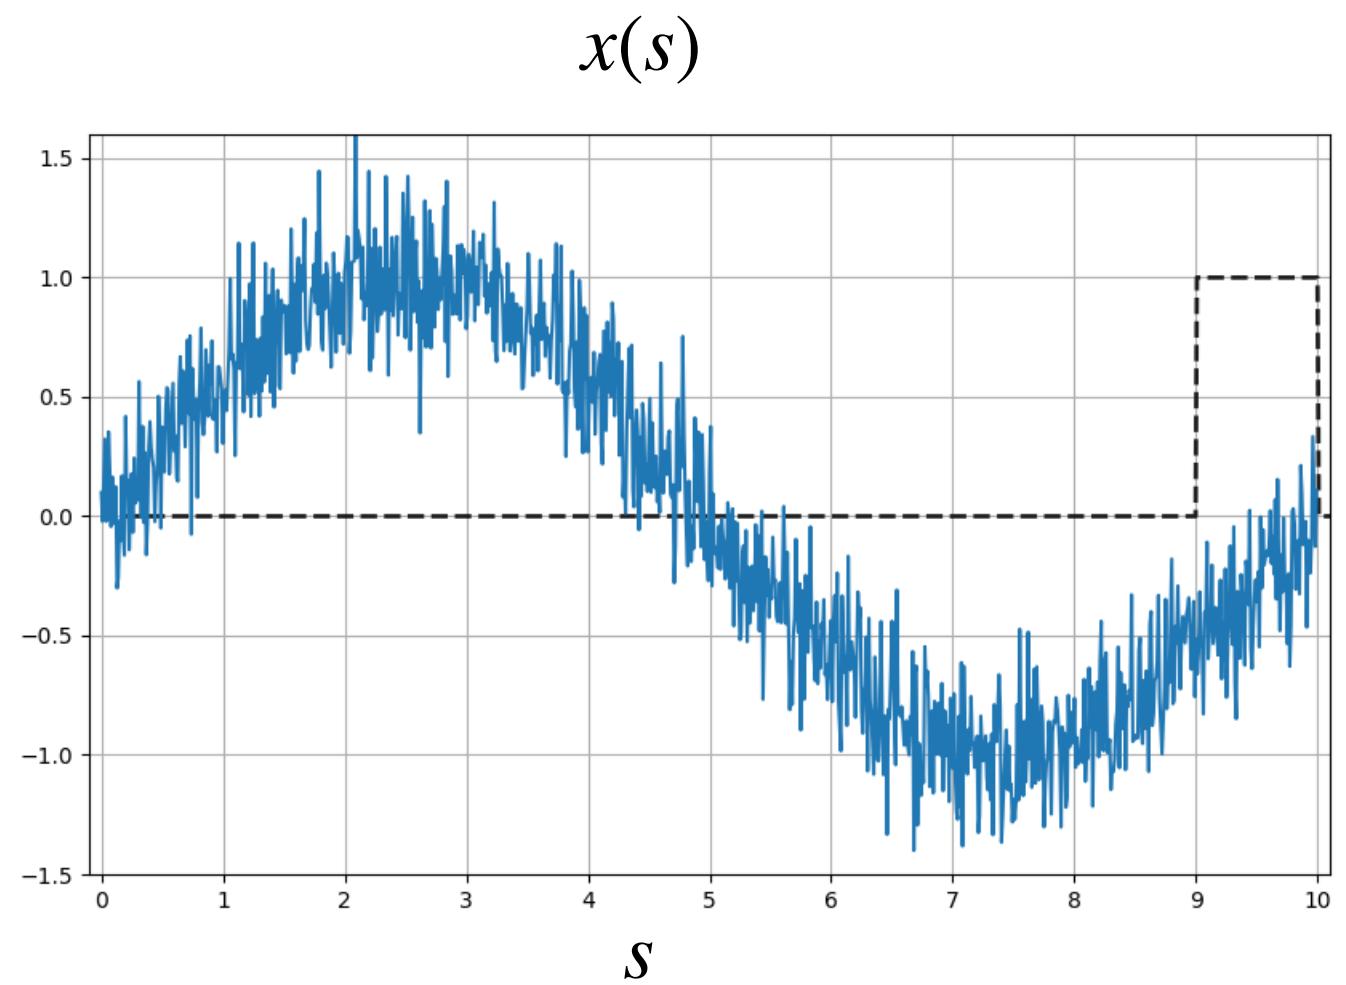
$$(x * k)(t) = \int x(s) k(t - s) ds$$



Convolutions

General definition of convolution between an input signal x and the kernel k

$$(x * k)(t) = \int x(s) k(t - s) ds$$



Convolutions

General definition of convolution between an input signal x and the kernel k

$$(x * k)(t) = \int x(s) k(t - s) ds$$

Discrete 1D convolution

$$(x * k)_i = \sum_n x_n k_{i-n}$$

Discrete input x

← $x = (x_1, \dots, x_N)$ is a 1D vector

Discrete 2D convolution

$$(x * k)_{i,j} = \sum_n \sum_m x_{n,m} k_{i-n, j-m}$$

← $x = \begin{pmatrix} x_{1,1} & \dots & x_{1,W} \\ x_{H,1} & \dots & x_{H,W} \end{pmatrix}$ is a 2D vector/matrix

Convolutions

General definition of convolution between an input signal x and the kernel k

$$(x * k)(t) = \int x(s) k(t - s) ds$$

Most ML libraries use *cross-correlation*

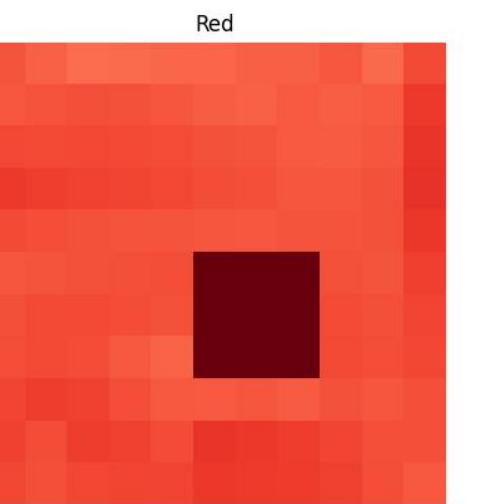
$$(x * k)_{i,j} = \sum_n \sum_m x_{i+n, j+m} k_{n,m}$$

Convolutions

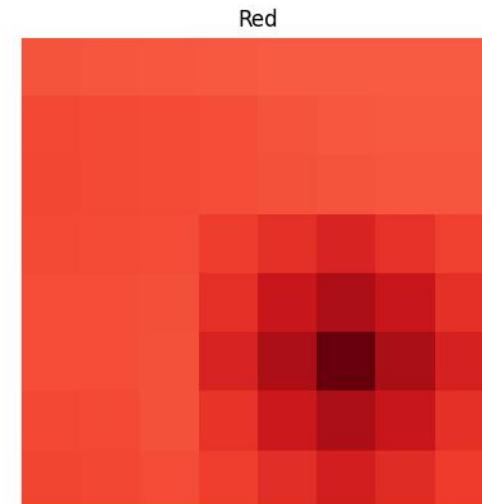
General definition of convolution between an input signal x and the kernel k

$$(x * k)(t) = \int x(s) k(t - s) ds$$

x



$x * k$



Most ML libraries use *cross-correlation*

$$(x * k)_{i,j} = \sum_n \sum_m x_{i+n, j+m} k_{n,m}$$

113	122	130	128	126	125	121	121	115	127
115	112	110	111	114	119	122	117	120	117
105	106	105	106	108	111	113	117	118	114
94	98	101	103	105	108	110	114	114	111
106	109	111	112	113	114	115	113	112	111
114	113	111	110	109	0	0	0	111	113
110	106	106	108	111	0	0	0	107	110
109	106	108	116	124	0	0	0	106	109
102	97	100	109	116	117	114	118	111	114
100	108	97	100	107	90	93	97	102	110

	113	114	115	116	118	118	118	118	
	105	106	107	109	112	115	116	116	
	104	106	107	109	111	113	114	114	
	106	108	108	97	86	75	88	100	
	110	110	110	86	62	38	62	86	
	109	109	111	75	38	0	36	73	
	105	106	111	89	65	39	62	86	
	103	105	109	98	85	70	82	96	

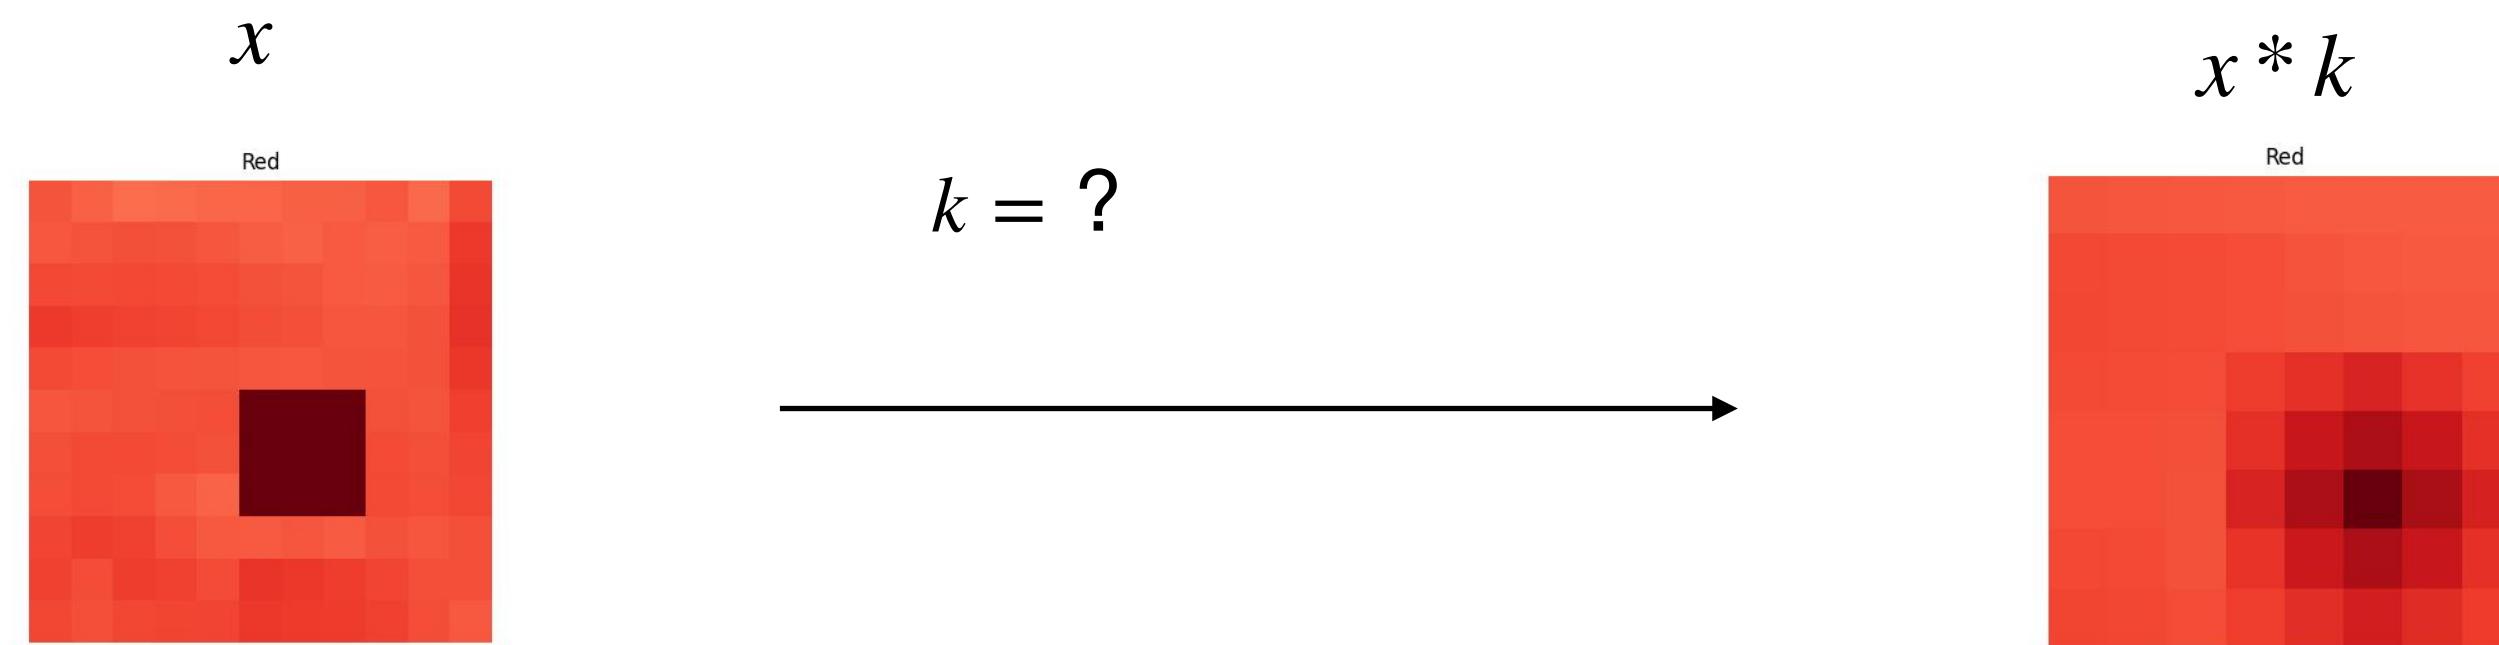
Convolutions

General definition of convolution between an input signal x and the kernel k

$$(x * k)(t) = \int x(s) k(t - s) ds$$

Most ML libraries use *cross-correlation*

$$(x * k)_{i,j} = \sum_n \sum_m x_{i+n, j+m} k_{n,m}$$



113	122	130	128	126	125	121	121	115	127
115	112	110	111	114	119	122	117	120	117
105	106	105	106	108	111	113	117	118	114
94	98	101	103	105	108	110	114	114	111
106	109	111	112	113	114	115	113	112	111
114	113	111	110	109	0	0	0	111	113
110	106	106	108	111	0	0	0	107	110
109	106	108	116	124	0	0	0	106	109
102	97	100	109	116	117	114	118	111	114
100	108	97	100	107	90	93	97	102	110

	113	114	115	116	118	118	118	118	
	105	106	107	109	112	115	116	116	
	104	106	107	109	111	113	114	114	
	106	108	108	97	86	75	88	100	
	110	110	110	86	62	38	62	86	
	109	109	111	75	38	0	36	73	
	105	106	111	89	65	39	62	86	
	103	105	109	98	85	70	82	96	

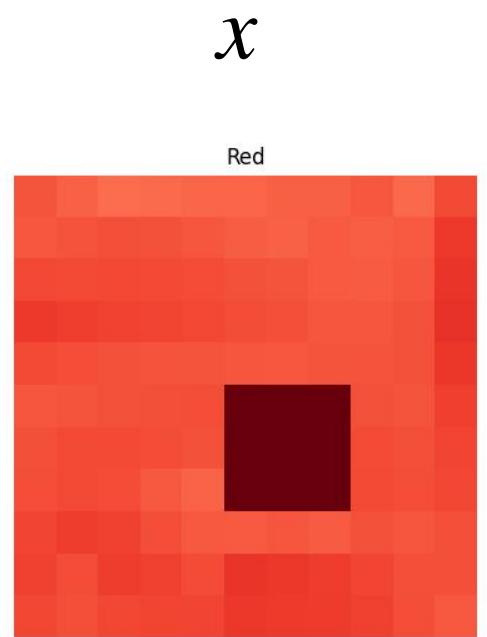
Convolutions

General definition of convolution between an input signal x and the kernel k

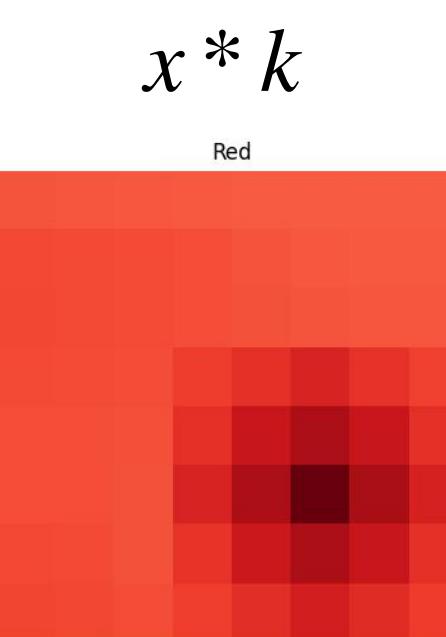
$$(x * k)(t) = \int x(s) k(t - s) ds$$

Most ML libraries use *cross-correlation*

$$(x * k)_{i,j} = \sum_n \sum_m x_{i+n, j+m} k_{n,m}$$



$$k = \begin{bmatrix} \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \end{bmatrix}^*$$



113	122	130	128	126	125	121	121	115	127
115	112	110	111	114	119	122	117	120	117
105	106	105	106	108	111	113	117	118	114
94	98	101	103	105	108	110	114	114	111
106	109	111	112	113	114	115	113	112	111
114	113	111	110	109	0	0	0	111	113
110	106	106	108	111	0	0	0	107	110
109	106	108	116	124	0	0	0	106	109
102	97	100	109	116	117	114	118	111	114
100	108	97	100	107	90	93	97	102	110

	113	114	115	116	118	118	118	118	
	105	106	107	109	112	115	116	116	
	104	106	107	109	111	113	114	114	
	106	108	108	97	86	75	88	100	
	110	110	110	86	62	38	62	86	
	109	109	111	75	38	0	36	73	
	105	106	111	89	65	39	62	86	
	103	105	109	98	85	70	82	96	

*More precisely $k_{i,j} = \begin{cases} \frac{1}{9} & \text{if } i \leq 2 \text{ and } j \leq 2 \\ 0 & \text{otherwise} \end{cases}$

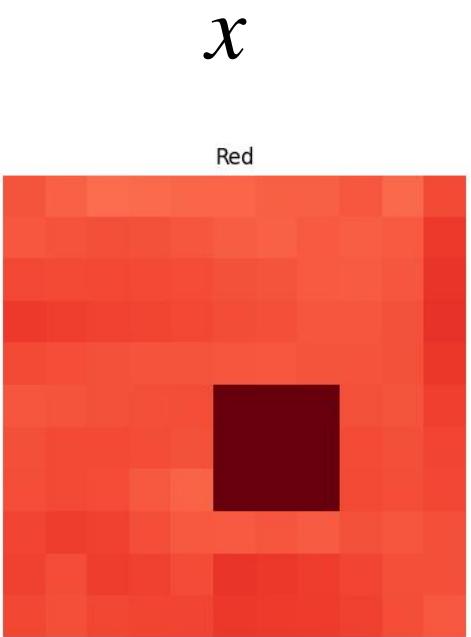
Convolutions

General definition of convolution between an input signal x and the kernel k

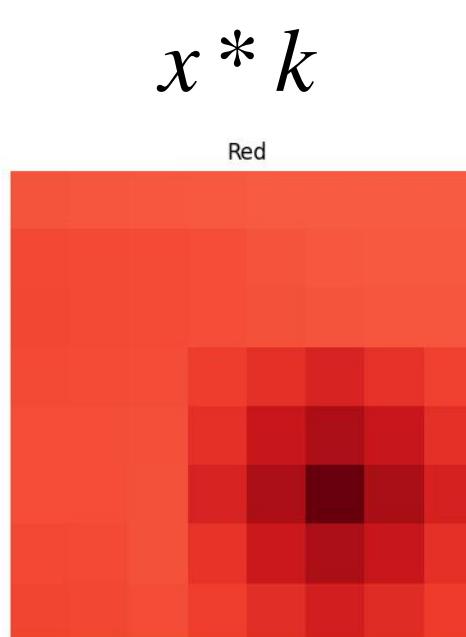
$$(x * k)(t) = \int x(s) k(t - s) ds$$

Most ML libraries use *cross-correlation*

$$(x * k)_{i,j} = \sum_n \sum_m x_{i+n, j+m} k_{n,m}$$



$$k = \begin{bmatrix} \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \end{bmatrix}^*$$



113	122	130	128	126	125	121	121	115	127
115	112	110	111	114	119	122	117	120	117
105	106	105	106	108	111	113	117	118	114
94	98	101	103	105	108	110	114	$\frac{1}{9}$	$\frac{1}{9}$
106	109	111	112	113	114	110	113	$\frac{1}{9}$	$\frac{1}{9}$
114	113	111	110	109	0	0	0	$\frac{1}{9}$	$\frac{1}{9}$
110	106	106	108	111	0	0	0	107	110
109	106	108	116	124	0	0	0	106	109
102	97	100	109	116	117	114	118	111	114
100	108	97	100	107	90	93	97	102	110

	113	114	115	116	118	118	118	118	
	105	106	107	109	112	115	116	116	
	104	106	107	109	111	113	114	114	
	106	108	108	97	86	75	88	100	
	110	110	110	86	62	38	62	86	
	109	109	111	75	38	0	36	73	
	105	106	111	89	65	39	62	86	
	103	105	109	98	85	70	82	96	

*More precisely $k_{i,j} = \begin{cases} \frac{1}{9} & \text{if } i \leq 2 \text{ and } j \leq 2 \\ 0 & \text{otherwise} \end{cases}$

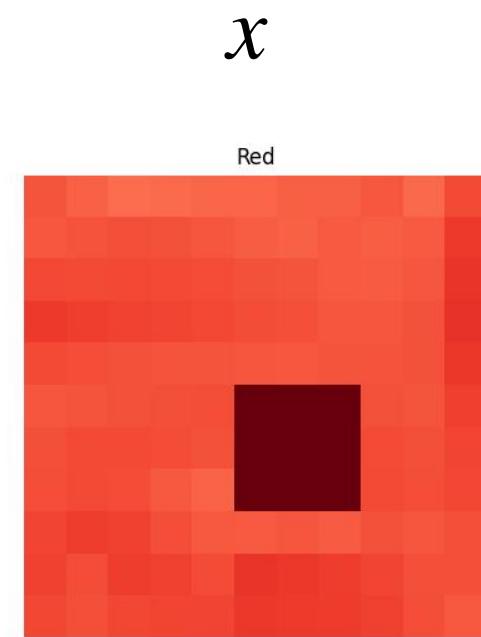
Convolutions

General definition of convolution between an input signal x and the kernel k

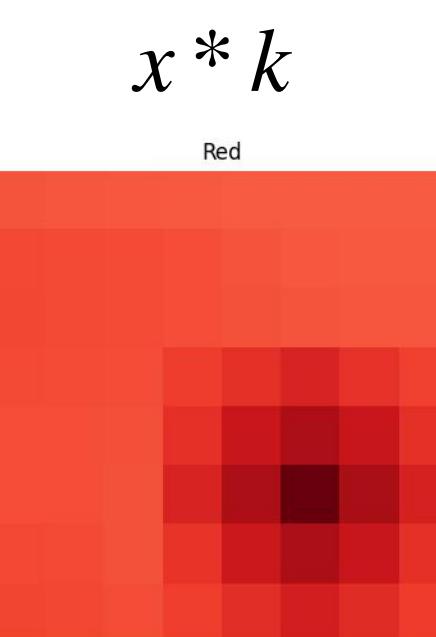
$$(x * k)(t) = \int x(s) k(t - s) ds$$

Most ML libraries use *cross-correlation*

$$(x * k)_{i,j} = \sum_n \sum_m x_{i+n, j+m} k_{n,m}$$



$$k = \begin{bmatrix} \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \end{bmatrix}^*$$



113	122	130	128	126	125	121	121	115	127
115	112	110	111	114	119	122	117	120	117
105	106	105	106	108	111	113	117	118	114
94	98	101	103	105	108	110	114	$\frac{1}{9}$	$\frac{1}{9}$
106	109	111	112	113	114	113	$\frac{1}{9}$	$\frac{1}{9}$	$\frac{1}{9}$
114	113	111	110	109	0	0	0	$\frac{1}{9}$	$\frac{1}{9}$
110	106	106	108	111	0	0	0	$\frac{1}{9}$	$\frac{1}{9}$
109	106	108	116	124	0	0	0	106	109
102	97	100	109	116	117	114	118	111	114
100	108	97	100	107	90	93	97	102	110

=

$\frac{103}{9}$	$\frac{105}{9}$	$\frac{108}{9}$
$\frac{112}{9}$	$\frac{113}{9}$	$\frac{114}{9}$
$\frac{110}{9}$	$\frac{109}{9}$	$\frac{0}{9}$

*More precisely $k_{i,j} = \begin{cases} \frac{1}{9} & \text{if } i \leq 2 \text{ and } j \leq 2 \\ 0 & \text{otherwise} \end{cases}$

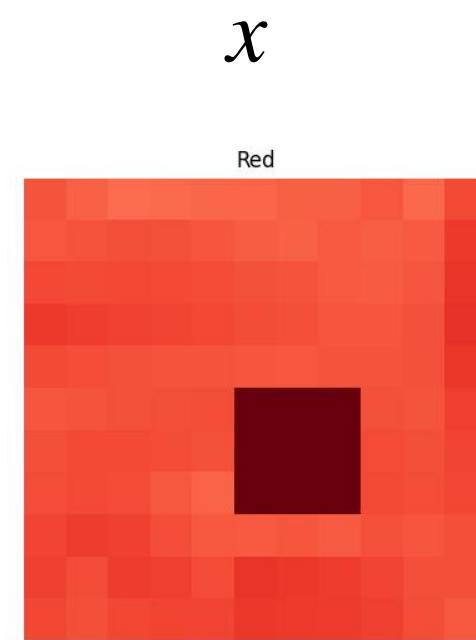
Convolutions

General definition of convolution between an input signal x and the kernel k

$$(x * k)(t) = \int x(s) k(t - s) ds$$

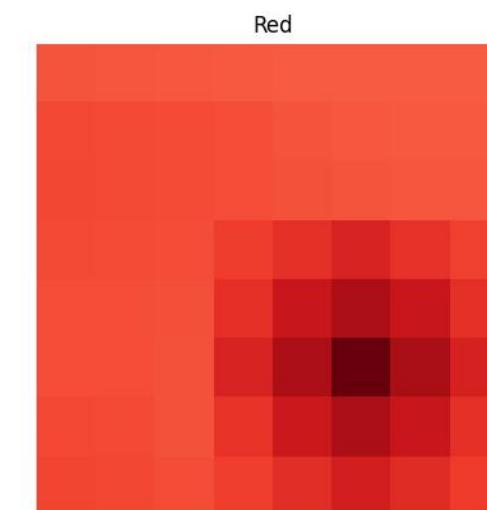
Most ML libraries use *cross-correlation*

$$(x * k)_{i,j} = \sum_n \sum_m x_{i+n, j+m} k_{n,m}$$



$$k = \begin{bmatrix} \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \end{bmatrix}^*$$

$x * k$



113	122	130	128	126	125	121	121	115	127
115	112	110	111	114	119	122	117	120	117
105	106	105	106	108	111	113	117	118	114
94	98	101	103	105	108	110	114	$\frac{1}{9}$	$\frac{1}{9}$
106	109	111	112	113	114	113	$\frac{1}{9}$	$\frac{1}{9}$	$\frac{1}{9}$
114	113	111	110	109	0	0	0	$\frac{1}{9}$	$\frac{1}{9}$
110	106	106	108	111	0	0	0	$\frac{1}{9}$	$\frac{1}{9}$
109	106	108	116	124	0	0	0	106	109
102	97	100	109	116	117	114	118	111	114
100	108	97	100	107	90	93	97	102	110

=

$\frac{103}{9}$	$\frac{105}{9}$	$\frac{108}{9}$
$\frac{112}{9}$	$\frac{113}{9}$	$\frac{114}{9}$
$\frac{110}{9}$	$\frac{109}{9}$	$\frac{0}{9}$

+

	113	114	115	116	118	118	118	118	
	105	106	107	109	112	115	116	116	
	104	106	107	109	111	113	114	114	
	106	108	108	97	86	75	88	100	
	110	110	110	86	62	38	62	86	
	109	109	111	75	38	0	36	73	
	105	106	111	89	65	39	62	86	
	103	105	109	98	85	70	82	96	

*More precisely $k_{i,j} = \begin{cases} \frac{1}{9} & \text{if } i \leq 2 \text{ and } j \leq 2 \\ 0 & \text{otherwise} \end{cases}$

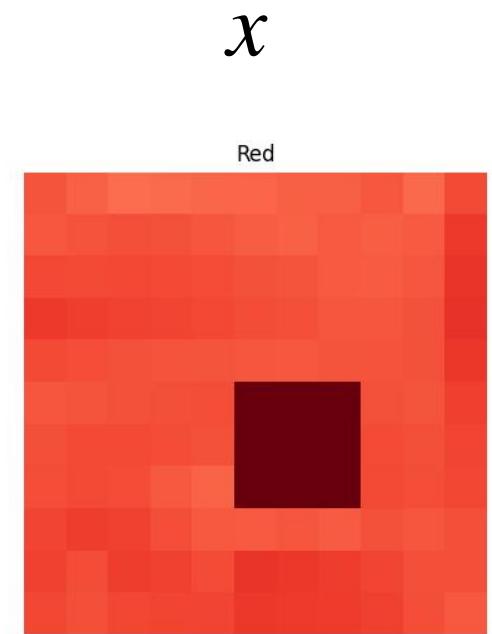
Convolutions

General definition of convolution between an input signal x and the kernel k

$$(x * k)(t) = \int x(s) k(t - s) ds$$

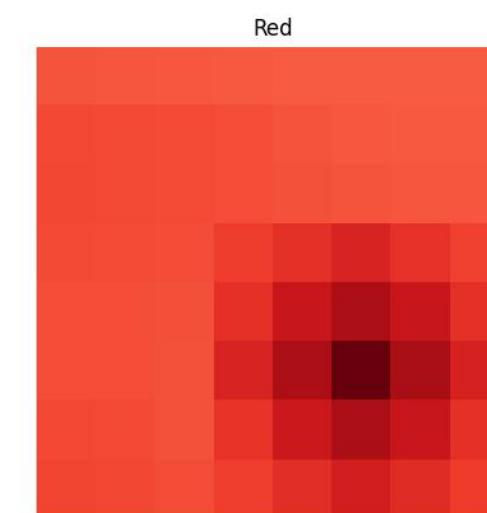
Most ML libraries use *cross-correlation*

$$(x * k)_{i,j} = \sum_n \sum_m x_{i+n, j+m} k_{n,m}$$



$$k = \begin{bmatrix} \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \end{bmatrix}^*$$

$x * k$



113	122	130	128	126	125	121	121	115	127
115	112	110	111	114	119	122	117	120	117
105	106	105	106	108	111	113	117	118	114
94	98	101	103	105	108	110	114	114	$\frac{1}{9}$
106	109	111	112	113	114	115	114	112	$\frac{1}{9}$
114	113	111	110	109	0	0	0	111	$\frac{1}{9}$
110	106	106	108	111	0	0	0	107	110
109	106	108	116	124	0	0	0	106	109
102	97	100	109	116	117	114	118	111	114
100	108	97	100	107	90	93	97	102	110

$$= \begin{array}{|c|c|c|} \hline & & \\ \hline & & \\ \hline & & \\ \hline \end{array}$$

+

*More precisely $k_{i,j} = \begin{cases} \frac{1}{9} & \text{if } i \leq 2 \text{ and } j \leq 2 \\ 0 & \text{otherwise} \end{cases}$

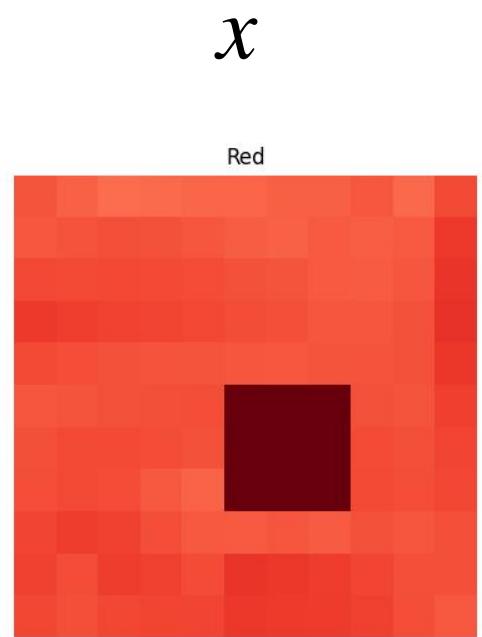
Convolutions

General definition of convolution between an input signal x and the kernel k

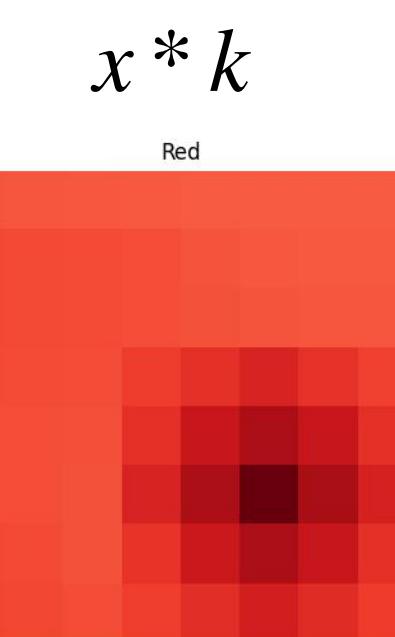
$$(x * k)(t) = \int x(s) k(t - s) ds$$

Most ML libraries use *cross-correlation*

$$(x * k)_{i,j} = \sum_n \sum_m x_{i+n, j+m} k_{n,m}$$



$$k = \begin{bmatrix} \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \end{bmatrix}^*$$



Computing cross-correlation/convolution:

1. Element-wise multiplication of kernel weights and pixels
2. Sum all the elements
3. Slide the kernel

113	122	130	128	126	125	121	121	115	127
115	112	110	111	114	119	122	117	120	117
105	106	105	106	108	111	113	117	118	114
94	98	101	103	105	108	110	114	114	111
106	109	111	112	113	114	115	113	111	111
114	113	111	110	109	0	0	0	111	113
110	106	106	108	111	0	0	0	107	110
109	106	108	116	124	0	0	0	106	109
102	97	100	109	116	117	114	118	111	114
100	108	97	100	107	90	93	97	102	110

$$\begin{array}{ccc} \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \times & & \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ = & & \end{array}$$

*More precisely $k_{i,j} = \begin{cases} \frac{1}{9} & \text{if } i \leq 2 \text{ and } j \leq 2 \\ 0 & \text{otherwise} \end{cases}$

Convolutional layer

Explore in practice the effects of different components of convolutional layers

1. Go to Brightspace → 5. Convolutional Networks → Convolutional Layer Practice
2. Follow instructions in the notebook
3. 10 mins
4. Discuss together

Note: it's not a programming exercise.

Convolutional layer: input



Shape of cat (1068, 1600, 3)

Maximum value: 255

Minimum value: 0



Shape of logo (166, 166)

Maximum value: 255

Minimum value: 0

Convolutional layer: input



Shape of cat (1068, 1600, 3)
Maximum value: 255
Minimum value: 0

Size of the image → (height, width, rgb channels)



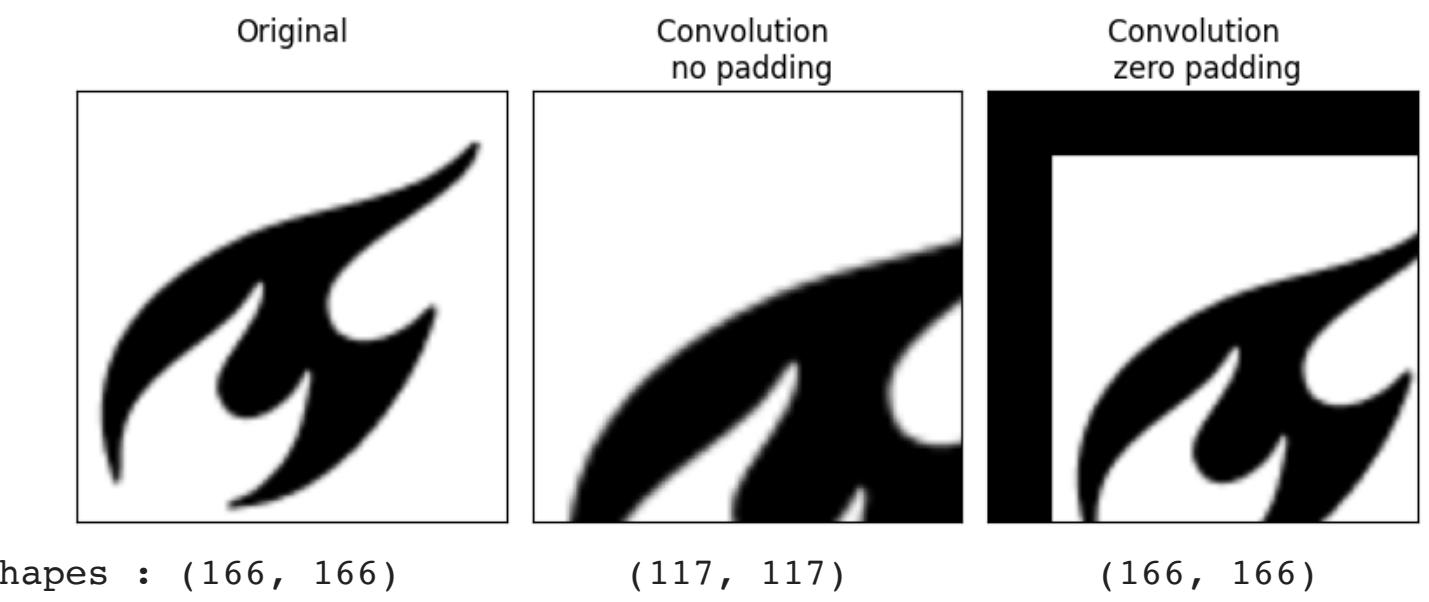
Shape of logo (166, 166)
Maximum value: 255
Minimum value: 0

Channel range [0, 255] integers

Grey scale black to white 0 to 255

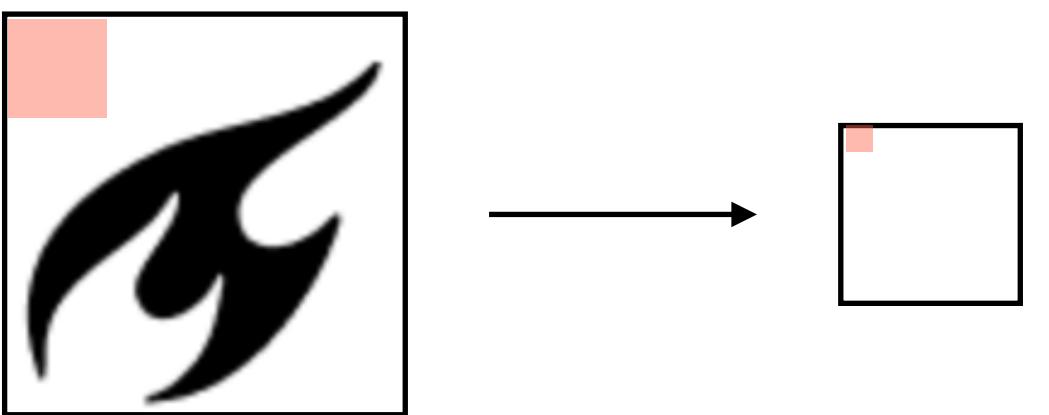
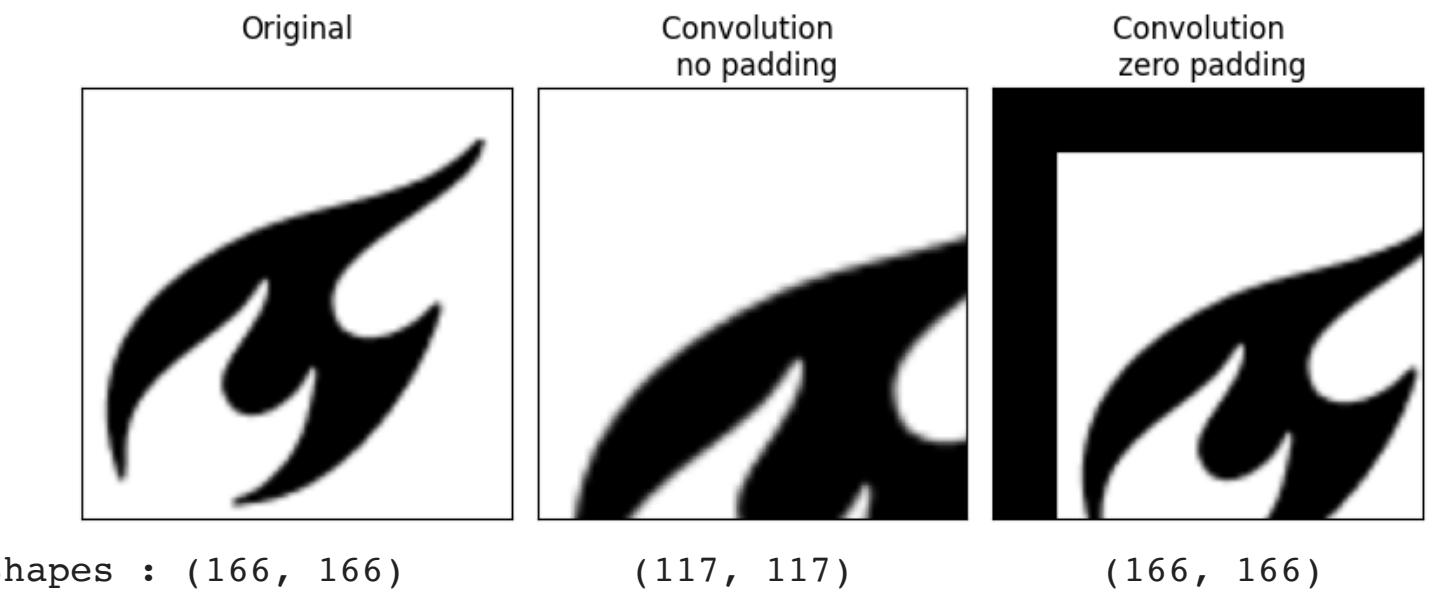
Convolutional layer: padding

- **Padding**



Convolutional layer: padding

- **Padding**

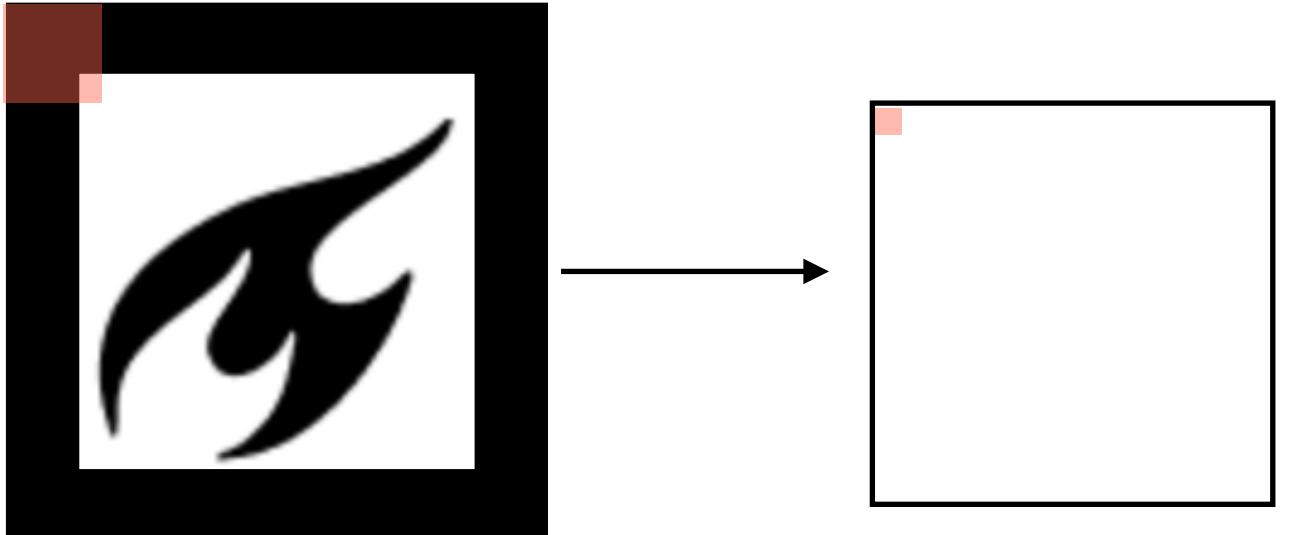
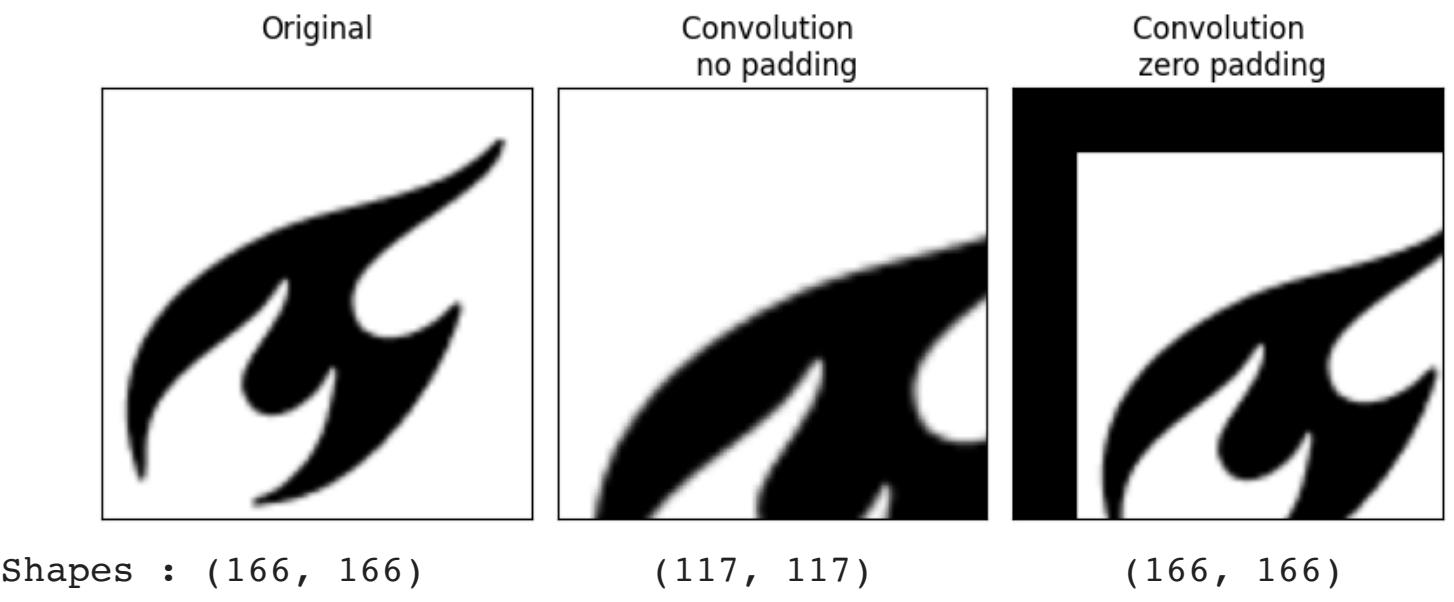


Convolution has a shrink effect*

*How much exactly? What is the shape of the output?

Convolutional layer: padding

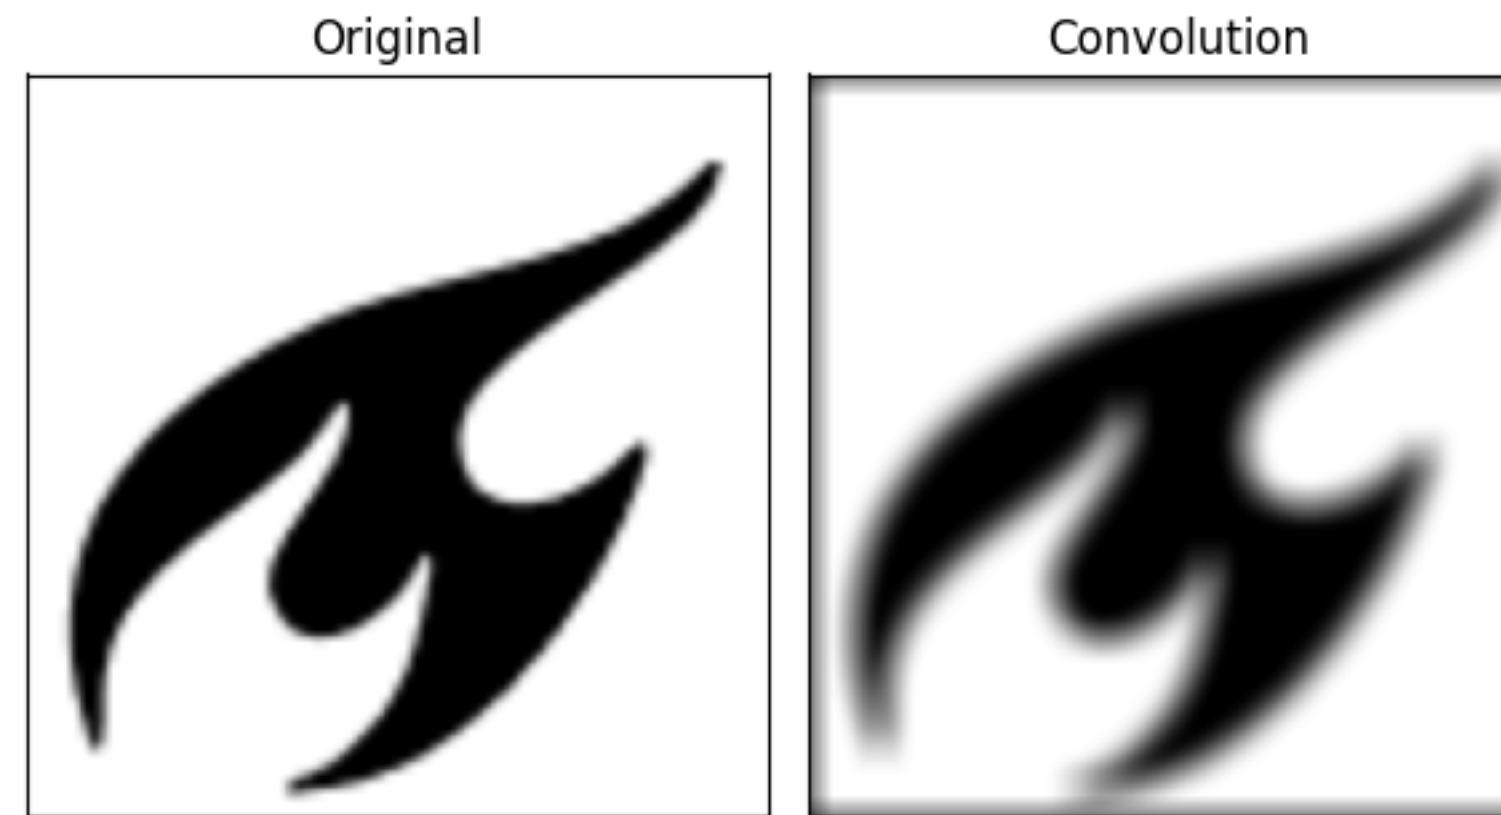
- **Padding**



Zero-padding → input and output have same shape

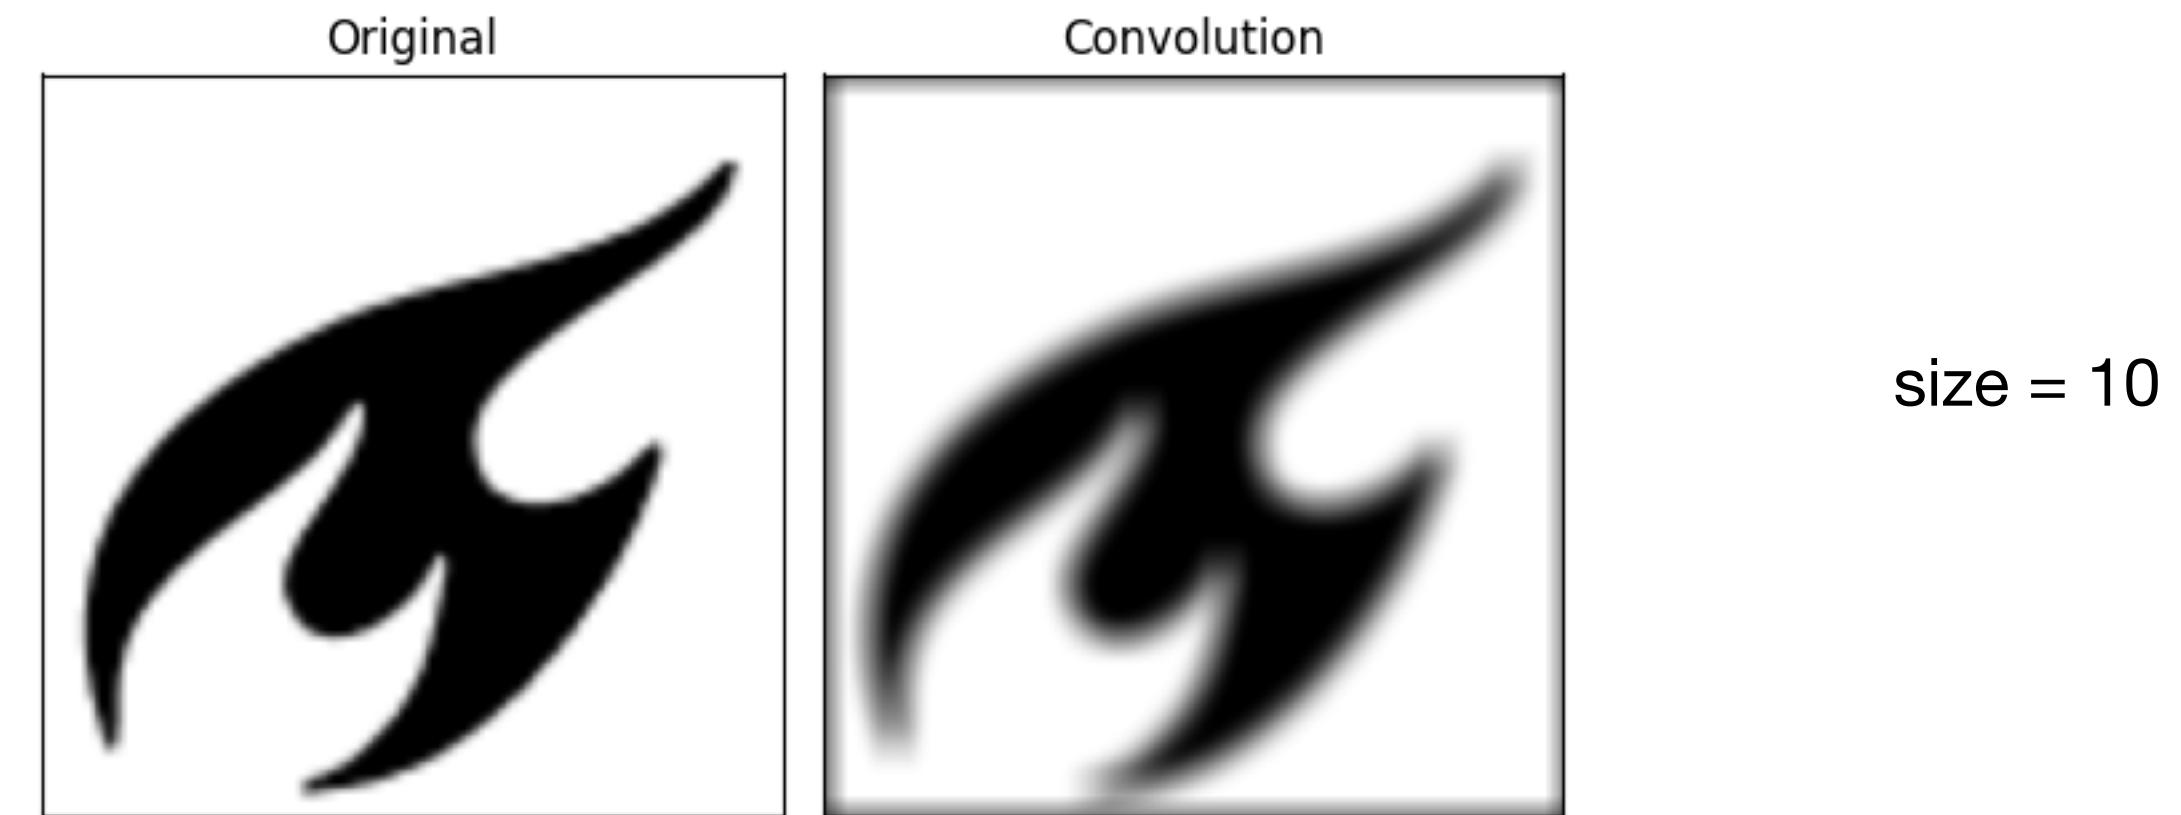
Convolutional layer: kernel effects

- Padding
- **Moving Average Filter**



Convolutional layer: kernel effects

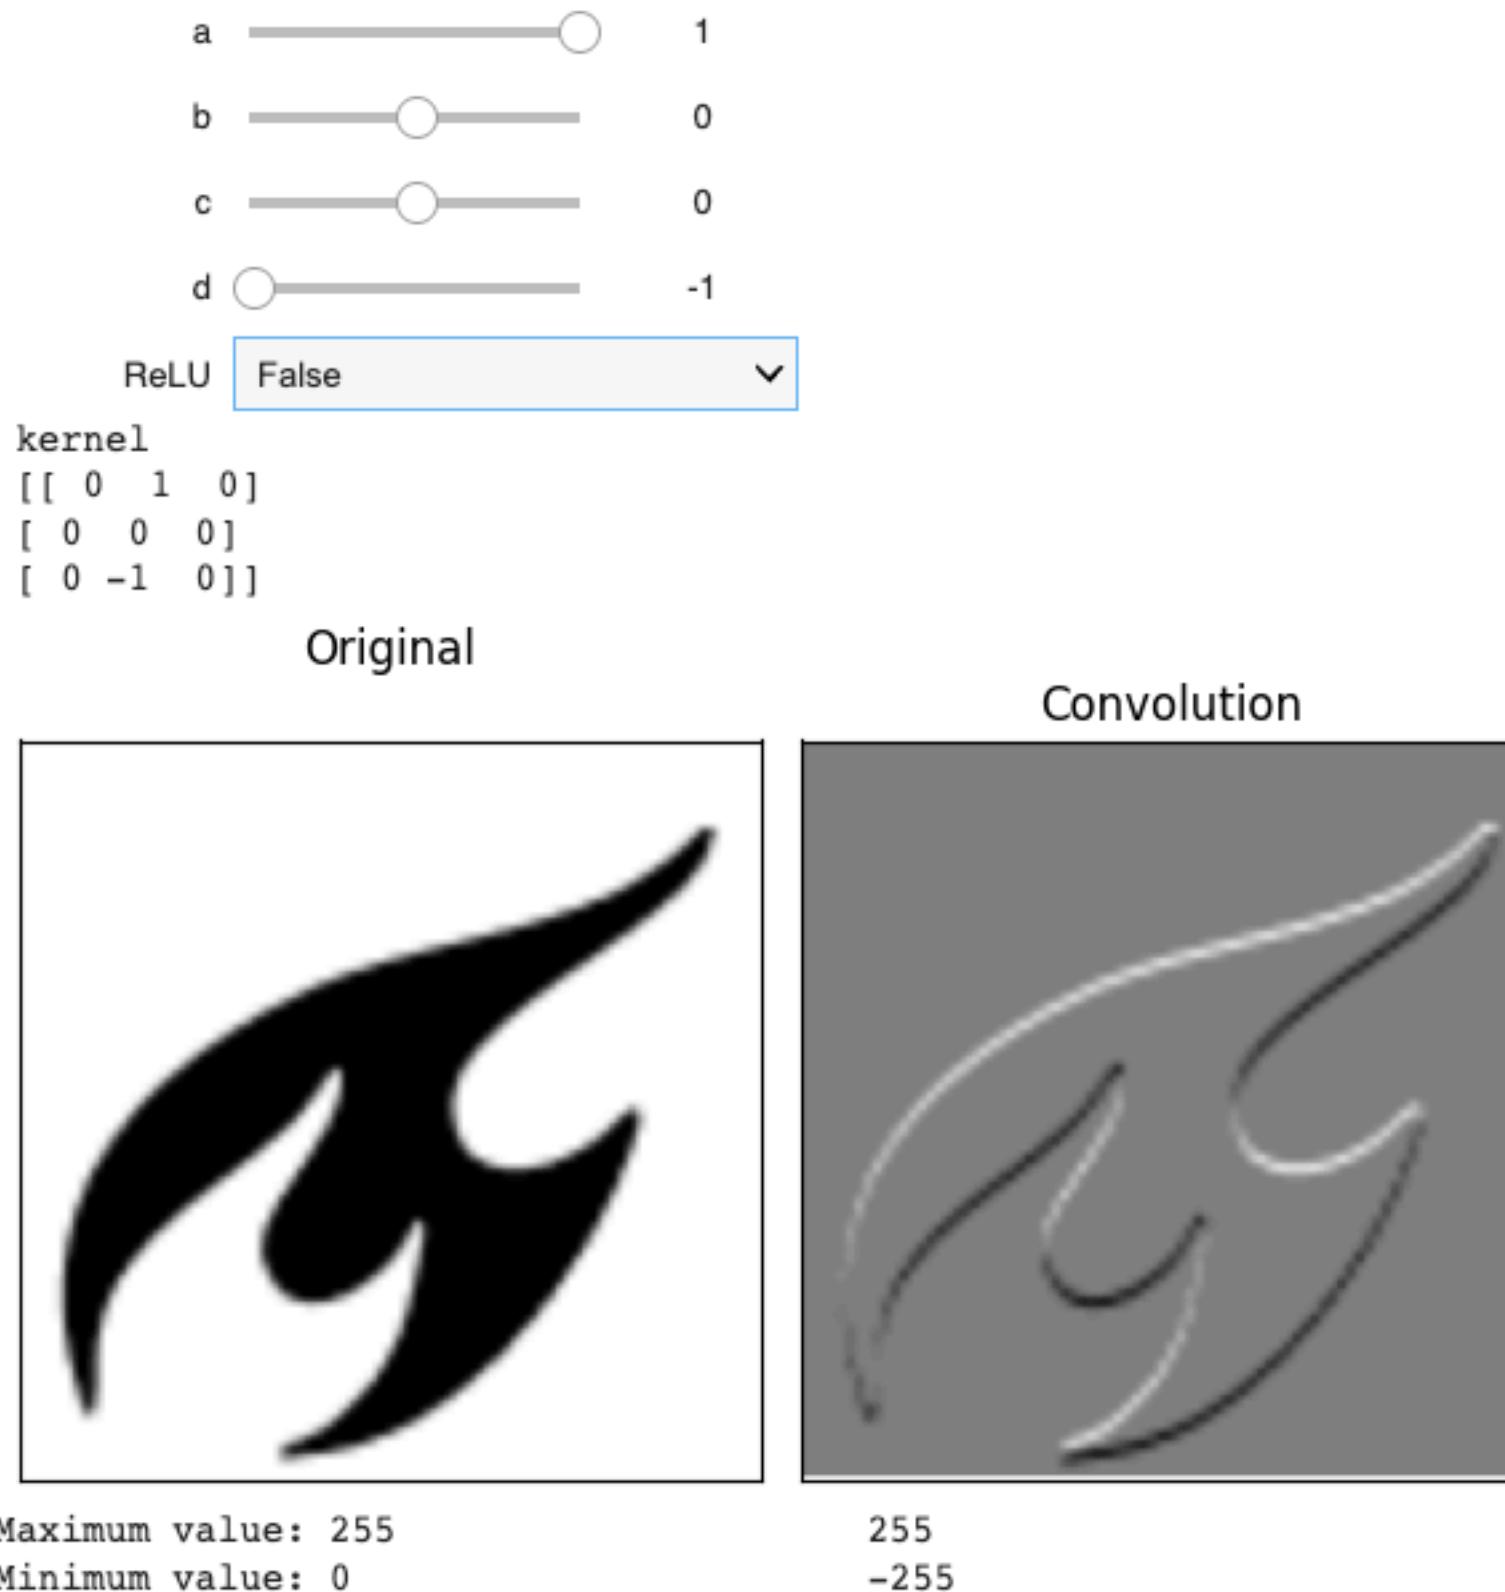
- Padding
- **Moving Average Filter**



Blur Filter

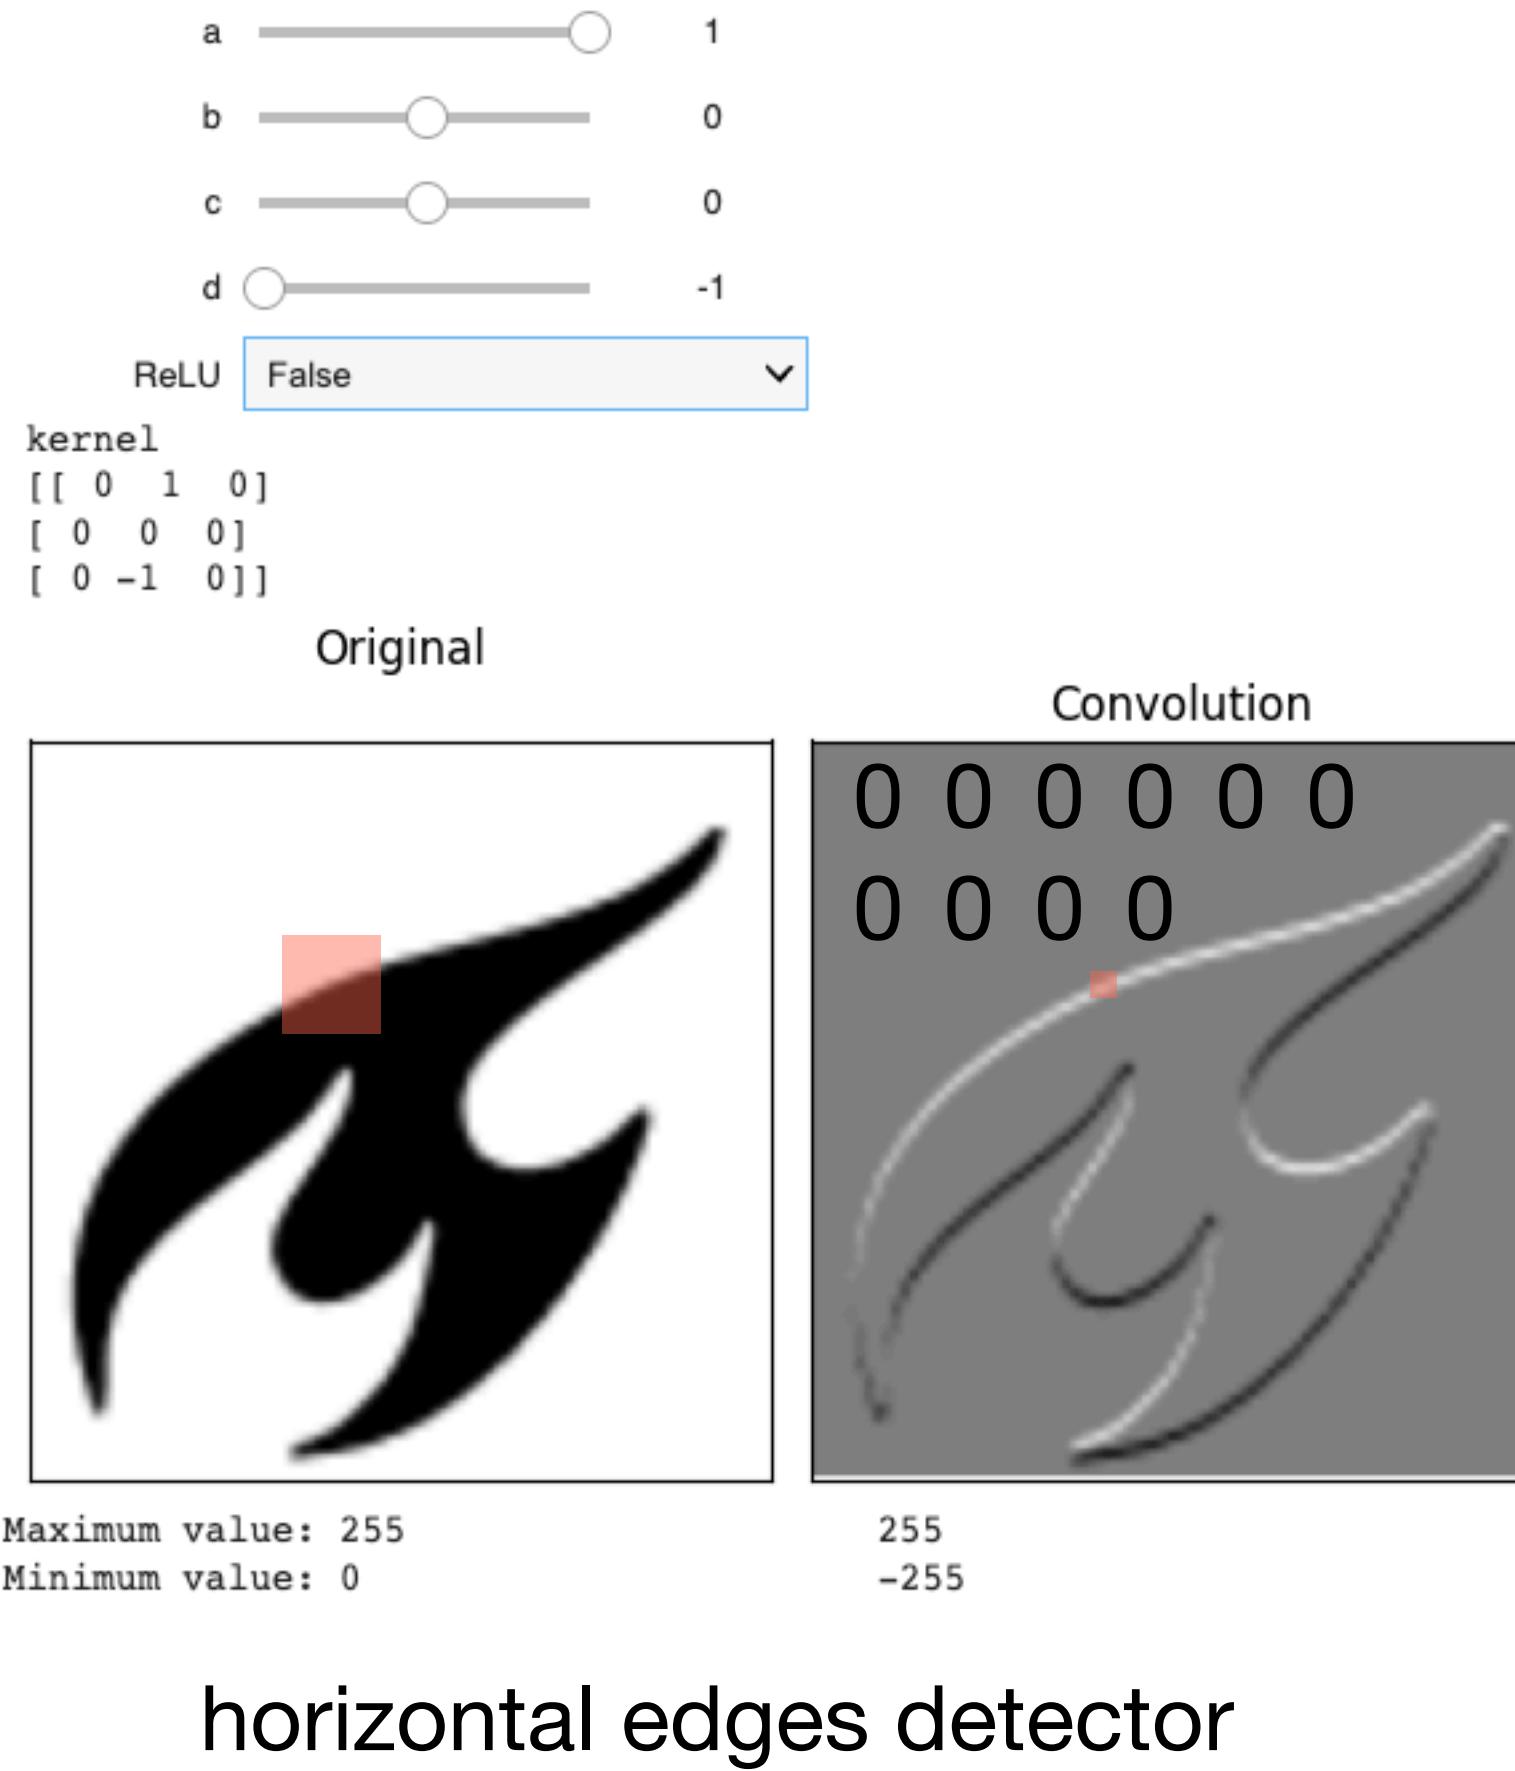
Convolutional layer: kernel effects

- Padding
- Moving Average Filter
- **Detector kernel**



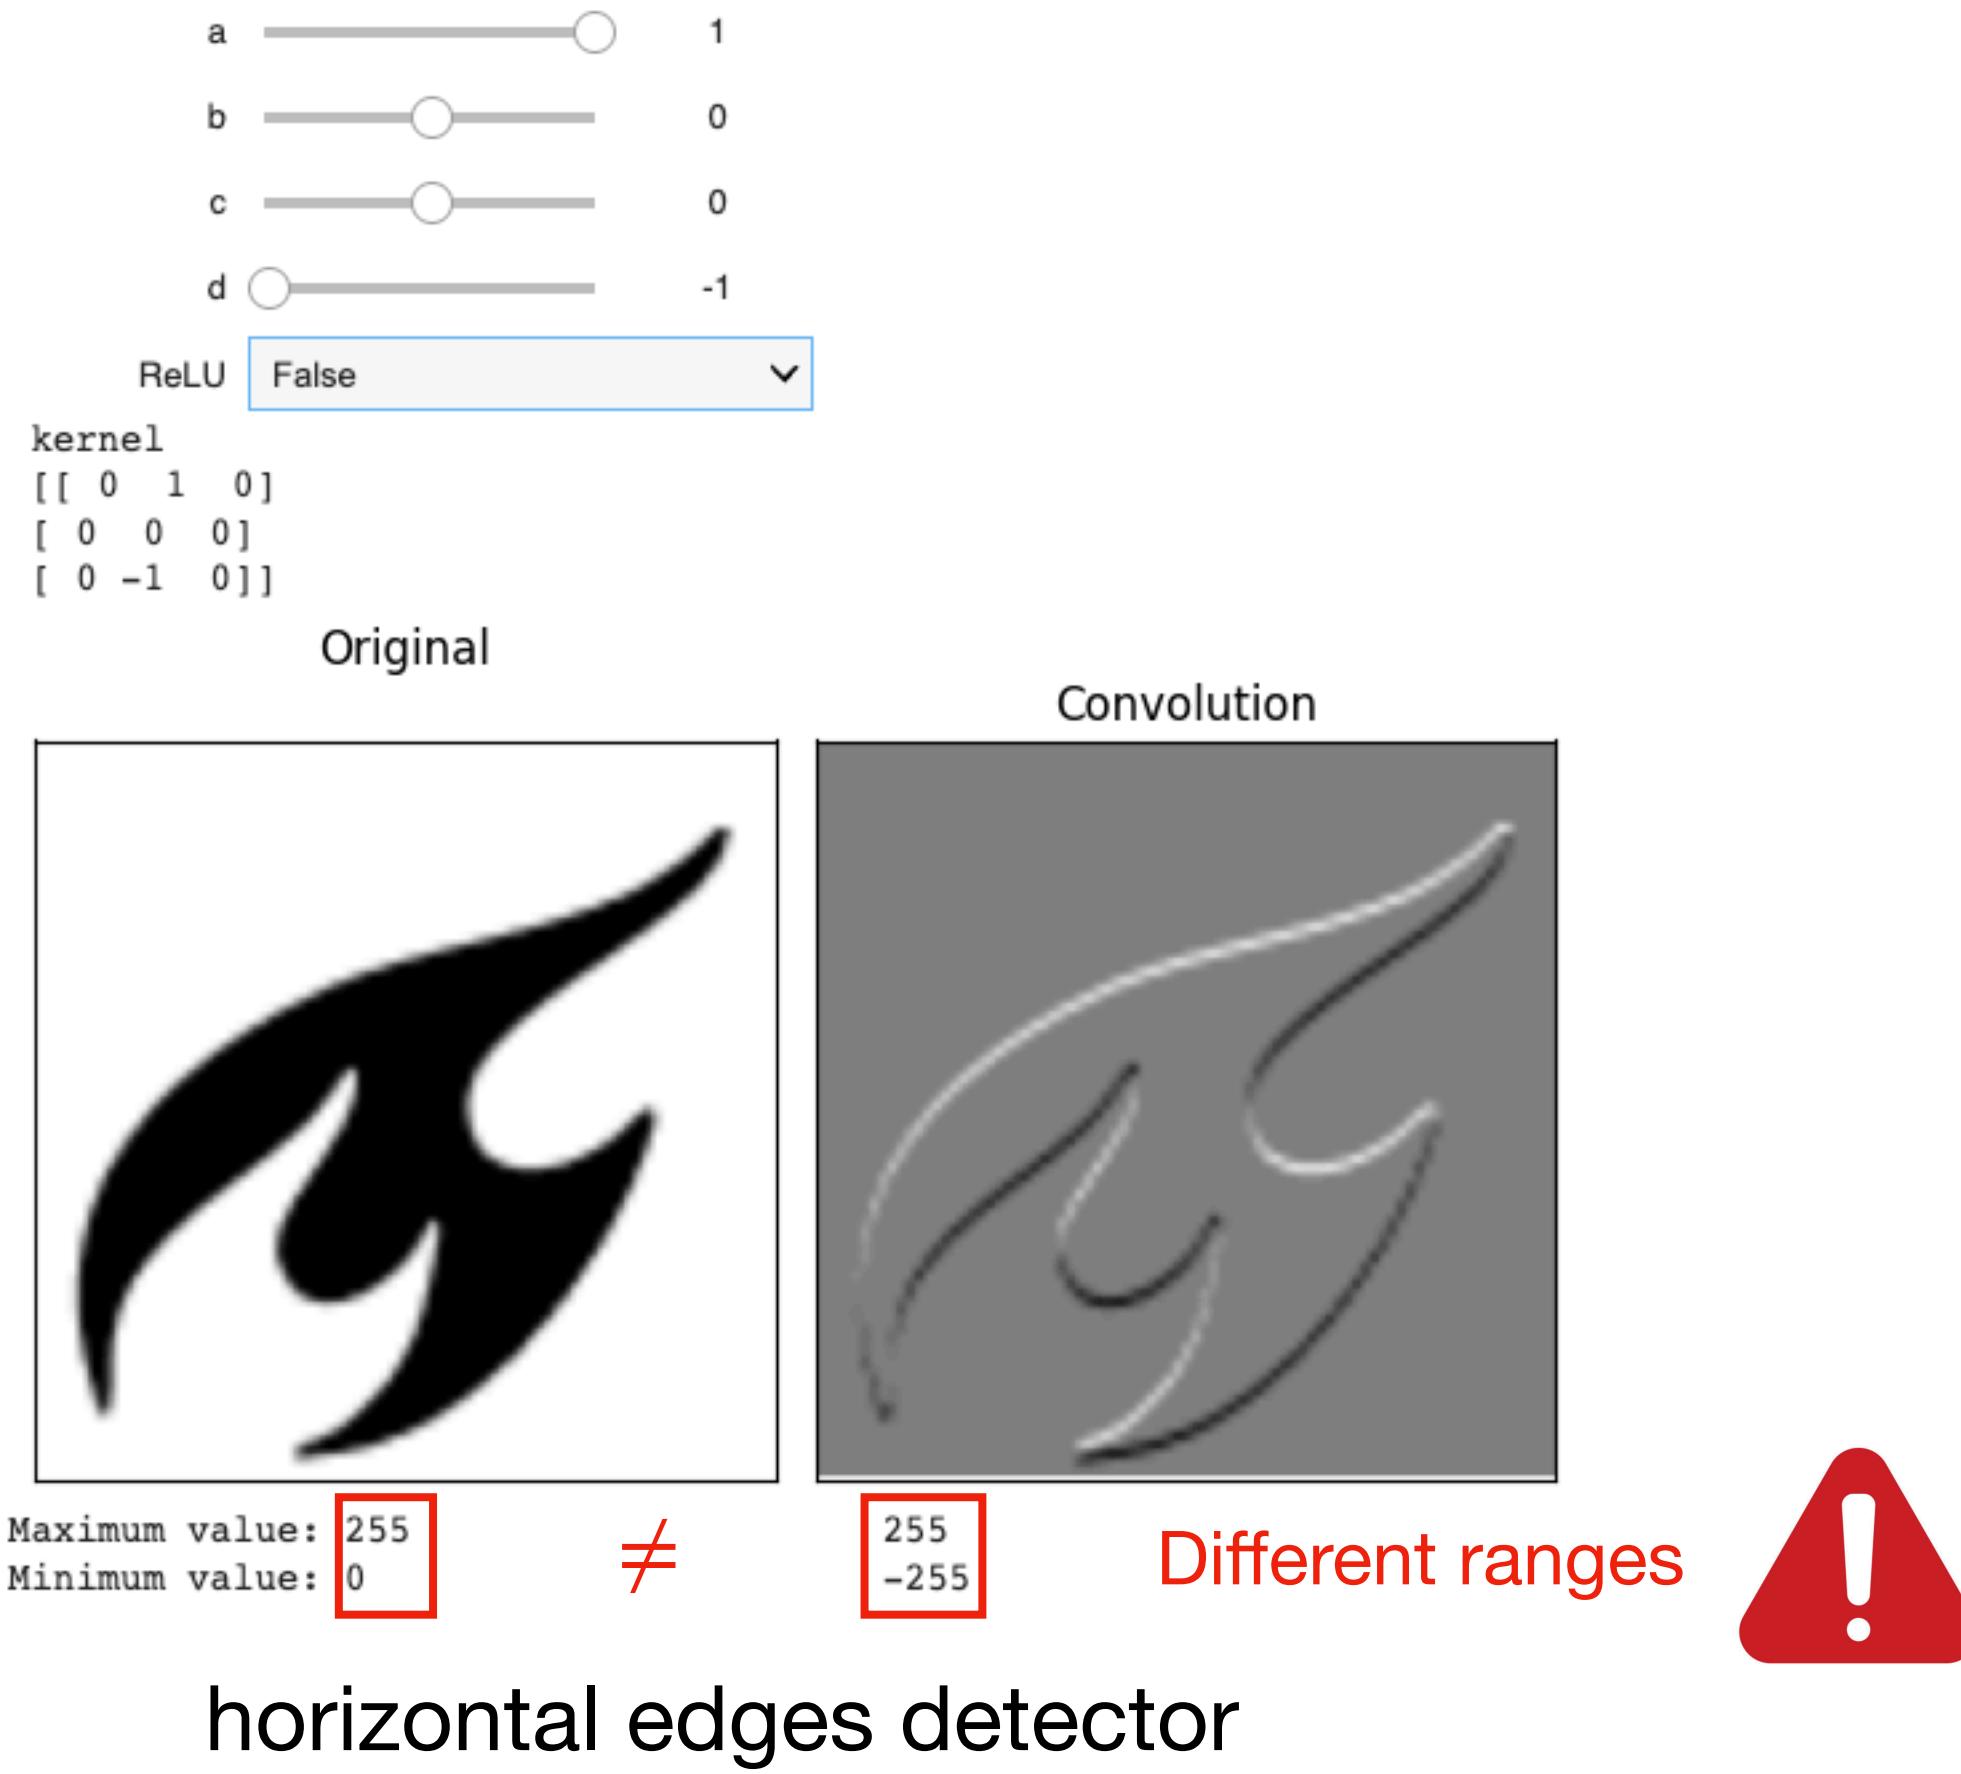
Convolutional layer: kernel effects

- Padding
- Moving Average Filter
- **Detector kernel**



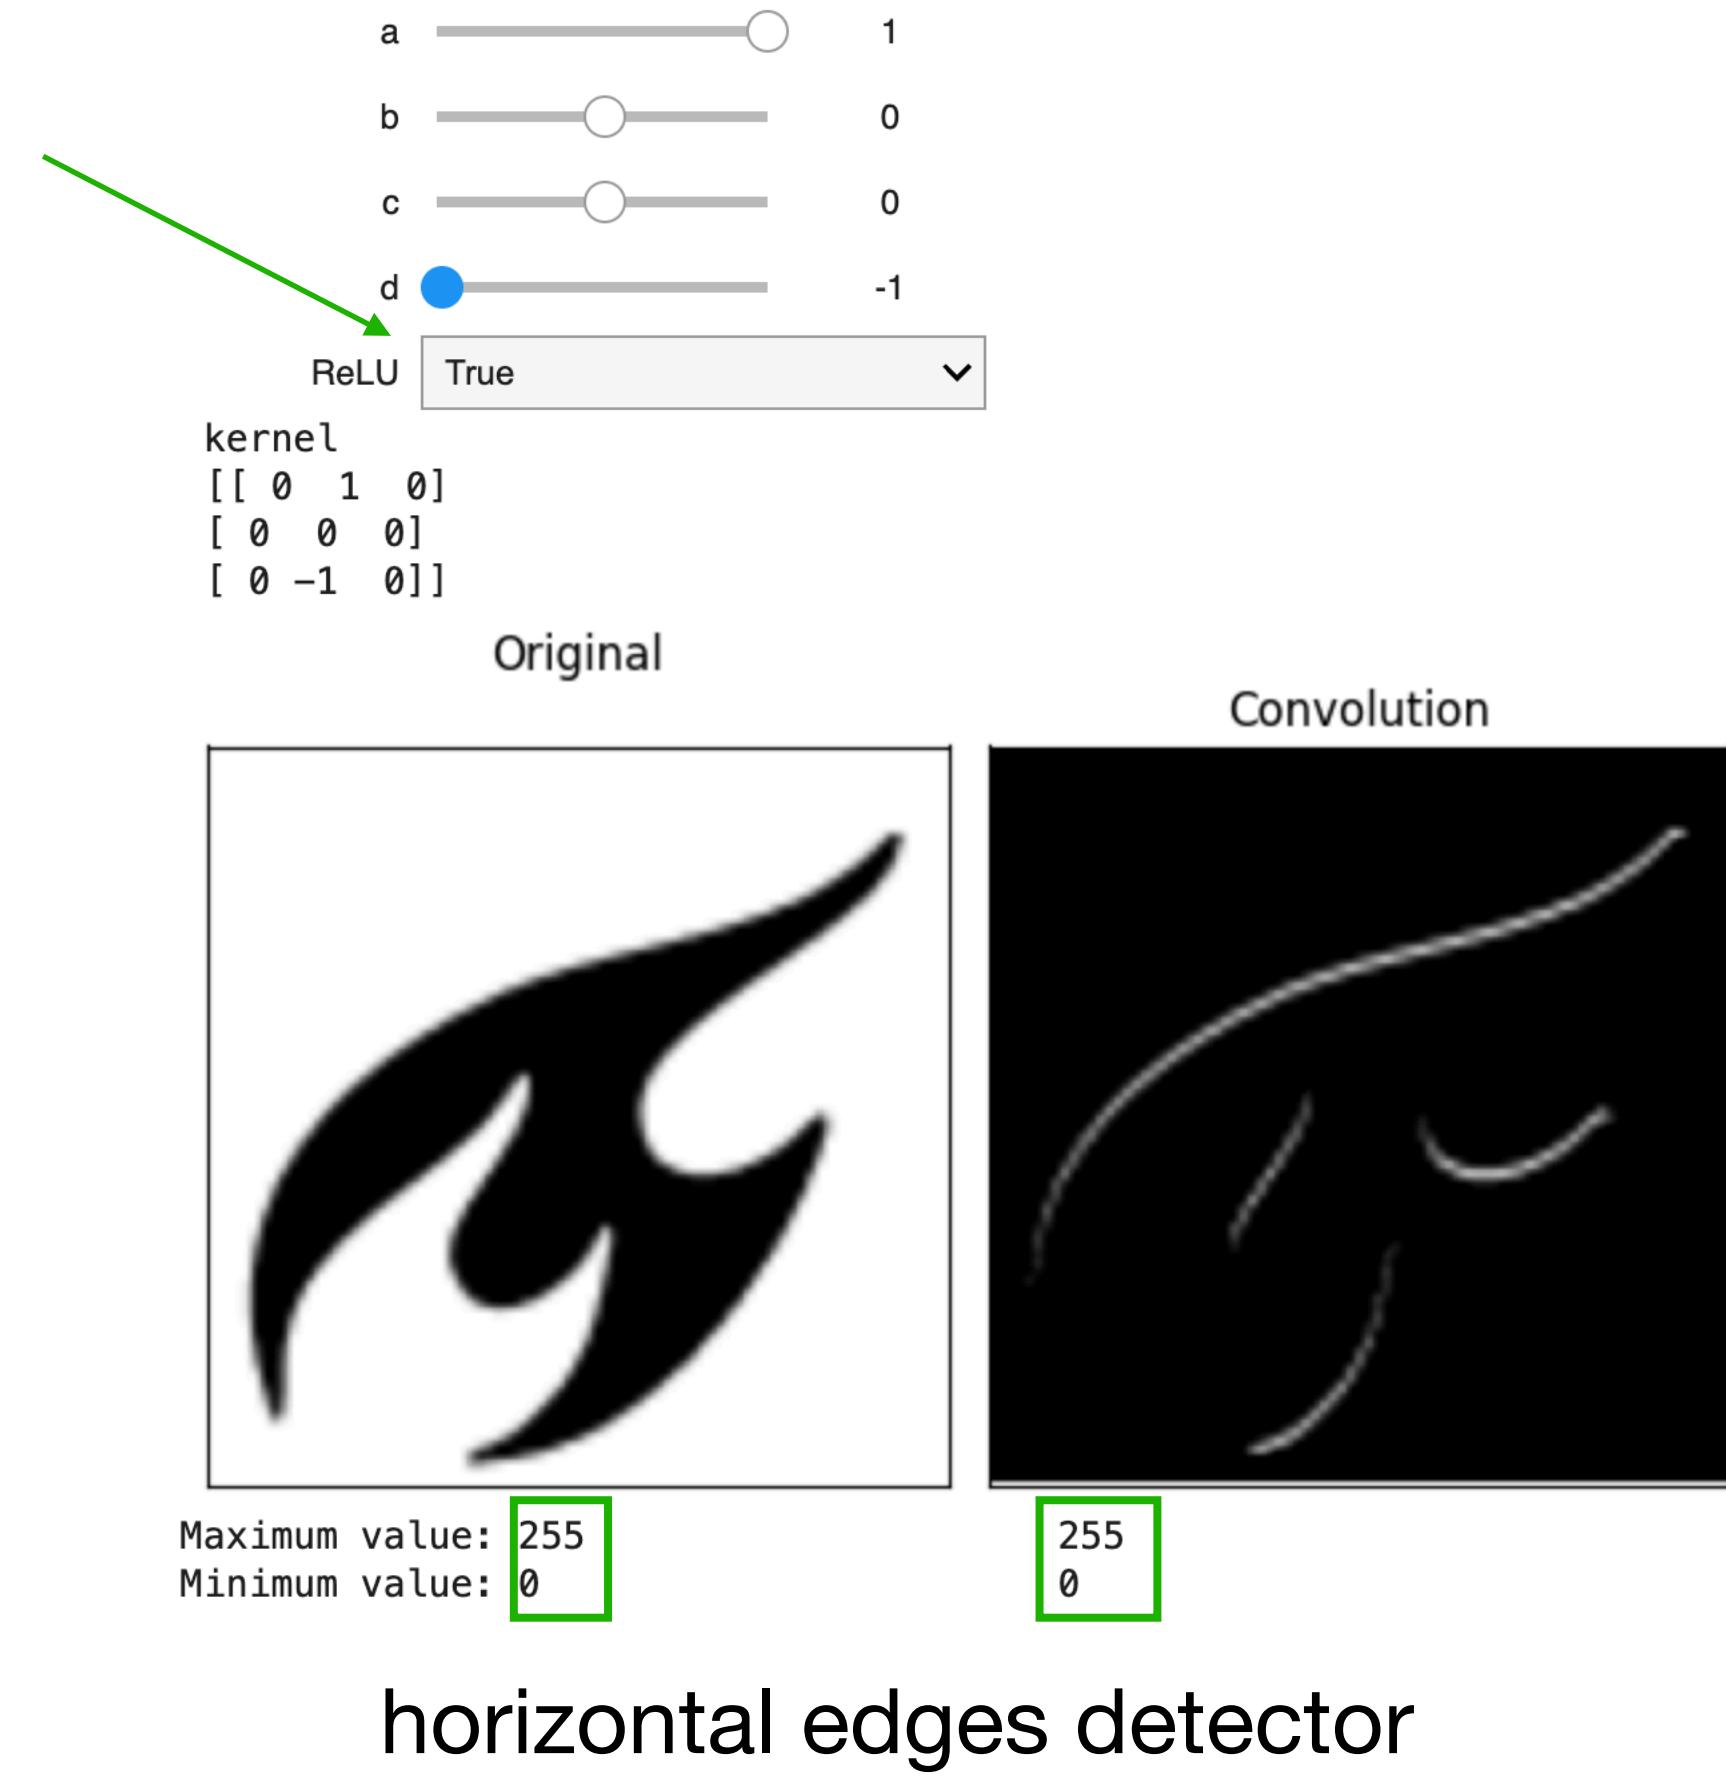
Convolutional layer: kernel effects

- Padding
- Moving Average Filter
- **Detector kernel**



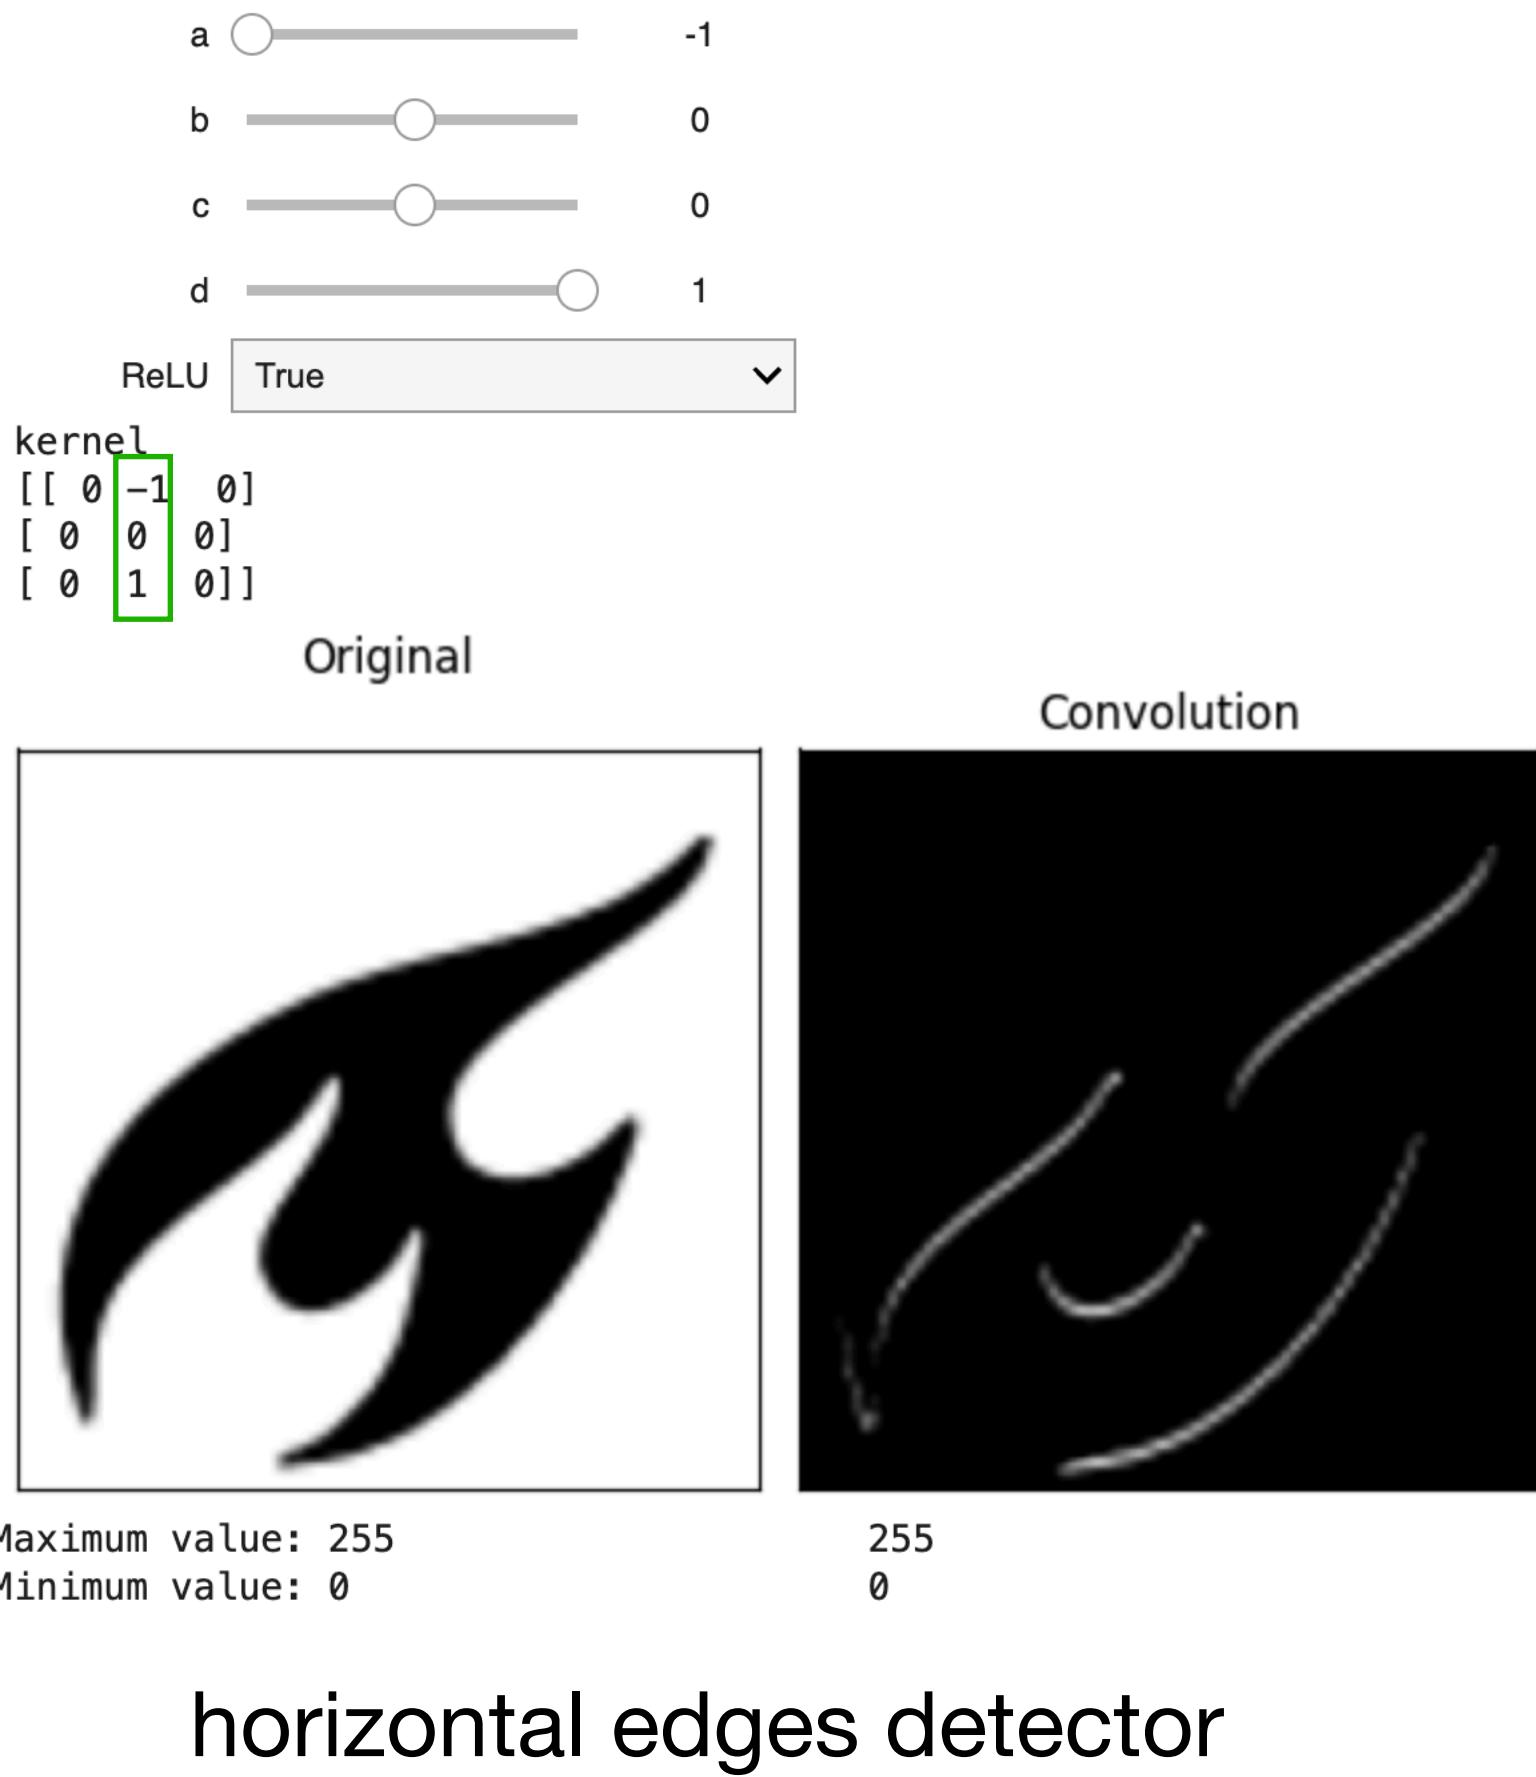
Convolutional layer: activation

- Padding
- Moving Average Filter
- Detector kernel
- **ReLU activation**



Convolutional layer: activation

- Padding
- Moving Average Filter
- Detector kernel
- **ReLU activation**



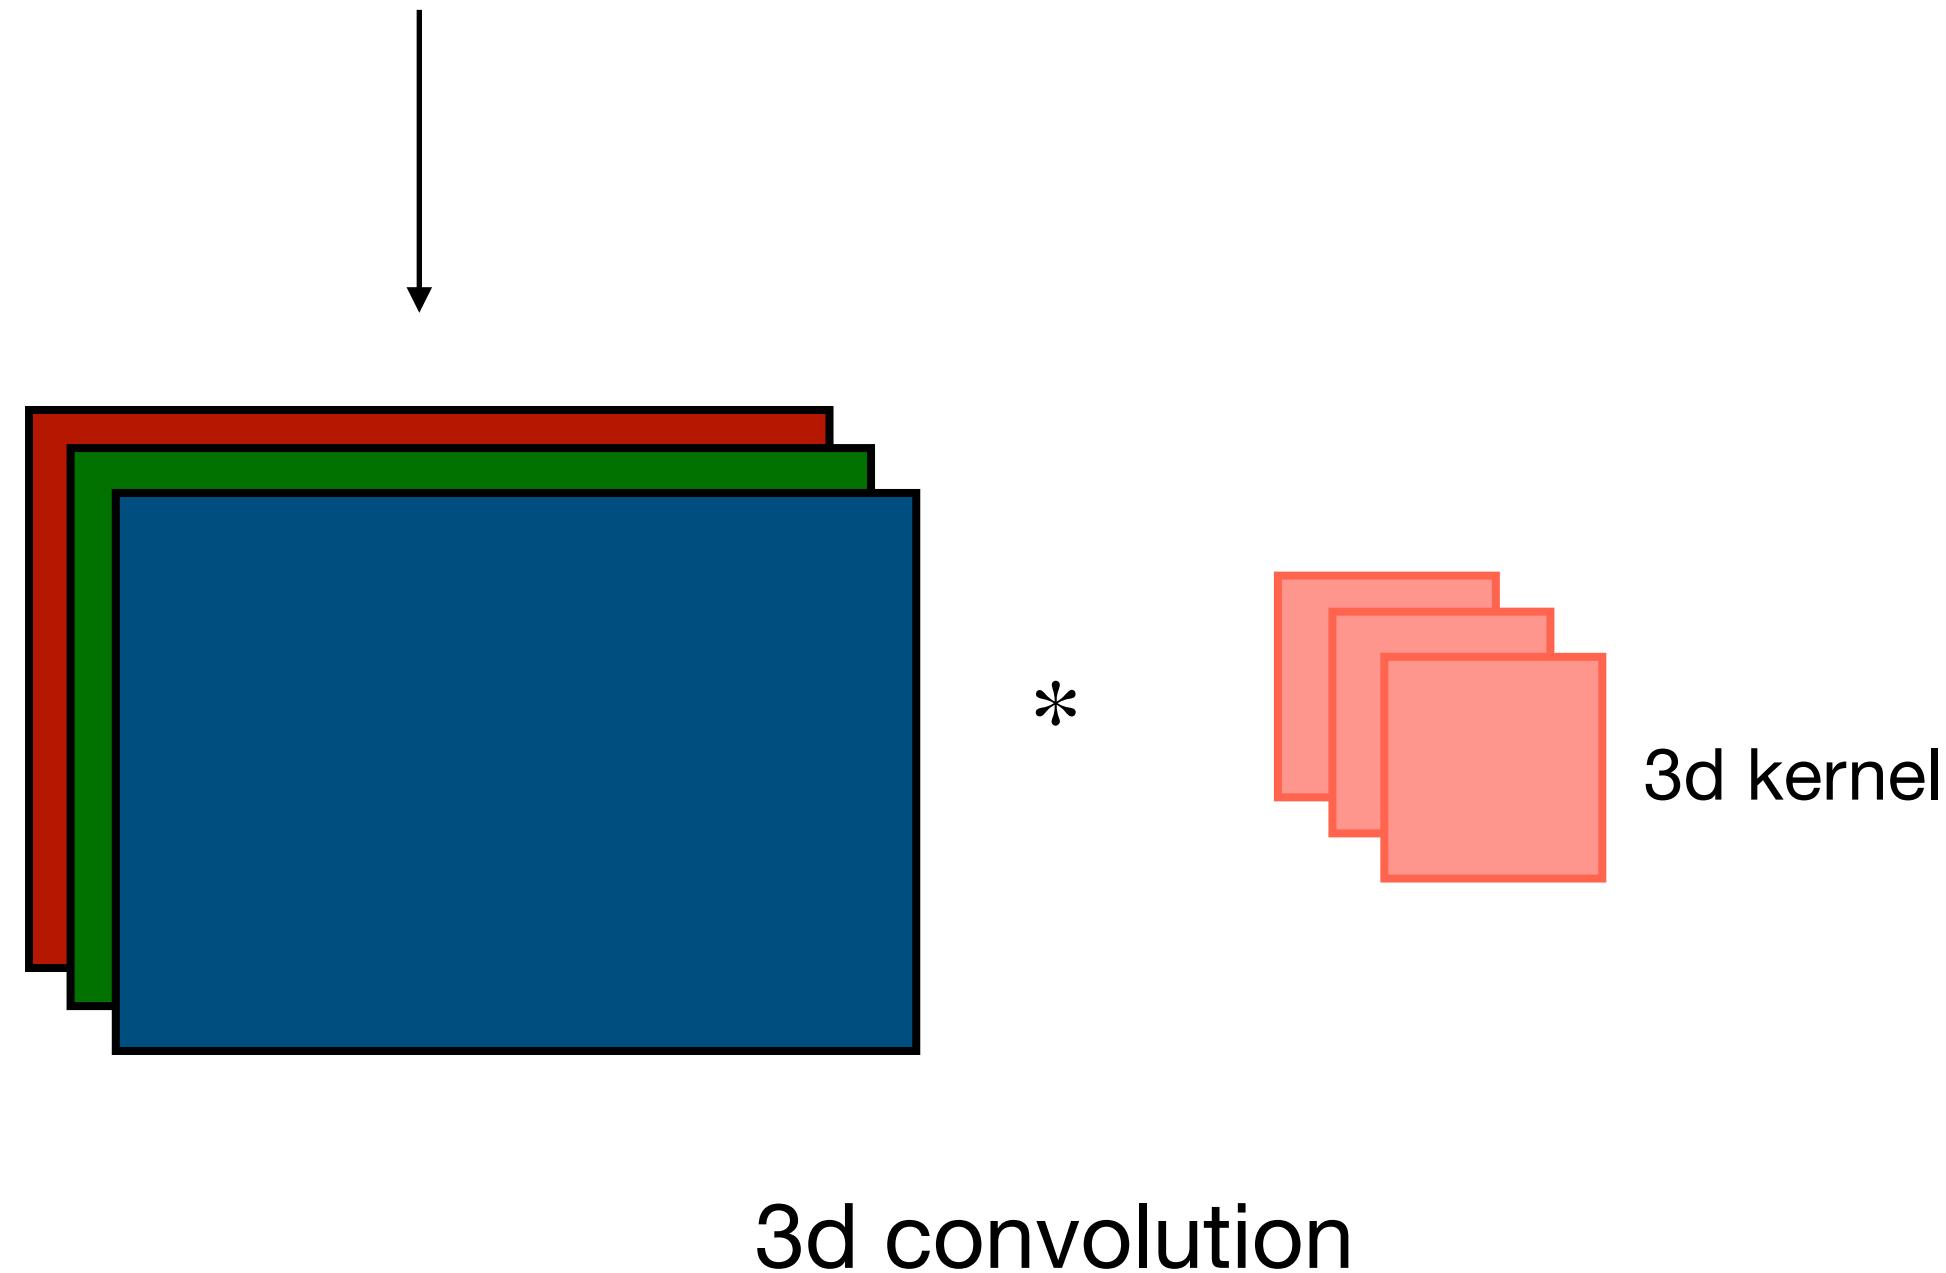
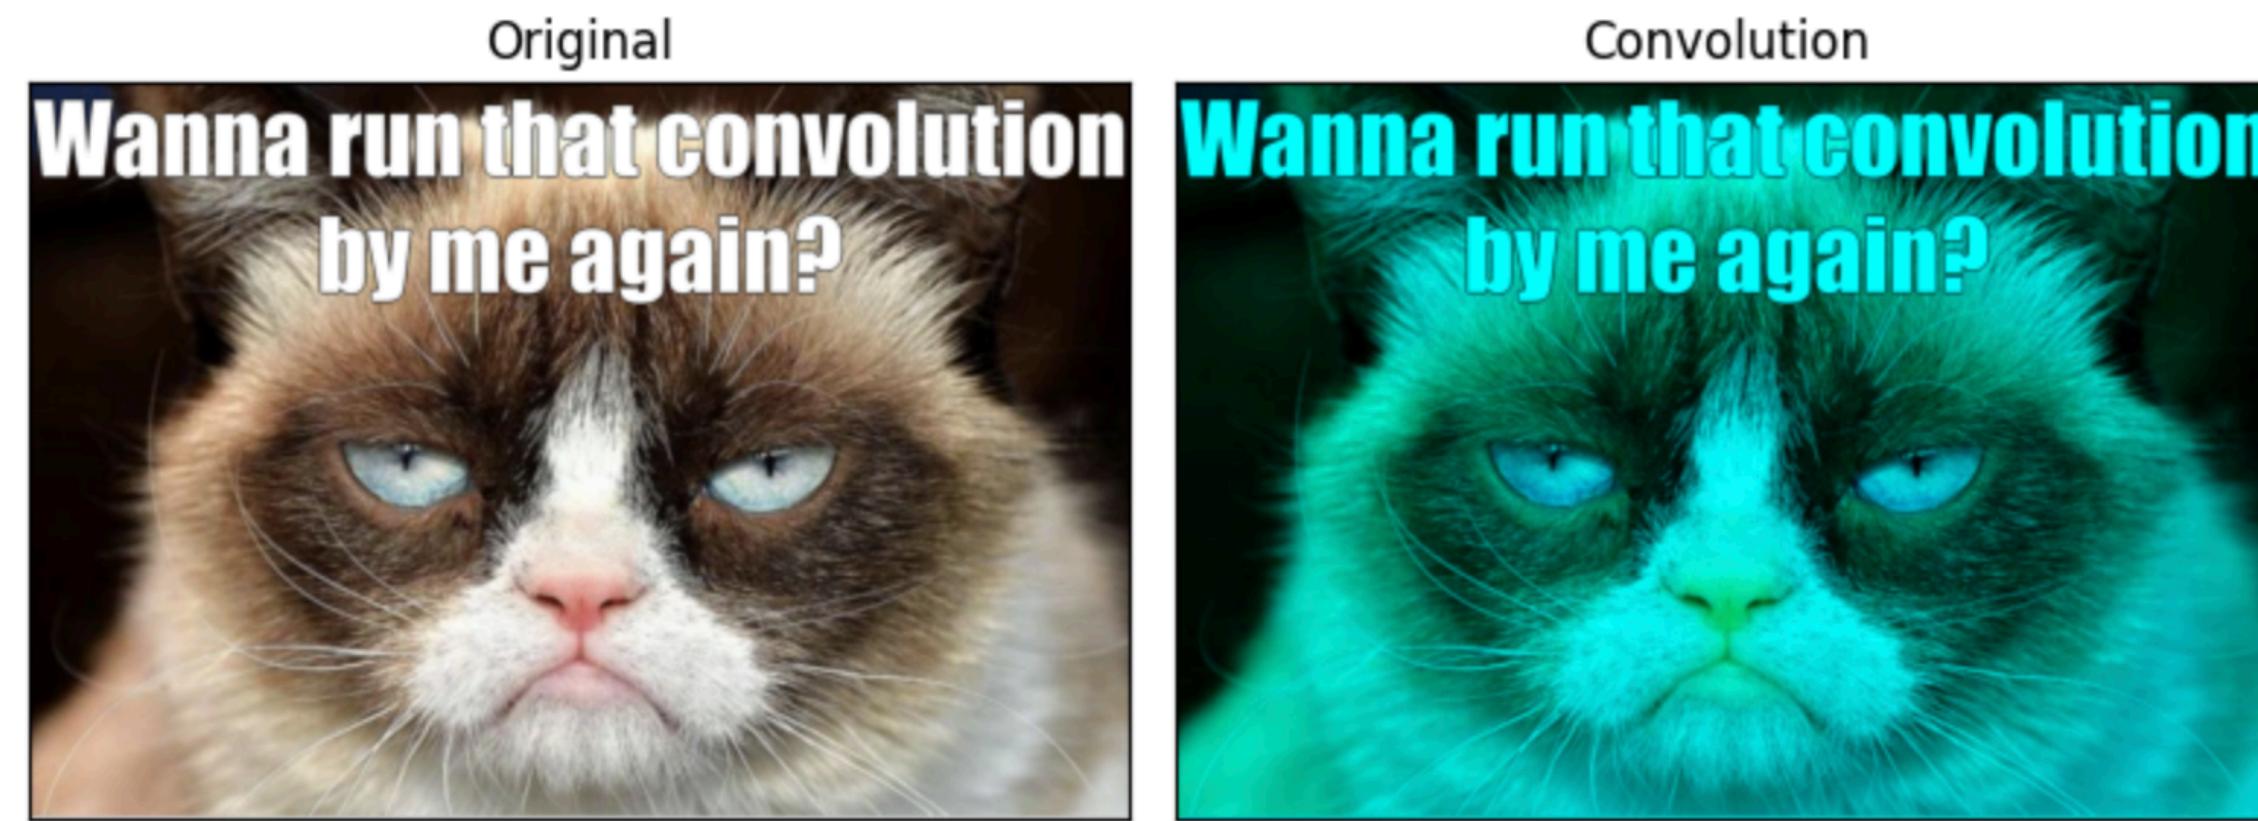
Convolutional layer: channels

- Padding
- Moving Average Filter
- Detector kernel
- ReLU activation
- **Input channels**



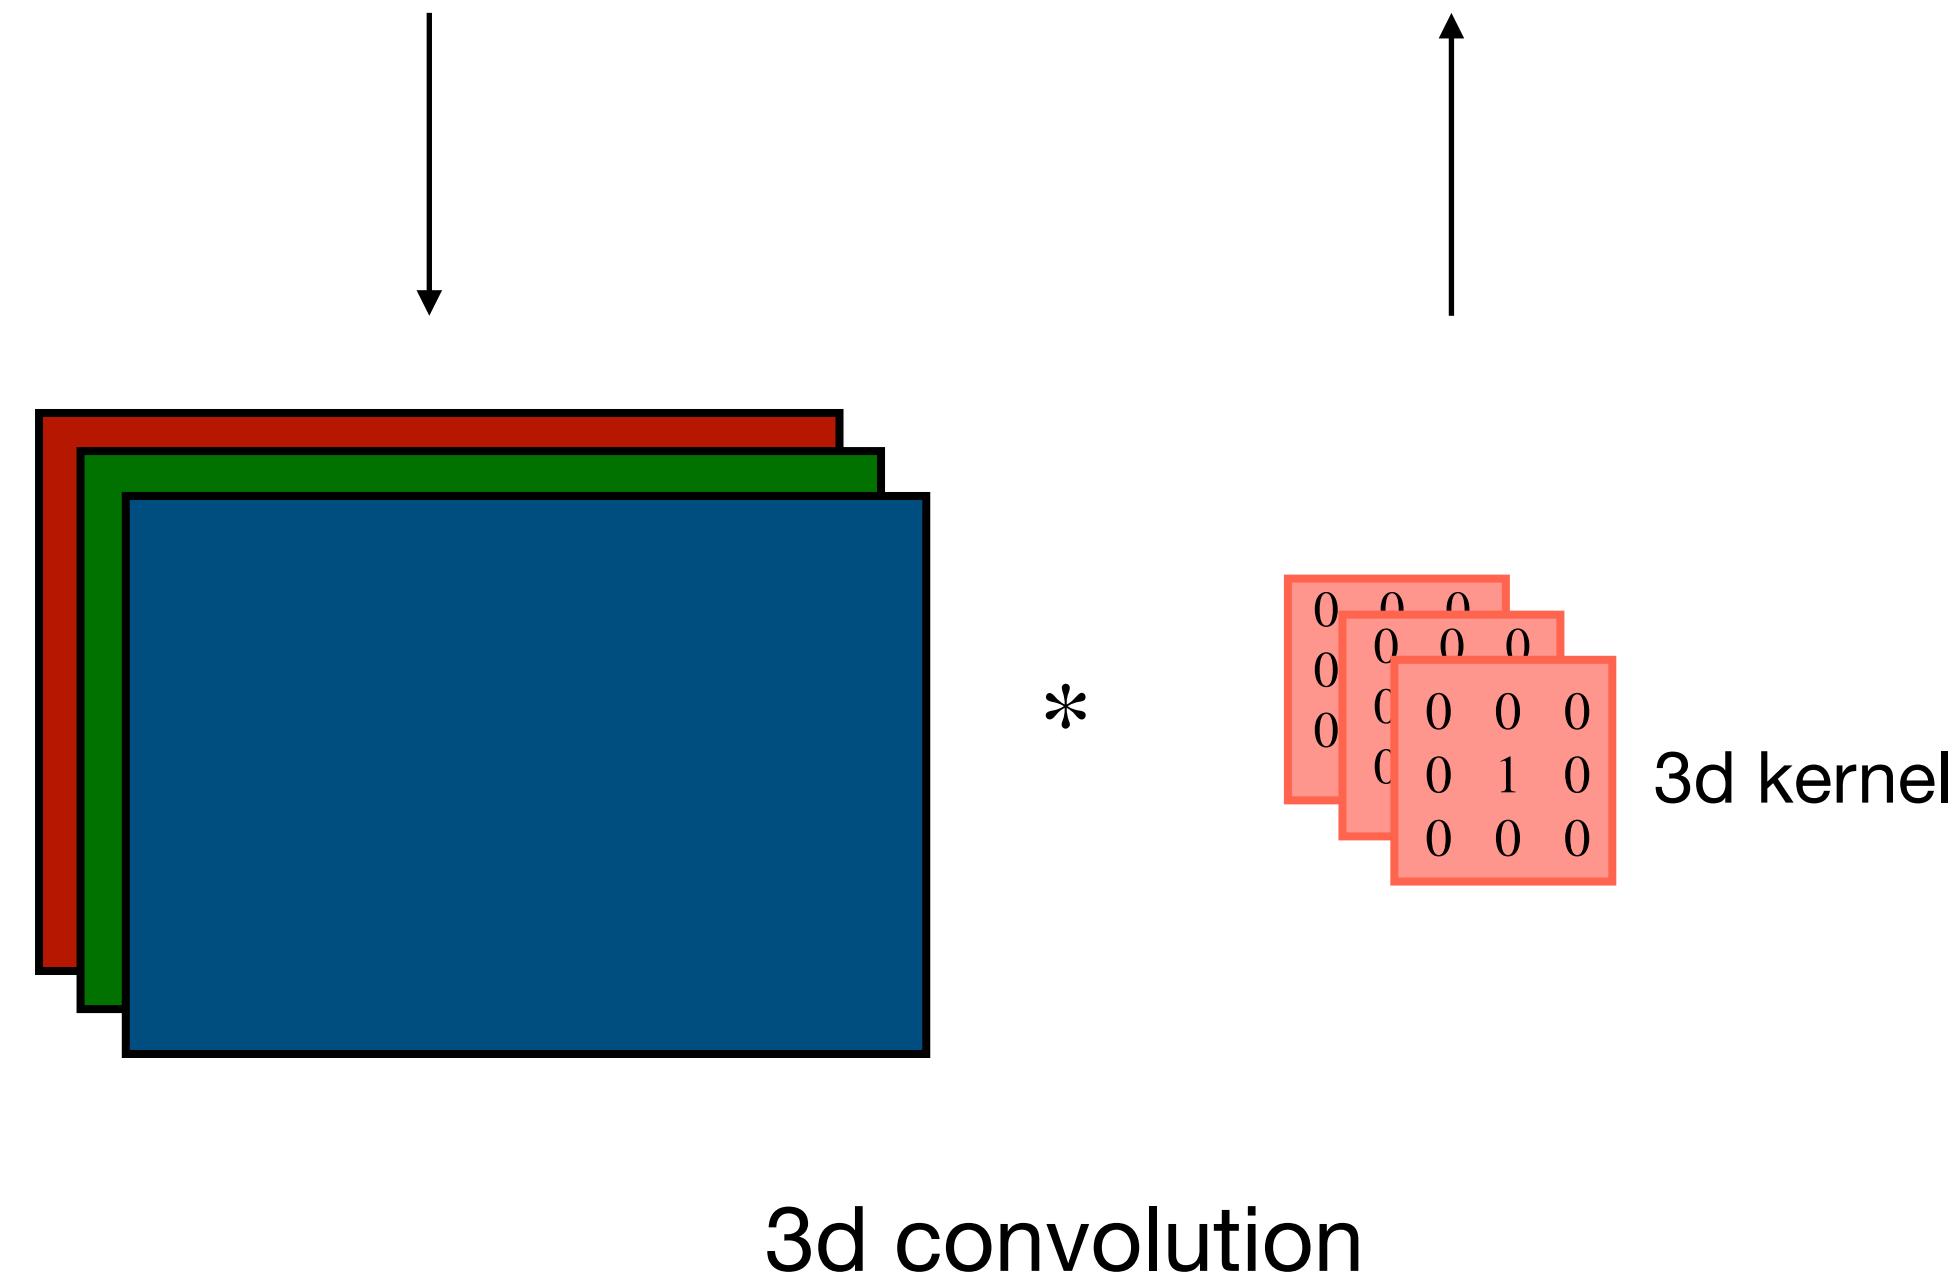
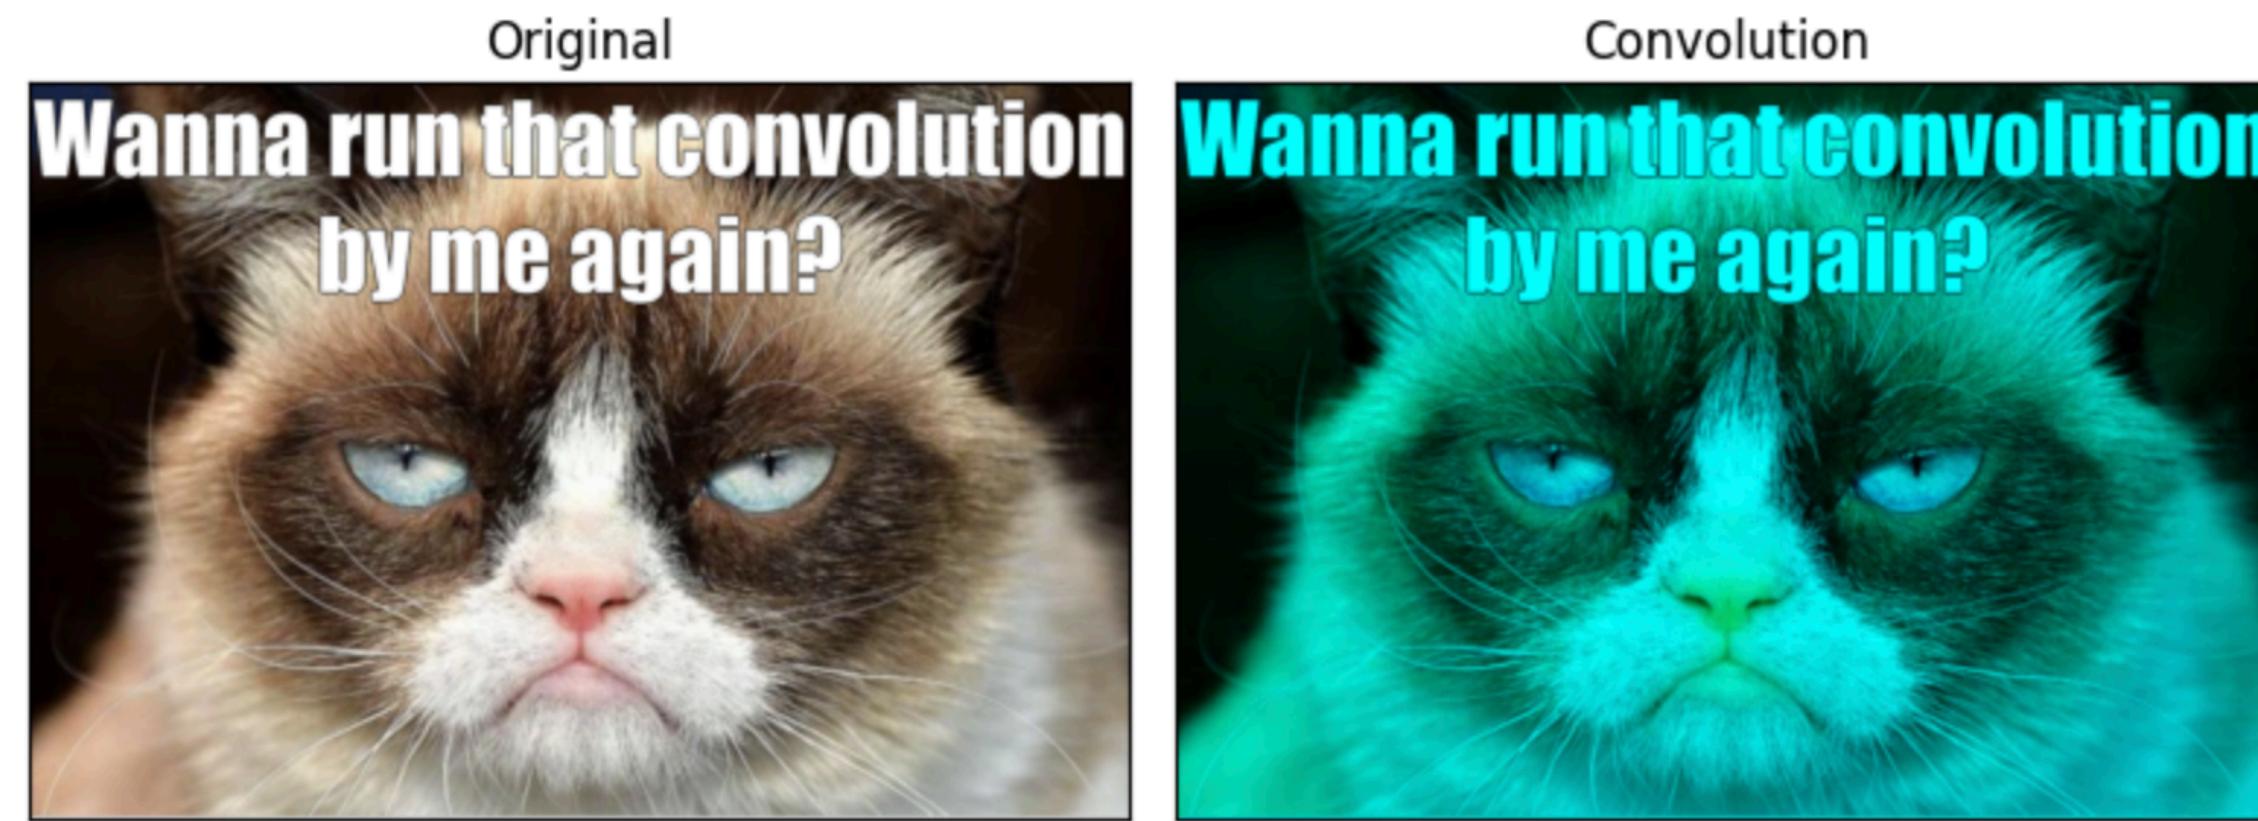
Convolutional layer: channels

- Padding
- Moving Average Filter
- Detector kernel
- ReLU activation
- **Input channels**



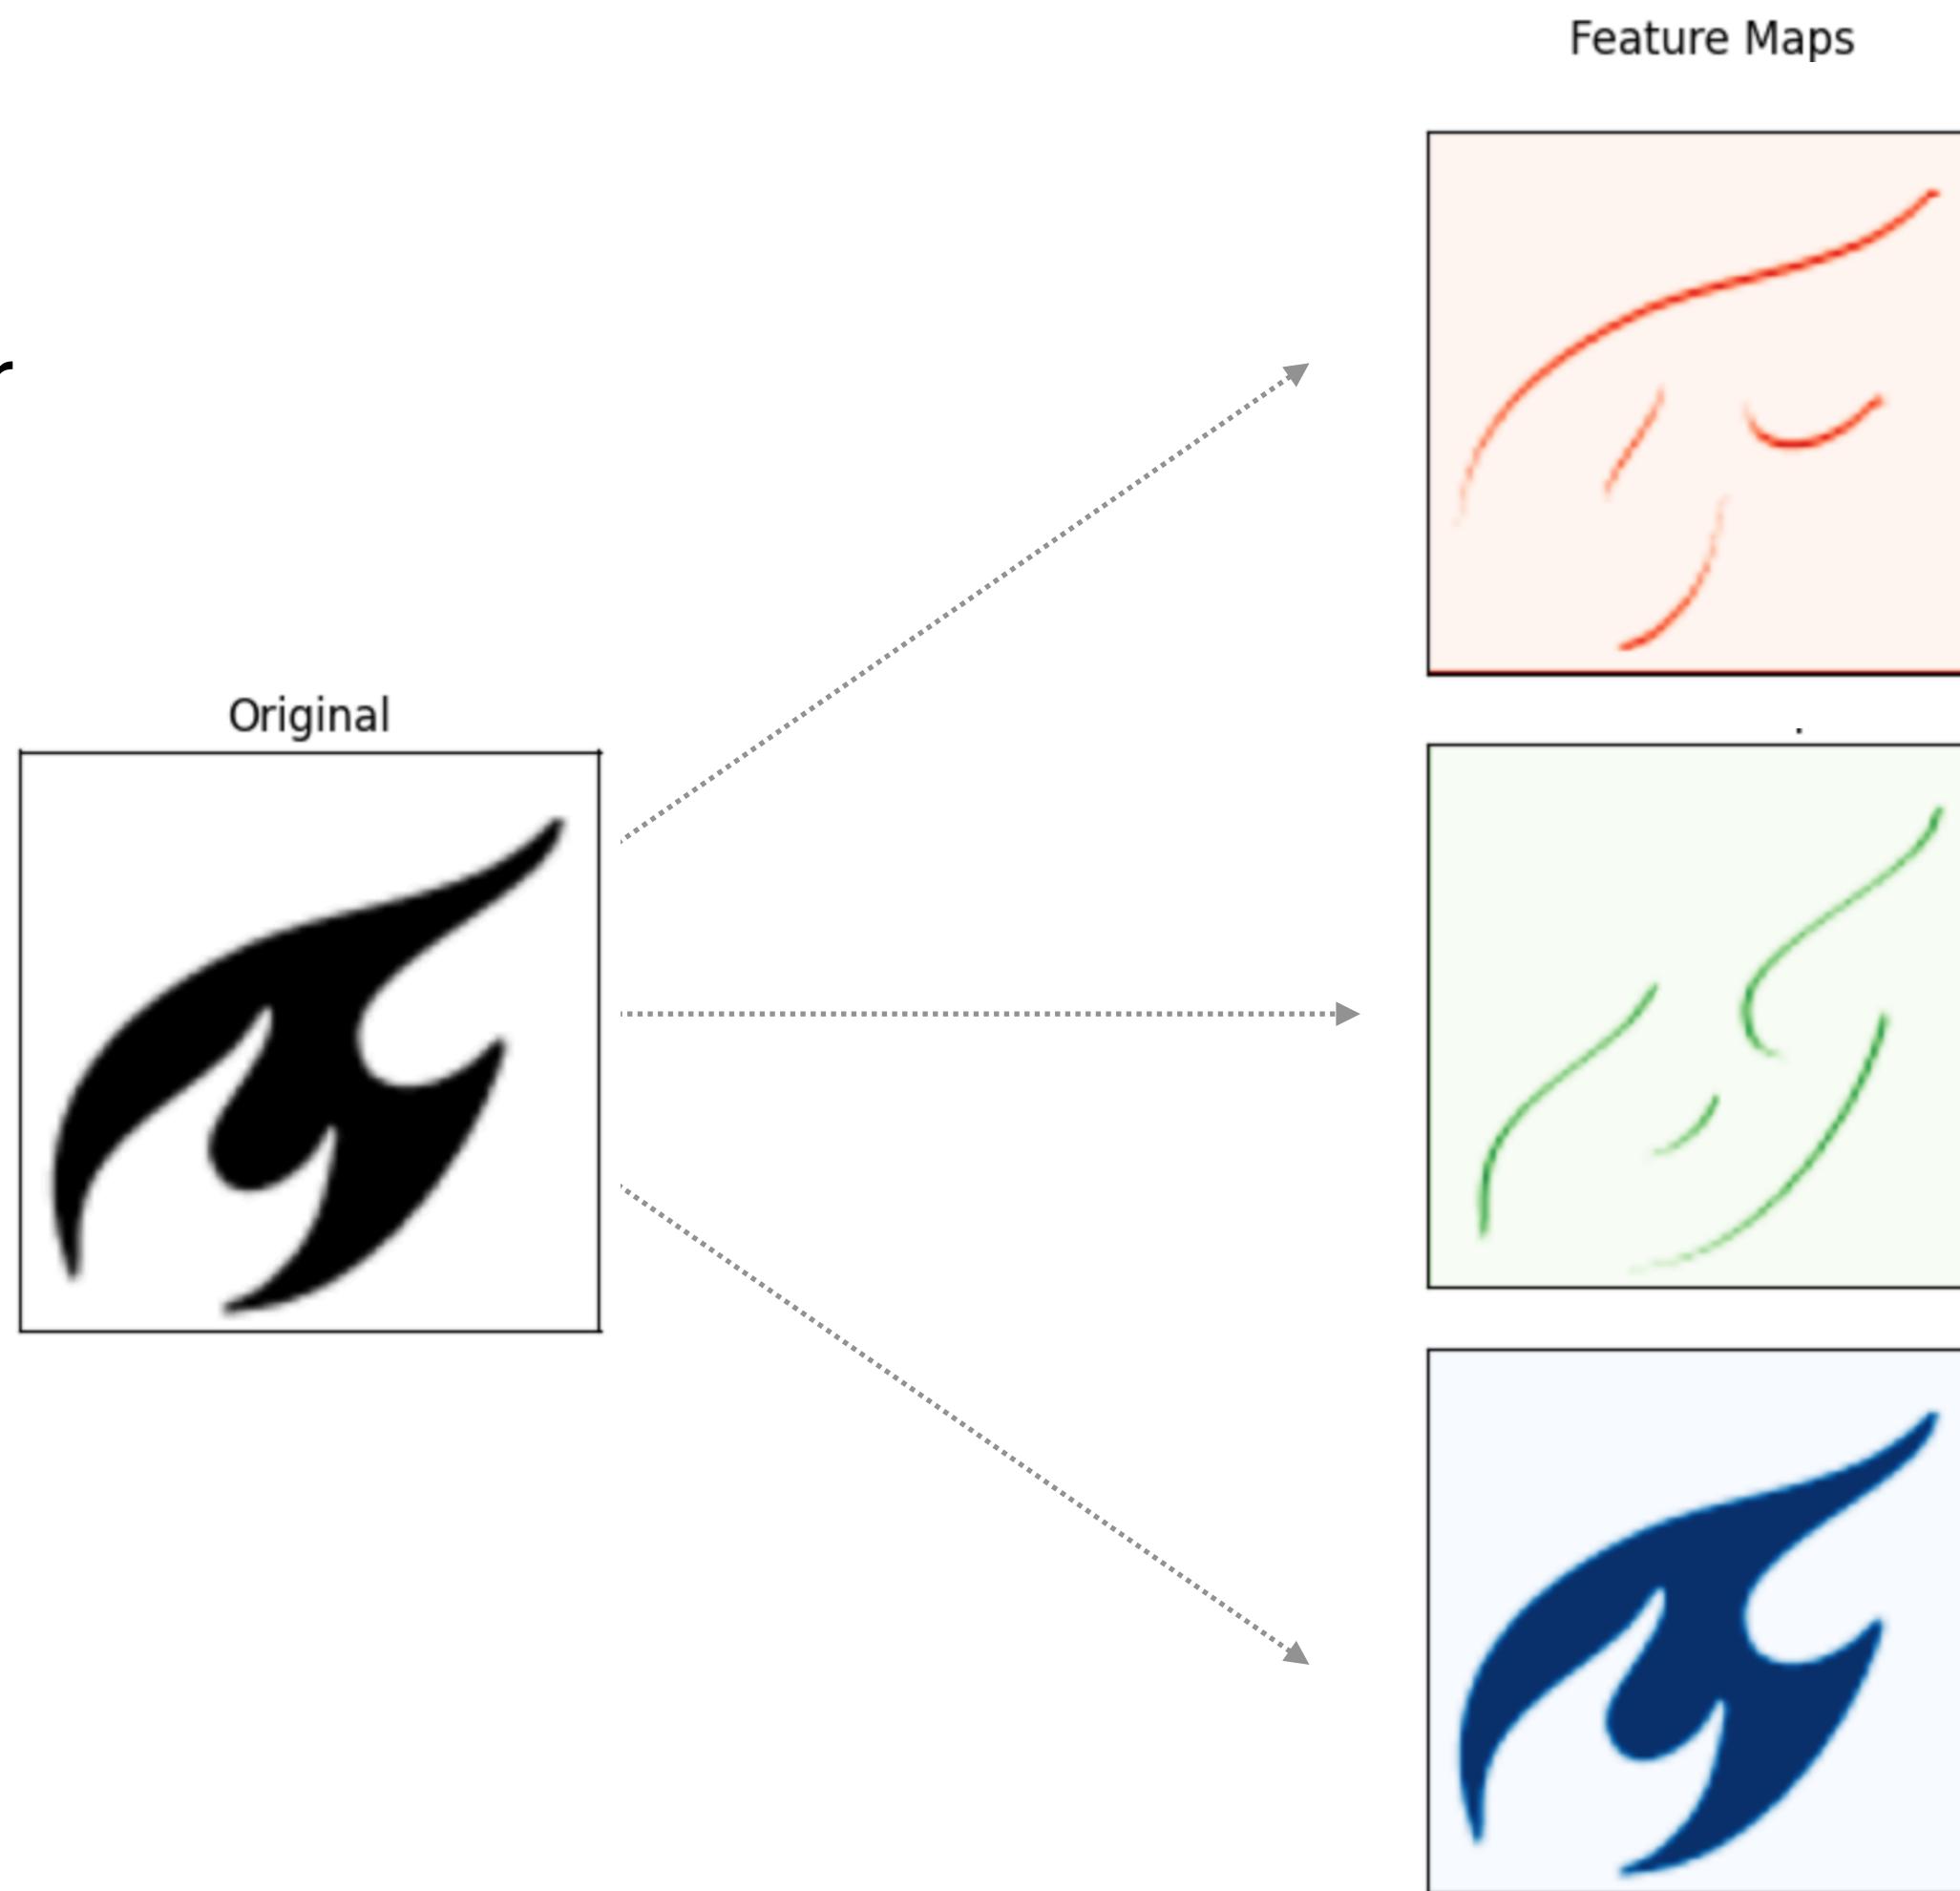
Convolutional layer: channels

- Padding
- Moving Average Filter
- Detector kernel
- ReLU activation
- **Input channels**



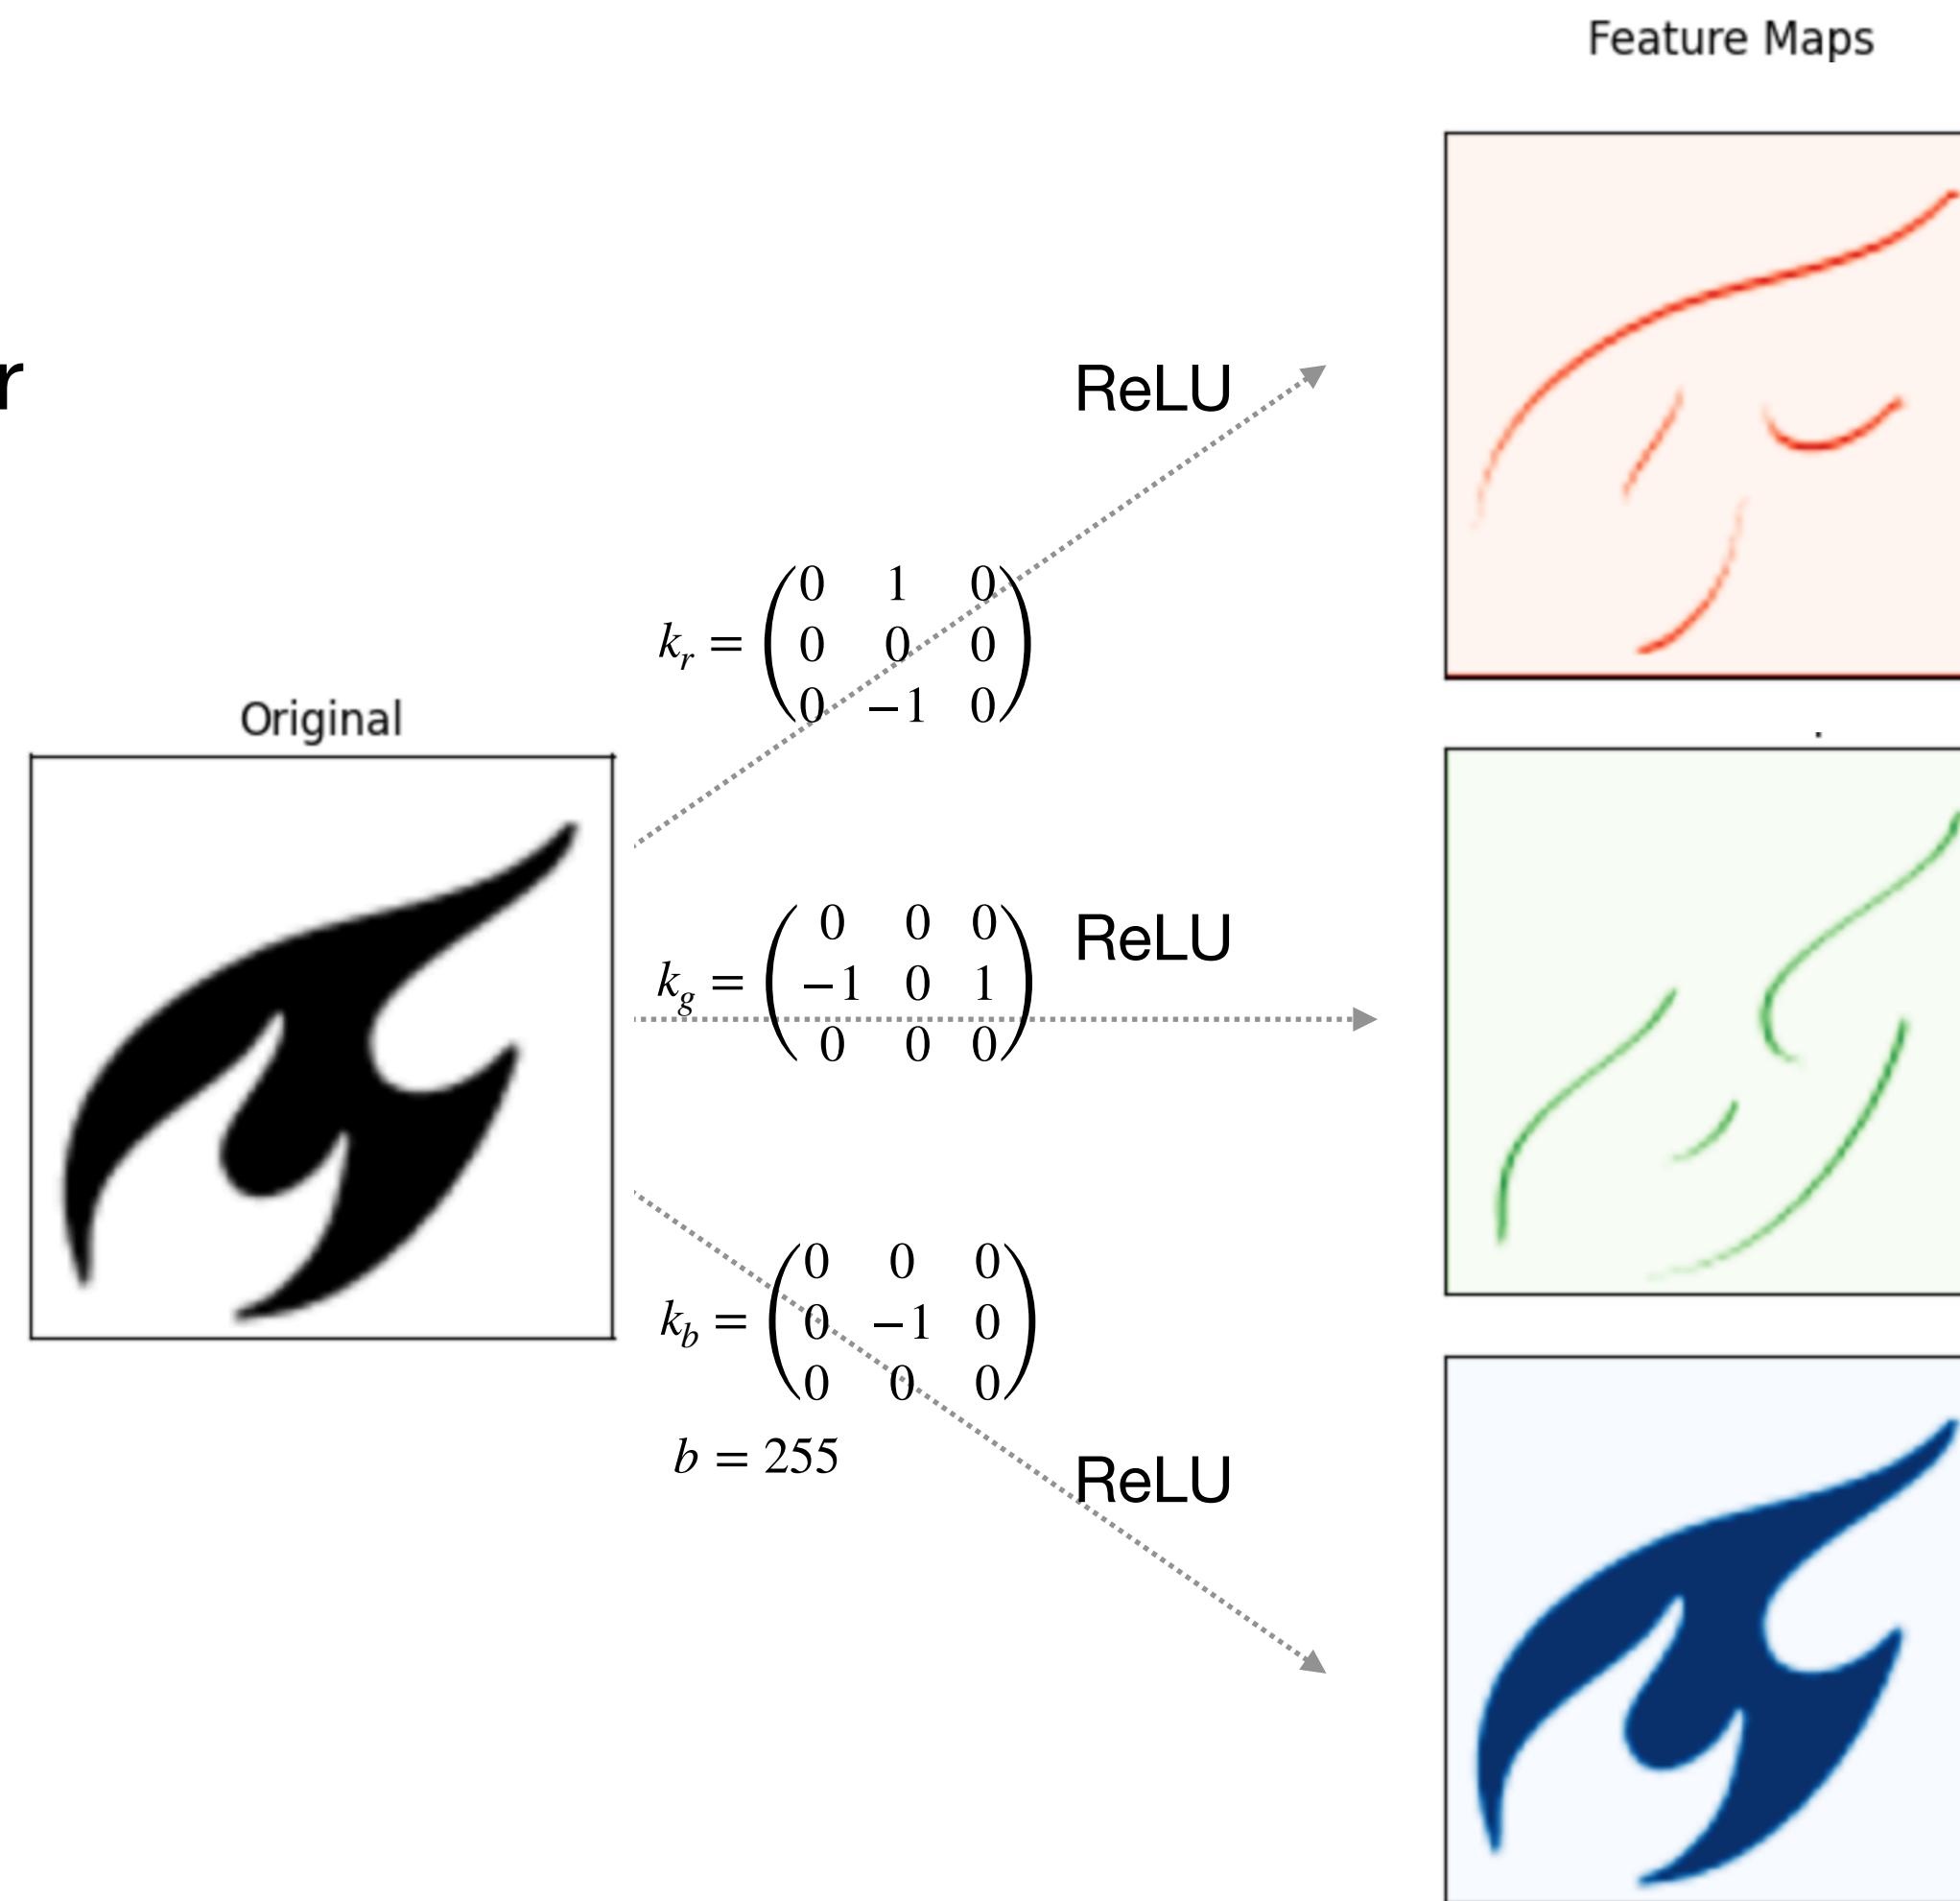
Convolutional layer: feature maps

- Padding
- Moving Average Filter
- Detector kernel
- ReLU activation
- Input channels
- **Feature Maps**



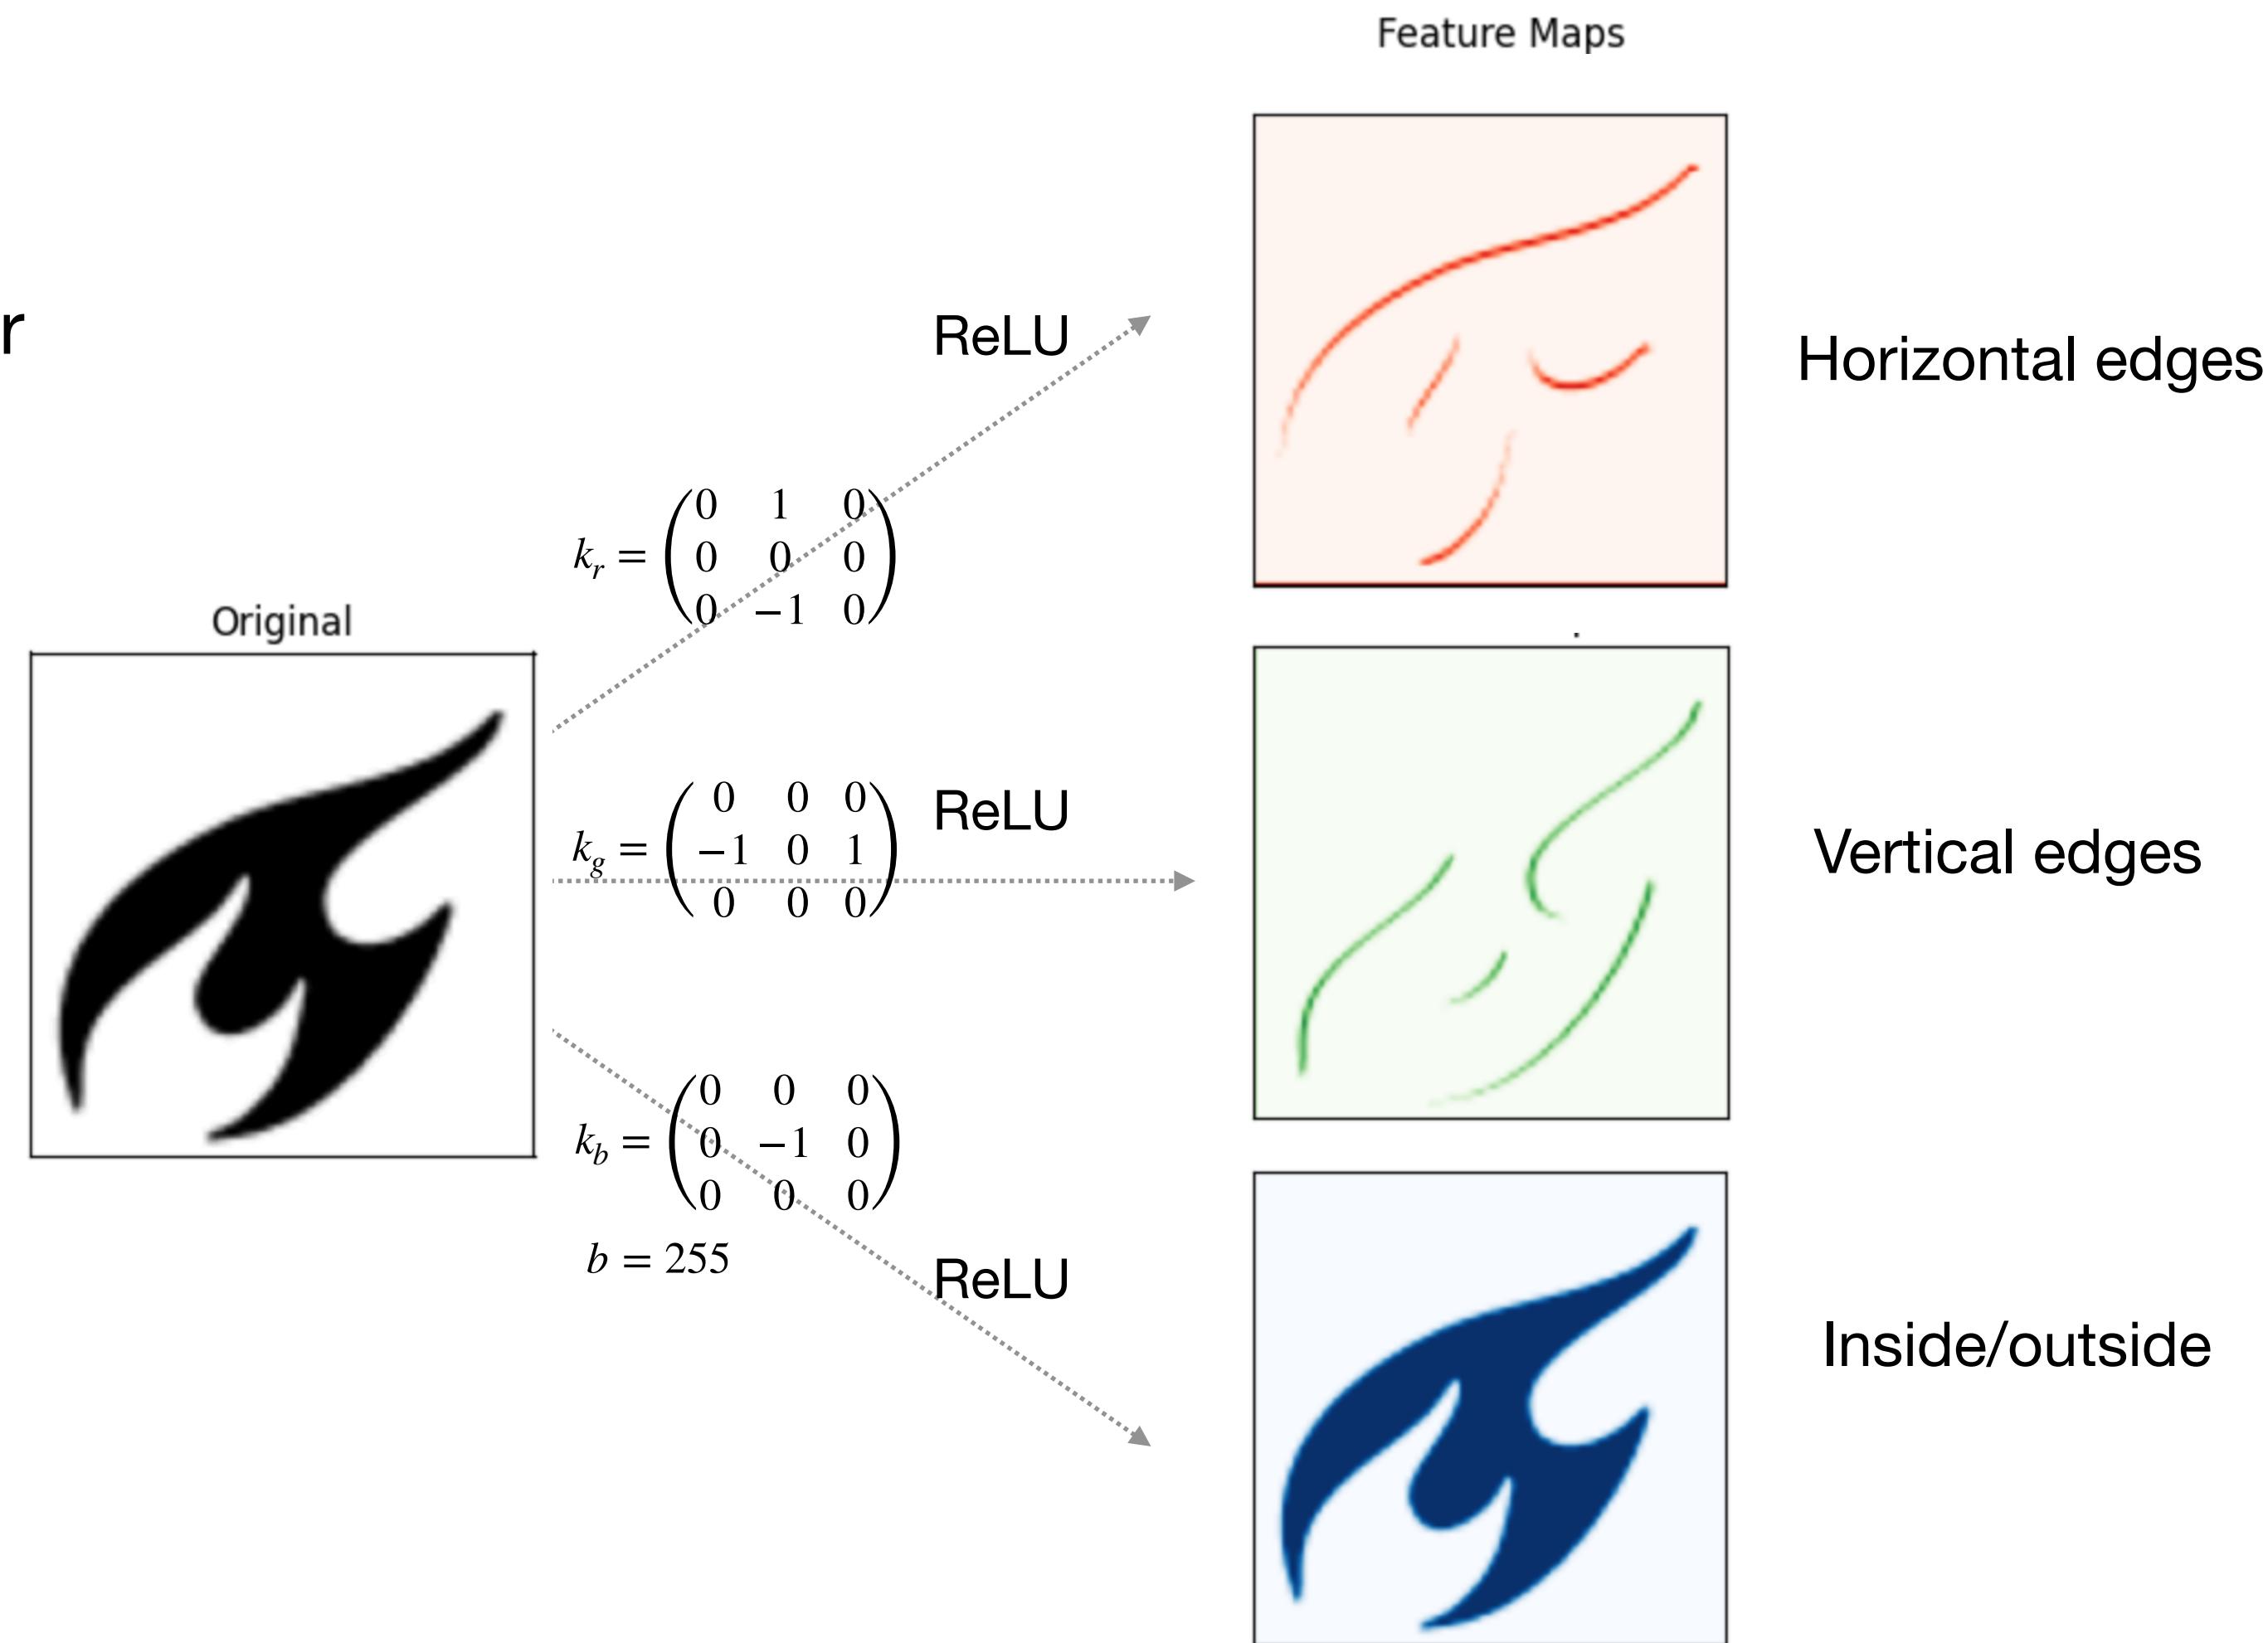
Convolutional layer: feature maps

- Padding
- Moving Average Filter
- Detector kernel
- ReLU activation
- Input channels
- **Feature Maps**



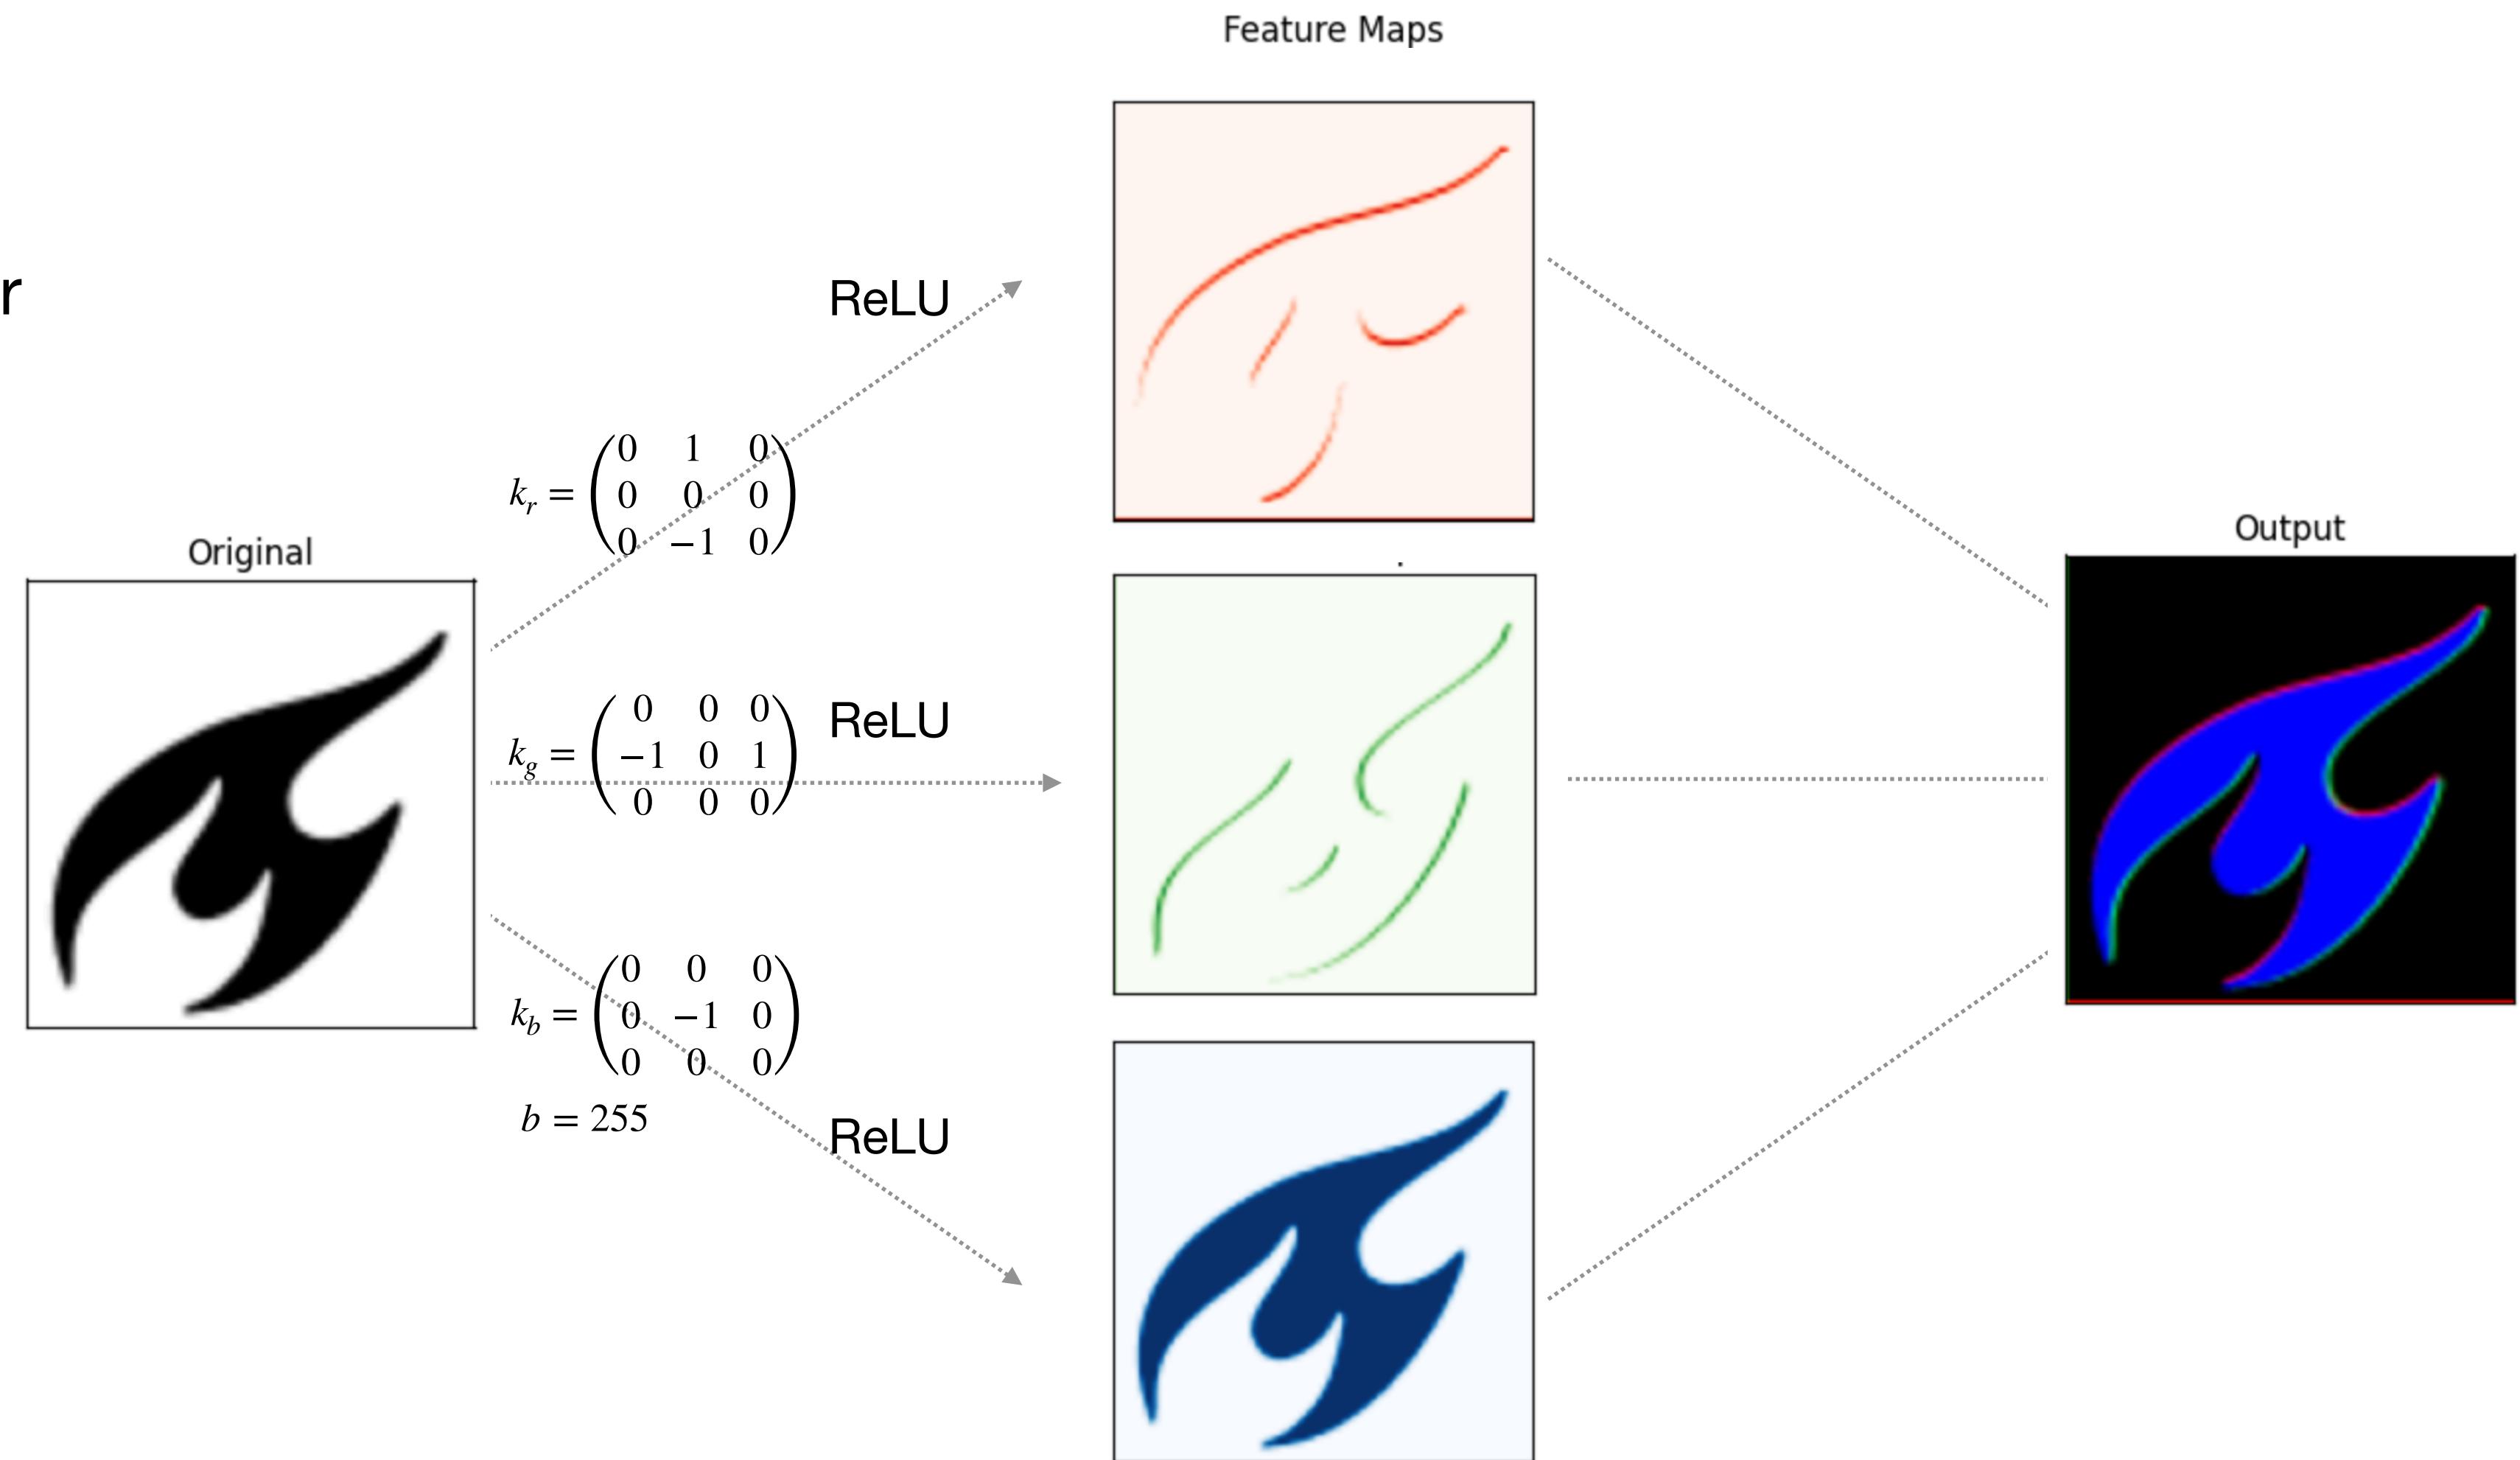
Convolutional layer: feature maps

- Padding
- Moving Average Filter
- Detector kernel
- ReLU activation
- Input channels
- **Feature Maps**



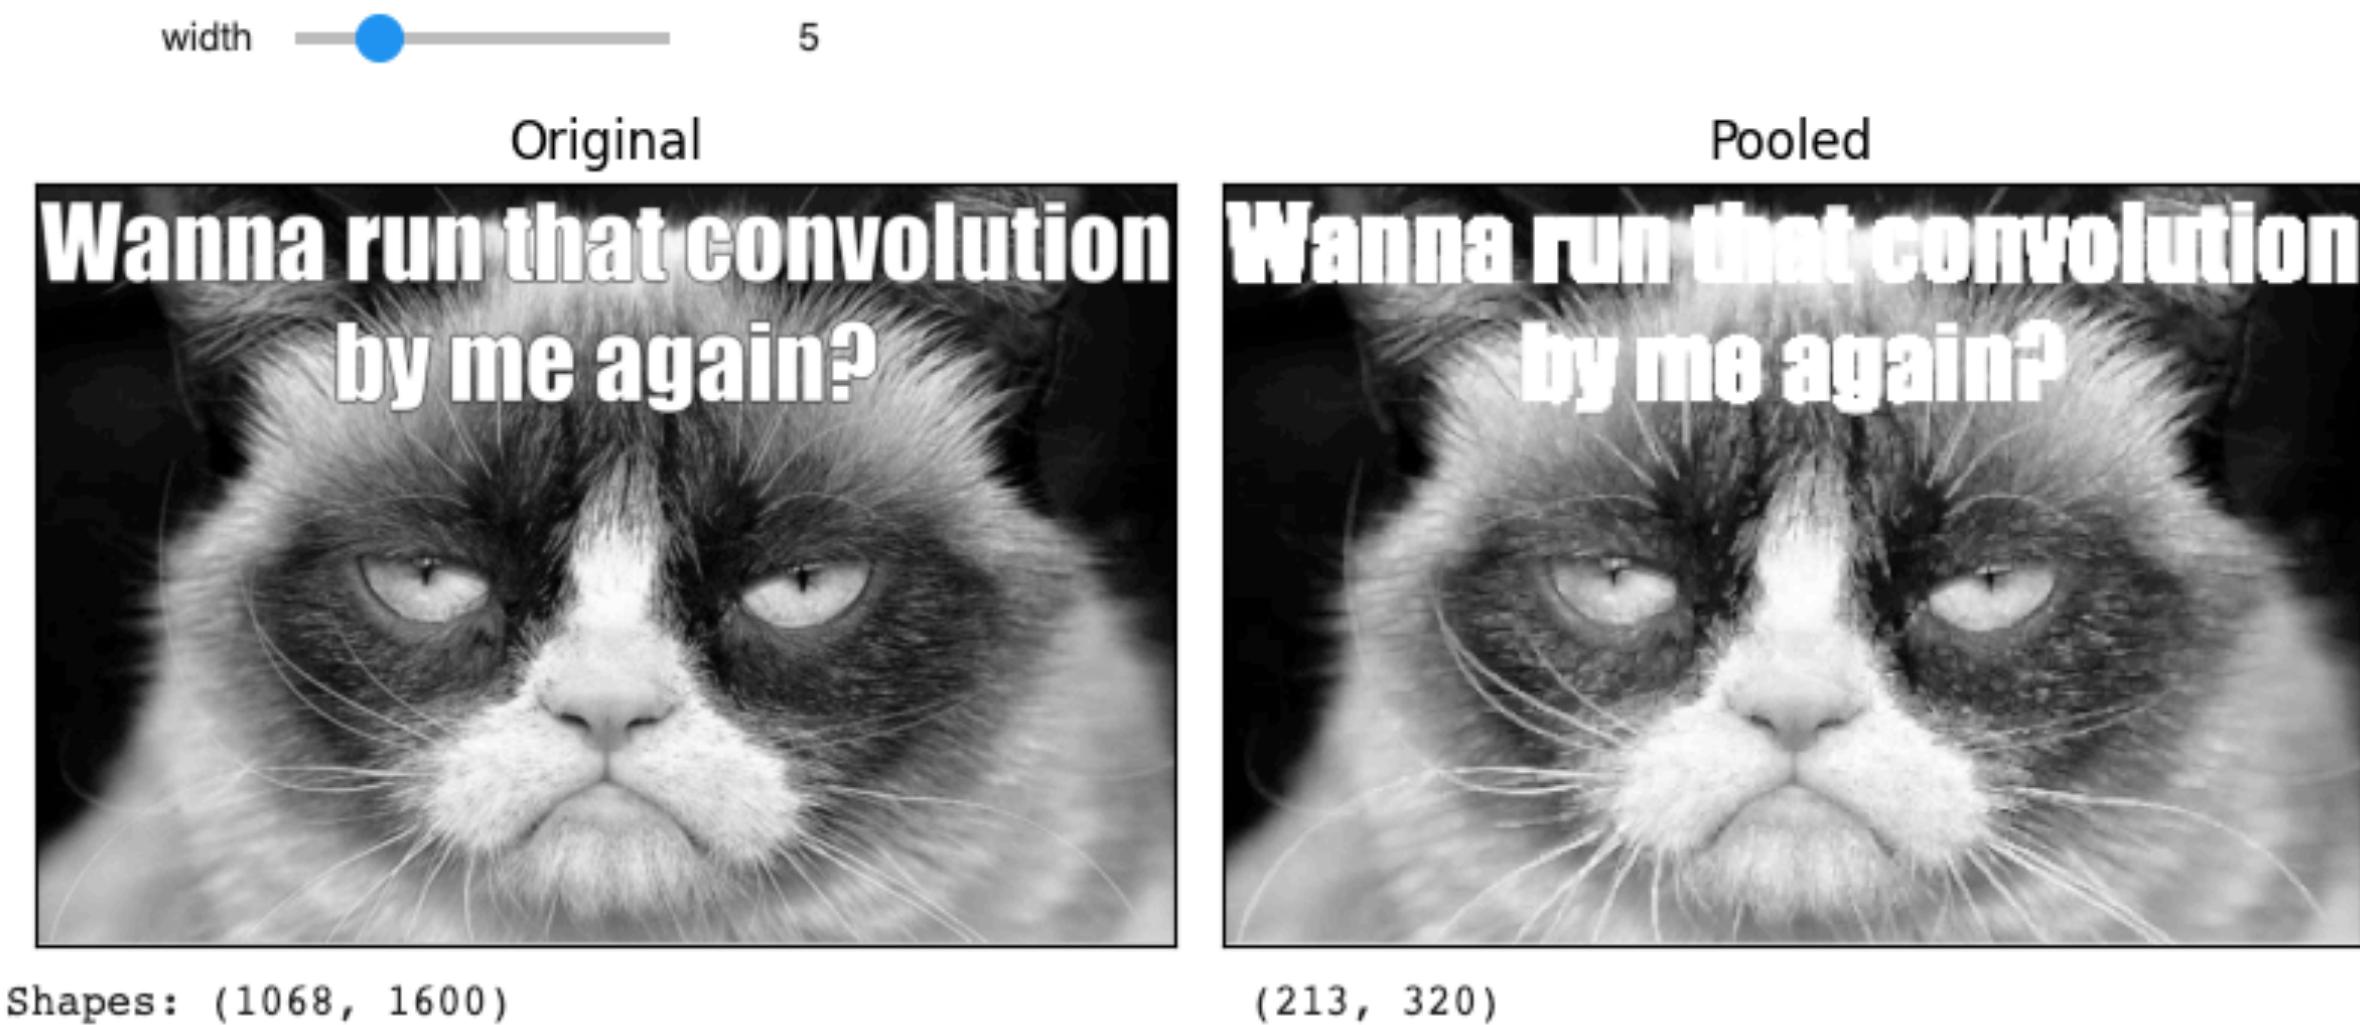
Convolutional layer: feature maps

- Padding
- Moving Average Filter
- Detector kernel
- ReLU activation
- Input channels
- **Feature Maps**



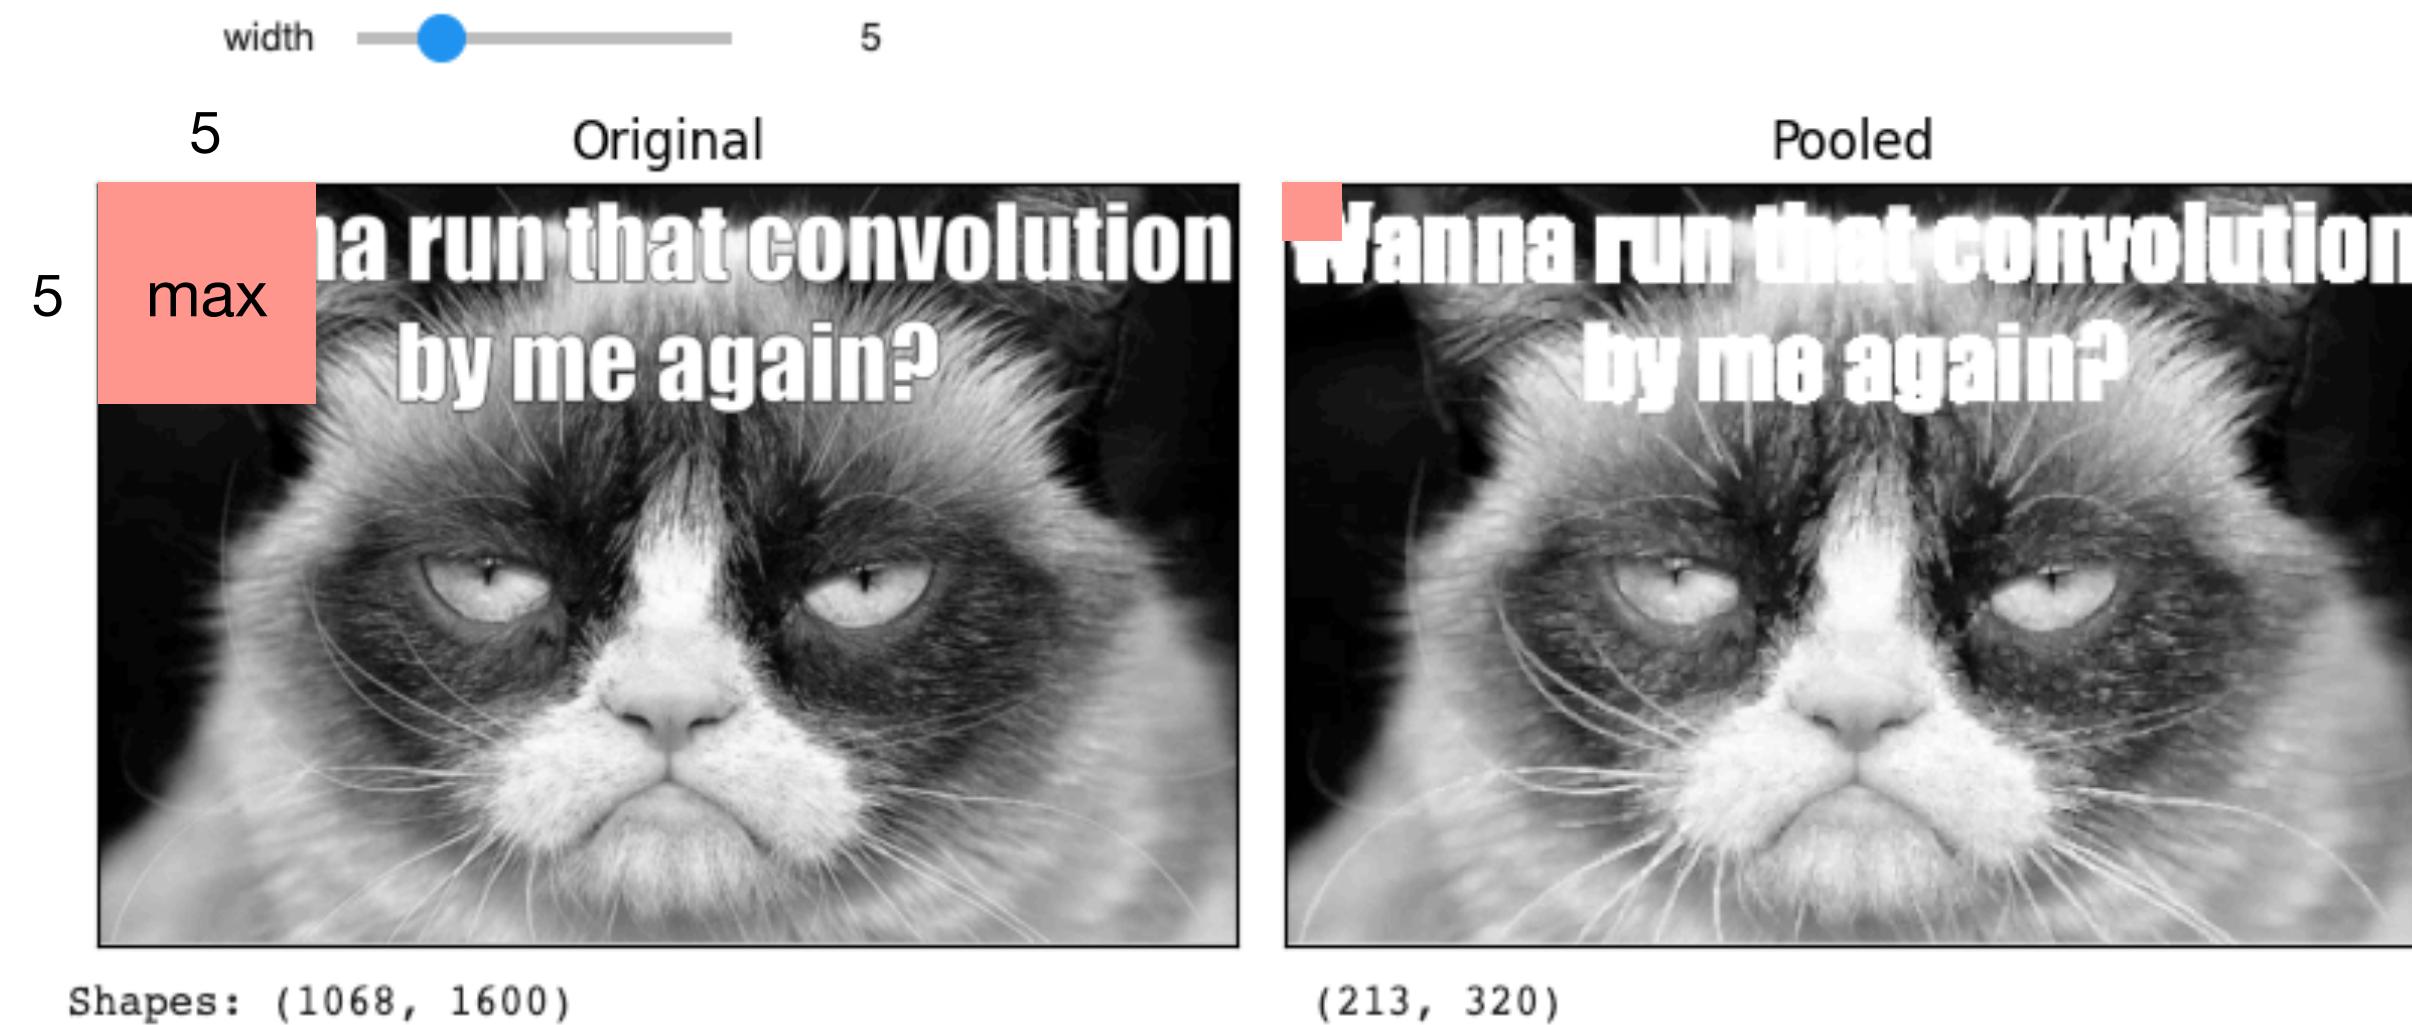
Convolutional layer: pooling

- Padding
- Moving Average Filter
- Detector kernel
- ReLU activation
- Input channels
- Feature Maps
- **Pooling**



Convolutional layer: pooling

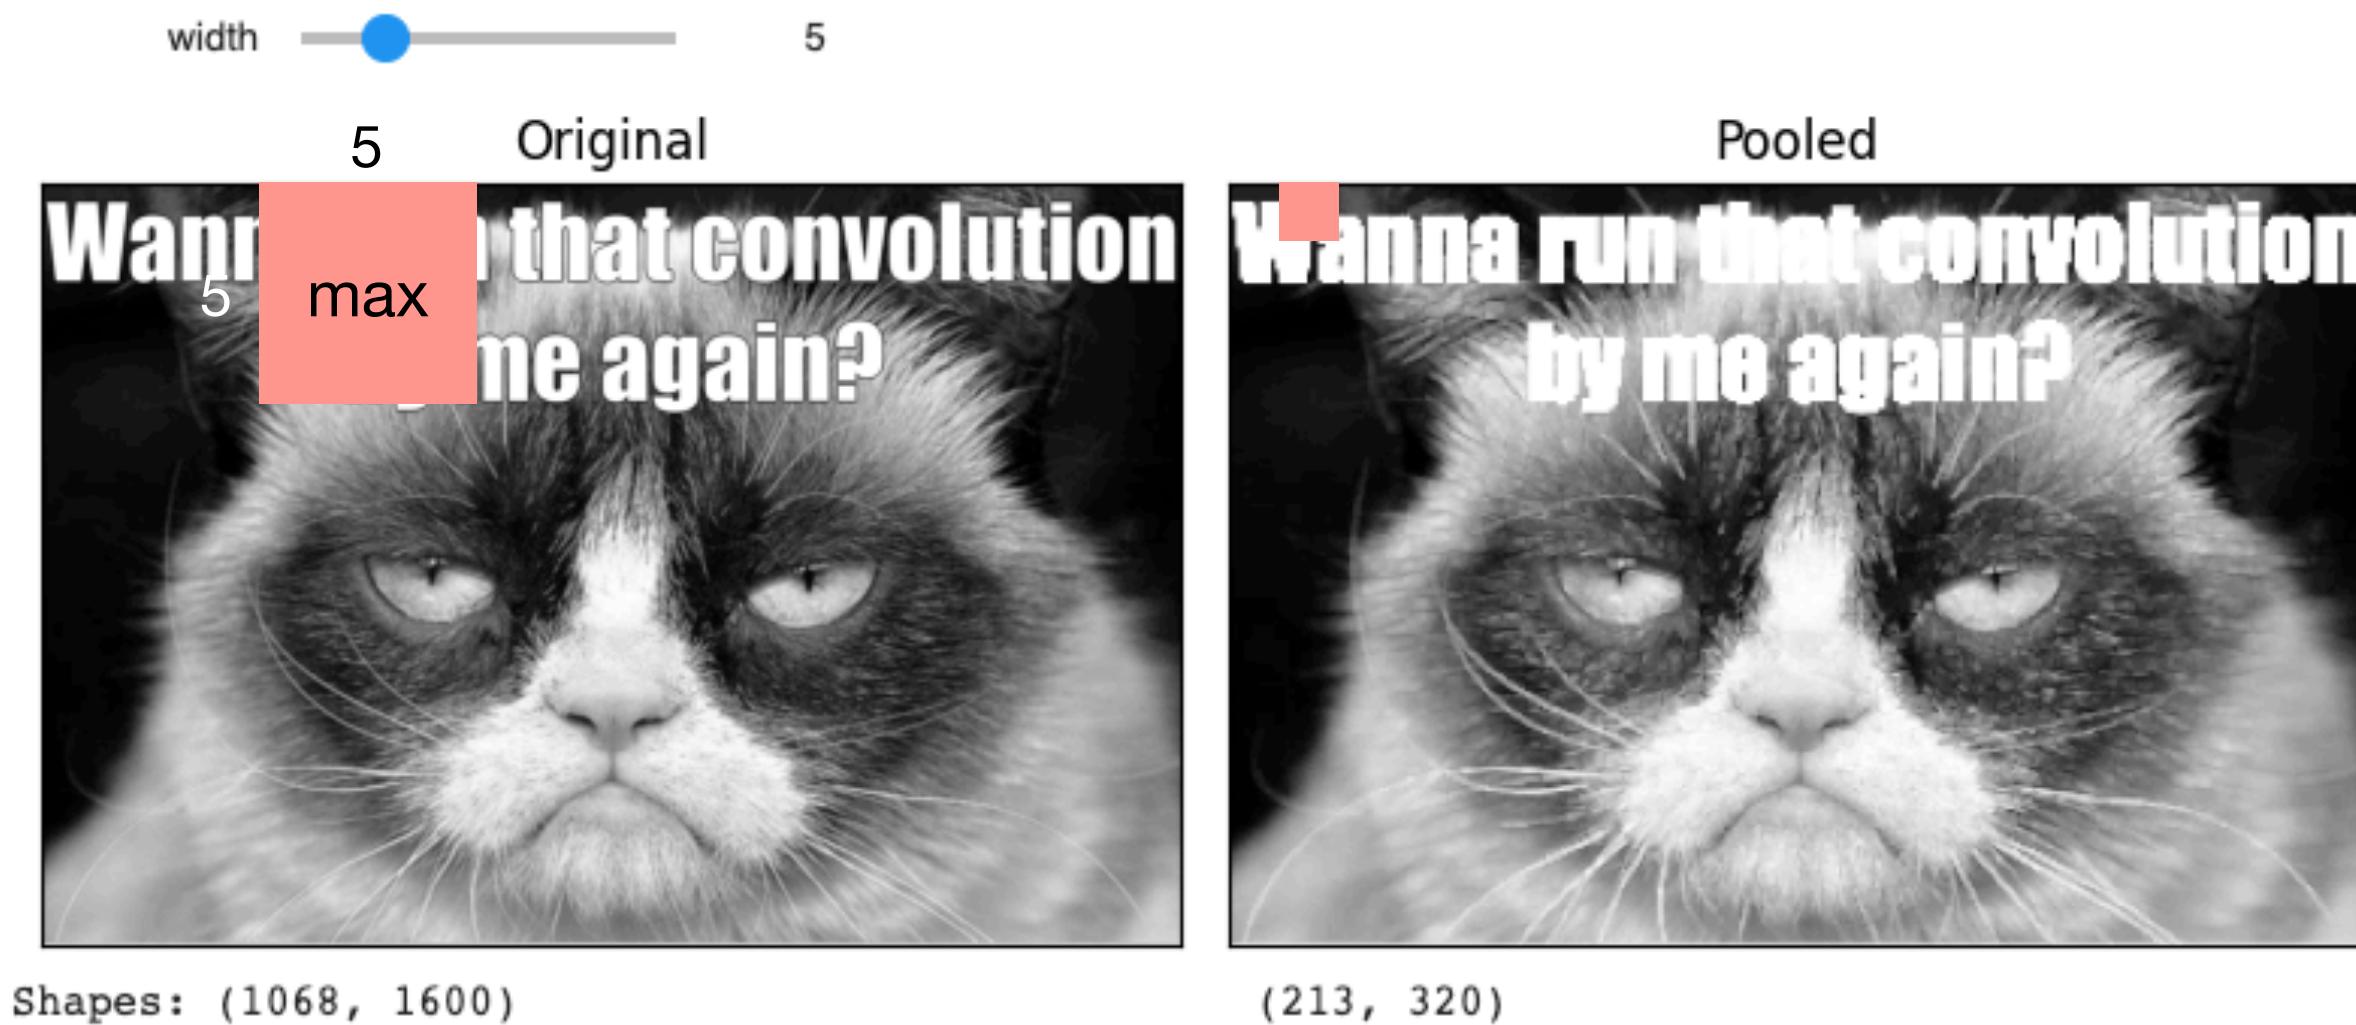
- Padding
- Moving Average Filter
- Detector kernel
- ReLU activation
- Input channels
- Feature Maps
- **Pooling**



Max Pooling preserves the features of the cat

Convolutional layer: pooling

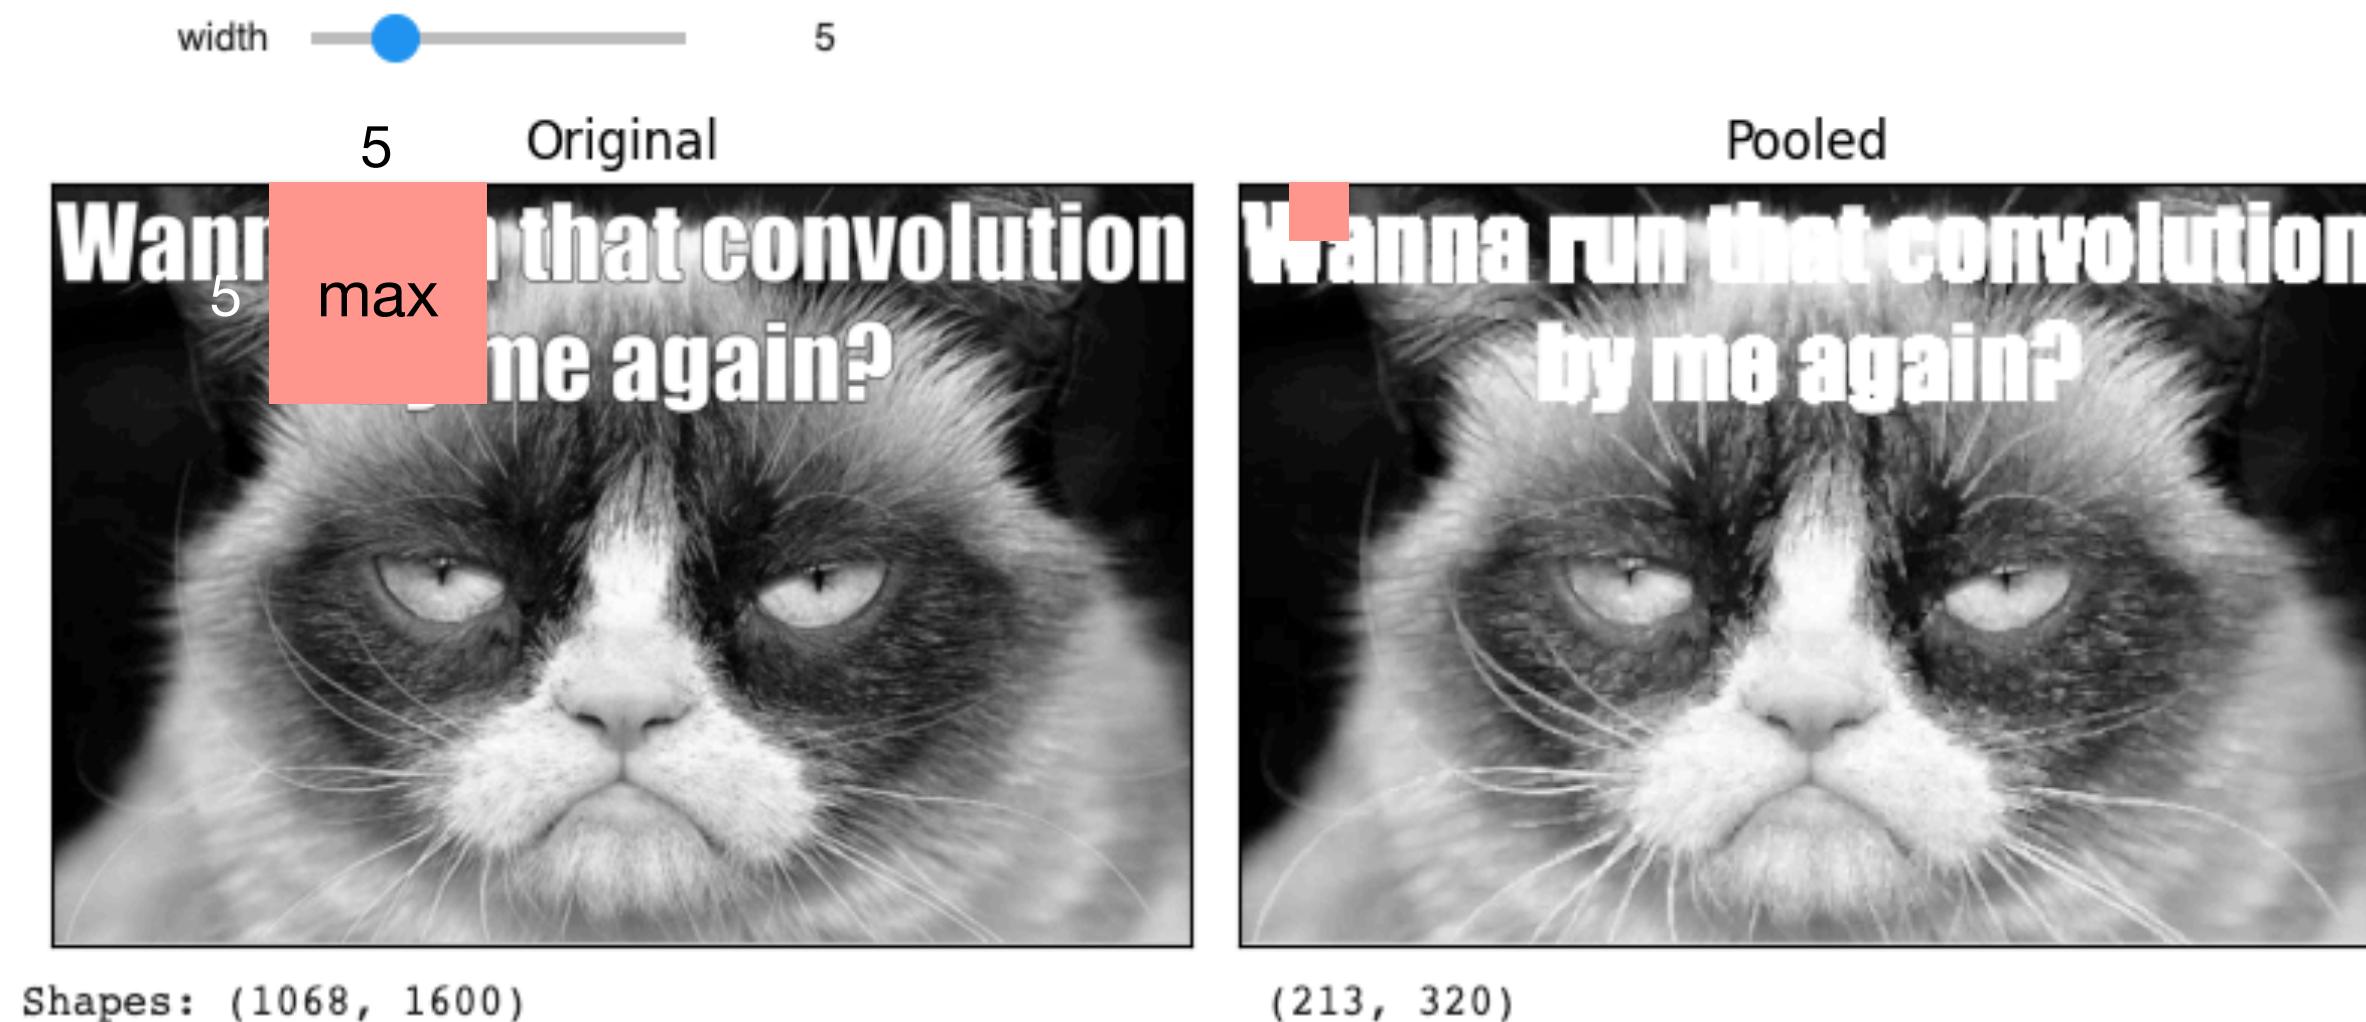
- Padding
- Moving Average Filter
- Detector kernel
- ReLU activation
- Input channels
- Feature Maps
- **Pooling**



Max Pooling preserves the features of the cat

Convolutional layer: pooling

- Padding
- Moving Average Filter
- Detector kernel
- ReLU activation
- Input channels
- Feature Maps
- **Pooling**

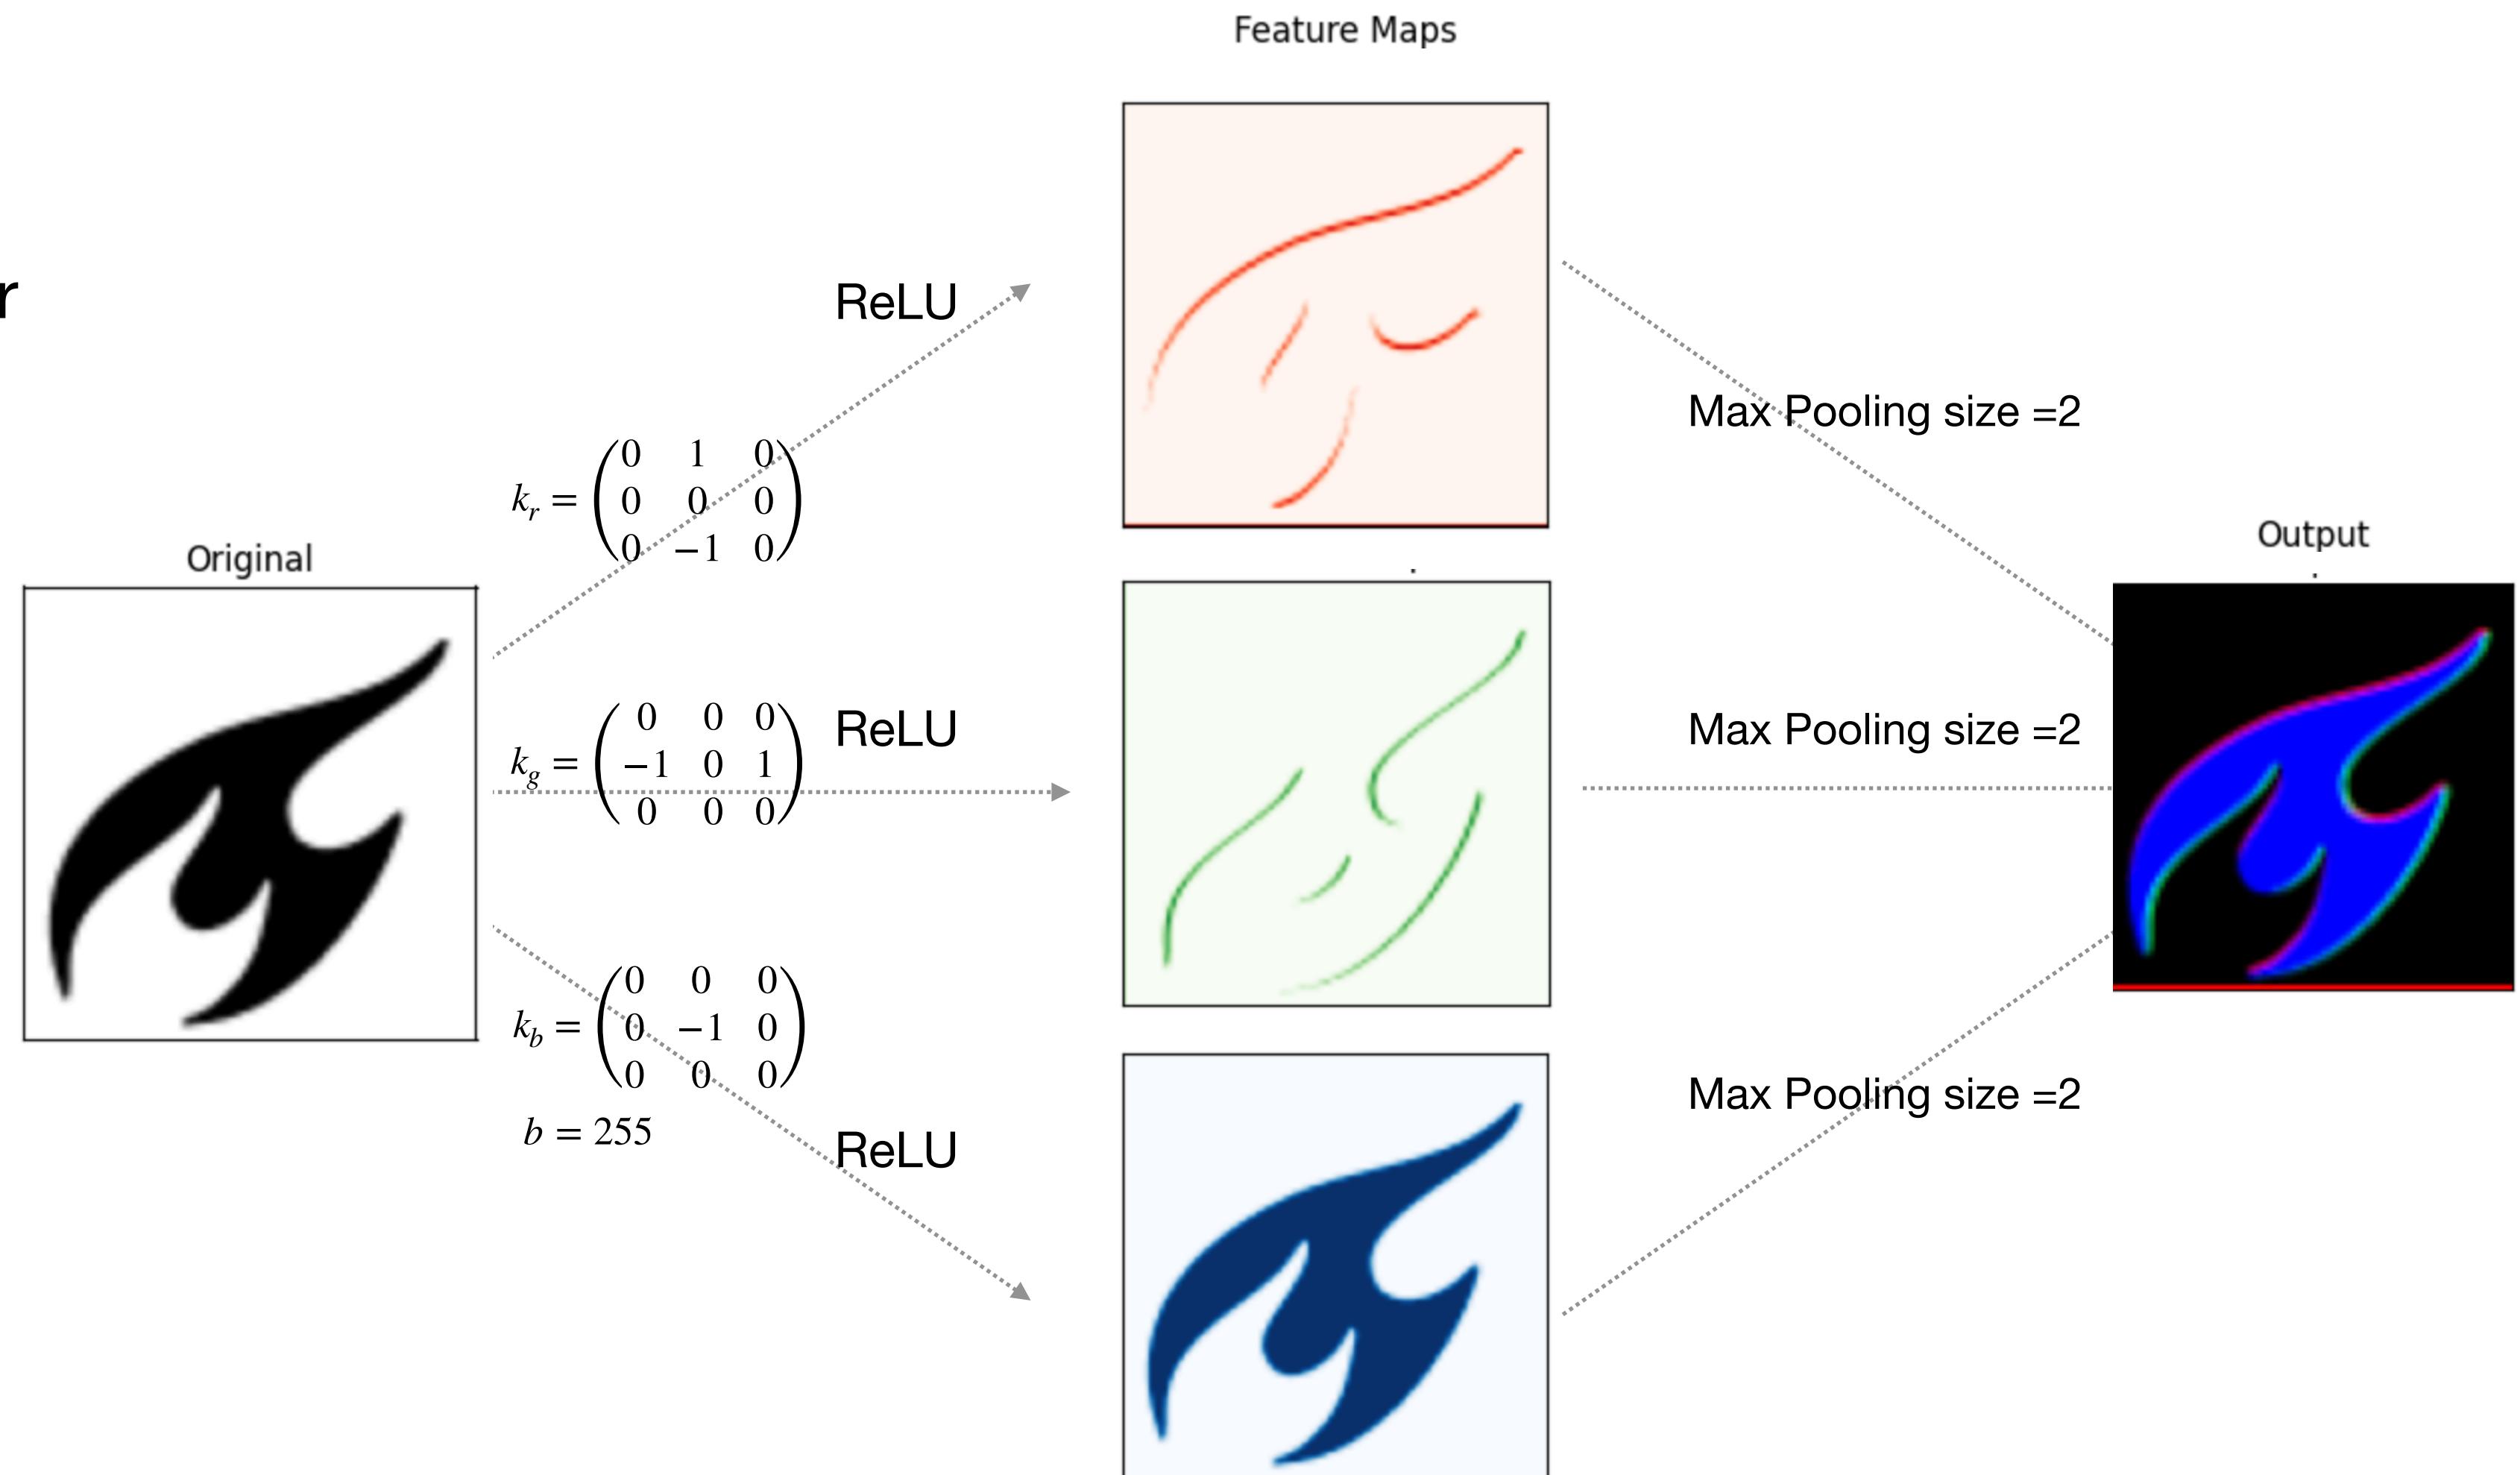


Max Pooling preserves the features of the cat

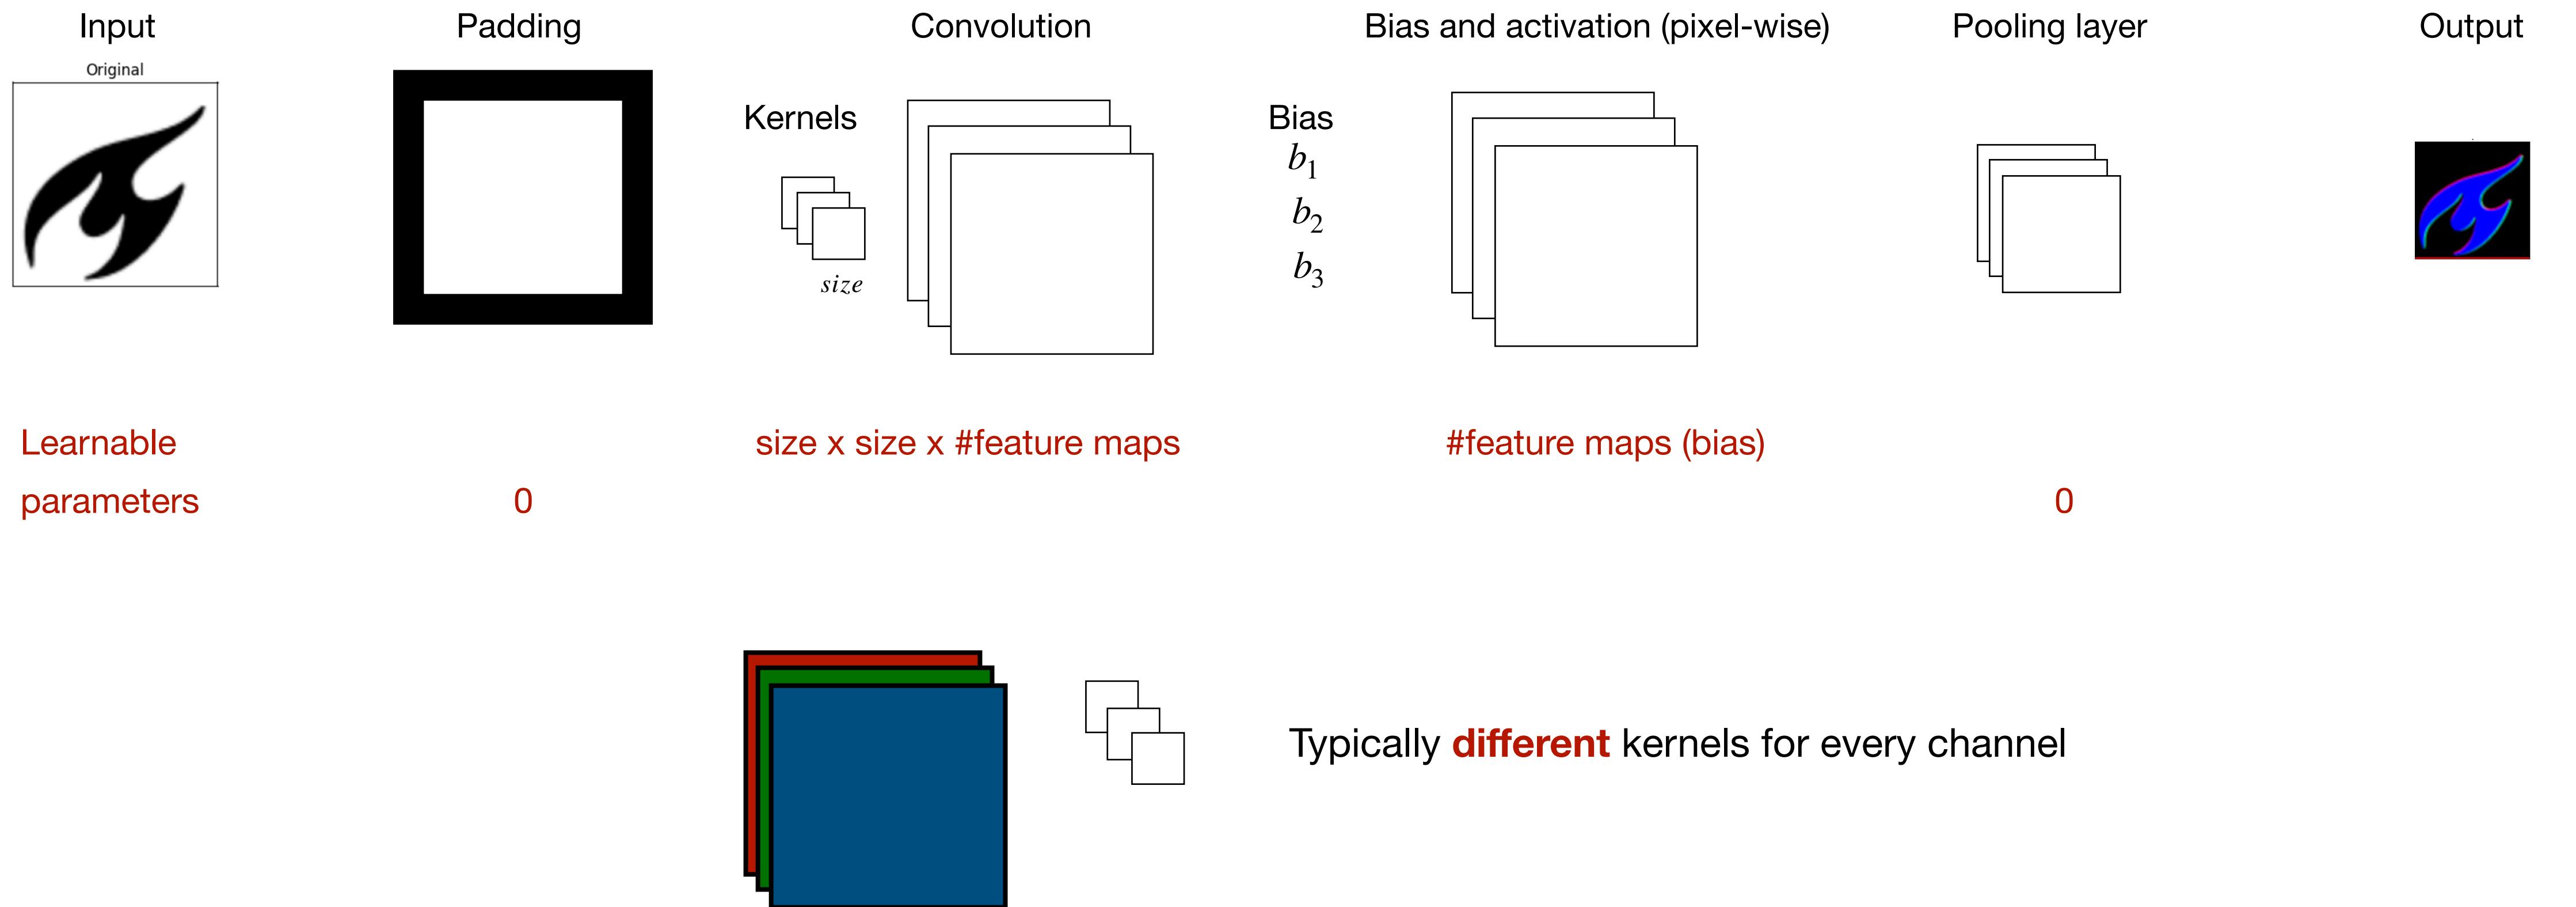
- Representation invariance to small translations
- Reduce memory

Convolutional layer

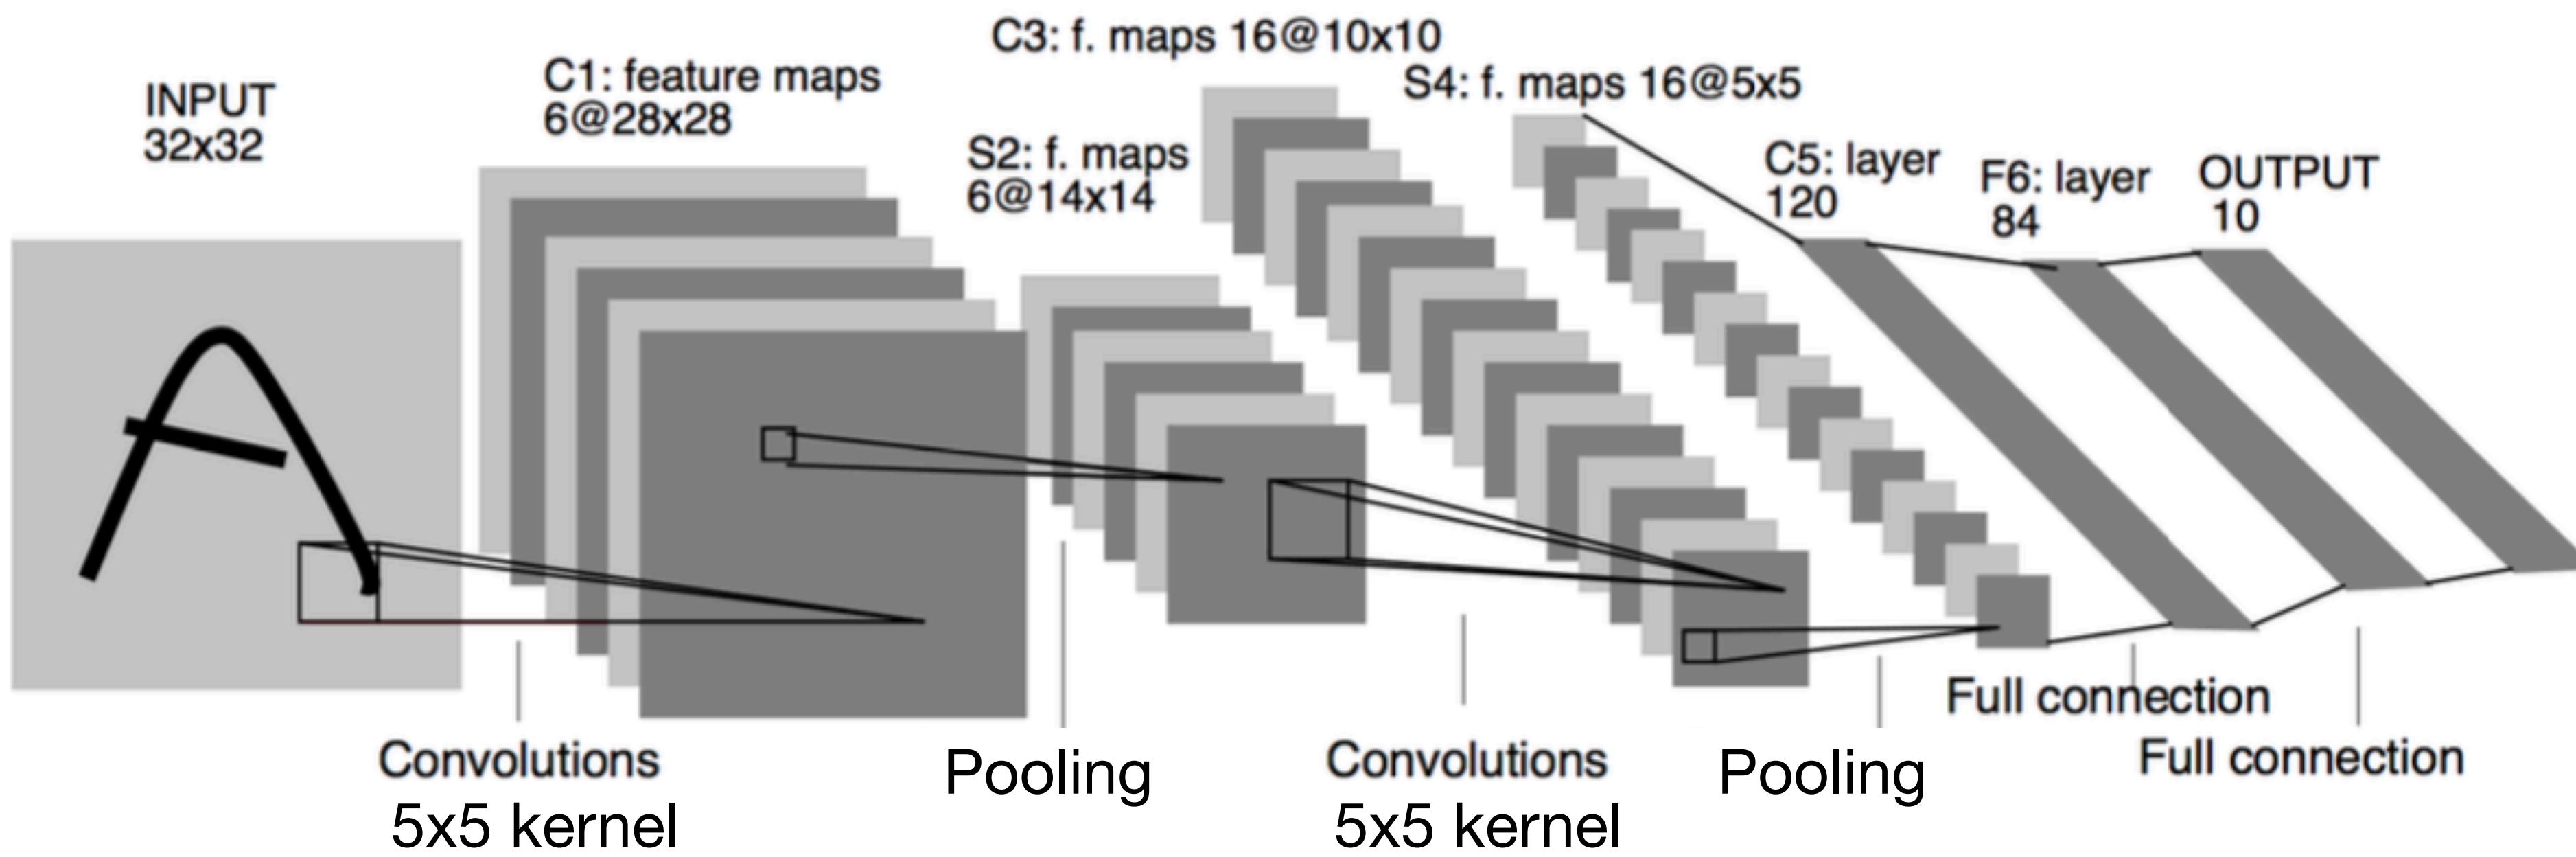
- Padding
- Moving Average Filter
- Detector kernel
- ReLU activation
- Input channels
- Feature Maps
- Pooling



Convolutional Layer



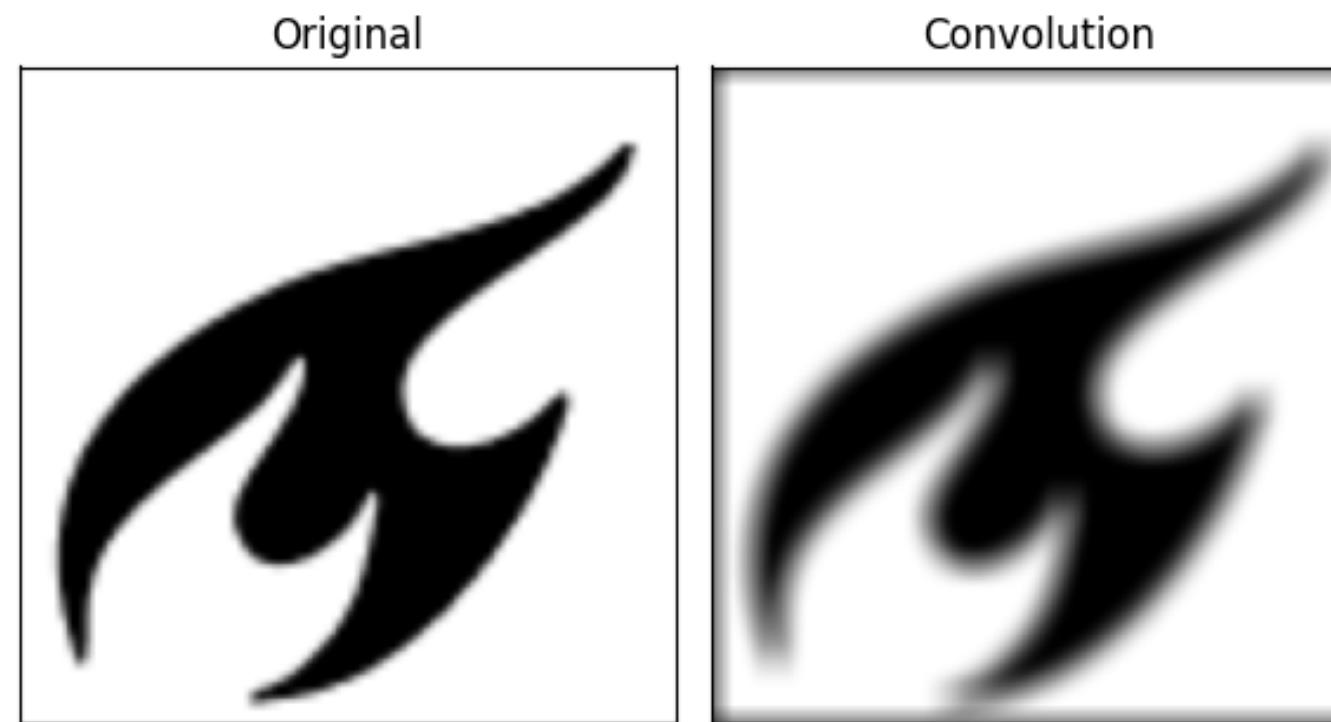
CNN Architecture



LeNet - 5 (1998)

Stride and Dilation

Bigger images require larger kernel size → more parameters to learn



kernel size = 10

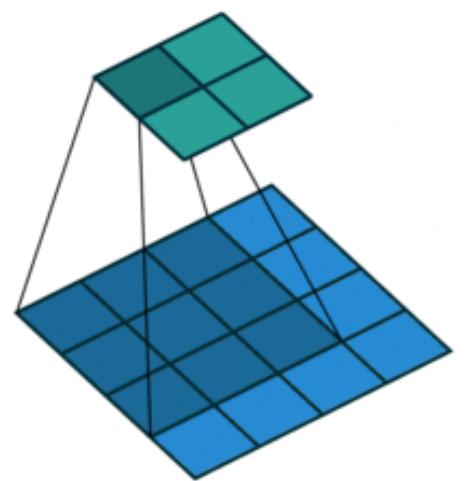


kernel size = 30

Stride and Dilation

Bigger images require larger kernel size → more parameters to learn

Output

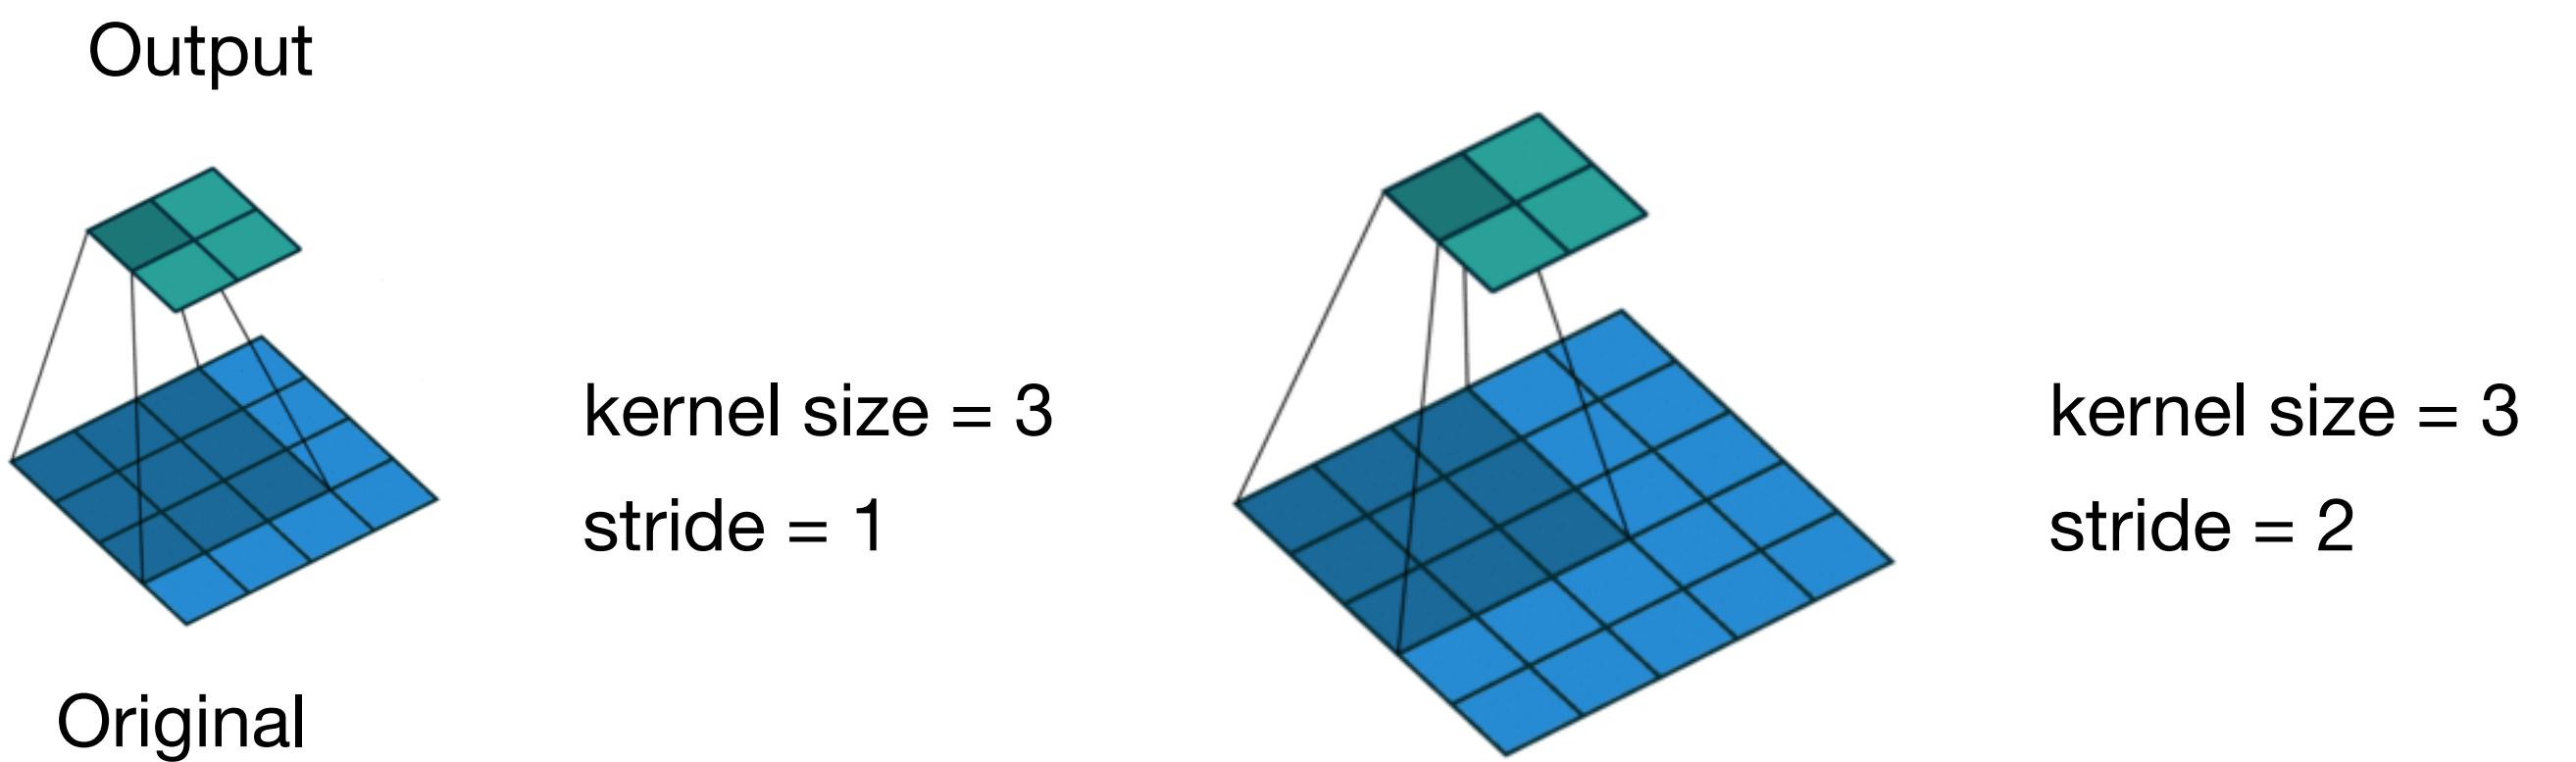


kernel size = 3 (3x3)
stride = 1

Original

Stride and Dilation

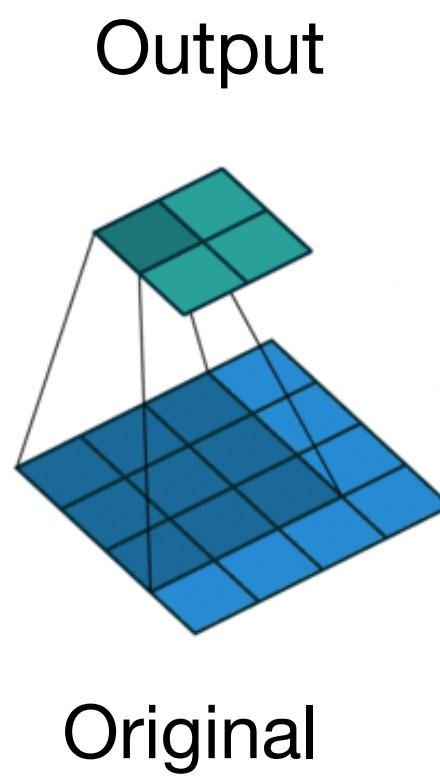
Bigger images require larger kernel size → more parameters to learn



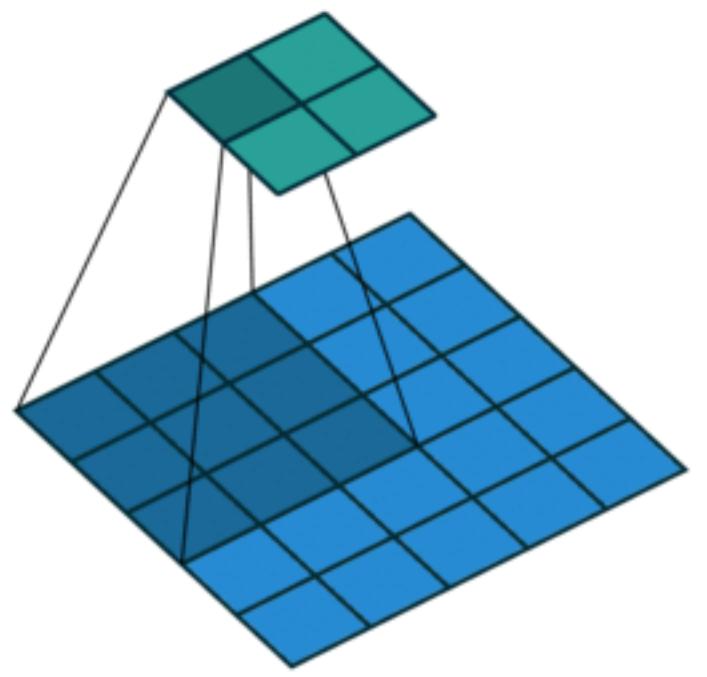
Remember, we have already seen stride for pooling!

Stride and Dilation

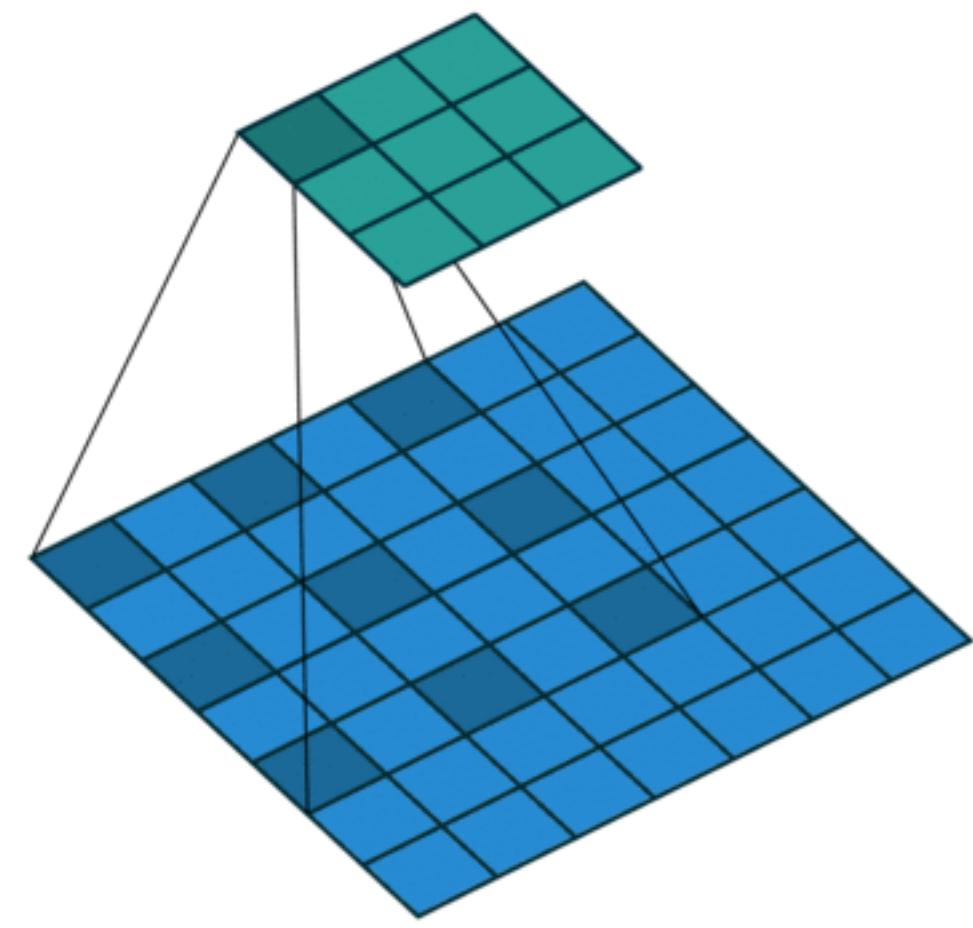
Bigger images require larger kernel size → more parameters to learn



kernel size = 3
stride = 1
dilation = 1



kernel size = 3
stride = 2
dilation = 1

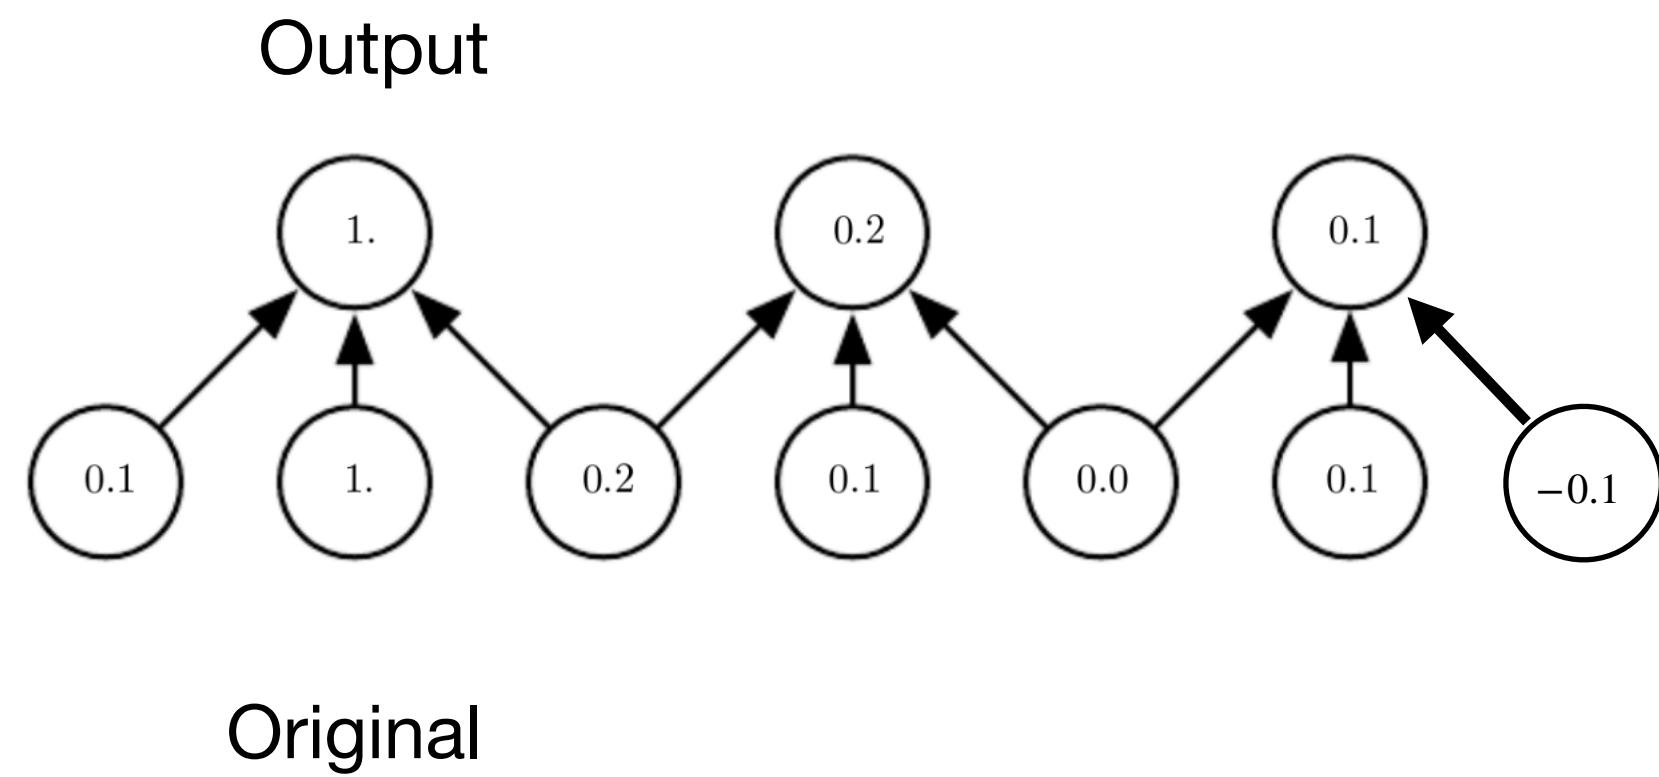


kernel size = 3
stride = 1
dilation = 2

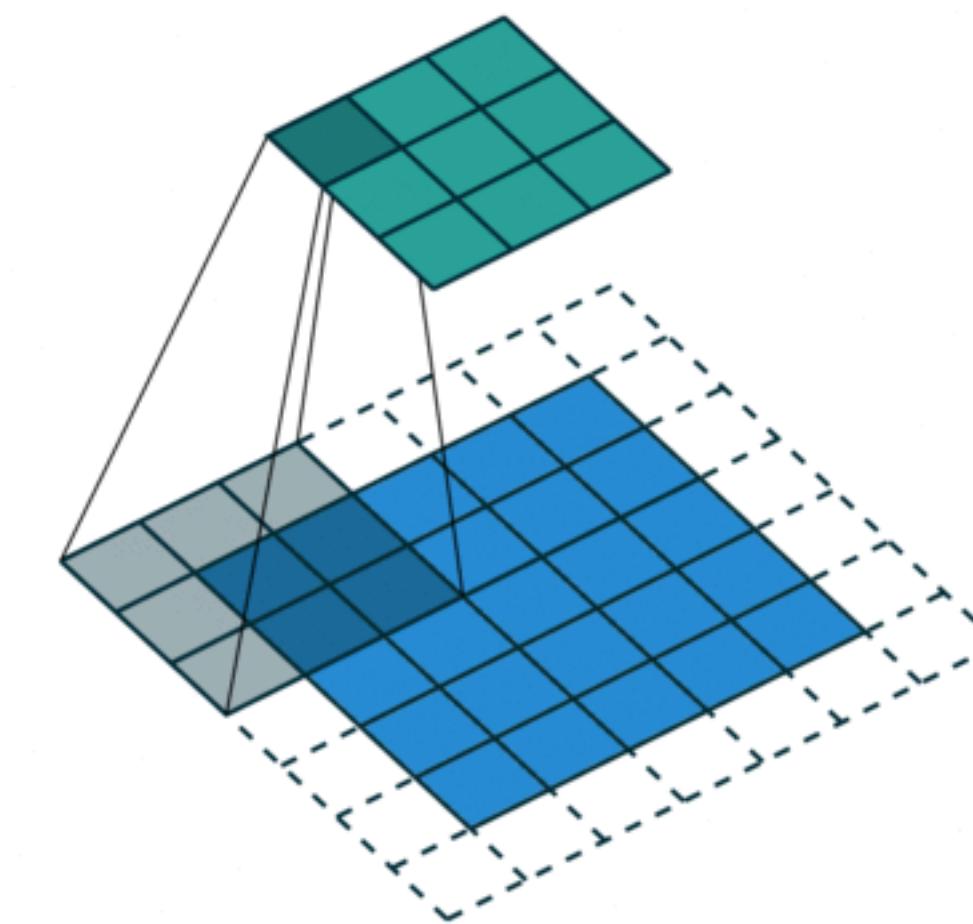
*Animation from Dumoulin, Vincent, and Francesco Visin. "A guide to convolution arithmetic for deep learning." (2016) https://github.com/vdumoulin/conv_arithmetic/blob/master/README.md

Stride and Dilation

Exercise



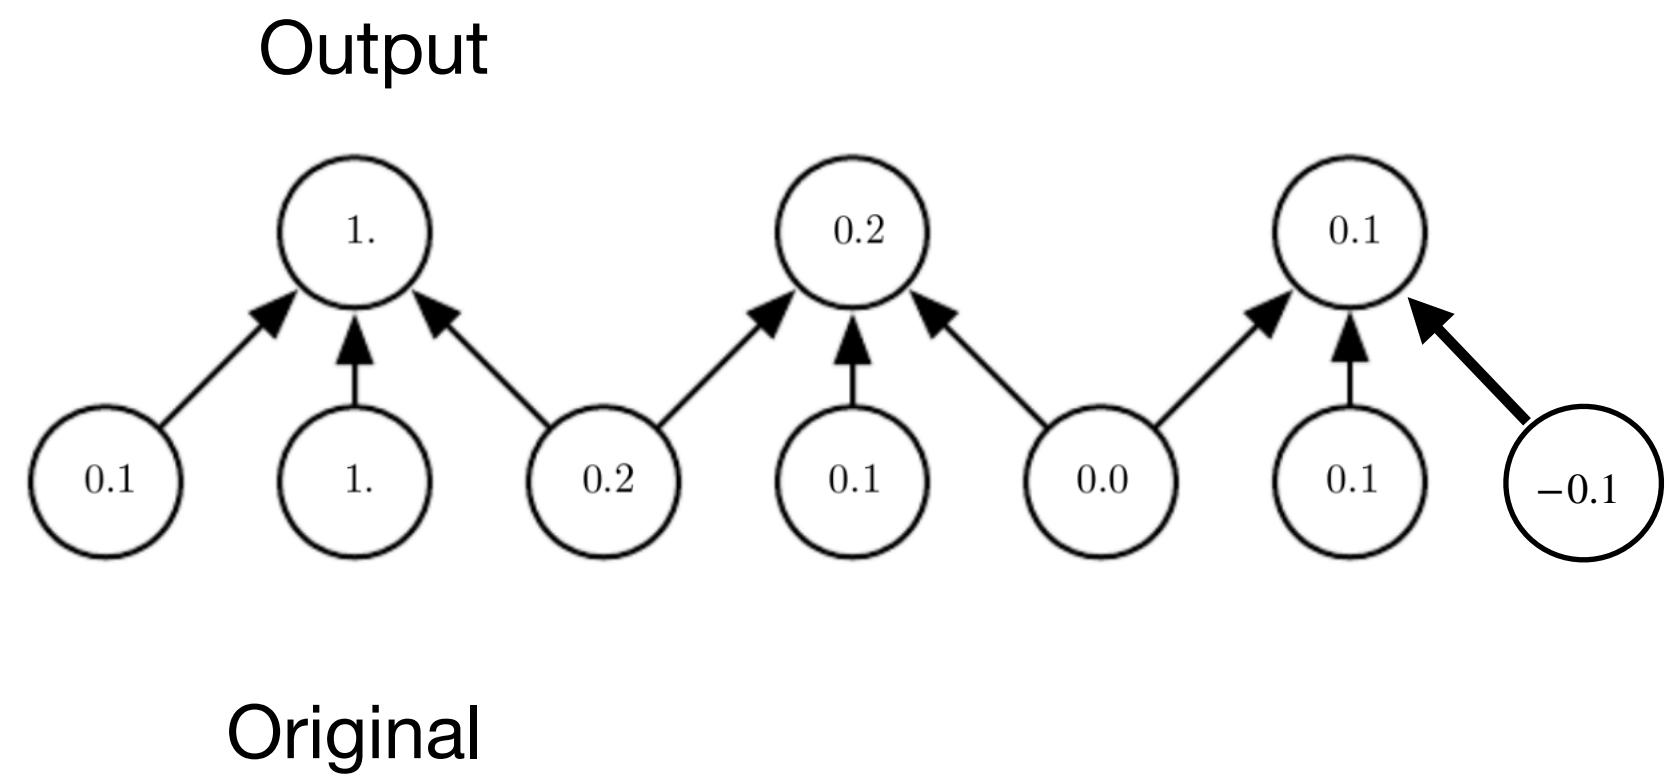
operation = ?
size = ?
stride = ?
dilation = ?
padding = ?



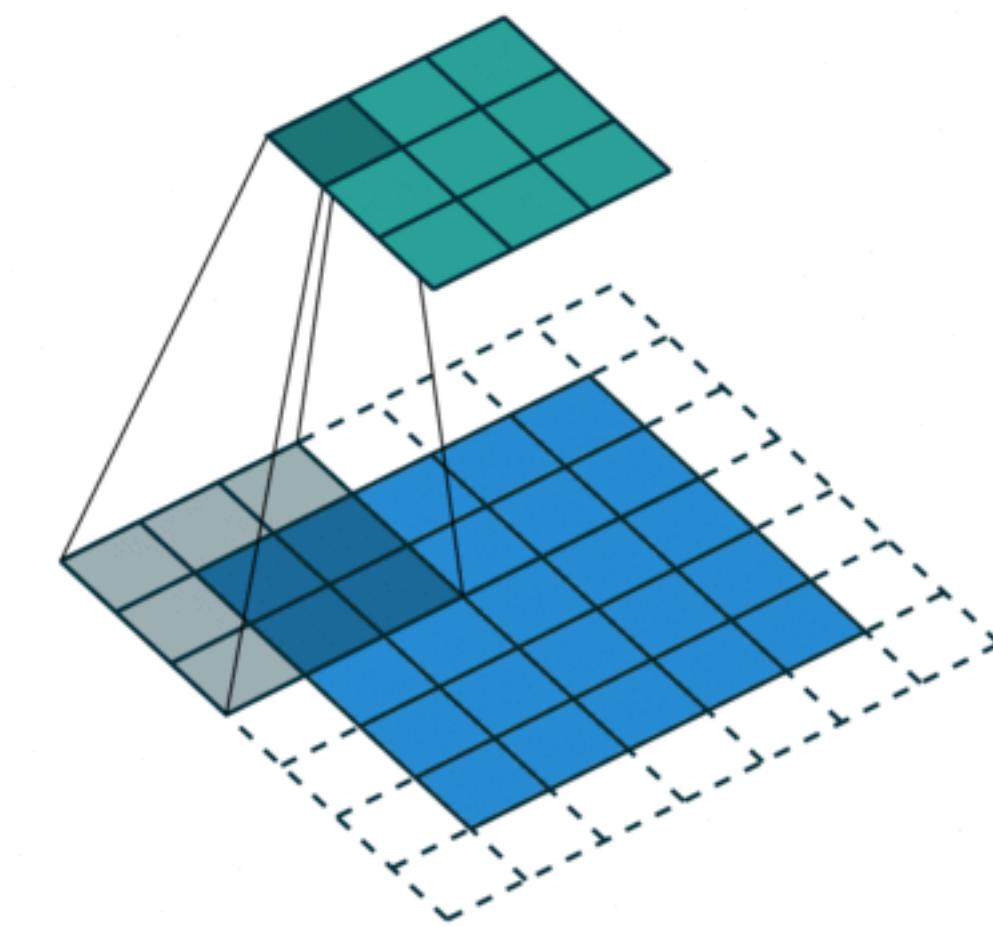
operation = ?
size=?
stride = ?
dilation = ?
padding = ?

Stride and Dilation

Solution



max pooling
size = 3
stride = 2
dilation = 1
padding = no



pooling or convolution
size = 3
stride = 2
dilation = 1
padding = yes

*Animation from Dumoulin, Vincent, and Francesco Visin. "A guide to convolution arithmetic for deep learning." (2016) https://github.com/vdumoulin/conv_arithmetic/blob/master/README.md

Convolutional vs fully connected layers

Images properties:

1. High dimensionality
2. Spacial structure
3. Stability under transformations

Convolutional vs fully connected layers

Images properties:

1. High dimensionality
2. Spacial structure
3. Stability under transformations

Convolution is a linear operation represented by a sparse matrix with shared parameters and a local structure

$$x * (-1,0,1) = x \cdot \begin{pmatrix} -1 & 0 & 0 & 0 \\ 0 & -1 & 0 & 0 \\ 1 & 0 & -1 & 0 & \dots \\ 0 & 1 & 0 & -1 \\ 0 & 0 & 1 & 0 \\ \dots \end{pmatrix}^*$$

Convolutional vs fully connected layers

Images properties:

1. High dimensionality
2. Spacial structure
3. Stability under transformations

Convolution is a linear operation represented by a sparse matrix with shared parameters and a local structure

$$x * (-1,0,1) = x \cdot \begin{pmatrix} -1 & 0 & 0 & 0 \\ 0 & -1 & 0 & 0 \\ 1 & 0 & -1 & 0 & \dots \\ 0 & 1 & 0 & -1 \\ 0 & 0 & 1 & 0 \\ \dots & & & \end{pmatrix}$$

#convolution parameters = 3

#fully connected parameters = $H \times W$

Convolutional vs fully connected layers

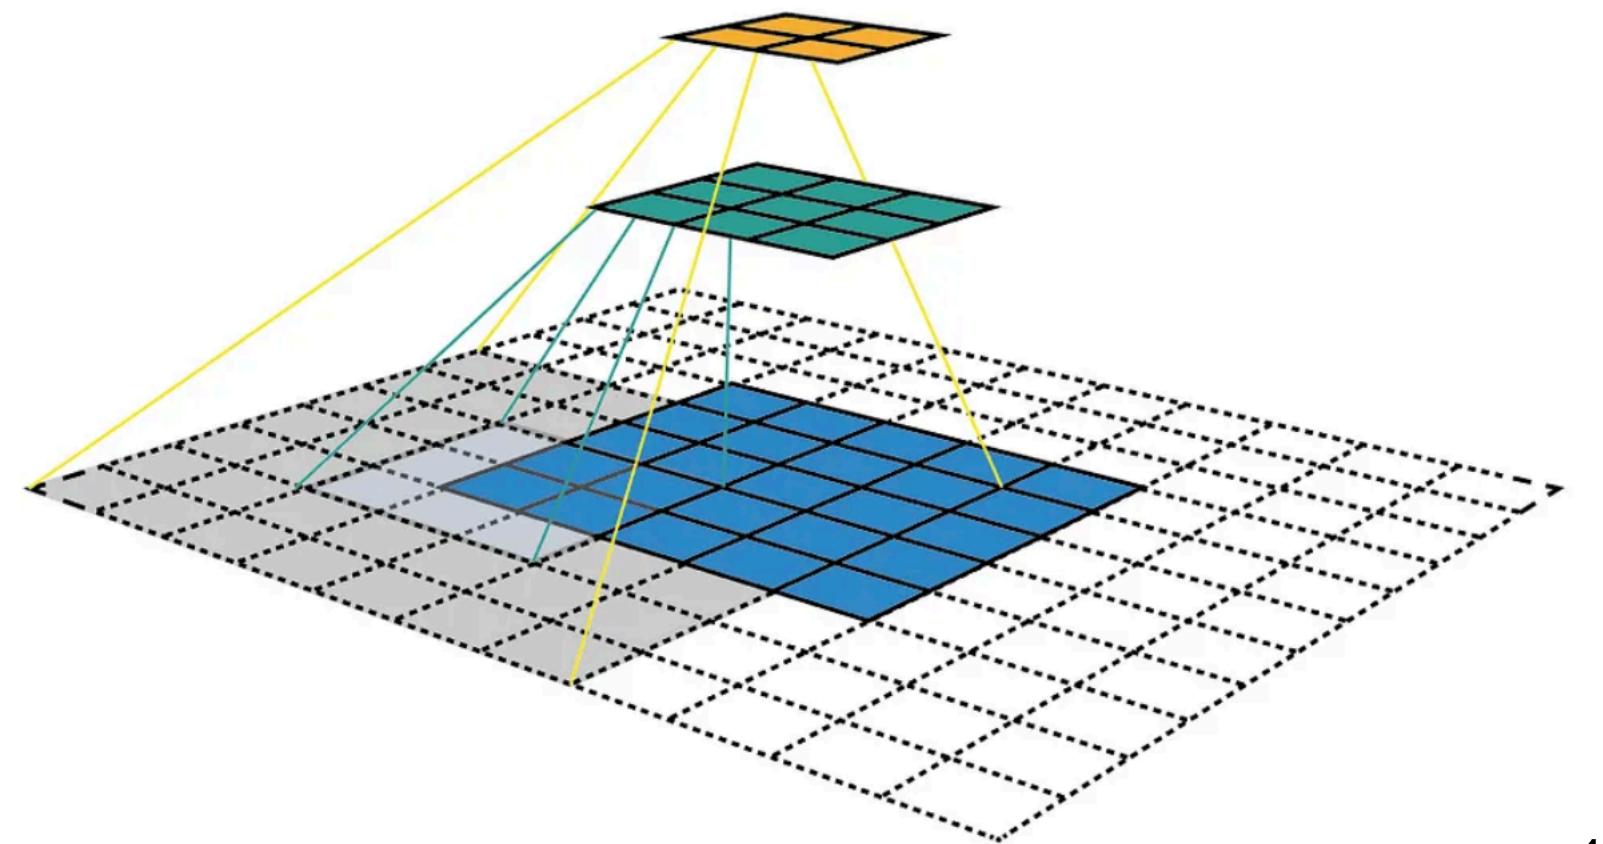
Images properties:

1. High dimensionality
2. Spacial structure
3. Stability under transformations

Convolution is a linear operation represented by a sparse matrix with shared parameters and a local structure

$$x * (-1,0,1) = x \cdot \begin{pmatrix} -1 & 0 & 0 & 0 \\ 0 & -1 & 0 & 0 \\ 1 & 0 & -1 & 0 & \dots \\ 0 & 1 & 0 & -1 \\ 0 & 0 & 1 & 0 \\ \dots & & & \end{pmatrix}$$

Receptive field, the region of the input which affects a specific output unit is **local**.



Convolutional vs fully connected layers

Images properties:

1. High dimensionality
2. Spacial structure
3. Stability under transformations

Convolution is a linear operation represented by a sparse matrix with shared parameters and a local structure

$$x * (-1,0,1) = x \cdot \begin{pmatrix} -1 & 0 & 0 & 0 \\ 0 & -1 & 0 & 0 \\ 1 & 0 & -1 & 0 & \dots \\ 0 & 1 & 0 & -1 \\ 0 & 0 & 1 & 0 \\ \dots & & & \end{pmatrix}$$

Convolution is a translation equivariant operation* and pooling enhances translation invariance

Exercise

Multi-armed bandit



Possible reward y : 0, +1 , +100

Bandit x : 0, 1

$$x^{(1)}, x^{(2)}, \dots x^{(10)} = 0, 1, 1, 1, 0, 1, 0, 1, 0, 0$$

$$y^{(1)}, y^{(2)}, \dots y^{(10)} = 1, 0, 0, 0, 1, 0, 0, 100, 0, 0$$

Neural network setup

1. What is the learning task? What is the input and output?
2. How many input, hidden and output units?
3. Which is the last layer activation function?
4. What is the loss function?

Multiclass Classification

Training set $\{x^{(i)}, y^{(i)}\}_{i=1,\dots,n}$ of i.i.d input/output pairs, $y^{(i)}$ is a class $\{1, \dots, K\}$

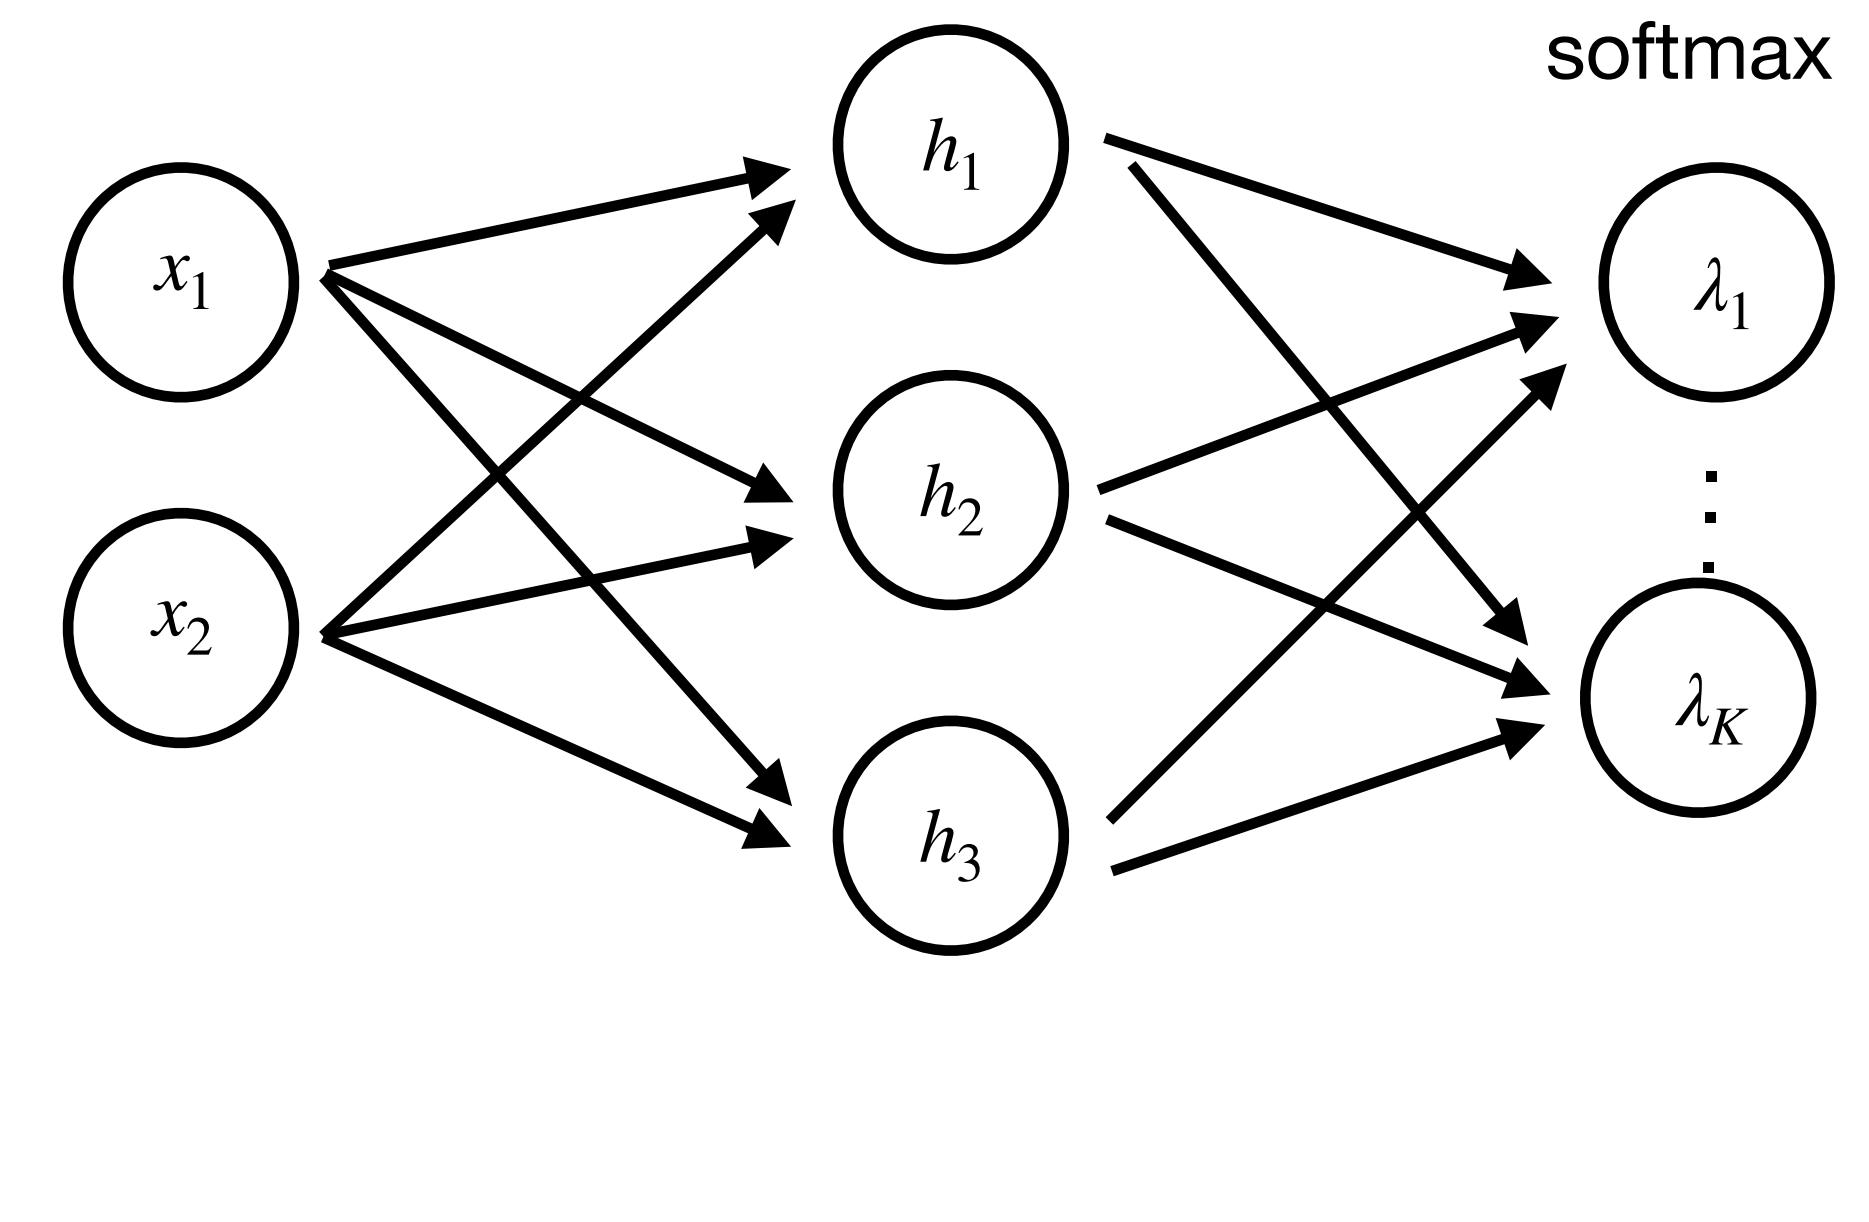
- One-hot vectors $y = (0, \dots, 0, 1, 0, \dots, 0)$
- K output features
- λ_k is the probability assigned to class k
- Softmax last layer:

$$\lambda_k = \text{softmax}(z_1, \dots, z_K) = \frac{e^{z_k}}{\sum_{k'} e^{z'_{k'}}}$$

- Cross entropy loss

$$L_{CE}(\theta) = \frac{1}{n} \sum_{i=1}^n \sum_{k=1}^K \underbrace{-y_k^{(i)} \log(\lambda_k^{(i)})}_{l^{(i)}}$$

- The predicted class predicted $\hat{y} = \text{argmax}_k \{\lambda_k\}$



Summary

Topics

- Structured Data
- Convolutions
- Convolutional modules: kernels, padding, feature maps, pooling, stride, dilatation.
- Convolutional vs fully connected networks

Reading material

- *Deep Learning book* - Chapter 9 (until 9.4 excluded)
- *Understanding Deep Learning* - Chapter 10 (10.4 excluded)

Additional resources

- *Neural Networks and Deep Learning*, Michael Nielsen - [Chapter 6 Introducing Convolutional Neural Networks](#) (up to “Convolutional neural network in practice” excluded)
- [Convolutional layer notebook](#)

