

# Introduction to Reinforcement Learning


Elena Congeduti, 09-12-2024



# Lecture's Agenda


- Sequential decision making
- Markov decision processes (MDPs)
- Value iteration
- Tabular Q-learning

# What is Reinforcement Learning?






# What is Reinforcement Learning?




Intelligent agent learns from experience how to make decisions that maximize its return in the face of uncertainty

# Robot Navigation




- Robot moving in a maze

# Robot Navigation




- Robot moving in a maze

# Robot Navigation



- Robot moving in a maze


# Robot Navigation



- Robot moving in a maze




# Robot Navigation




- Robot moving in a maze

# Robot Navigation




- Robot moving in a maze

# Robot Navigation




- Robot moving in a maze

# Robot Navigation




- Robot moving in a maze

# Robot Navigation



- Robot moving in a maze
- Goal: find the exit

# Robot Navigation



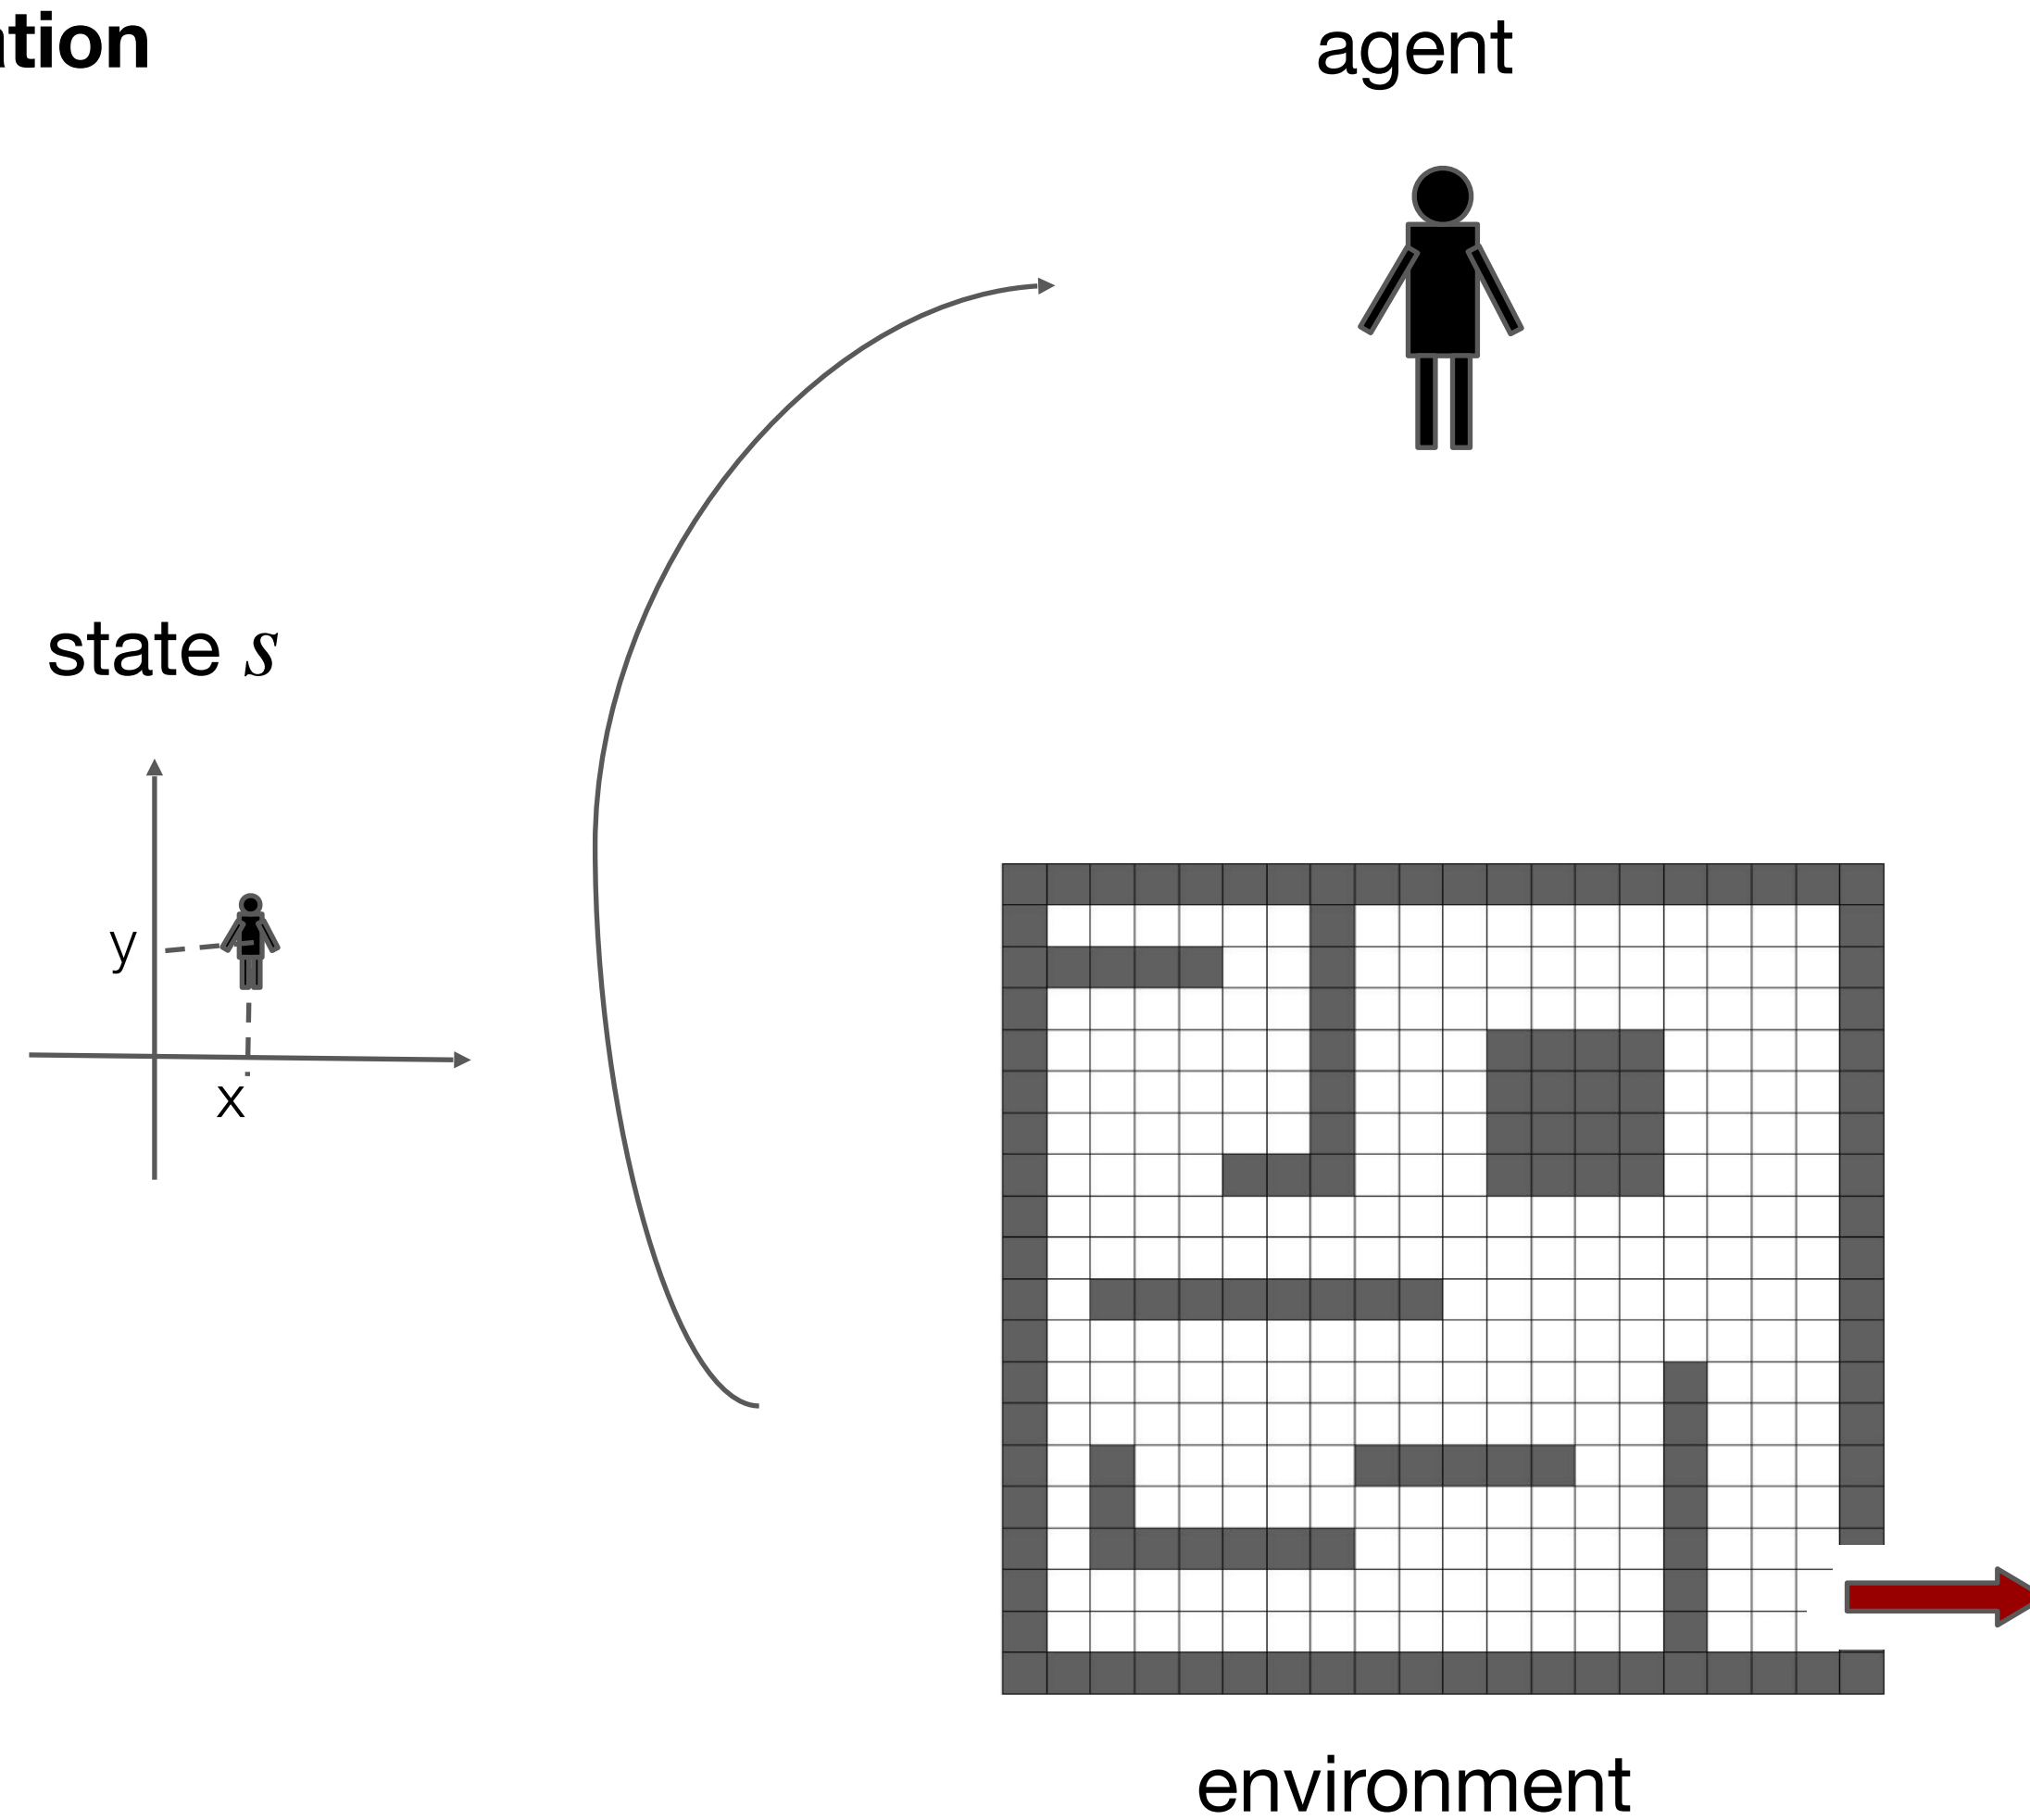
- Robot moving in a maze
- Goal: find the exit
- Chance of failing





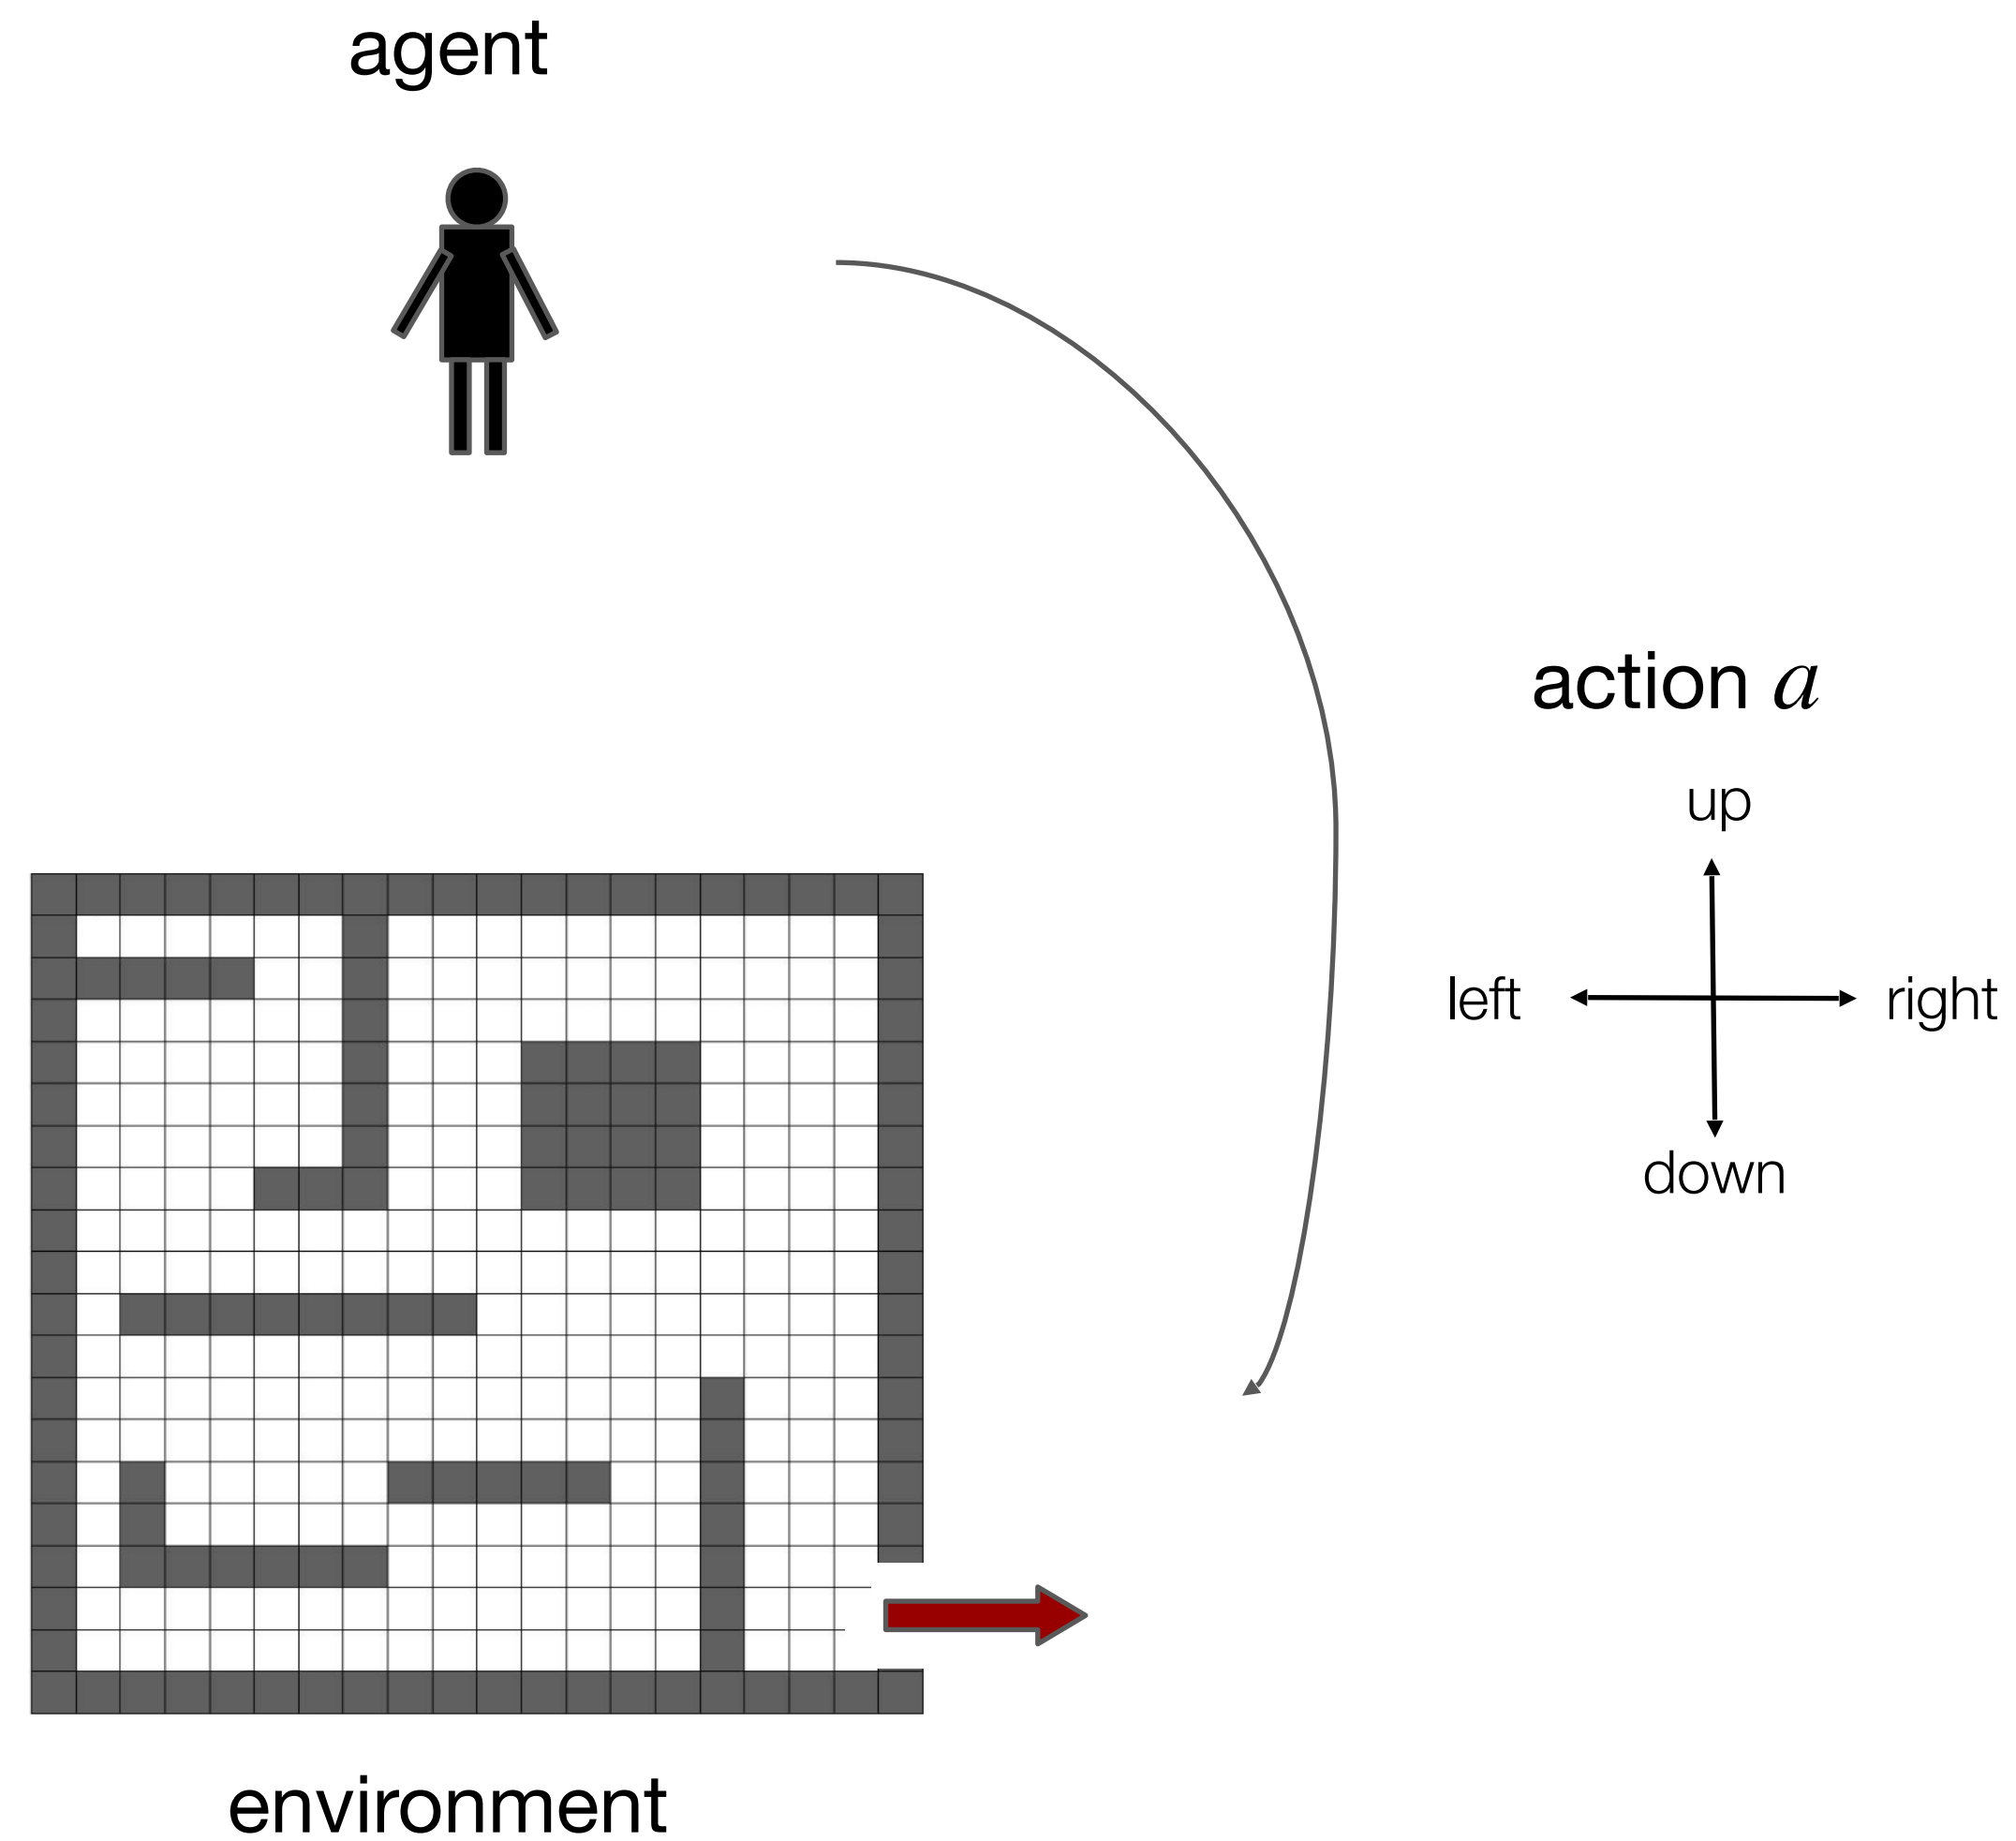
# Sequential Decision Making

## Robot Navigation



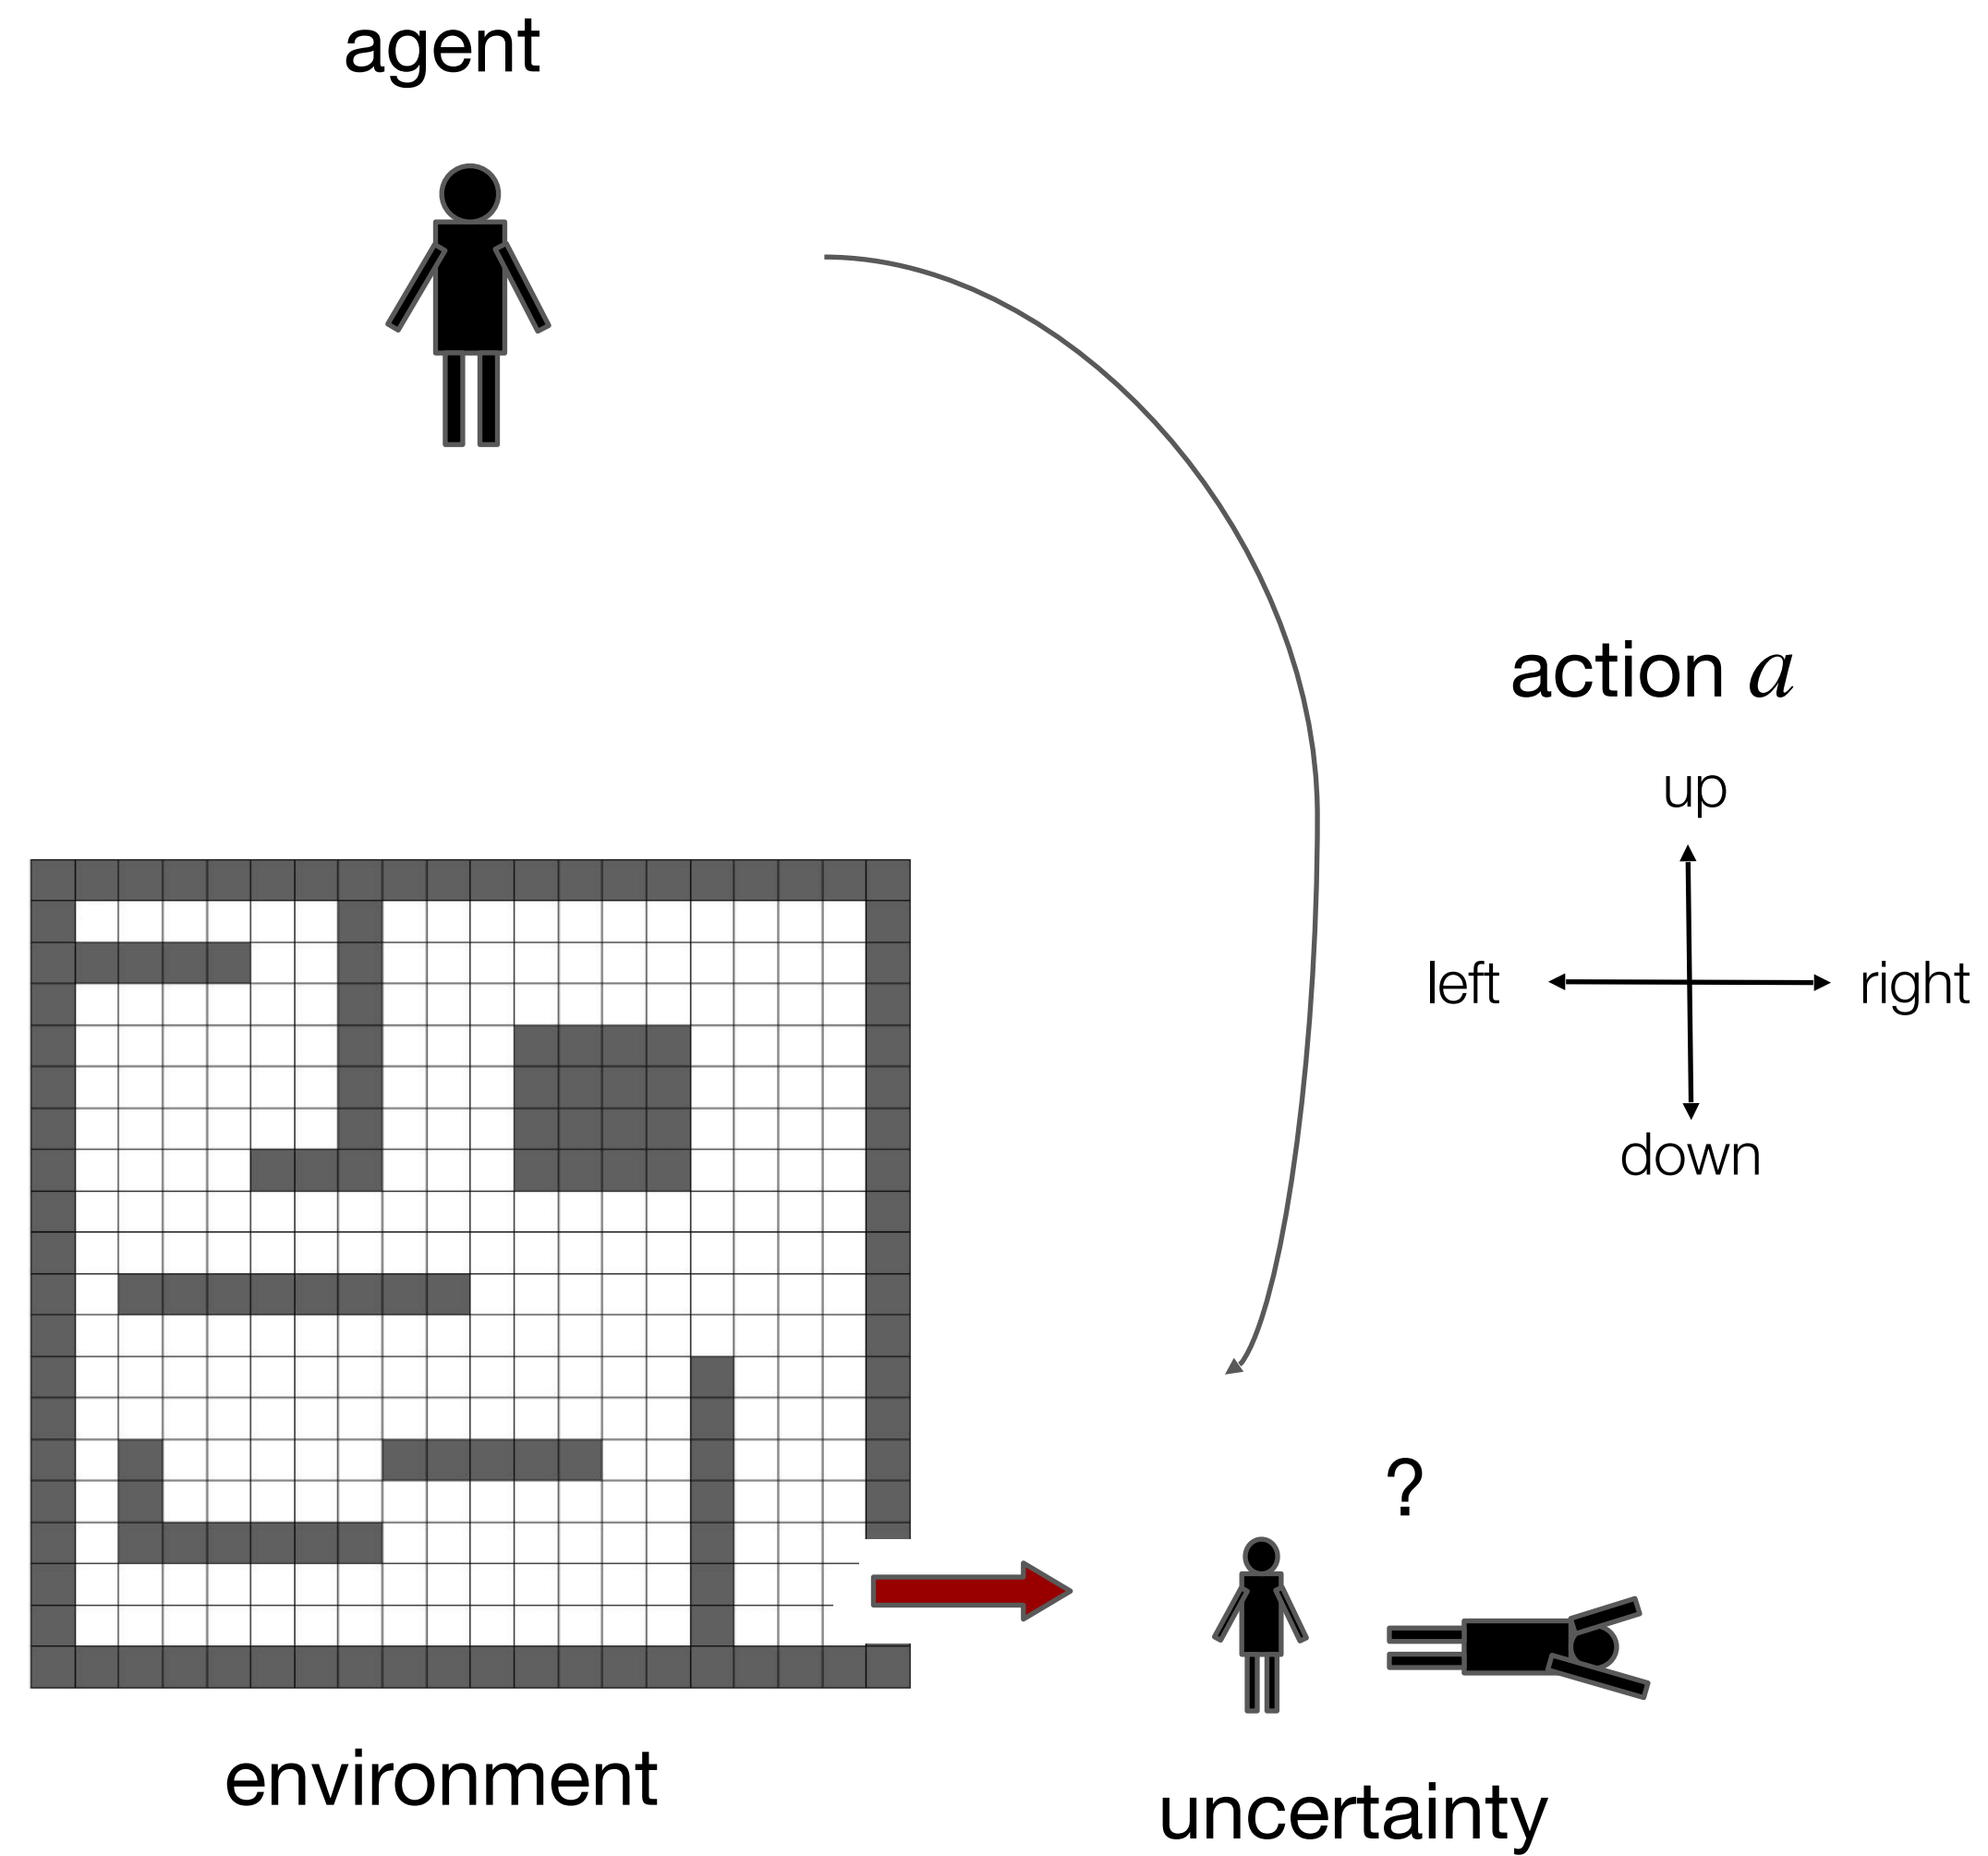
# Sequential Decision Making

## Robot Navigation




# Sequential Decision Making

## Robot Navigation



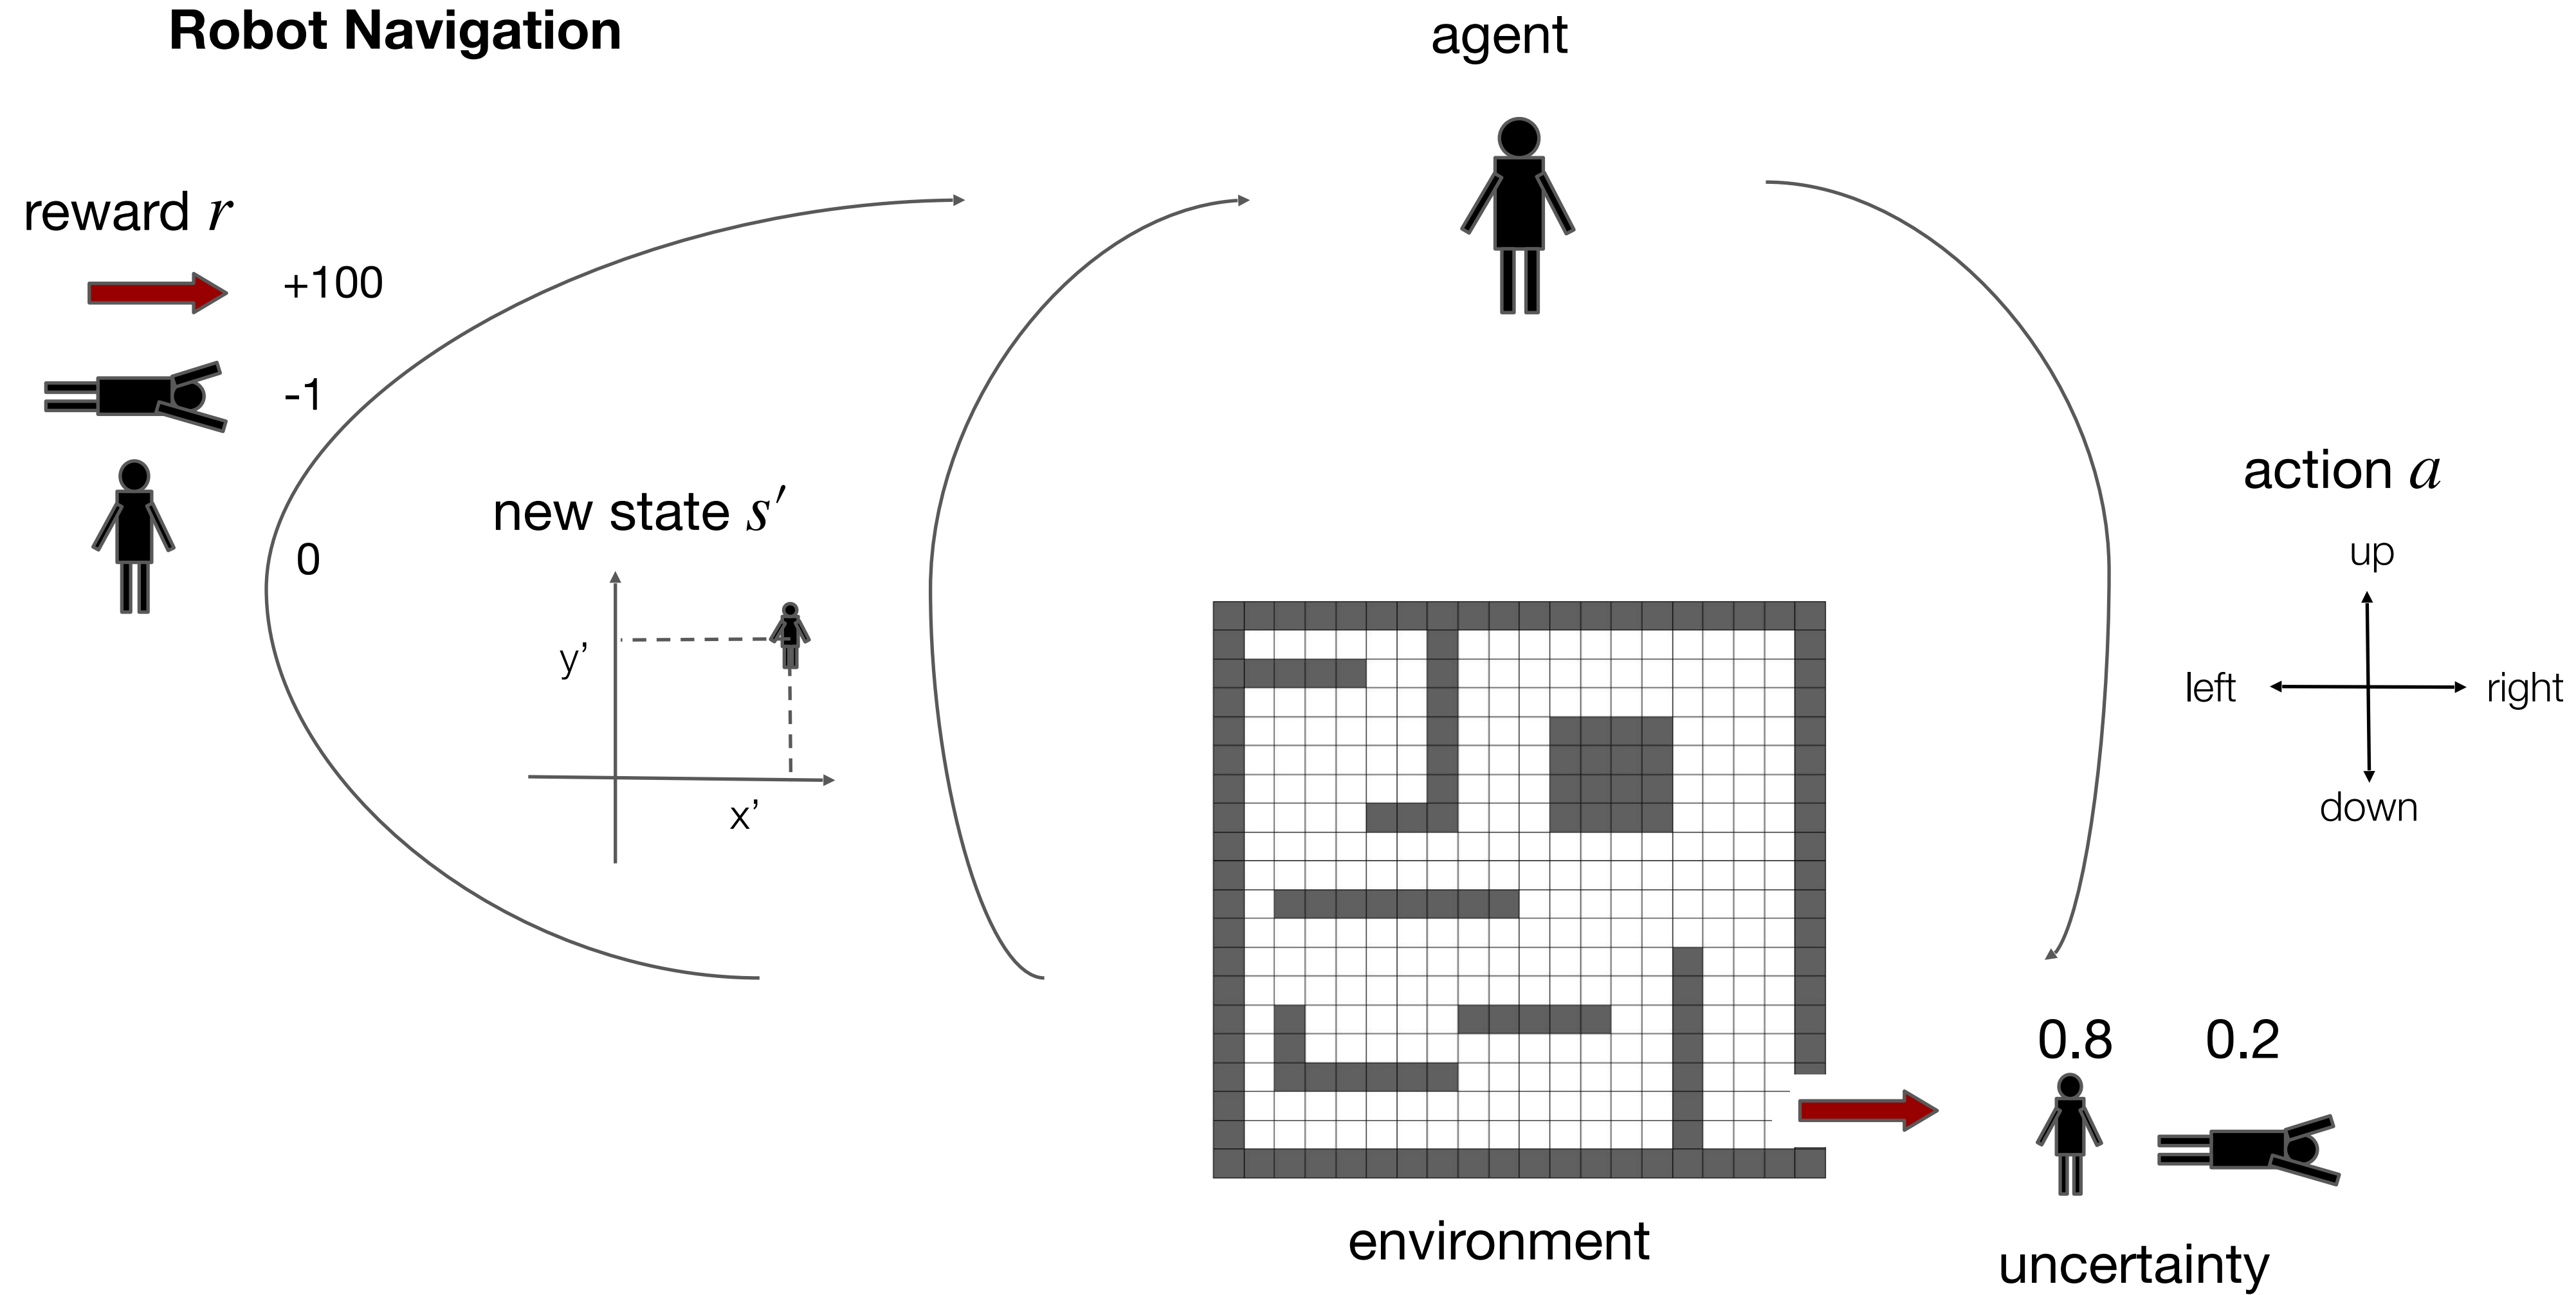
# Sequential Decision Making

## Robot Navigation



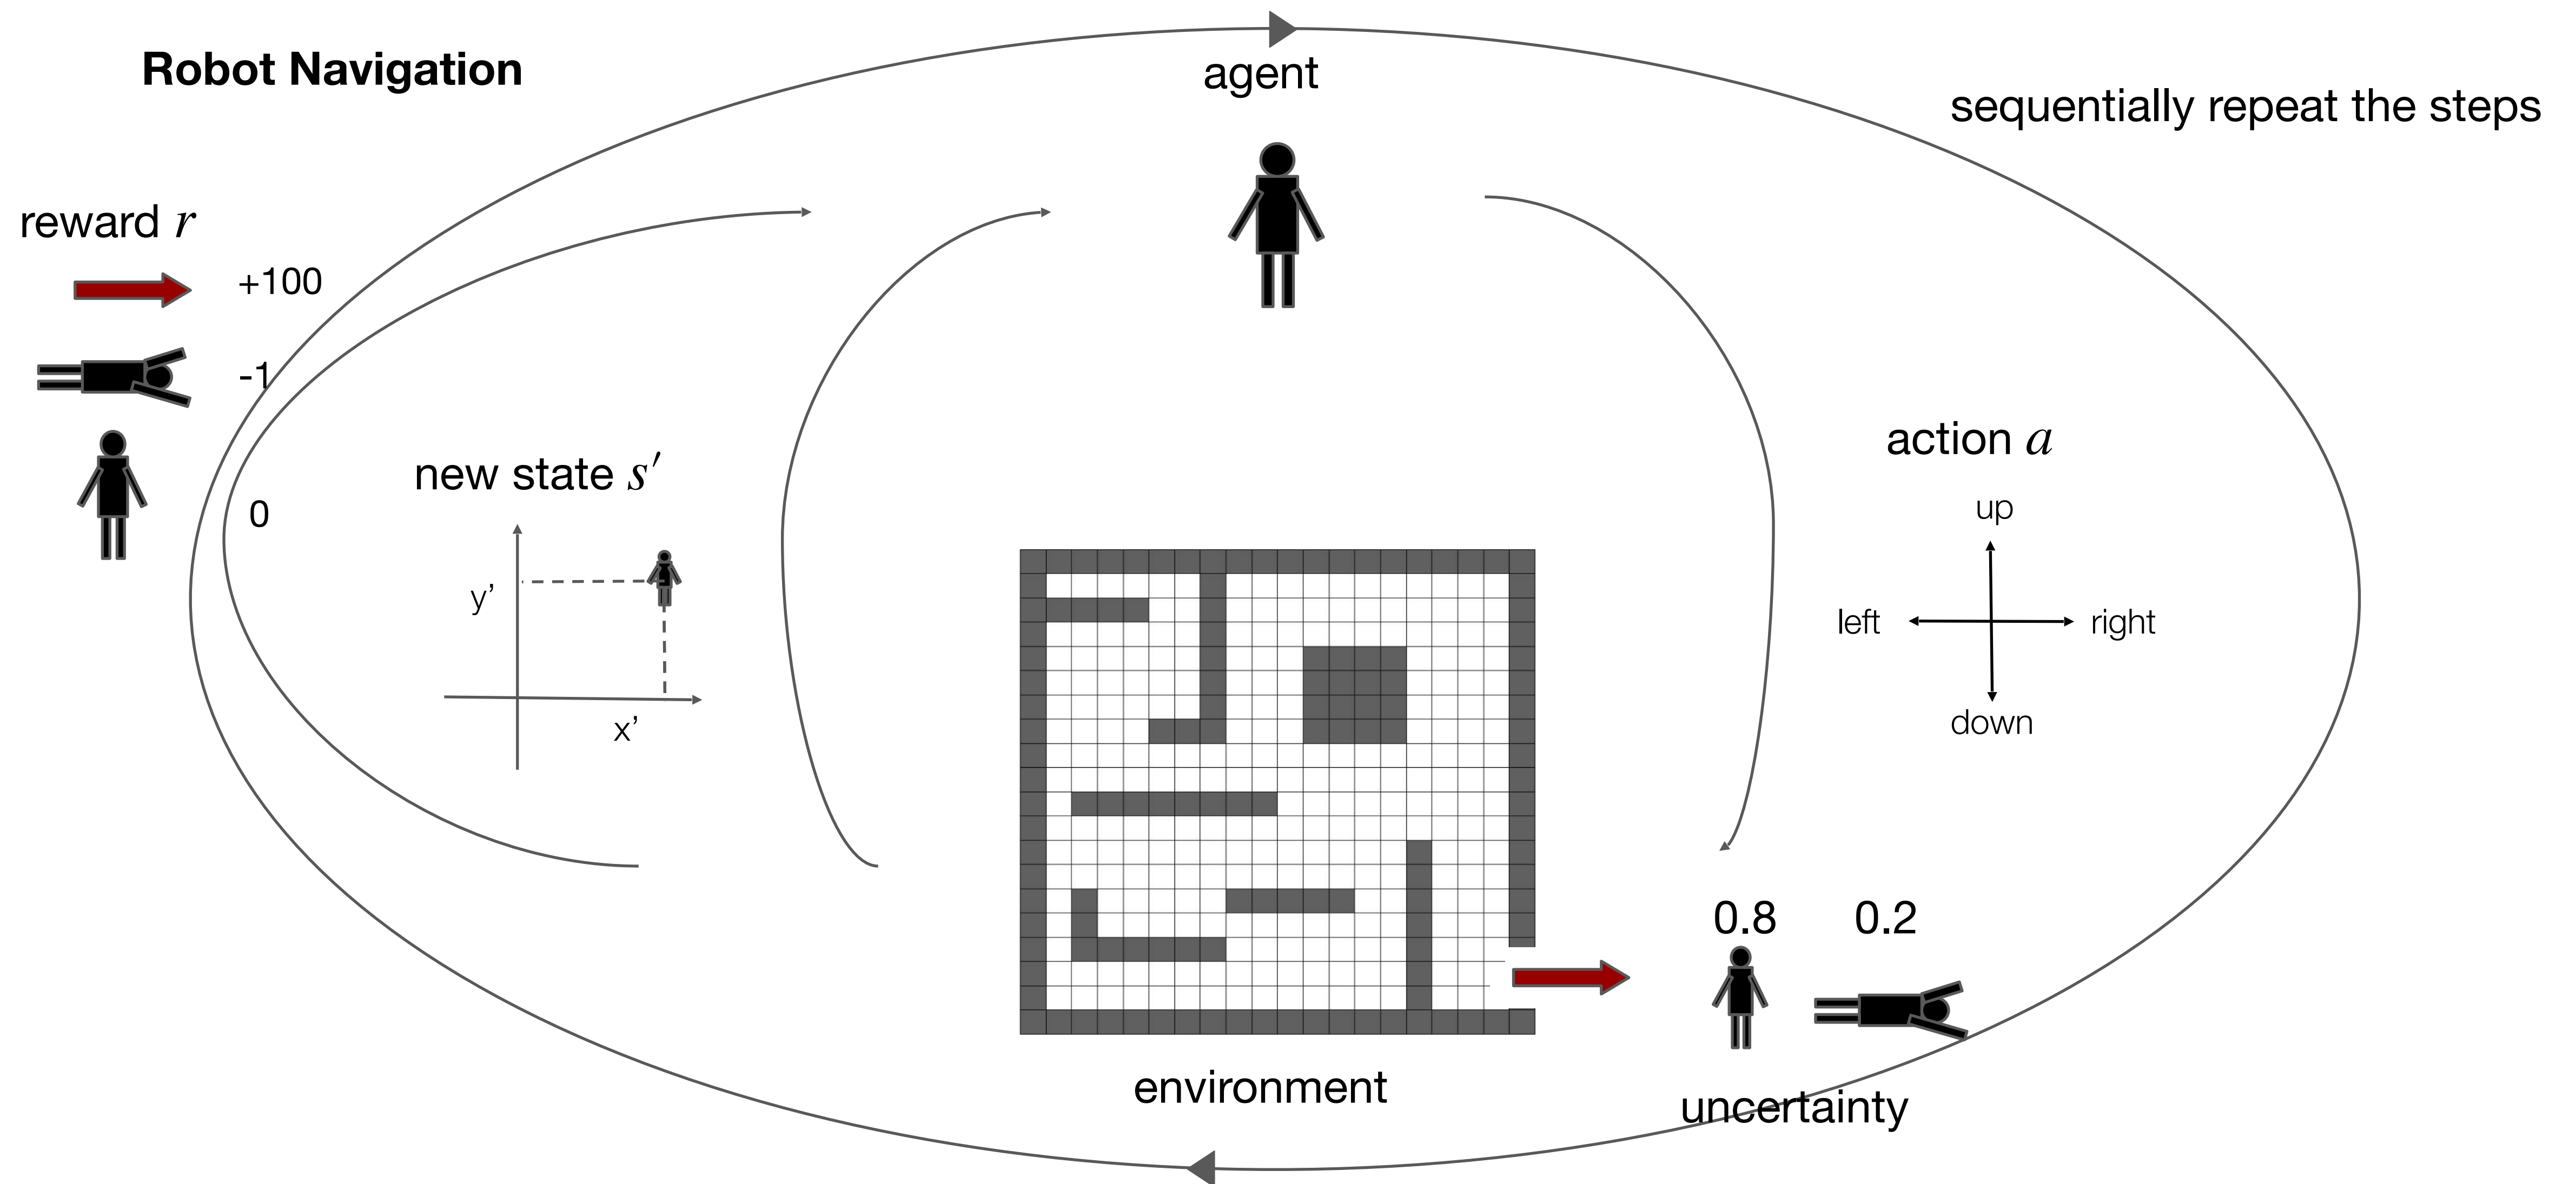
# Sequential Decision Making

## Robot Navigation

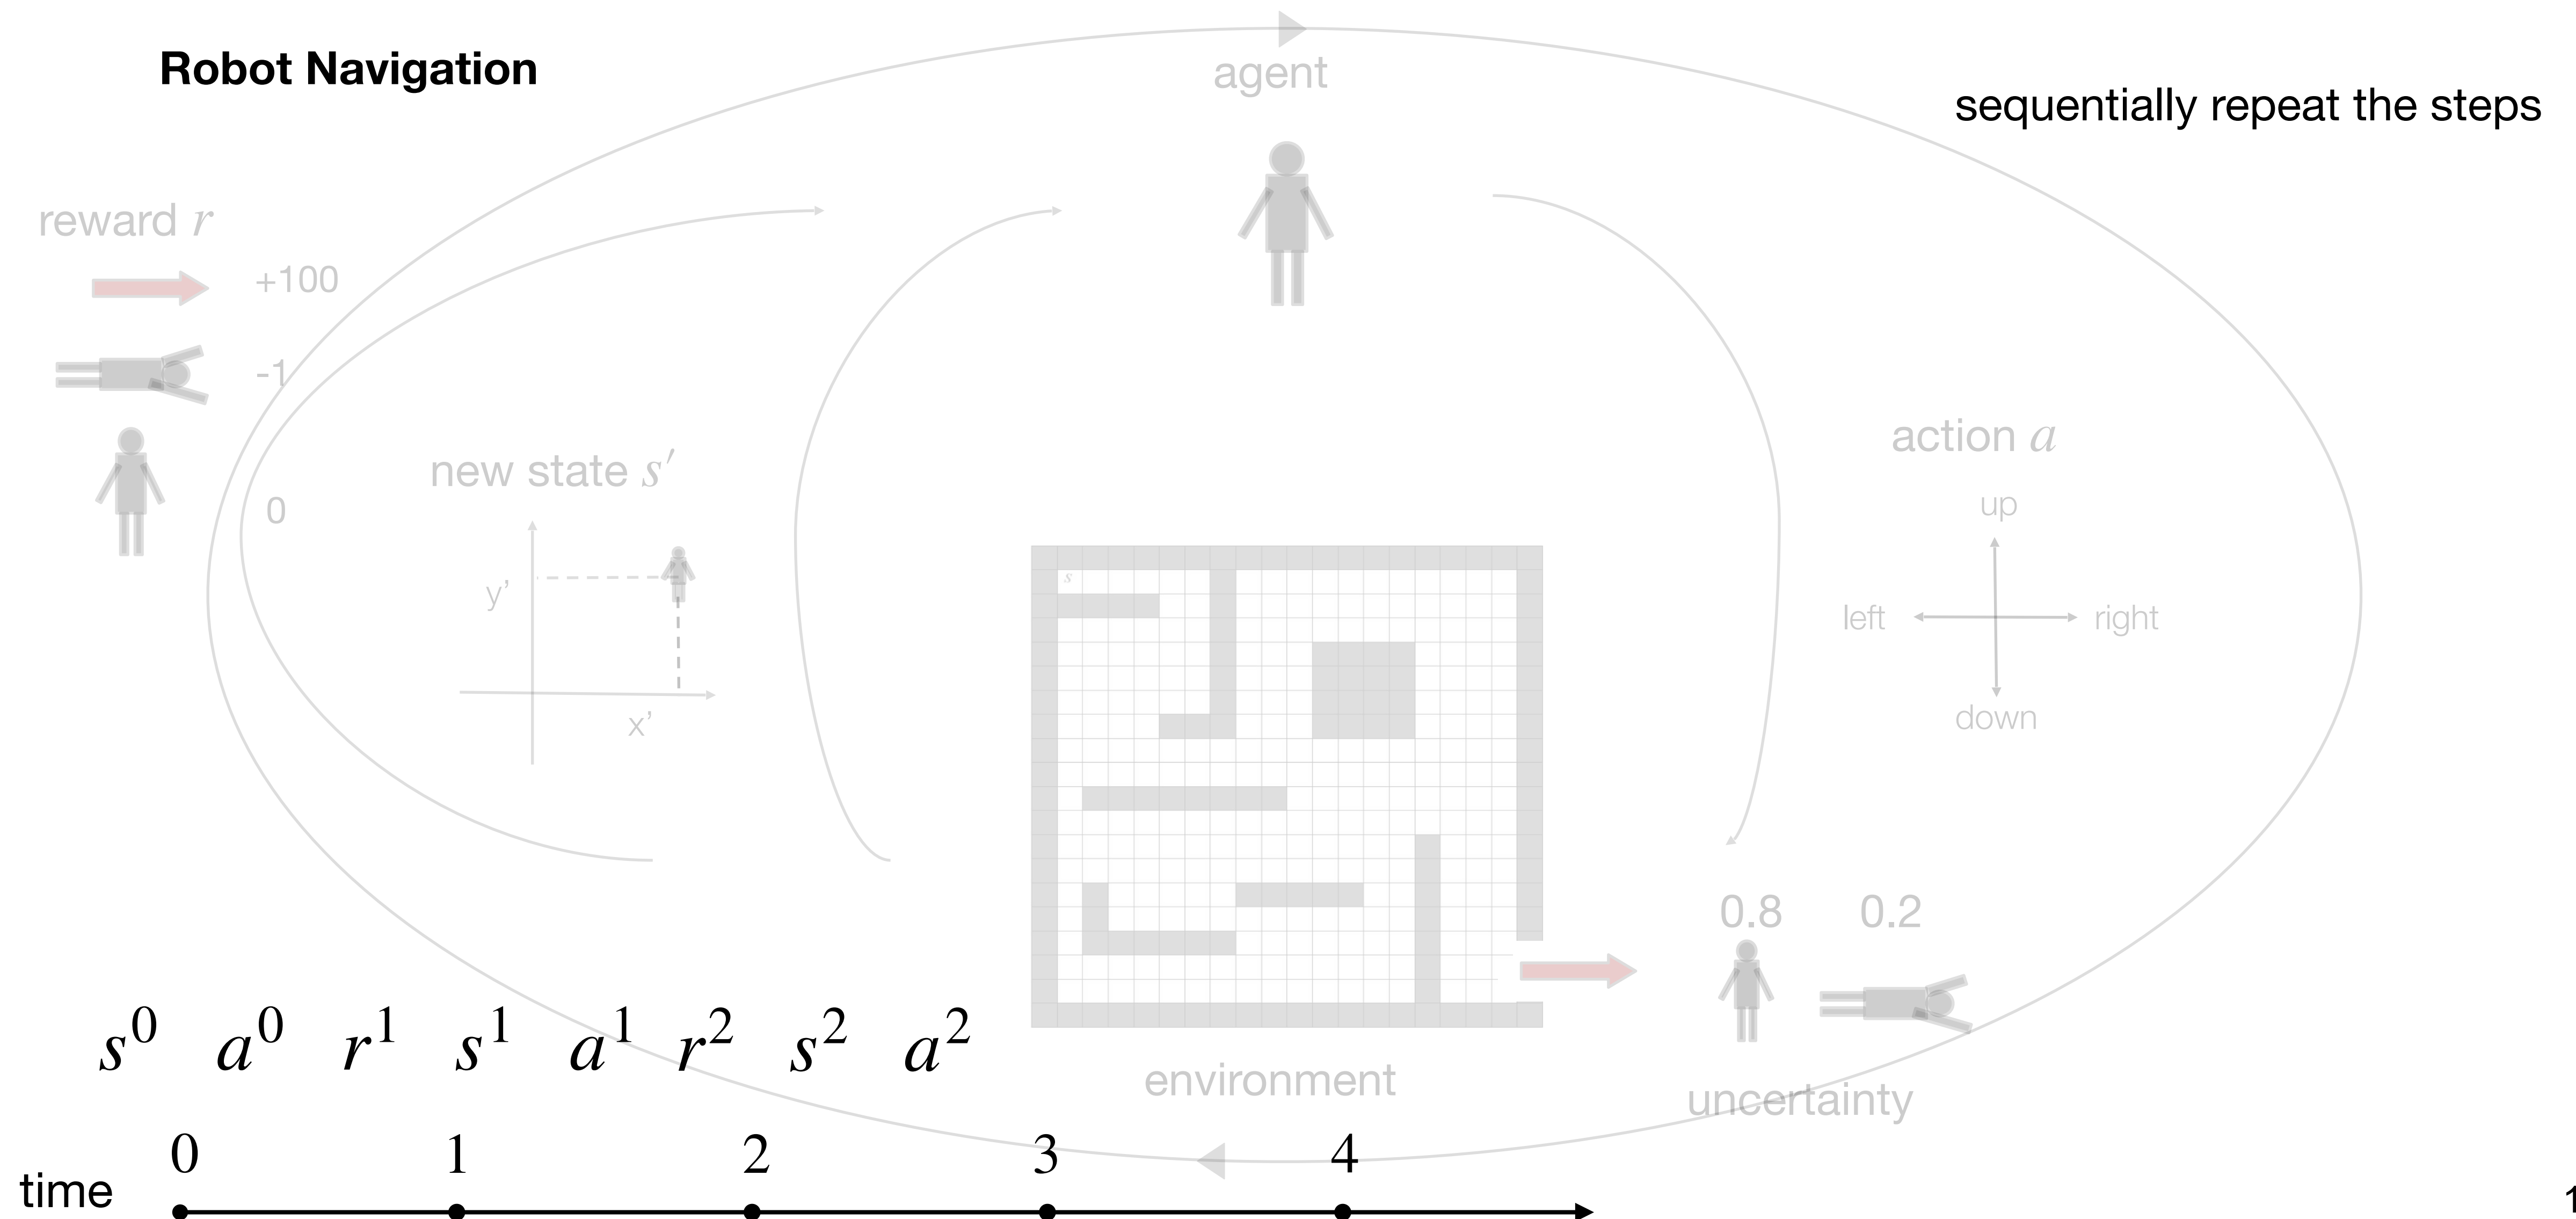




# Sequential Decision Making




# Sequential Decision Making



# Markov Decision Processes


Markov decision process (MDP)  $M = (S, A, s^0, P, R)$



# Markov Decision Processes

Markov decision process (MDP)  $M = (S, A, s^0, P, R)$

- State space  $S$  of the environment
- Initial state  $s^0 \in S$




$$S = \{\text{feasible } (x, y)\}$$

$$s^0 = (x^0, y^0) \quad \text{initial robot coordinates}$$

# Markov Decision Processes

Markov decision process (MDP)  $M = (S, A, s^0, P, R)$

- State space  $S$  of the environment
- Initial state  $s^0 \in S$
- Action space  $A$  of possible agent actions




$$A = \{right, down, left, up\}$$

# Markov Decision Processes

Markov decision process (MDP)  $M = (S, A, s^0, P, R)$

- State space  $S$  of the environment
- Initial state  $s^0 \in S$
- Action space  $A$  of possible agent actions
- Transition probabilities  $P(s' | s, a)$  from state  $s$  to  $s'$  after choosing  $a$

Transitions






$$P(s' | a = \textit{right}, s = (x, y)) = \begin{cases} 0.8 & \text{if } s' = (x + 1, y) \\ 0.2 & \text{if } s' = (x, y) \\ 0 & \text{otherwise} \end{cases}$$



# Markov Decision Processes

Markov decision process (MDP)  $M = (S, A, s^0, P, R)$

- State space  $S$  of the environment
- Initial state  $s^0 \in S$
- Action space  $A$  of possible agent actions
- Transition probabilities  $P(s' | s, a)$  from state  $s$  to  $s'$  after choosing  $a$
- Reward  $R(s', a, s)$  or  $R(s, a)$  for resulting in state  $s'$  from state  $s$  after choosing action  $a$

reward	
	+100
	-1
	0


$$R(s', a, s) = \begin{cases} 100 & \text{if } s' = s_{goal} \\ -1 & \text{if } s = s' \\ 0 & \text{otherwise} \end{cases}$$

# Policy

- A policy  $\pi$  encodes the agent's behaviour
- A map from states to actions,  $\pi(s) = a$  is the agent's action when the environment state is  $s$


# Policy

- A policy  $\pi$  encodes the agent's behaviour
- A map from states to actions,  $\pi(s) = a$  is the agent's action when the environment state is  $s$



# Policy


- A policy  $\pi$  encodes the agent's behaviour
- A map from states to actions,  $\pi(s) = a$  is the agent's action when the environment state is  $s$



**How good is this policy?**

# Value


$V^\pi(s)$  is the expected sum of discounted future rewards for employing a policy  $\pi$  starting from an initial state  $s$



# Value

$V^\pi(s)$  is the expected sum of discounted **future rewards** for employing a policy  $\pi$  starting from an initial state  $s$

$$r^1, r^2, r^3 \dots$$






# Value

$V^\pi(s)$  is the expected sum of discounted **future rewards** for employing a policy  $\pi$  starting from an initial state  $s$

0,  $r^2$ ,  $r^3$  ...

$r^1 = 0$




# Value

$V^\pi(s)$  is the expected sum of discounted **future rewards** for employing a policy  $\pi$  starting from an initial state  $s$

0, -1,  $r^3$  ...

$$r^2 = -1$$




# Value

$V^\pi(s)$  is the expected sum of discounted **future rewards** for employing a policy  $\pi$  starting from an initial state  $s$

0, -1, 0 ...

$$r^3 = 0$$

Initial state  $s^0$



goal, +100

failing, -1


# Value

$V^\pi(s)$  is the expected sum of discounted **future rewards** for employing a policy  $\pi$  starting from an initial state  $s$

0, -1, 0 ... 100

$$r^T = 100$$

Initial state  $s^0$




# Value

$V^\pi(s)$  is the expected **sum of discounted** future rewards for employing a policy  $\pi$  starting from an initial state  $s$

$$\gamma^1 0 + \gamma^2 (-1) + \gamma^3 0 + \dots \gamma^T (100)$$

$0 < \gamma < 1$  discount factor




# Value

$V^\pi(s)$  is the expected **sum of discounted** future rewards for employing a policy  $\pi$  starting from an initial state  $s$

$$\gamma^1 0 + \gamma^2 (-1) + \gamma^3 0 + \dots \gamma^T (100) = \sum_t \gamma^t r^t$$


$0 < \gamma < 1$  discount factor



# Value

$V^\pi(s)$  is the **expected** sum of discounted future rewards for employing a policy  $\pi$  starting from an initial state  $s$


$$V^\pi(s) = \underset{\substack{\downarrow \\ \text{Average over all possible futures}}}{\mathbb{E}} \left[ \sum_t \gamma^t r^t \mid s^0 = s, \pi \right]$$



# Value

$V^\pi(s)$  is the expected sum of discounted future rewards for employing a policy  $\pi$  starting from an initial state  $s$

$$V^\pi(s) = \mathbb{E} \left[ \sum_t \gamma^t r^t \mid s^0 = s, \pi \right]$$



**How good is this policy?**

Deterministic transitions - no chance of failing,  $\gamma = 0.9$


$V^\pi(s^0) = ?$



# Value

$V^\pi(s)$  is the expected sum of discounted future rewards for employing a policy  $\pi$  starting from an initial state  $s$

$$V^\pi(s) = \mathbb{E} \left[ \sum_t \gamma^t r^t \mid s^0 = s, \pi \right]$$



## How good is this policy?


Deterministic transitions - no chance of failing,  $\gamma = 0.9$

$$V^\pi(s^0) = 0 + 0 + \dots + 0 + \gamma^{11} 100 \approx 31$$

# Value

$V^\pi(s)$  is the expected sum of discounted future rewards for employing a policy  $\pi$  starting from an initial state  $s$

$$V^\pi(s) = \mathbb{E} \left[ \sum_t \gamma^t r^t \mid s^0 = s, \pi \right]$$



**How good is this policy?**


Deterministic transitions - no chance of failing,  $\gamma = 0.9$

$$V^\pi(s^0) = \gamma^{11} 100 \approx 31 \quad V^{\pi}(s^0) = ?$$

# Value

$V^\pi(s)$  is the expected sum of discounted future rewards for employing a policy  $\pi$  starting from an initial state  $s$

$$V^\pi(s) = \mathbb{E} \left[ \sum_t \gamma^t r^t \mid s^0 = s, \pi \right]$$



## How good is this policy?


Deterministic transitions - no chance of failing,  $\gamma = 0.9$

$$V^\pi(s^0) = \gamma^{11} 100 \approx 31 \quad V^{\pi}(s^0) = \gamma^9 100 \approx 39$$

# Value

$V^\pi(s)$  is the expected sum of discounted future rewards for employing a policy  $\pi$  starting from an initial state  $s$

$$V^\pi(s) = \mathbb{E} \left[ \sum_t \gamma^t r^t \mid s^0 = s, \pi \right]$$



**How good is this policy?**

Deterministic transitions - no chance of failing,  $\gamma = 0.9$


$$V^\pi(s^0) = \gamma^{11} 100 \approx 31 < V^{\pi^*}(s^0) = \gamma^9 100 \approx 39$$

**Value measures the quality of a policy: best policy gives the highest value**

# Value

$V^\pi(s)$  is the expected sum of discounted future rewards for employing a policy  $\pi$  starting from an initial state  $s$

$$V^\pi(s) = \mathbb{E} \left[ \sum_t \gamma^t r^t \mid s^0 = s, \pi \right]$$



**Goal:** find the optimal policy  $\pi^*$  that achieves the highest value

$$V^{\pi^*}(s) = \max_{\pi} V^\pi(s) = V^*(s)$$

**Value measures the quality of a policy: best policy gives the highest value**

# Quiz Show

1. Go to brightspace → 9. Reinforcement Learning → Quiz Show
2. Read the description and play the quiz show
3. Do not compute anything just follow your intuition
4. 3 mins then discussion





# Quiz Show: discussion

**How do you formulate this game as a Markov decision process?**

1. Who is the agent?
2. What are the states and initial state?
3. What are the actions?
4. What is the reward?



# Quiz Show: discussion

How do you formulate this game as a Markov decision process?

1. Who is the agent?

1. You, the player

2. What are the states and initial state?

2.  $S = \{Q_1, Q_2, Q_3, Q_4, home\}$  and  $s^0 = Q_1$

3. What are the actions?

3.  $A = \{quit, answer\}$

4. What is the reward?

4. Reward  $R(home, ...) = 0\$, R(Q_2, quit) = 100\$, R(Q_3, quit) = 1100\$ \dots,$




Optimal policy  $\begin{cases} \pi^*(Q_4) &= quit \\ \pi^*(Q) &= answer \end{cases}$  for  $Q \neq Q_4$  with optimal value  $V^*(Q_1) = 3746$




# Quiz Show: intuition

Assume we are at question 4, what is the optimal action?



# Quiz Show: intuition


Assume we are at question 4, what is the optimal action?





# Quiz Show: intuition


Assume we are at question 4, what is the optimal action?



$$\pi^*(Q_4) = \textit{quit} \quad V^*(Q_4) = 11.100\$$$

# Quiz Show: intuition


Assume we are at question 3, what is the optimal action?






# Q-Value

- Expected value for executing an action in a given state assuming that you will act optimally in the future



# Q-Value

- Expected value for executing an action in a given state assuming that you will act optimally in the future



$$Q(Q_3, \text{quit}) = R(Q_3, \text{quit}) + 0 = 1100$$

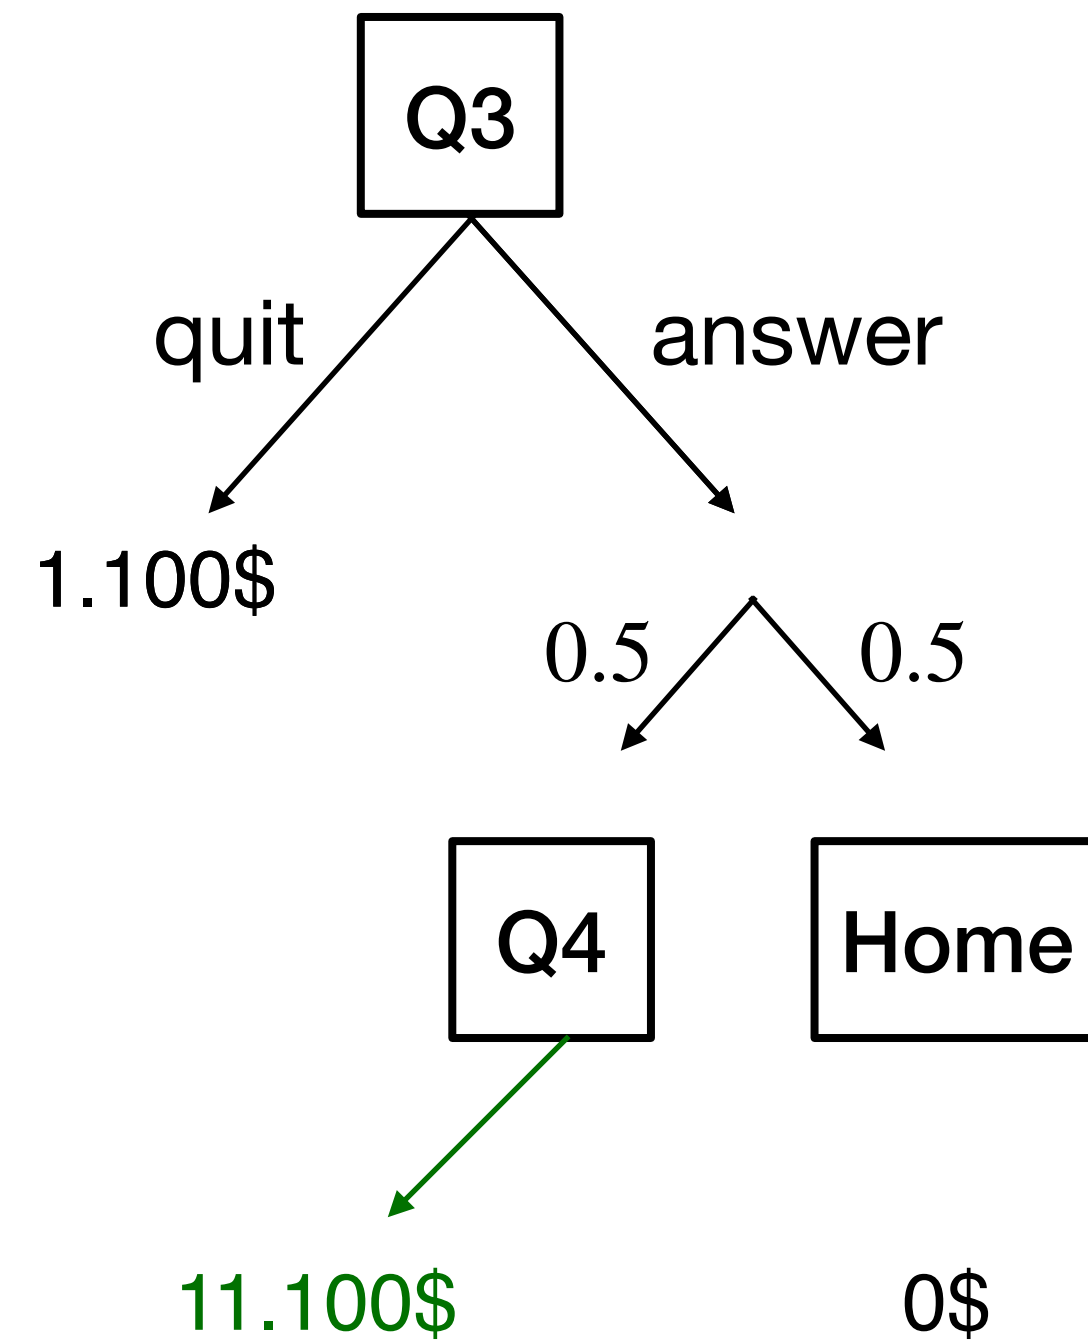
Immediate reward for  
quitting at question 3

There's no future state as  
your game is terminated



# Q-Value

- Expected value for executing an action in a given state assuming that you will act optimally in the future

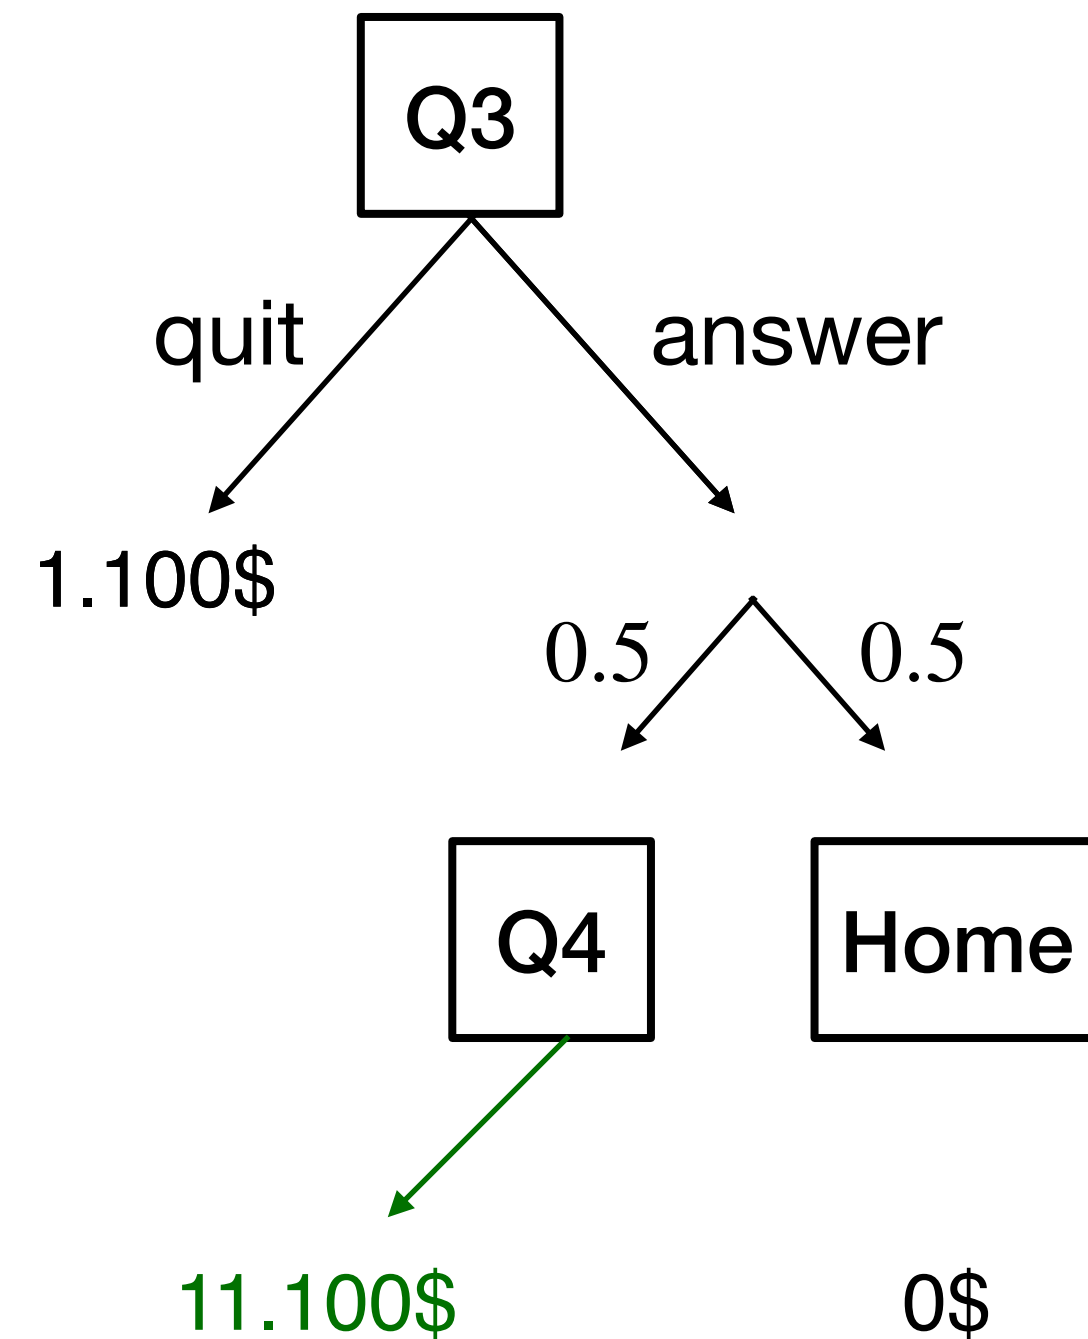


$$Q(Q_3, \text{answer}) =$$



# Q-Value

- Expected value for executing an action in a given state assuming that you will act optimally in the future



$$Q(Q_3, \text{answer}) = 0 + (0.5 \times 11.100 + 0.5 \times 0) = 5.550$$


↓  
Immediate reward for  
answering at question 3

↓  
Expected value for acting optimally in the future



# Q-Value

- Expected value for executing an action in a given state assuming that you will act optimally in the future



$$Q(Q_3, \text{answer}) = 0 + (0.5 \times 11.100 + 0.5 \times 0) = 5.550$$


$$R(Q_3, \text{answer}) + (P(Q_4 | \text{answer}, Q_3) \times V^*(Q_4) + P(\text{home} | \text{answer}, Q_3) \times V^*(\text{home}))$$

↓  
Immediate reward for  
answering at question 3

↓  
Expected value for acting optimally in the future

# Q-Value

- Expected value for executing an action in a given state assuming that you will act optimally in the future



$$Q(Q_3, \text{answer}) = 0 + (0.5 \times 11.100 + 0.5 \times 0) = 5.550$$

$$R(Q_3, \text{answer}) + (P(Q_4 | \text{answer}, Q_3) \times V^*(Q_4) + P(\text{home} | \text{answer}, Q_3) \times V^*(\text{home}))$$

↓  
Immediate reward for  
answering at question 3

↓  
Expected value for acting optimally in the future



# Q-Value

- Expected value for executing an action in a given state assuming that you will act optimally in the future

$$Q(s, a) = R(s, a) + \gamma \sum_{s'} P(s' | s, a) V^*(s')$$

$\downarrow$                        $\downarrow$   
Immediate reward      Discounted expected value for acting optimally in the future



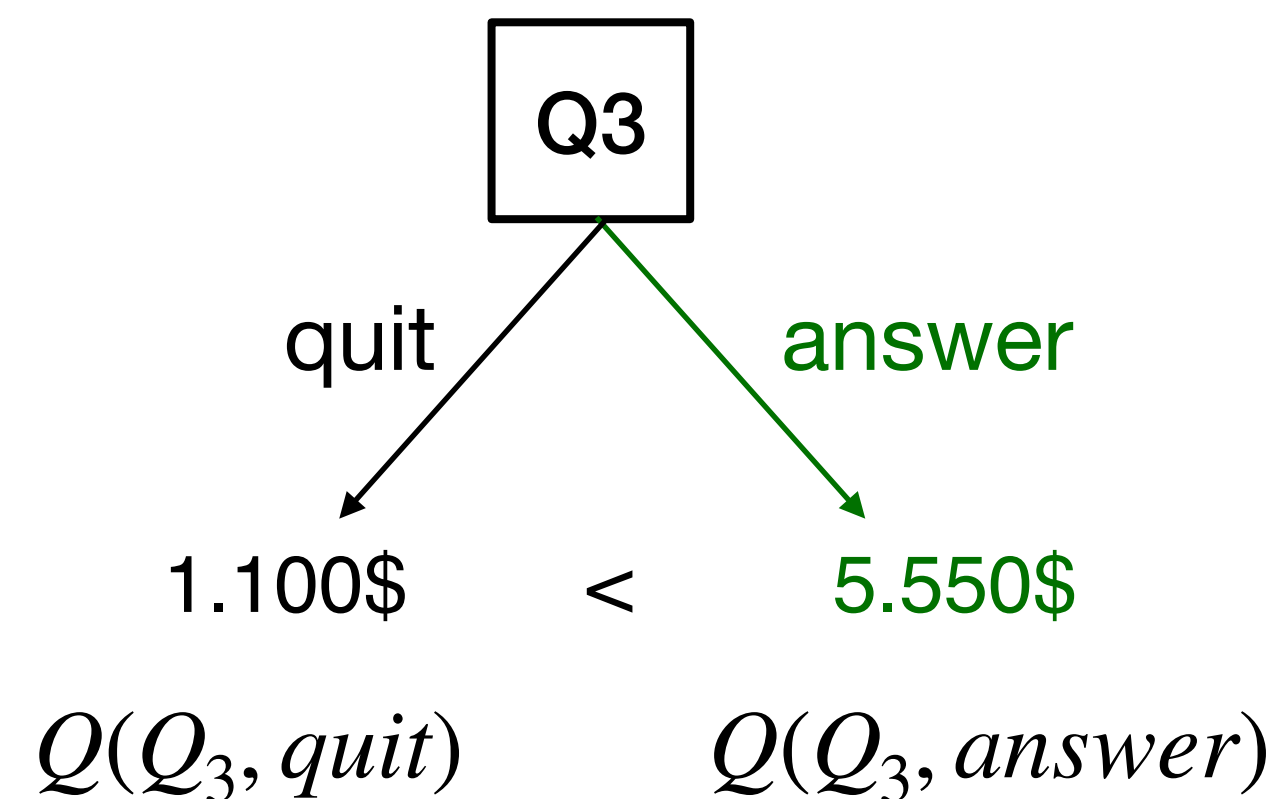
$$Q(Q_3, \text{answer}) = R(Q_3, \text{answer}) + P(Q_4 | \text{answer}, Q_3) V^*(Q_4) + P(\text{home} | \text{answer}, Q_3) V^*(\text{home})$$

# Q-Value

- Expected value for executing an action in a given state assuming that you will act optimally in the future

$$Q(s, a) = R(s, a) + \gamma \sum_{s'} P(s' | s, a) V^*(s')$$

- Select the optimal action as the one that maximizes the Q-value



$$\pi^*(Q_3) = \operatorname{argmax}_a Q(Q_3, a) = answer \quad V^*(Q_3) = \max_a Q(s, a) = 5500$$

# Value Iteration

- Compute approximation of  $V^*(s)$  for every state  $s$
- Compute Q-values for every state  $s$  and action  $a$

$$Q(s, a) = R(s, a) + \gamma \sum_{s'} P(s' | s, a) V^*(s')$$

- Build the optimal policy as  $\pi^*(s) = \operatorname{argmax}_a Q(s, a)$  that maximizes the Q-value

# Value Iteration

## Algorithm

Initialize random values  $V^0$

$$\Delta \leftarrow 0$$

Repeat:

For  $s \in S$  :

For  $a \in A$

$$Q^k(s, a) \leftarrow R(s, a) + \gamma \sum_{s'} P(s' | s, a) V^{k-1}(s')$$

$$V^k(s) = \max_a Q^k(s, a)$$

$$\Delta \leftarrow \max(\Delta, |V^k(s) - V^{k-1}(s)|)$$

Until  $\Delta < \epsilon$

Return  $V^k \approx V^* \in \mathbb{R}^{|S|}$

Iterative algorithm that converges to the optimal value  $V^*$



# Value Iteration

## Algorithm

Initialize random values  $V^0$

$\Delta \leftarrow 0$

Repeat:

For  $s \in S$  :

For  $a \in A$

$$Q^k(s, a) \leftarrow R(s, a) + \gamma \sum_{s'} P(s' | s, a) V^{k-1}(s')$$

$$V^k(s) = \max_a Q^k(s, a)$$

$$\Delta \leftarrow \max(\Delta, |V^k(s) - V^{k-1}(s)|)$$

Until  $\Delta < \epsilon$

Return  $V^k \approx V^* \in \mathbb{R}^{|S|}$

Vector of values  $V \in \mathbb{R}^{|S|}$ , one value for each state

# Value Iteration

## Algorithm

Initialize random values  $V^0$

$$\Delta \leftarrow 0$$

Repeat:

For  $s \in S$  :

For  $a \in A$

$$Q^k(s, a) \leftarrow R(s, a) + \gamma \sum_{s'} P(s' | s, a) V^{k-1}(s')$$

$$V^k(s) = \max_a Q^k(s, a)$$

$$\Delta \leftarrow \max(\Delta, |V^k(s) - V^{k-1}(s)|)$$

Until  $\Delta < \epsilon$

Return  $V^k \approx V^* \in \mathbb{R}^{|S|}$

Updates the Q-values for each action  $a$  and state  $s$



# Value Iteration

## Algorithm

Initialize random values  $V^0$

$$\Delta \leftarrow 0$$

Repeat:

For  $s \in S$  :

For  $a \in A$

$$Q^k(s, a) \leftarrow R(s, a) + \gamma \sum_{s'} P(s' | s, a) V^{k-1}(s')$$

$$V^k(s) = \max_a Q^k(s, a)$$

$$\Delta \leftarrow \max(\Delta, |V^k(s) - V^{k-1}(s)|)$$

Until  $\Delta < \epsilon$

Return  $V^k \approx V^* \in \mathbb{R}^{|S|}$

Retain the maximum Q-value

# Value Iteration

## Algorithm

Initialize random values  $V^0$

$$\Delta \leftarrow 0$$

Repeat:

For  $s \in S$  :

For  $a \in A$

$$Q^k(s, a) \leftarrow R(s, a) + \gamma \sum_{s'} P(s' | s, a) V^{k-1}(s')$$

$$V^k(s) = \max_a Q^k(s, a)$$

$$\Delta \leftarrow \max(\Delta, |V^k(s) - V^{k-1}(s)|)$$

Until  $\Delta < \varepsilon$

Return  $V^k \approx V^* \in \mathbb{R}^{|S|}$

Repeat until the updates are small enough

# Value Iteration

## Algorithm

Initialize random values  $V^0$

$$\Delta \leftarrow 0$$

Repeat:

For  $s \in S$  :

For  $a \in A$

$$Q^k(s, a) \leftarrow R(s, a) + \gamma \sum_{s'} P(s' | s, a) V^{k-1}(s')$$

$$V^k(s) = \max_a Q^k(s, a)$$

$$\Delta \leftarrow \max(\Delta, |V^k(s) - V^{k-1}(s)|)$$

Until  $\Delta < \varepsilon$

Return  $V^k \approx V^* \in \mathbb{R}^{|S|}$

Return the optimal value  $V^*$

# Value Iteration

## Algorithm

Initialize random values  $V^0$   
 $\Delta \leftarrow 0$

Repeat:

  For  $s \in S$  :

    For  $a \in A$


$Q^k(s, a) \leftarrow R(s, a) + \gamma \sum_{s'} P(s' | s, a) V^{k-1}(s')$

$V^k(s) = \max_a Q^k(s, a)$

$\Delta \leftarrow \max(\Delta, |V^k(s) - V^{k-1}(s)|)$

Until  $\Delta < \epsilon$

Return  $V^k \approx V^* \in \mathbb{R}^{|S|}$



For  $s \in S$  :

  For  $a \in A$

$Q(s, a) = R(s, a) + \gamma \sum_{s'} P(s' | s, a) V^*(s')$

$\pi^*(s) = \operatorname{argmax}_a Q(s, a)$

From the optimal value  $V^*$  to the optimal policy  $\pi^*$

# Value Iteration

## Algorithm

Initialize random values  $V^0$

$$\Delta \leftarrow 0$$

Repeat:

For  $s \in S$ :

For  $a \in A$

$$Q^k(s, a) \leftarrow R(s, a) + \gamma \sum_{s'} P(s' | s, a) V^{k-1}(s')$$

$$V^k(s) = \max_a Q^k(s, a)$$

$$\Delta \leftarrow \max(\Delta, |V^k(s) - V^{k-1}(s)|)$$

Until  $\Delta < \epsilon$

Return  $V^k \approx V^* \in \mathbb{R}^{|S|}$

For  $s \in S$ :

For  $a \in A$

$$Q(s, a) = R(s, a) + \gamma \sum_{s'} P(s' | s, a) V^*(s')$$

$$\pi^*(s) = \operatorname{argmax}_a Q(s, a)$$

Assumption that transitions and rewards are known



# Planning vs Reinforcement Learning

## Planning

- The transitions and reward model are known
- Quiz show, robot navigation
- Use value iteration to find optimal policy to deploy afterwards



## Reinforcement learning

- No prior knowledge of transition or reward model
- Robot parkour, Atari Breakout
- Find an optimal policy **while** interacting with an unfamiliar environment



# Supervised Learning vs Reinforcement Learning

## Supervised Learning

- Examples of correct or incorrect behaviour
- Independent sample of experience
- Learn a model and use it afterwards

## Reinforcement learning

- Only reward for a sequence of actions tried
- Agent has partial control over the training data
- Agent learns on-line: it must maximize performance during learning





# Tabular Q-Learning

Estimate  $Q$ -value while interacting with the environment: MDP is unknown but can be sampled!

Table of q-values  $Q$  with shape  $|S| \times |A|$  (one  $Q(s, a)$  for every  $s, a$ )



# Tabular Q-Learning

Estimate  $Q$ -value while interacting with the environment: MDP is unknown but can be sampled!

Table of q-values  $Q$  with shape  $|S| \times |A|$  (one  $Q(s, a)$  for every  $s, a$ )

After observing a transition  $s, a, s', r$  update the table  $Q$

$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + \alpha(r + \gamma \max_{a'} Q(s', a'))$$

# Tabular Q-Learning

Estimate  $Q$ -value while interacting with the environment: MDP is unknown but can be sampled!

Table of q-values  $Q$  with shape  $|S| \times |A|$  (one  $Q(s, a)$  for every  $s, a$ )

After observing a transition  $s, a, s', r$  update the table  $Q$

$$Q(s, a) \leftarrow (1 - \alpha) \underbrace{Q(s, a)}_{\text{old value}} + \alpha \underbrace{(r + \gamma \max_{a'} Q(s', a'))}_{\text{update}}$$

# Tabular Q-Learning

Estimate  $Q$ -value while interacting with the environment: MDP is unknown but can be sampled!

Table of q-values  $Q$  with shape  $|S| \times |A|$  (one  $Q(s, a)$  for every  $s, a$ )

After observing a transition  $s, a, s', r$  update the table  $Q$

$$Q(s, a) \longleftarrow (1 - \alpha)Q(s, a) + \alpha(r + \gamma \max_{a'} Q(s', a'))$$

↓  
update

$$Q(s, a) = R(s, a) + \gamma \sum_{s'} P(s' | s, a) \max_{a'} Q(s', a')$$

# Tabular Q-Learning

Estimate  $Q$ -value while interacting with the environment: MDP is unknown but can be sampled!

Table of q-values  $Q$  with shape  $|S| \times |A|$  (one  $Q(s, a)$  for every  $s, a$ )

After observing a transition  $s, a, s', r$  update the table  $Q$

$$Q(s, a) \longleftarrow (1 - \alpha)Q(s, a) + \alpha(r + \gamma \max_{a'} Q(s', a')) = Q(s, a) + \underbrace{\alpha(r + \gamma \max_{a'} Q(s', a') - Q(s, a))}_{\text{Temporal difference TD error}}$$

Temporal average between the old value and the update weighted by the learning rate  $\alpha$

# Tabular Q-Learning

## Algorithm

Initialize constant Q-table  $Q(s, a) \leftarrow c$

Repeat

    Initialize the state  $s^0$

    Repeat:

        Learning agent chooses  $a^t = \max_a Q(s^t, a)$

        Perform one env step  $s^{t+1}, r^{t+1} \leftarrow \text{env}(s^t, a^t)$

        Update Q-table

$$Q(s^t, a^t) \leftarrow Q(s^t, a^t) + \alpha \left( r^{t+1} + \gamma \max_{a'} Q(s^{t+1}, a') - Q(s^t, a^t) \right)$$

    Until the episode is done

Until max number of steps

If all state-action pairs are visited often enough, tabular Q-learning converges to  $Q$ -value

# From Tabular to Deep RL

- The size of the Q table grows exponentially with the state and action dimensions
- Tabular RL can only work for small state-action spaces

# From Tabular to Deep RL

- The size of the Q table grows exponentially with the state and action dimensions
- Tabular RL can only work for small state-action spaces
- How can the robot learn parkour?
  - No transitions and reward models available
  - High dimensional state space



Robot Parkour vision of an obstacle

$$|S| = 48 \times 64 \times 255 \approx 10^6$$



# From Tabular to Deep RL

- The size of the Q table grows exponentially with the state and action dimensions
- Tabular RL can only work for small state-action spaces
- How can the robot learn parkour?
  - No transitions and reward models available
  - High dimensional state space
- Idea: use a neural network to approximate  $Q$

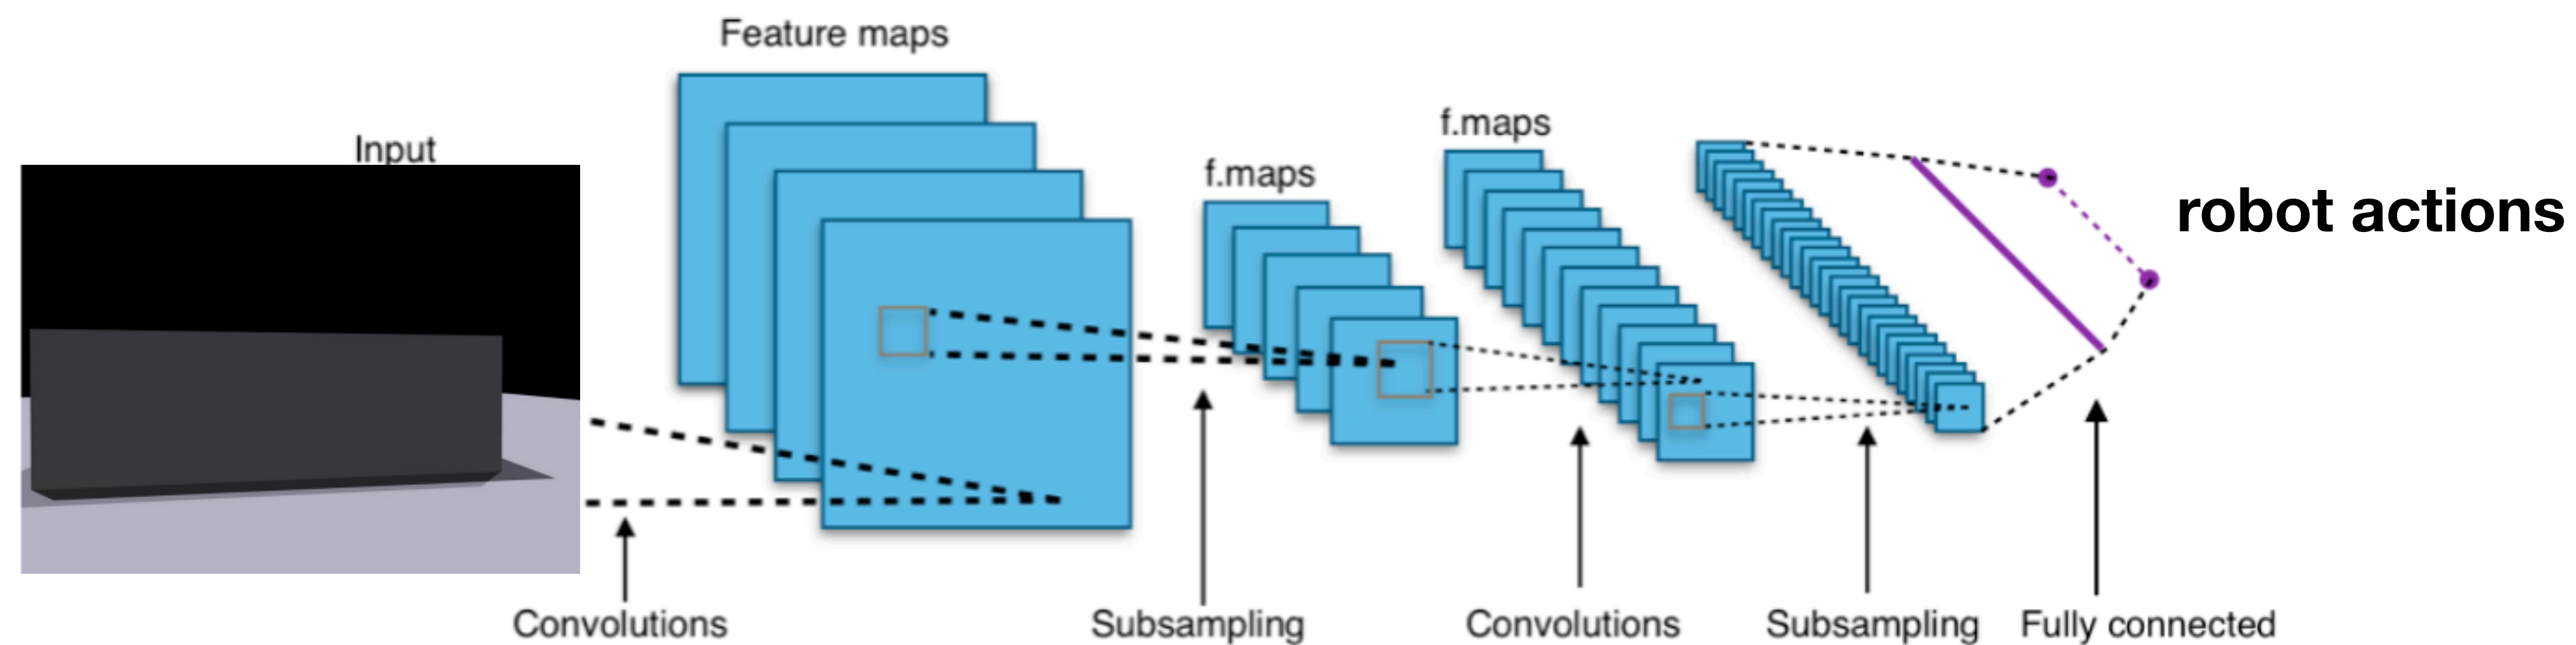


Robot Parkour vision of an obstacle

$$|S| = 48 \times 64 \times 255 \approx 10^6$$

# Deep Q-Network

Convolutional neural networks  $q(s; \theta) \in \mathbb{R}^{|A|}$ : output Q-value approximation per action



Gradient descent on mean-squared TD error

$$\theta \longleftarrow \theta - \alpha \nabla \left( \underbrace{r + \gamma \max_{a'} q(s'; \theta) - q(s; \theta)_a}_{\text{TD-error}} \right)^2$$

# Summary

## Topics

- Sequential decision making: agents, environment, reward...
- MDPs
- Planning and Value Iteration
- Reinforcement learning and Tabular Q-learning

## Reading material

- Reinforcement Learning: An Introduction - Chapter 3
- Wendelin Böhmer - Learning to interact, tabular Q-learning

